

## Introduction

This document details the instruction set for the PIC32A family of MCU devices. This is intended to guide development in the native assembly language for optimization and direct control over instruction execution. Assembly code can be used to simplify and accelerate time-sensitive applications including embedded control loops and data processing. More information about these devices and products, along with corresponding technical documentation, is available on the Microchip website ([www.microchip.com](http://www.microchip.com)).

This manual is a software developer's reference that describes the instruction set in detail and also provides general information to assist the development of software for the PIC32A device families.

This manual does not include detailed information about the core, peripherals, system integration or device-specific information. The user should refer to the device specific data sheet for information on the core, peripherals and system integration. The data sheet will also include information on:

- Device memory map
- Device pinout and packaging details
- Device electrical specifications
- List of peripherals included on the device

Code examples/snippets are given throughout this manual. These examples are valid for the devices from the PIC32A families.

## Development Support

Microchip offers a wide range of development tools that allow users to efficiently develop and debug application code. Microchip's development tools can be broken down into four categories:

- Code Generation
- Hardware/Software Debug
- Device Programmer
- Product Evaluation Boards

Information about the latest tools, product briefs and user guides can be obtained from the Microchip website ([www.microchip.com](http://www.microchip.com)) or from a local Microchip sales office.

Microchip offers additional resources to speed up the development cycle. These include:

- Application Notes
- Reference Designs
- Microchip Website
- Local Sales Offices with Field Application Support
- Corporate Support Line

The Microchip website also lists other sites that may serve as useful references.

## Style and Symbol Conventions

Throughout this document, certain style and font format conventions are used. [Table 1](#) provides a description of the conventions used in this document.

**Table 1.** Document Conventions

Symbol or Term	Description
set	To force a bit/register to a value of logic '1'.
clear	To force a bit/register to a value of logic '0'.
Reset	1. To force a register/bit to its default state. 2. A condition in which the device places itself after a device reset occurs. Some bits will be forced to '0' (such as Interrupt Enable bits), while others will be forced to '1' (such as the I/O Data Direction bits).
0xnnnn	Designates the number 'nnnn' in the hexadecimal number system. These conventions are used in the code examples. For example, 0x013F or 0xA800.
: (colon)	Used to specify a range or the concatenation of registers/bits/pins. One example is ACCAU:ACCAH:ACCAL, which is the concatenation of three registers to form the 72-bit Accumulator. Concatenation order (left-right) usually specifies a positional relationship (MSb to LSb, higher to lower).
[]	Specifies bit locations in a particular register. One example is SR[7:5] (or IPL[2:0]), which specifies the register and associated bits or bit locations.
LSb, MSb	Indicates the Least Significant or Most Significant bit in a field.
LSB, MSB	Indicates the Least/Most Significant Byte in a field of bits.
lsw, msw	Indicates the least/most significant word in a field of bits
Courier New Font	Used for code examples, binary numbers and for instruction mnemonics in the text.
<i>Times New Roman Font, Italic</i>	Used for equations and variables.
<b><i>Times New Roman Font, Bold Italic</i></b>	Used in explanatory text for items called out from a figure, equation or example.
<b>Note:</b>	A Note presents information that we want to re-emphasize, either to help the user avoid a common pitfall or to make the user aware of operating differences between some device family members. A note can be in a box, or when used in a table or figure, it is located at the bottom of the table or figure.

## Instruction Set Symbols

The summary tables in [Instruction Set Overview](#) and [Instruction Set Summary Tables](#), and the instruction descriptions in [Instruction Descriptions](#) utilize the symbols shown in [Table 2](#).

**Table 2.** Symbols Used in Instruction Summary Tables and Descriptions

Symbol <sup>(1)</sup>	Description
{ }	Optional field or operation
[text]	The location addressed by text
(text)	The contents of text
#text	The literal defined by text
{label:}	Optional label name
[n:m]	Register bit field
.l	32-bit Long Word mode selection
.b	8-bit Byte mode selection
.sl	24-bit (literal) Word mode selection
.v	Destination data value select (MAXABW, MINABW and FLIMW)
.w	16-bit Word mode selection (default)
AWB	Accumulator write back destination address register
bit3	3-bit bit selection field (used in byte addressed instructions) (0:7)
bit4	4-bit bit selection field (used in word addressed instructions) (0:15)
C, N, OV, Z	ALU status bits: Carry, Digit Carry, Negative, Overflow, Zero
d	File register destination (W0, none)
Expr	Absolute address, label or expression (resolved by the linker)
f	File register address (0x0000:0xFFFF) or (0x00000:0xFFFF) (addressable space varies depending upon instruction class)
Fd <sup>(2)</sup>	One of 32 FPU destination data registers (F0:F31) (Register Direct)
Fs <sup>(2)</sup>	One of 32 FPU source data registers (F0:F31) (Register Direct)
FSR, FSRH, FCR, FEAR <sup>(2)</sup>	FPU special (control & status) coprocessor registers (Register Direct)
label	Translates to a literal representing the location of the label name
lit1	1-bit unsigned literal (0:1)
lit3	3-bit unsigned literal (0:7)
lit5	5-bit unsigned literal (0:31)
lit6	6-bit unsigned literal (0:63)
lit8	8-bit unsigned literal (0:255)
lit16	16-bit unsigned literal (0:65535)
lit24	24-bit unsigned literal (0:1677215; LSB must be 0 if an address)
lit32	32-bit unsigned literal (0:4294967295)
none	Field does not require an entry and may be blank
OA, OB, SA, SB	DSC status bits: ACCA Overflow, ACCB Overflow, ACCA Saturate, ACCB Saturate
PC	Program Counter
Wd	Destination Working register
Rnd	Instruction rounding mode [E, Z, P, N]
Ws	Source Working register
Slit6	Signed 6-bit literal (-32:31)
Slit7	Signed 7-bit literal (-64:63)
Slit8	Signed 8-bit literal (-128:127)
Slit20	Signed 20-bit literal (-524288:524287)

**Table 2.** Symbols Used in Instruction Summary Tables and Descriptions (continued)

Symbol <sup>(1)</sup>	Description
SR	Status Register
text1 ∈ {text2, {text3,...}}	text1 must be in the set of text2, text3, ...
v	Selects MULxxx operand data types
Wb	Base Working register
Wd	Destination Working register
Wm	One of 16 Working registers (W0:W15)
Wn	Both source and destination Working register (W0:W15)
Wnd	One of 16 destination Working registers
Wns	One of 16 source Working registers
Wm * Wm	Multiplicand and Multiplier Working register for Square instructions
Wm * Wn	Multiplicand and Multiplier W register for DSP instructions
Ws	Source Working register
Wx	X data space fetch address register for DSP instructions
Wy	Y data space fetch address register for DSP instructions

**Notes:**

1. The range of each symbol is instruction-dependent. Refer to [Instruction Descriptions](#) for the specific instruction range.
2. Only applicable when the FPU coprocessor is present.

# Table of Contents

Introduction.....	1
Development Support.....	2
Style and Symbol Conventions.....	2
Instruction Set Symbols.....	3
1. PIC32A Core Architecture Overview.....	7
1.1. Features Specific to the PIC32A Core.....	7
1.2. Floating Point Unit (FPU) Overview.....	8
1.3. Programmer's Model.....	12
1.4. Working Register Array.....	13
1.5. Software Stack Frame Pointer.....	13
1.6. Software Stack Pointer.....	13
1.7. Stack Pointer Limit Register (SPLIM).....	13
1.8. Accumulator A and Accumulator B.....	13
1.9. Program Counter.....	13
1.10. RCOUNT Register.....	13
1.11. STATUS Register.....	13
1.12. Core Control Register.....	15
1.13. CPU STATUS Register.....	16
1.14. Core Control Register.....	19
2. Instruction Set Overview.....	21
2.1. Multicycle Instructions.....	21
2.2. Multiword Instructions.....	22
2.3. Instruction Set Summary Tables.....	22
3. Instruction Set Details.....	33
3.1. Data Addressing Modes.....	33
3.2. Data Addressing Mode Tree.....	40
3.3. Program Addressing Modes.....	40
3.4. Instruction Stalls.....	41
3.5. Byte Operations.....	42
3.6. Word Move Operations.....	44
3.7. Using 16-Bit Literal Operands.....	47
3.8. Bit Field Insert/Extract Instructions.....	47
3.9. Software Stack Pointer and Frame Pointer.....	48
3.10. Conditional Branch Instructions.....	52
3.11. Z Status Bit.....	54
3.12. DSP Data Formats.....	54
3.13. Accumulator Usage.....	56
3.14. Accumulator Access.....	57
3.15. DSP MAC Instructions.....	58
3.16. DSP Accumulator Instructions.....	60
3.17. Scaling Data with the FBCL Instruction.....	61
3.18. Data Range Limit Instructions.....	62
3.19. Normalizing the Accumulator with the NORM Instruction.....	62

4.	Instruction Descriptions.....	63
4.1.	Instruction Symbols.....	63
4.2.	Instruction Encoding Field Descriptors Introduction.....	63
4.3.	Instruction Description Example.....	67
4.4.	Instruction Descriptions (A to BZ).....	68
4.5.	Instruction Descriptions (C to DTB).....	146
4.6.	Instruction Descriptions (E to MULUU).....	175
4.7.	Instruction Descriptions (N to XORWF).....	268
4.8.	FPU Instruction Encoding and Opcode Field Description .....	347
4.9.	Floating-Point Instruction Description.....	348
5.	Built-In Functions.....	385
5.1.	Introduction.....	385
5.2.	Built-In Functions List.....	386
6.	Reference.....	403
6.1.	Instruction Bit Map.....	403
6.2.	Instruction Set Summary Table.....	403
7.	Revision History.....	424
	Microchip Information.....	425
	Trademarks.....	425
	Legal Notice.....	425
	Microchip Devices Code Protection Feature.....	425

# 1. PIC32A Core Architecture Overview

This section provides an overview of the features and capabilities of the PIC32A family of devices.

## 1.1. Features Specific to the PIC32A Core

The PIC32A devices are 32-bit (data) modified Harvard architecture, with a 5-stage instruction pipeline and a single-phase clock design, featuring an enhanced instructions set. The core has a 32-bit instruction word with an 8-bit opcode field. The Program Counter (PC) is 24 bits wide and addresses up to 4M x 32 bits of user program memory space. An instruction prefetch mechanism is used to help maintain throughput and provides predictable execution. The majority of instructions execute in a single cycle.

### 1.1.1. Registers

The PIC32A devices have sixteen 32-bit Working registers. Each of the Working registers can act as a data, address or offset register. The 16th Working register (W15) operates as a Software Stack Pointer (SSP) for interrupts and function calls.

### 1.1.2. Instruction Set

The instruction set is almost identical for the 16-bit MCU and DSC architectures. The instruction set includes many addressing modes and was designed for optimum C compiler efficiency.

### 1.1.3. Addressing Modes

The core supports Inherent (no operand), Relative, Literal, Memory Direct, Register Direct, Register Indirect and Register Offset Addressing modes. Each instruction is associated with a predefined addressing mode group, depending upon its functional requirements. As many as seven addressing modes are supported for each instruction.

For most instructions, the CPU is capable of executing a data (or program data) memory read, a Working register (data) read, a data memory write and a program (instruction) memory read per instruction cycle. As a result, 3-operand instructions can be supported, allowing  $A + B = C$  operations to be executed in a single cycle.

### 1.1.4. Arithmetic and Logic Unit

A high-speed, 33-bit by 33-bit multiplier is included to significantly enhance the core's arithmetic capability and throughput. The multiplier supports signed and unsigned, as well as 32-bit by 32-bit, or 16-bit by 16-bit integer multiplication. All multiply instructions execute in a single cycle.

The Arithmetic Logic Unit (ALU) is enhanced with integer divide assist hardware that supports an iterative non-restoring divide algorithm. It operates in conjunction with the REPEAT instruction looping mechanism and a selection of iterative divide instructions to support 32/32-bit, 32/16-bit, and 16/16-bit signed and unsigned, integer and fractional divide operations.

### 1.1.5. Exception Processing

The PIC32A devices have a vectored exception scheme with support for up to eight sources of non-maskable traps and up to 502 interrupt sources. Each interrupt source can be assigned to one of seven priority levels.

### 1.1.6. MCU Multiplications With 64-Bit Result

32x32-bit MUL instructions include an option to store the product in a single 32-bit Working register rather than a pair of registers. This feature helps free up a register for other purposes in cases where the numbers being multiplied are small in magnitude and therefore expected to provide a 16-bit result. See the individual MUL instruction descriptions in [Instruction Descriptions](#) for more details.

### 1.1.7. DSP Context Switch Support

DSP Overflow and Saturation Status bits are writable. This allows the state of the DSP engine to be efficiently saved and restored while switching between DSP tasks. See [DSP ALU Status Bits](#) for more details on DSP Status bits. There are also seven additional sets of CORCON, DSP Accumulators A and B for fast context switching; each set is inherently assigned to a respective IPL.

Any writes to CORCON at context-0 will be replicated across all DSP-related bits within CORCON registers mapped to contexts 1–7, until a write is made to the CORCON within one of the contexts 1–7 for the first time. A reset will re-enable CORCON write replication.

For instance, if the CPU IPL changes to IPL3 and the user writes to the associated CORCON (that is, CORCON of context-3), any subsequent writes to CORCON from context-0 will no longer affect CORCON from context-3. This applies to any CORCON register that was written from within its context. In contrast, further writes to CORCON at context-0 will continue to replicate over all CORCON registers that haven't been written to from within their contexts.

The code can assume that the DSP CORCON state remains consistent across all CPU register contexts following a reset and background initialization. In many applications, all DSP functions will utilize the same DSP engine configuration, making consistent initialization beneficial by eliminating the need to set up CORCON for each context individually. However, if a specific context requires a unique DSP engine configuration, it can initialize it as needed and subsequently prevent any background changes to CORCON from affecting the local CORCON configuration.

### 1.1.8. DSP Instruction Class

The DSP class of instructions are seamlessly integrated into the architecture and will execute from a single execution unit.

### 1.1.9. Data Space Addressing

The data space is split into two blocks, referred to as X and Y data memory. Each memory block has its own independent Address Generation Unit (AGU). The MCU class of instructions operates solely through the X memory AGU, which accesses the entire memory map as one linear data space. The DSP dual source class of instructions operates through the X and Y AGUs, which splits the data address space into two parts. The X and Y data space boundary is arbitrary and device-specific.

### 1.1.10. Modulo and Bit-Reversed Addressing

Overhead-free circular buffers (Modulo Addressing) are supported in both X and Y address spaces. The Modulo Addressing removes the software boundary checking overhead for DSP algorithms. Furthermore, the X AGU Circular Addressing can be used with any of the MCU class of instructions. The X AGU also supports Bit-Reversed Addressing to greatly simplify input or output data reordering for radix-2 FFT algorithms.

### 1.1.11. DSP Engine

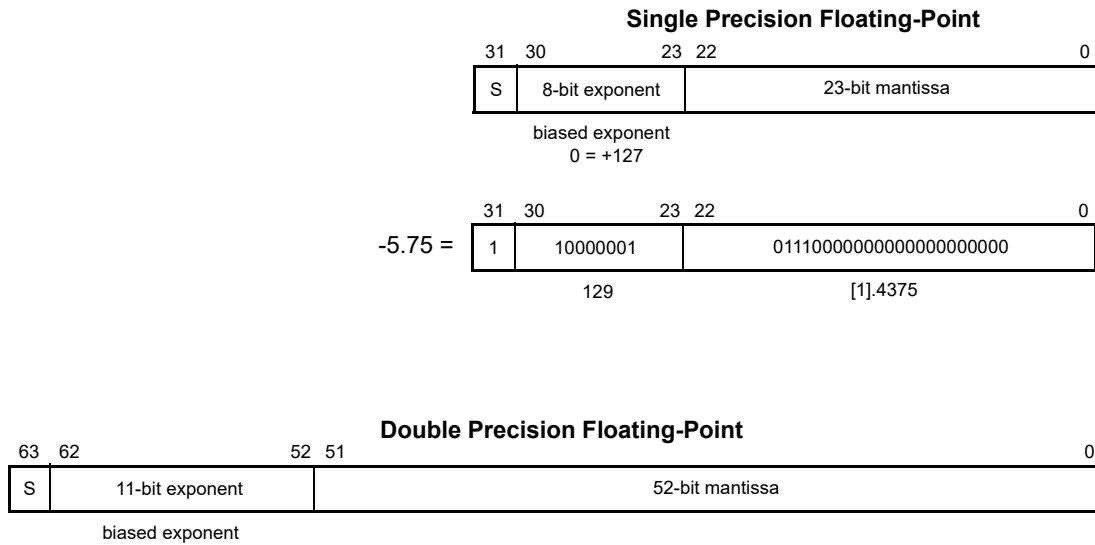
The DSP engine features a high-speed 33-bit by 33-bit multiplier, a 72-bit ALU, two 72-bit saturating accumulators and a 72-bit bidirectional barrel shifter. The barrel shifter is capable of shifting a 72-bit value up to 32 bits right or up to 32 bits left in a single cycle. The DSP instructions operate seamlessly with all other instructions and have been designed for optimal real-time performance. The MAC instruction and other associated instructions can concurrently fetch two data operands from memory while multiplying two Working registers. This requires that the data space be split for these instructions and linear for all others. This is achieved in a transparent and flexible manner through dedicating certain Working registers to each address space.

## 1.2. Floating Point Unit (FPU) Overview

The IEEE standard for floating-point arithmetic (IEEE 754-2008) specifies the floating-point data formats as shown in [Figure 1-1](#), which are comprised of a Sign bit, an exponent value and a (fractional) mantissa value. The PIC32A Floating-Point Unit (FPU) supports both single precision (32-bit, SP) and double precision (64-bit, DP) operations for most instructions.

To avoid the need for another Sign bit in the exponent, the IEEE floating-point format exponent is biased by 127 (SP) or 1023 (DP). Consequently, for any datum, the required IEEE exponent value = datum exponent + bias. In addition, the '1' to the left of the Most Significant bit (MSb) of the mantissa is implied for all numbers except subnormal numbers and is consequently referred to as the leading bit convention or "hidden bit." Therefore, the mantissa is a fractional value with an implied integer value of [1].

**Figure 1-1.** IEEE Floating-Point Data Formats and Single Precision Example



$$(-1)^S \times [1].m_{\text{base}2} \times 2^{(e-\text{bias})}$$

where:

- 'S' indicates the sign of the number (same values as a signed integer value)
- 'e' represents the exponent value
- 'm' represents the fractional mantissa value
- 'bias' is 127 (SP) or 1023 (DP)

For example,  $-5.75 = -(1.4375 \times 2^2)$ . In IEEE SP format, this would be represented as:

$$(-1)^1 \times [1].4375 \times 2^{(129-127)}$$

or (as shown in [Figure 1-1](#)):

$$S = 1, \text{exponent} = 129_{10}, \text{mantissa} = [1].4375_{10}$$

or: 0xC0B8 0000

### 1.2.1. Floating-Point Unit Registers

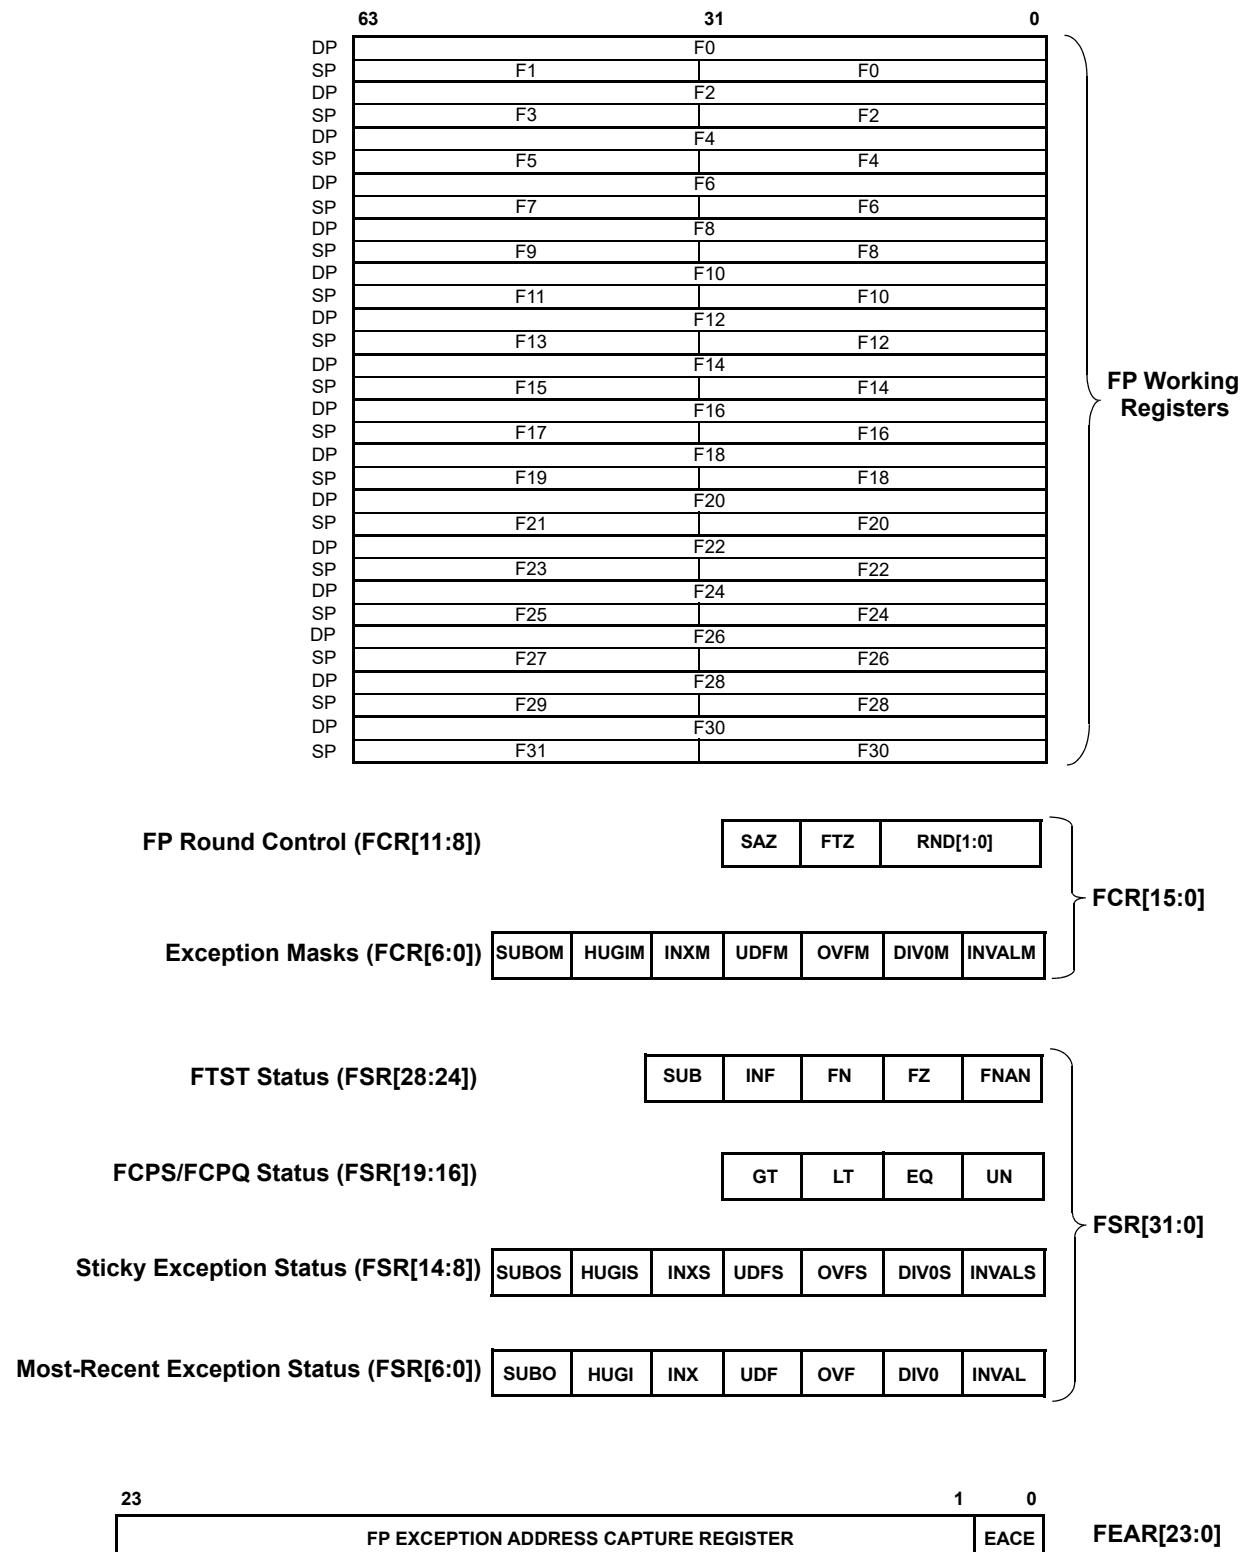
The PIC32A Floating-Point Unit (FPU) provides a large set of Working registers (F-reg):

- 32 x 32-bit (Single Precision, F0 ... F31) or
- 16 x 64-bit (Double Precision, F0, F2 ... F28, F30) or
- A mix of the two sizes aligned as shown in [Figure 1-2](#).

In addition to the F-reg, status (FSR) and control (FCR) registers are also supported as shown in [Figure 1-2](#):

- FSR (FPU Status Register, 32-bit): Holds the status of retired floating-point instructions:

- FSR[6:0]: Instruction "most-recent" exception status
- FSR[14:8]: Instruction "sticky" exception status
- FSR[19:16]: CPS/CPQ instruction status
- FSR[28:24]: FTST instruction status
- FCR (FPU Control Register, 16-bit):
  - FCR[6:0]: Exception mask control
  - FCR[9:8]: Rounding mode control
  - FCR[10]: Subnormal result "Flush to Zero" (FTZ) control
  - FCR[11]: Subnormal operand "Subnormals are Zeros" (SAZ) control
- FEAR: (FPU Exception Address Capture Register, 24-bit):
  - Holds the address of the first instruction encountered that causes an exception. All subsequent instructions in the FPU pipeline that subsequently retire will not affect the FEAR, even if they too generate exceptions.

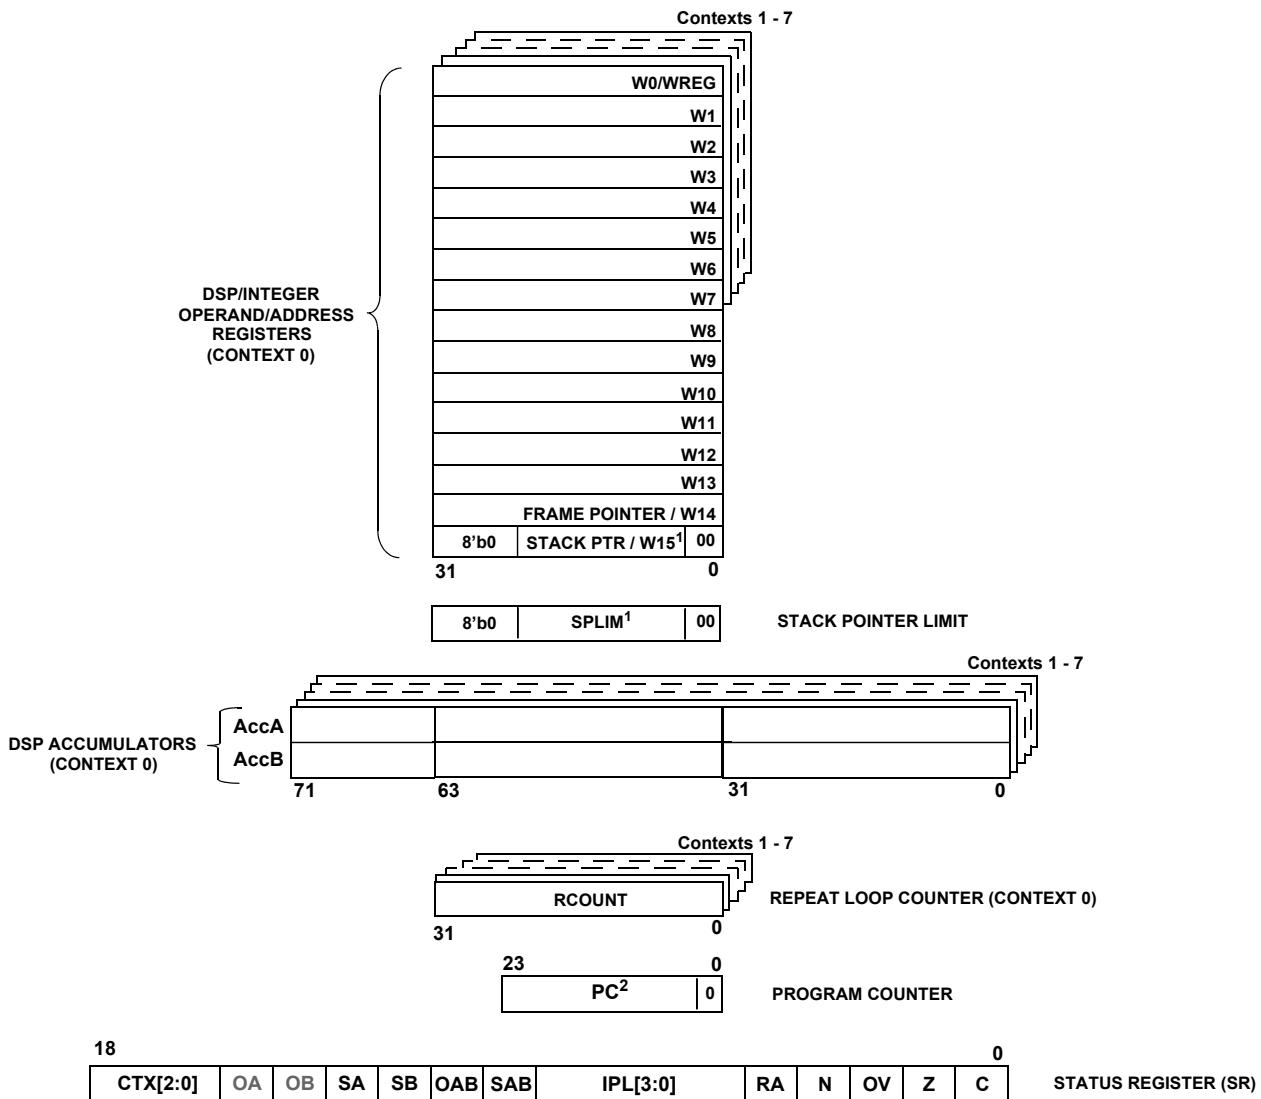
**Figure 1-2.** FPU Programmer's Model

**Note:** Only a single register context shown.

### 1.3. Programmer's Model

Figure 1-3 shows the programmer's model diagrams for the PIC32A family of devices.

**Figure 1-3. PIC32A Programmer's Model**



**Notes:**

1. W15[1:0] and SPLIM[1:0] always = 0b00.
2. PC[0] always = 0b0.

**Table 1-1. Programmer's Model Register Descriptions**

Register	Description
CORCON	CPU Core Configuration register
PC	24-Bit Program Counter
RCOUNT	REPEAT Loop Counter register
SPLIM	Stack Pointer Limit Value register
SR	ALU and DSP Engine STATUS Register
W0-W15	Working register array
ACCA, ACCB	72-Bit DSP Accumulators

## 1.4. Working Register Array

The 16 Working (W) registers can function as data, address or offset registers. The function of a W register is determined by the instruction that accesses it.

Byte instructions, which target the Working register array, only affect the Least Significant Byte (LSB) of the target register.

## 1.5. Software Stack Frame Pointer

A frame is a user-defined section of memory in the stack, used by a function to allocate memory for local variables. W14 has been assigned for use as a Stack Frame Pointer with the link (`LNK`) and unlink (`ULNK`) instructions. However, if a Stack Frame Pointer and the `LNK` and `ULNK` instructions are not used, W14 can be used by any instruction in the same manner as all other W registers. See [Software Stack Frame Pointer](#) for detailed information about the Frame Pointer.

## 1.6. Software Stack Pointer

W15 serves as a dedicated Software Stack Pointer and will be automatically modified by function calls, exception processing and returns. However, W15 can be referenced by any instruction in the same manner as all other W registers. This simplifies reading, writing and manipulating the Stack Pointer. Refer to [Software Stack Pointer](#) for detailed information about the Stack Pointer.

## 1.7. Stack Pointer Limit Register (SPLIM)

The SPLIM is a 32-bit register associated with the Stack Pointer. It is used to prevent the Stack Pointer from overflowing and accessing memory beyond the user allocated region of stack memory. Refer to [Stack Pointer Overflow](#) for detailed information about the SPLIM.

## 1.8. Accumulator A and Accumulator B

Accumulator A (ACCA) and Accumulator B (ACCB) are 72-bit wide registers utilized by DSP instructions to perform mathematical and shifting operations.

Accumulator A and Accumulator B can also be used as destination registers in MCU MUL.xx instructions. This helps reduce the execution time of extended precision arithmetic operations.

Refer to [Figure 3-13](#) for details on using ACCA and ACCB.

## 1.9. Program Counter

The Program Counter (PC) is 24 bits wide. Instructions are addressed in the 4M x 32-bit user program memory space by PC[23:1], where PC[0] is always set to '0' to maintain instruction word alignment and provide compatibility with Data Space Addressing. This means that during normal instruction execution, the PC increments by two, four or eight during sequential 16-bit, 32-bit or 64-bit instruction execution, respectively.

.

## 1.10. RCOUNT Register

The 32-bit RCOUNT register contains the loop counter for the REPEAT instruction. When a REPEAT instruction is executed, RCOUNT is loaded with the repeat count of the instruction, either "lit20" for the "REPEAT #lit20" instruction, "lit5" for the "REPEAT #lit5" instruction or Wn register for the "REPEAT Wn" instruction. The REPEAT loop will be executed RCOUNT + 1 time.

**Note:** If a REPEAT loop is executing and gets interrupted, RCOUNT may be cleared by the Interrupt Service Routine (ISR) to break out of the REPEAT loop when the foreground code is re-entered.

## 1.11. STATUS Register

The 32-bit STATUS Register maintains status information for the instructions which have been executed most recently. Operation Status bits exist for MCU operations, loop operations and DSP operations. Additionally, the STATUS Register contains the CPU Interrupt Priority Level bits, IPL[2:0],

which are used as a context identifier and for interrupt processing. See [SR](#) for more detailed information.

### 1.11.1. MCU ALU Status Bits

The MCU operation Status bits are either affected or used by the majority of instructions in the instruction set. Most of the logic, math, rotateshift and bit instructions modify the MCU Status bits after execution, and the conditional branch instructions use the state of individual Status bits to determine the flow of program execution. All conditional branch instructions are listed in [Conditional Branch Instructions](#).

The Carry (C), Zero (Z), Overflow (OV) and Negative (N) bits show the immediate status of the MCU ALU by indicating whether an operation has resulted in a Carry, Zero, Overflow or Negative result. When a subtract operation is performed, the C flag is used as a Borrow flag.

The Z Status bit is useful for extended precision arithmetic. The Z Status bit functions like a normal Z flag for all instructions except those that use a carry or borrow input (ADDC, CPB, SUBB and SUBBR). See [Z Status Bit](#) for more detailed information.

**Note:**

1. All MCU bits are stacked during exception processing (see [Software Stack Pointer](#)).

### 1.11.2. REPEAT Loop Active (RA) Status Bit

The REPEAT Loop Active bit (RA) is used to indicate when looping is active. The RA flag indicates that a REPEAT instruction is being executed, and it is only affected by the REPEAT instructions. The RA flag is set to '1' when the instruction being repeated begins execution, and it is cleared when the instruction being repeated completes execution for the last time.

Since the RA flag is also read-only, it may not be directly cleared. However, if a REPEAT or its target instruction is interrupted, the Interrupt Service Routine may clear the RA flag, which resides on the stack. This action will disable looping once program execution returns from the Interrupt Service Routine because the restored RA will be '0'.

### 1.11.3. DSP ALU Status Bits

The high byte of the STATUS Register is used by the DSP class of instructions and it is modified when data passes through one of the adders. It provides status information about overflow and saturation for both accumulators. The Saturate A, Saturate B, Overflow A and Overflow B (SA, SB, OA, OB) bits provide individual accumulator status, while the Saturate AB and Overflow AB (SAB, OAB) bits provide combined accumulator status. The SAB and OAB bits provide an efficient method for the software developer to check the register for saturation or overflow.

The OA and OB bits are used to indicate when an operation has generated an overflow into the Guard bits (bits 63 through 71) of the respective accumulator. This condition can only occur when the processor is in Super Saturation mode or if saturation is disabled. It indicates that the operation has generated a number which cannot be represented with the lower 62 bits of the accumulator.

The SA and SB bits are used to indicate when an operation has generated an overflow out of the MSb of the respective accumulator. The SA and SB bits are active, regardless of the Saturation mode (Disabled, Normal or Super) and may be considered "sticky." Namely, once the SA or SB bit is set to '1', it can only be cleared by software, regardless of subsequent DSP operations. When it is required, the `BCLR` instruction can be used to clear the SA or SB bit.

In addition, the SA and SB bits can be set by software, enabling efficient context state switching.

For convenience, the OA and OB bits are logically ORed together to form the OAB flag, and the SA and SB bits are logically ORed to form the SAB flag. These cumulative Status bits provide efficient overflow and saturation checking when an algorithm is implemented. Instead of interrogating the OA and OB bits independently for arithmetic overflows, a single check of OAB can be performed.

Likewise, when checking for saturation, SAB may be examined instead of checking both the SA and SB bits. Note that clearing the SAB flag will clear both the SA and SB bits.

#### 1.11.4. Interrupt Priority Level Status Bits

The three Interrupt Priority Level (IPL) bits of the SRL (SR[7:5]) and the IPL3 bit (SR[8]) set the CPU's IPL, which is used for exception processing. Exceptions consist of interrupts and hardware traps. Interrupts have a user-defined priority level between 0 and 7, while traps have a fixed priority level between 8 and 15. The fourth Interrupt Priority Level bit, IPL3, is a special IPL bit that may only be read or cleared by the user. This bit is only set when a hardware trap is activated, and it is cleared after the trap is serviced.

The CPU's IPL identifies the lowest level exception which may interrupt the processor. The interrupt level of a pending exception must always be greater than the CPU's IPL for the CPU to process the exception. This means that if the IPL is 0, all exceptions at Priority Level 1 and above may interrupt the processor. If the IPL is 7, only hardware traps may interrupt the processor.

When an exception is serviced, the IPL is automatically set to the priority level of the exception being serviced, which will disable all exceptions of equal and lower priority. However, since the IPL field is read/write, one may modify the lower three bits of the IPL in an Interrupt Service Routine to control which exceptions may preempt the exception processing. Since the SRL is stacked during exception processing, the original IPL is always restored after the exception is serviced. If required, one may also prevent exceptions from nesting by setting the NSTDIS bit (INTCON1[15]).

### 1.12. Core Control Register

The Core Control register (CORCON) is used to set the configuration of the CPU.

In addition to setting CPU modes, the following features are available through the CORCON register:

- Sets the ACCA and ACCB saturation enable
- Sets the Data Space Write Saturation mode
- Sets the Accumulator Saturation and Rounding modes
- Sets the Multiplier mode for DSP operations

## 1.13. CPU STATUS Register

**Name:** SR

**Notes:**

1. This bit may be read or cleared (not set). Clearing this bit will clear SA and SB irrespective of the value simultaneously written to SA and/or SB.
2. IPL[2:0] become read only bits if NSTDIS(INTCON1[15])=1 (nesting disabled)..

**Legend:** C = Clearable bit

Bit	31	30	29	28	27	26	25	24
Access								
Reset								
Bit	23	22	21	20	19	18	17	16
Access	VF						CTX[2:0]	
Reset	R					R	R	R
Bit	15	14	13	12	11	10	9	8
Access	OA	OB	SA	SB	OAB	SAB		IPL3
Reset	R/W	R/W	R/W	R/W	R	R/C		R/C
Bit	7	6	5	4	3	2	1	0
Access		IPL[2:0]		RA	N	OV	Z	C
Reset	R/W	R/W	R/W	R	R/W	R/W	R/W	R/W

**Bit 23 – VF** Vector (Fetch) Fail Status

Value	Description
1	Indicates to the bus error handler that the source of the bus error is a vector fetch. The vector data read will be substituted with the contents of the Vector Fail Address (VFA) SFR.
0	Indicates to the bus error handler that the source of the bus error is not a vector fetch.

**Bits 18:16 – CTX[2:0]** Current (W register) Context Identifier

Identifies which W register context is currently in use by the CPU

Value	Description
111	Context 7 is currently in use
110	Context 6 is currently in use
101	Context 5 is currently in use
100	Context 4 is currently in use
011	Context 3 is currently in use
010	Context 2 is currently in use
001	Context 1 is currently in use
000	Context 0 is currently in use

**Bit 15 – OA** Accumulator A Fractional Overflow Status

Value	Description
1	Accumulator A fractional overflow has occurred (its contents can no longer be represented as a 1.31 fractional value)

Value	Description
0	Accumulator A not overflowed

**Bit 14 - OB** Accumulator B Fractional Overflow Status

Value	Description
1	Accumulator A fractional overflow has occurred (its contents can no longer be represented as a 1.31 fractional value)
0	Accumulator A not overflowed

**Bit 13 - SA** Accumulator A Saturation/Sign Overflow 'Sticky' Status

Value	Description
1	Accumulator A is saturated or has been saturated at some time, or has overflowed into bit 71 (if saturation is disabled)
0	Accumulator A is not saturated or has not overflowed into bit 71 (if saturation is disabled)

**Bit 12 - SB** Accumulator B Saturation/Sign Overflow 'Sticky' Status

Value	Description
1	Accumulator B is saturated or has been saturated at some time, or has overflowed into bit 71 (if saturation is disabled)
0	Accumulator B is not saturated or has not overflowed into bit 71 (if saturation is disabled)

**Bit 11 - OAB** OA || OB Combined Accumulator Fractional Overflow Status

Value	Description
1	Accumulators A or B fractional overflow has occurred (one or both of their contents can no longer be represented as a 1.31 fractional value)
0	Neither Accumulators A nor B have overflowed

**Bit 10 - SAB** SA || SB Combined Accumulator 'Sticky' Status<sup>(1)</sup>

Value	Description
1	Accumulators A or B are saturated or have been saturated at some time, or have overflowed into bit 71 (if saturation is disabled)
0	Neither Accumulator A nor B are saturated or have overflowed into bit 71 (if saturation is disabled)

**Bit 8 - IPL3** MSb (Most Significant bit) of CPU Priority Level Nibble

Value	Description
1	CPU Priority $\geq 8$ (trap exception underway)
0	CPU Priority $< 8$ (no trap exception underway)

**Bits 7:5 - IPL[2:0]** CPU Interrupt Priority Level Status<sup>(2)</sup>

Value	Description
111	All interrupts disabled
110	Level 7 interrupts enabled
101	Level 6 and 7 interrupts enabled
100	Level 5 through 7 interrupts enabled
011	Level 4 through 7 interrupts enabled
010	Level 3 through 7 interrupts enabled
001	Level 2 through 7 interrupts enabled
000	Level 1 through 7 interrupts enabled

**Bit 4 - RA** REPEAT Loop Active

Value	Description
1	REPEAT loop in progress
0	REPEAT loop not in progress

**Bit 3 - N** MCU ALU Negative

Value	Description
1	Result was negative
0	Result was non-negative (zero or positive)

**Bit 2 - OV** MCU ALU Overflow

This bit is used for signed arithmetic (2's complement). It indicates an overflow of the magnitude that causes the Sign bit to change state.

Value	Description
1	Overflow occurred for signed arithmetic (in this arithmetic operation)
0	No overflow occurred

**Bit 1 - Z** MCU ALU 'Sticky' Zero

Value	Description
1	An operation which effects the Z bit has set it at some time in the past
0	The most recent operation which effects the Z bit has cleared it (i.e., a non-zero result)

**Bit 0 - C** MCU ALU Carry/Borrow

Value	Description
1	A carry-out from the MSb of the result occurred
0	No carry-out from the MSb of the result occurred

## 1.14. Core Control Register

**Name:** CORCON

**Note:**

1. This bit has no effect if US = 1 (Unsigned mode).

Bit	31	30	29	28	27	26	25	24
Access								
Reset								
Bit	23	22	21	20	19	18	17	16
Access								
Reset								
Bit	15	14	13	12	11	10	9	8
Access				US				
Reset				R/W				
Bit	7	6	5	4	3	2	1	0
Access	SATA	SATB	SATDW	ACCSAT			RND	IF
Reset	R/W	R/W	R/W	R/W			R/W	R/W
	0	0	1	0			0	0

**Bit 12 - US** DSP Multiply Unsigned/Signed Control

Value	Description
1	DSP Engine Unsigned mode enabled
0	DSP Engine Signed mode enabled

**Bit 7 - SATA** ACCA Saturation Enable<sup>(1)</sup>

Value	Description
1	Accumulator A saturation is enabled
0	Accumulator A saturation is disabled

**Bit 6 - SATB** ACCB Saturation Enable<sup>(1)</sup>

Value	Description
1	Accumulator B saturation is enabled
0	Accumulator B saturation is disabled

**Bit 5 - SATDW** Data Space Write from DSP Engine Saturation Enable

Value	Description
1	Data Space write saturation is enabled
0	Data Space write saturation is disabled

**Bit 4 - ACCSAT** Accumulator Saturation Mode Select

Value	Description
1	9.31 Saturation (super saturation)
0	1.31 Saturation (normal saturation)

**Bit 1 – RND** Rounding Mode Select

Value	Description
1	Biased (conventional) rounding is enabled
0	Unbiased (convergent) rounding is enabled

**Bit 0 – IF** Integer or Fractional Multiplier Mode Select

Value	Description
1	Integer mode is enabled for DSP multiply
0	Fractional mode is enabled for DSP multiply

## 2. Instruction Set Overview

The PIC32A instruction set provides a broad suite of instructions that includes traditional microcontroller applications and DSP-based math-intensive applications. Since almost all of the functionality of the 16-bit MCU and DSC instruction set has been maintained, this hybrid instruction set allows an easy 32-bit migration path for users already familiar with the PIC microcontroller architecture and instructions.

Instructions can be grouped into the functional categories, as shown in [Table 2-1](#). [Table 2](#) defines the symbols used in the instruction summary tables. [Table 2-1](#) through [Table 2-13](#) define the syntax, description, storage and execution requirements for each instruction. Storage requirements are represented in 32-bit instruction words and execution requirements are represented in instruction cycles.

**Table 2-1.** Instruction Groups

Functional Group	Summary Table
Move Instructions	<a href="#">Table 2-3</a>
Math Instructions	<a href="#">Table 2-4</a>
Logic Instructions	<a href="#">Table 2-5</a>
Rotate/Shift Instructions	<a href="#">Table 2-6</a>
Bit Instructions	<a href="#">Table 2-7</a>
Compare/Skip and Compare/Branch Instructions	<a href="#">Table 2-8</a>
Program Flow Instructions	<a href="#">Table 2-9</a>
Shadow/Stack Instructions	<a href="#">Table 2-10</a>
Control Instructions	<a href="#">Table 2-11</a>
DSP Instructions	<a href="#">Table 2-12</a>
FPU Instructions	<a href="#">Table 2-13</a>

Several instructions have different addressing modes and execution flows, which require different instruction variants. For instance, there are up to seven unique `ADD` instructions and each instruction variant has its own instruction encoding. Instruction format descriptions and specific instruction operations are provided in [Instruction Descriptions](#). Additionally, a composite alphabetized instruction set table is provided in the [Reference](#) section.

### 2.1. Multicycle Instructions

As shown in the instruction summary tables, most instructions execute in a single cycle with the following exceptions:

- ASRM, LSRM, ED, EDAC and a few other instructions require two cycles to execute.
- Instructions, `DIVF`, `DIVFL`, `DIVU`, `DIVUL`, `DIVS`, and `DIVSL` are single-cycle instructions, which should be executed consecutive times as the target of a `REPEAT` instruction.
- Instructions that change the Program Counter require two cycles to execute. Instructions such as `CALL` also require two cycles to execute.
- `RETFIE`, `RETLW` and `RETURN` are a special case of an instruction that changes the Program Counter. These execute in three to four cycles, unless an exception is pending, and then they execute in two cycles.

The cycle count of program flow change instructions also depend on the status of PBU cache. The table below shows the cycle counts of these instructions under different scenarios.

**Table 2-2.** Program Flow Control and Instruction Execution Times

Instruction	Cycle Counts		
	Cache/ISB Hit	Cache/ISB Miss	Execution from RAM
Conditional Branch <sup>1</sup>	1 (3)	1 (6 to 9)	1 (3)
BRA	1	4-7	1
BRA W	5-8	5-8	2
CALL	1	4-7	1
CALL W	5-8	5-8	2
GOTO	1	4-7	1
GOTO W	5-8	5-8	2
RCALL	1	4-7	1
RCALL W	5-8	5-8	2
RETFIE	7-10	7-10	4
RETLW	6-9	6-9	3
RETURN	6-9	6-9	3

**Note:**

1. Branch taken execution times in ( ) brackets.

## 2.2. Multiword Instructions

As defined by [Table 2-3](#) through [Table 2-13](#), almost all instructions consume one instruction word (32 bits), with the exception of the `CALL` and `GOTO` instructions, which are program flow instructions listed in [Table 2-9](#). These instructions require two words of memory because their opcodes embed large literal operands.

## 2.3. Instruction Set Summary Tables

**Table 2-3.** Move Instructions

Assembled Syntax	Description	Words	Cycles
EXCH Wns, Wnd	Swap Wns with Wnd	1	2
MOV Ws, Wd	Move Ws to Wd	0.5/1	1
MOV.l lit32,Wnd	Move 32-bit unsigned literal to Wnd	2	2
MOV.sl lit24,Wnd	Move 24-bit unsigned literal to Wnd; 0 extend to 32-bits	1	1
MOV.w lit16,Wnd	Move 16-bit unsigned literal to Wnd; 0 extend to 32-bits	1	1
MOV.bz lit8,Wnd	Move 8-bit unsigned literal to Wnd; 0 extend to 32-bits	1	1
MOV.l [W15-lit7], Wnd [W14+slit7], Wnd	Move from system stack with literal offset to Wnd using SP or FP	0.5	1
MOV.l Wns, [W15-lit7] Wns, [W14+slit7]	Move from Wns to system stack with literal off-set using SP or FP	0.5	1
MOV.l f, Wnd	Move f to Wnd (Word or Long Word)(f < ~1MB)	1	1
MOV.b f, Wnd	Move f to Wnd (Byte)	1	1
MOV.l Wns, f	Move Wns to f (Word or Long Word)(f < ~1MB)	1	1
MOV.w Wns, f	Move Wns to f (Word or Long Word)(f > ~1MB)	2	2
MOV.b Wns, f	Move Wns to f (Byte)	1	1
MOV [Wns+Slit12], Wnd	Move [Wns+Slit12] to Wnd	1	1

**Table 2-3.** Move Instructions (continued)

Assembled Syntax	Description	Words	Cycles
MOV Wns, [Wnd+Slit12]	Move Wns to [Wnd+Slit12]	1	1
MOVIF.l CC, Wb, Wns, Wd	If CC == True Move W1 to [W15++] Else Move W2 to [W15++] where, CC -> Z, N, C, OV, GT, LT, GTU	1	1
MOVR.l	Move Ws to Wd with destination Bit Reversed addressing		
MOVS.l slit16, Wd	Move signed extended 16-bit literal to Wd	1	1
MOVS.b slit8, Wnd	Move 8-bit literal to Wd; no extension.	1	1
SWAP Wn	Wn = Word or byte swap Wn	1	1
TST f	Test f	1	1
TST f,Wnd	Test f and move f to Wnd	1	1

**Table 2-4.** Math Instructions

Assembled Syntax	Description	Words	Cycles
ADD f,Wn	f = f + Wn	1	1
ADD f,Wn,Wn	Wn = f + Wn	1	1
ADD.l lit5,Wn	Wn = Wn + lit5	0.5	1
ADD lit16,Wn	Wn = Wn + lit16	1	1
ADD Wb,Ws,Wd	Wd = Wb + Ws	0.5/1	1
ADD Wb,lit7,Wd	Wd = Wb + lit7(literal zero-extended)	1	1
ADDC f,Wn	f = f + Wn + (C)	1	1
ADDC f,Wn,Wn	Wn = f + Wn + (C)	1	1
ADDC lit16,Wn	Wn = Wn + lit16 + (C)	1	1
ADDC Wb,Ws,Wd	Wd = Wb + Ws + (C)	0.5/1	1
ADDC Wb,lit7,Wd	Wd = Wb + lit7 + (C)(literal zero-extended)	1	1
DEC f	f = f - 1	1	1
DEC f,Wd	W5 = f - 1	1	1
DEC Ws,Wd	Wd = Ws - 1	1	1
DEC2 f	f = f - 2	1	1
DEC2 f,Wd	W5 = f - 2	1	1
DEC2 Ws,Wd	Wd = Ws - 2	1	1
DIVF Wm/Wn	Interruptible Signed 16/16 or 32/16 Fractional Divide	1	1
DIVFL Wm/Wn	Interruptible Signed 32/32 Fractional Divide	1	1
DIVS.w Wm/Wn	Interruptible Signed 16/16-bit Integer Divide	1	1
DIVS.l Wm/Wn	Interruptible Signed 32/16-bit Integer Divide	1	1
DIVSL Wm/Wn	Interruptible Signed 32/32 Integer Divide	1	1
DIVU.w Wm/Wn	Interruptible Unsigned 16/16-bit Integer Divide	1	1
DIVU.l Wm/Wn	Interruptible Unsigned 32/16-bit Integer Divide	1	1
DIVUL Wm/Wn	Interruptible Unsigned 32/32 Integer Divide	1	1

**Table 2-4.** Math Instructions (continued)

Assembled Syntax	Description	Words	Cycles
FLIM Wb, Ws	Force Data (Upper and Lower) Range Limit without Limit Excess Result	1	1
FLIM Wb, Ws, Wd	Force Data (Upper and Lower) Range Limit with Limit Excess Flag (Wd=-1)	1	2
FLIM.V Wb, Ws, Wd	Force Data (Upper and Lower) Range Limit with Limit Excess Result	1	2
INC f	$f = f + 1$	1	1
INC f,Wd	$W5 = f + 1$	1	1
INC Ws,Wd	$Wd = Ws + 1$	1	1
INC2 f	$f = f + 2$	1	1
INC2 f,Wd	$W5 = f + 2$	1	1
INC2 Ws,Wd	$Wd = Ws + 2$	1	1
MULSS Wb,Ws,Wnd	{Wd}=signed(Wb) * signed(Ws)	0.5/1	1
MULSU Wb,Ws,Wnd	{Wd}=signed(Wb) * unsigned(Ws)	0.5/1	1
MULUS Wb,Ws,Wnd	{Wd}=unsigned(Wb) * signed(Ws)	0.5/1	1
MULUU Wb,Ws,Wnd	{Wd}=unsigned(Wb) * unsigned(Ws)	0.5/1	1
MULSU Wb,lit8,Wnd	{Wd}=signed(Wb) * unsigned(lit8)	1	1
MULUU Wb,lit8,Wnd	{Wd}=unsigned(Wb) * unsigned(lit8)	1	1
MULSS Wb,slit8,Wnd	{Wd}=signed(Wb) * signed(slit8)	1	1
MULUS Wb,slit8,Wnd	{Wd}=unsigned(Wb) * signed(slit8)	1	1
MUL f, Wn	$W2 = f * Wn$	1	1
SE Ws,Wnd	$Wd = \text{sign-extended } Ws$	0.5/1	1
SUB f,Wn	$f = f - Wn$	1	1
SUB f,Wn,Wn	$Wn = f - Wn$	1	1
SUB.l lit5,Wn	$Wn = Wn - \text{lit5}$	0.5	1
SUB lit16,Wn	$Wn = Wn - \text{lit16}$	1	1
SUB Wb,Ws,Wd	$Wd = Wb - Ws$	0.5/1	1
SUB Ws,lit7,Wd	$Wd = Ws - \text{lit7} \text{ (literal zero-extended)}$	1	1
SUBB f,Wn	$f = f - Wn - (C)$	1	1
SUBB f,Wn,Wn	$Wn = f - Wn - (C)$	1	1
SUBB lit16,Wn	$Wn = Wn - \text{lit16} - (C)$	1	1
SUBB Wb,Ws,Wd	$Wd = Wb - Ws - (C)$	0.5/1	1
SUBB Ws,lit7,Wd	$Wd = Ws - \text{lit7} - \text{ (literal zero-extended)}$	1	1
SUBR f,Wn	$f = Wn - f$	1	1
SUBR f,Wn,Wn	$Wn = Wn - f$	1	1
SUBR Wb,Ws,Wd	$Wd = Ws - Wb$	0.5/1	1
SUBR Ws,lit7,Wd	$Wd = \text{lit7} - Ws \text{ (literal zero-extended)}$	0.5/1	1
SUBBR f,Wn	$f = Wn - f - (C)$	1	1
SUBBR f,Wn,Wn	$Wn = Wn - f - (C)$	1	1
SUBBR Wb,Ws,Wd	$Wd = Ws - Wb - (C)$	0.5/1	1
SUBBR Ws,lit7,Wd	$Wd = \text{lit7} - Ws - (C) \text{ (literal zero-extended)}$	1	1
ZE Ws,Wnd	$Wd = \text{Zero-extend } Ws$	0.5/1	1

**Table 2-5.** Logic Instructions

Assembled Syntax	Description	Words	Cycles
AND f,Wn	$f = f .AND. Wn$	1	1

**Table 2-5.** Logic Instructions (continued)

Assembled Syntax	Description	Words	Cycles
AND f,Wn,Wn	W0 = f .AND. Wn	1	1
AND lit16,Wn	Wn = Wn .AND. lit16	1	1
AND Wb,Ws,Wd	Wd = Wb .AND. Ws	0.5/1	1
AND Wb,lit7,Wd	Wd = Wb .AND. Lit7 (literal zero-extended)	1	1
AND1 Wb,lit7,Wd	Wd = Wb .AND. Lit7 (literal zero-extended)	1	1
CLR f	f = 0x0000	1	1
CLR Wd	Wd = 0x0000	1	1
COM f	f = f	1	1
COM f,Wd	Wd = f	1	1
COM Ws,Wd	Wd = Ws	0.5/1	1
IOR f,Wn	f = f .IOR. Wn	1	1
IOR f,Wn,Wn	Wn = f .IOR. Wn	1	1
IOR lit16,Wn	Wn = Wn .IOR. lit16	1	1
IOR Wb,Ws,Wd	Wd = Wb .IOR. Ws	0.5/1	1
IOR Wb,lit7,Wd	Wd = Wb .IOR. lit7	1	1
NEG f	f = f + 1	1	1
NEG f,Wd	Wd = f + 1	1	1
NEG Ws,Wd	Wd = Ws + 1	0.5/1	1
SETM f	f = 0xFFFF	1	1
SETM Wd	Wd = 0xFFFF	1	1
XOR f,Wn	f = f .XOR. Wn	1	1
XOR f,Wn,Wn	Wn = f .XOR. Wn	1	1
XOR lit16,Wn	Wn = Wn .XOR. lit16	1	1
XOR Wb,Ws,Wd	Wd = Wb .XOR. Ws	0.5/1	1
XOR Wb,lit7,Wd	Wd = Wb .XOR. Lit7 (literal zero-extended)	1	1

**Table 2-6.** Rotate/Shift Instructions

Assembled Syntax	Description	Words	Cycles
ASR f	f = Arithmetic Right Shift f by 1	1	1
ASR f,Wn	Wn = Arithmetic Right Shift f by 1	1	1
ASR Ws,Wd	Wd = Arithmetic Right Shift Ws by 1	0.5/1	1
ASR Ws,Wb,Wd	Wnd = Arithmetic Right Shift Ws by Wb	0.5/1	1
ASR Ws,lit5,Wd	Wnd = Arithmetic Right Shift Ws by lit5	0.5/1	1
ASRM Ws, lit5, Wnd	Wnd = Arithmetic Right Shift Ws by lit5, then logically OR with next lsw	1	2
ASRM Ws, Wb, Wnd	Wnd = Arithmetic Right Shift Ws by Wb, then logically OR with next lsw	1	2
LSR f	f = Logical Right Shift f by 1	1	1
LSR f,Wd	Wd = Logical Right Shift f by 1	1	1
LSR Ws,Wd	Wd = Logical Right Shift Ws by 1	0.5/1	1
LSR Ws,Wb,Wd	Wnd = Logical Right Shift Ws by Wns	0.5/1	1
LSR Ws,lit5,Wd	Wnd = Logical Right Shift Ws by lit5	0.5/1	1
LSRM Ws, lit5, Wnd	Wnd = Logical Right Shift Ws by lit5, then logically OR with next lsw	1	2
LSRM Ws, Wb, Wnd	Wnd = Logical Right Shift Ws by Wb, then logically OR with next lsw	1	2
RLC f	f = Rotate Left through Carry f	1	1

**Table 2-6.** Rotate/Shift Instructions (continued)

Assembled Syntax	Description	Words	Cycles
RLC f,Wd	Wd = Rotate Left through Carry f	1	1
RLC Ws,Wd	Wd = Rotate Left through Carry Ws	0.5/1	1
RLNC f	f = Rotate Left (No Carry) f	1	1
RLNC f,Wd	Wd = Rotate Left (No Carry) f	1	1
RLNC Ws,Wd	Wd = Rotate Left (No Carry) Ws	0.5/1	1
RRC f	f = Rotate Right through Carry f	1	1
RRC f,Wd	Wd = Rotate Right through Carry f	1	1
RRC Ws,Wd	Wd = Rotate Right through Carry Ws	0.5/1	1
RRNC f	f = Rotate Right (No Carry) f	1	1
RRNC f,Wd	Wd = Rotate Right (No Carry) f	1	1
RRNC Ws,Wd	Wd = Rotate Right (No Carry) Ws	0.5/1	1
SL f	f = Left Shift f by 1	1	1
SL f,Wd	Wd = Left Shift f by 1	1	1
SL Ws,Wd	Wd = Left Shift Ws by 1	0.5/1	1
SL Ws,Wb,Wnd	Wnd = Left Shift Wb by Wns	0.5/1	1
SL Ws,lit5,Wnd	Wnd = Left Shift Ws by lit5	0.5/1	1
SLM Ws, lit5, Wnd	Wnd = Left Shift Wb by lit5, then logically OR with next msw	1	2
SLM Ws, Wb, Wnd	Wnd = Left Shift Wb by Wb, then logically OR with next msw	1	2

**Table 2-7.** Bit Instructions

Assembled Syntax	Description	Words	Cycles
BCLR.b f,bit3	Bit Clear f	1	1
BCLR Ws,bit5	Bit Clear Ws	0.5/1	1
BFEXT bit5,wid6,Ws,Wb	Bit Field Extract from Ws to Wb	1	1
BFEXT bit5,wid6,f,Wb	Bit Field Extract from f to Wb	2	2
BFINS bit5,wid6,Wb,Ws	Bit Field Insert from Wb into Ws	1	1
BFINS bit5,wid6,Wb,f	Bit Field Insert from Wb into f	2	2
BFINS bit5, wid6, lit16, Ws	Bit Field Insert lit8 into Ws	2	2
BSET.b f,bit3	Bit Set f	1	1
BSET Ws,bit5	Bit Set Ws	0.5/1	1
BSW.C Ws,Wb	Write C or Z bit to Ws<Wb>	1	1
BSW.Z Ws,Wb	Write C or Z bit to Ws<Wb>	0.5/1	1
BTG.b f,bit3	Bit Toggle f	1	1
BTG Ws,bit5	Bit Toggle Ws	0.5/1	1
BTST.b f,bit3	Bit Test f	1	1
BTST.C Ws,bit5	Bit Test Ws to C	0.5/1	1
BTST.Z Ws,bit5	Bit Test Ws to Z	1	1
BTST.C Ws,Wb	Bit Test Ws<Wb> to C	0.5/1	1
BTST.Z Ws,Wb	Bit Test Ws<Wb> to Z	1	1
BTSTS.b f,bit3	Bit Test then Set f	1	1
BTSTS.C Ws,bit5	Bit Test Ws to C then Set	0.5/1	1
BTSTS.Z Ws,bit5	Bit Test Ws to Z then Set	1	1
FBCL Ws,Wnd	Find Bit Change from Left (MSb) Side	1	1
FF1L Ws,Wnd	Find First One from Left (MSb) Side	1	1

**Table 2-7.** Bit Instructions (continued)

Assembled Syntax	Description	Words	Cycles
FF1R Ws,Wnd	Find First One from Right (LSb) Side	1	1

**Table 2-8.** Compare/Skip and Compare/Branch Instructions

Assembled Syntax	Description	Words	Cycles
CP f,Ws	Compare f with Ws	1	1
CP Ws,lit13	Compare Ws with lit13 (literal zero-extended)	1	1
CP Wb,lit16	Compare Wb with lit16 (literal zero-extended)	1	1
CP Wb, Ws	Compare Wb with Ws	0.5/1	1
CPO f	Compare f with 0x0000	1	1
CPO Ws	Compare Ws with 0x0000 (substitute CPLS Ws ,#0)	1	1
CPB f,Ws	Compare f with Ws, with borrow	1	1
CP Wb,lit13	Compare Wb with lit13, with borrow (literal zero-extended)	1	1
CP Wb,lit16	Compare Wb with lit16, with borrow (literal zero-extended)	1	1
CPB Wb,Ws	Compare Borrow Wb with Ws	0.5/1	1
DTB Wn,Label	Decrement Wn, then branch if not zero	1	1(2/3)

**Table 2-9.** Program Flow Instructions

Assembled Syntax	Description	Words	Cycles
BRA Label	Branch Unconditionally	1	1
BRA Wn	Computed Branch	1	2
BRA C,Label	Branch if Carry	1	1(2/3)
BRA GE,Label	Branch if greater than or equal	1	1(2/3)
BRA GEU,Label	Branch if unsigned greater than or equal	1	1(2/3)
BRA GT,Label	Branch if greater than	1	1(2/3)
BRA GTU,Label	Branch if unsigned greater than	1	1(2/3)
BRA LE,Label	Branch if less than or equal	1	1(2/3)
BRA LEU,Label	Branch if unsigned less than or equal	1	1(2/3)
BRA LT,Label	Branch if less than	1	1(2/3)
BRA LTU,Label	Branch if unsigned less than	1	1(2/3)
BRA N,Label	Branch if Negative	1	1(2/3)
BRA NC,Label	Branch if Not Carry	1	1(2/3)
BRA NN,Label	Branch if Not Negative	1	1(2/3)
BRA NOV,Label	Branch if Not Overflow	1	1(2/3)
BRA NZ,Label	Branch if Not Zero	1	1(2/3)
BRA Z,Label	Branch if Zero	1	1(2/3)
BREAK	Stop user code execution	0.5/1	1
CALL Label	Call subroutine (label < ~ 16MB)	1	1
CALL Wns	Call indirect subroutine at address [W11]	1	2
GOTO Label	Goto address (address < ~ 16MB)	1	1
GOTO Wn	Go to indirect address at [W11]	1	2
RCALL Label	Relative Call	1	1
RCALL Wns	Computed Call	1	2

**Table 2-9.** Program Flow Instructions (continued)

Assembled Syntax	Description	Words	Cycles
REPEAT lit15	Repeat Next Instruction lit15+1 times	1	1
REPEAT lit5	Repeat Next Instruction lit5+1 times	0.5	1
REPEAT Wn	Repeat Next Instruction (Wn)+1 times	1	1
RETFILE	Return from interrupt enable	0.5	4
RETLW lit16,Wn	Return from Subroutine with literal in Wn	1	3
RETURN	Return from Subroutine	0.5	3

**Table 2-10.** Shadow/Stack/Context Instructions

Assembled Syntax	Description	Words	Cycles
BOOTSWP	Swap Active and Inactive address space	0.5	2
CTXTSWP lit3	Swap to CPU register context #2	0.5	2
CTXTSWP Wn	Swap to CPU register context defined in Wn[2:0]	1	2
LNK lit16	Link frame pointer	1	1
LNK lit7	Link frame pointer (literal < 128)	0.5	1
POP f	Pop f from top of stack (TOS)	1	1
POP { [--Ws], } Wnd	Pop Wnd Register from system stack.	0.5	1
POP Fd	Pop Fd Register from system stack.	0.5	1
PUSH f	Push f to top of stack (TOS)	1	1
PUSH Wns, { [Wd++ ] }	Push Wns Register to system stack	0.5	1
PUSH Fs	Push Fs Register to system stack	0.5	1
ULNK	Unlink frame pointer	0.5	1

**Table 2-11.** Control Instructions

Assembled Syntax	Description	Words	Cycles
CLRWDT	Clear Watchdog Timer	0.5	1
DISICCTL lit3 {,Wd}	Disable interrupts at IPL <= lit3 Optionally save prior IPL threshold to Wd	1	1
DISICCTL Wns {,Wd}	Disable interrupts at IPL <= Wns[2:0] Optionally save prior IPL threshold to Wd	1	1
NEOP	None executable NOP (16-bit instruction pad)	0.5	0
NOP	No Operation	1	1
NOPR	No Operation	1	1
PWRSAV mode	Go into standby mode	0.5	2
RESET	Software device RESET	1	1

**Table 2-12.** DSP Instructions

Assembled Syntax	Description	Words	Cycles
ADD A	Add Accumulators	0.5	1
ADD Ws,Slit6, A	Signed Add to Accumulator	1	1

**Table 2-12.** DSP Instructions (continued)

Assembled Syntax	Description	Words	Cycles
BRA OA,Label	Branch if accumulator A overflow	1	1(2/3)
BRA OB,Label	Branch if accumulator B overflow	1	1(2/3)
BRA OV,Label	Branch if Overflow	1	1(2/3)
BRA SA,Label	Branch if accumulator A saturated	1	1(2/3)
BRA SB,Label	Branch if accumulator B saturated	1	1(2/3)
CLR A	Clear Accumulator	0.5	1
ED Wx, Wy, A, AWB	Euclidean Distance	1	2
EDAC Wx, Wy, A, AWB	Euclidean Distance Accumulate	1	2
LAC Ws,Slit6, A	Load Accumulator (16/32-bit), literal shift	1	1
LLAC.l Ws Slit6, A	Load Lower (LS-word of) Accumulator (32-bit), literal shift	1	1
LUAC.l Ws, Slit6, A	Load Upper (LS-byte) of Accumulator (32-bit), literal shift	1	1
MAC Wx, Wy, A, AWB	Multiply and Accumulate	1	1
MAX Wb, Ws	Force Data Maximum Range Limit	1	1
MAX A	Force Data Maximum Range Limit	0.5	1
MAX.V A, Wd	Force Data Maximum Range Limit with Result	1	2
MIN Wb, Ws	Force Data Minimum Range Limit	1	1
MIN A	Force Data Minimum Range Limit	0.5	1
MIN.V A, Wd	Force Data Minimum Range Limit with Result	1	2
MULISS Wb,Ws,A	Integer: Acc(A or B) = signed(Wb) * signed(Ws)	1	1
MULFSS Wb,Ws,A	Fractional: Acc(A or B) = signed(Wb) * signed(Ws)	1	1
MULISU Wb,Ws,A	Integer: Acc(A or B) = signed(Wb) * unsigned(Ws)	1	1
MULFSU Wb,Ws,A	Fractional: Acc(A or B) = signed(Wb) * unsigned(Ws)	1	1
MULIUS Wb,Ws,A	Integer: Acc(A or B) = unsigned(Wb) * signed(Ws)	1	1
MULFUS Wb,Ws,A	Fractional: Acc(A or B) = unsigned(Wb) * signed(Ws)	1	1
MULIUU Wb,Ws,A	Integer: Acc(A or B) = unsigned(Wb) * unsigned(Ws)	1	1
MULFUU Wb,Ws,A	Fractional: Acc(A or B) = unsigned(Wb) * unsigned(Ws)	1	1
MULISS Wb,slit8,A	Integer: Acc(A or B) = signed(Wb) * signed(slit8)	1	1

**Table 2-12.** DSP Instructions (continued)

Assembled Syntax	Description	Words	Cycles
MULFSS Wb,slit8,A	Integer: Acc(A or B) = signed(Wb) * signed(slit8)	1	1
MULISU Wb,lit8,A	Integer: Acc(A or B) = signed(Wb) * unsigned(lit8)	1	1
MULFSU Wb,lit8,A	Integer: Acc(A or B) = signed(Wb) * unsigned(lit8)	1	1
MULIUS Wb,slit8,A	Integer: Acc(A or B) = signed(Wb) * signed(slit8)	1	1
MULFUS Wb,slit8,A	Integer: Acc(A or B) = signed(Wb) * signed(slit8)	1	1
MULIUU Wb,lit8,A	Integer: Acc(A or B) = signed(Wb) * unsigned(lit8)	1	1
MULFUU Wb,lit8,A	Integer: Acc(A or B) = signed(Wb) * unsigned(lit8)	1	1
MPY Wx, Wy, A, AWB	Multiply Wm by Wn to Accumulator	1	1
MPYN Wx, Wy, A, AWB	-(Multiply Wm by Wn) to Accumulator	1	1
MSC Wx, Wy, A, AWB	Multiply and Subtract from Accumulator	1	1
NEG A	Negate Accumulator	0.5	1
NORM A, Wd	Normalize Accumulator	1	1
SAC A,Slit6, Wd	Store Accumulator (16/32-bit)	1	1
SACR A,Slit6,Wd	Store Rounded Accumulator (16/32-bit), literal shift	1	1
SACR A,Ws,Wd	Store Rounded Accumulator (16/32-bit), Wb shift	1	1
SLAC.1 A,Slit6,Wd	Store Lower (LS-Word of) Accumulator (32-bit), literal shift	1	1
SUAC.1 A,Slit6,Wd	Store sign extended Upper (MS-Byte of) Accumulator (32-bit), literal shift	1	1
SFTAC A,Wn	Arithmetic Shift by (Wn) Accumulator	1	1
SFTAC A,Slit7	Arithmetic Shift by Slit7 Accumulator	1	1
SQR Wx, A, AWB	Square to Accumulator	1	1
SQRAC Wx, A, AWB	Square and Accumulate	1	1
SQRN Wx, A, AWB	Negated Square to Accumulator	1	1
SQRSC Wx, A, AWB	Square and Subtract from Accumulator	1	1
SUB A	Subtract Accumulators	0.5	1
SUB Ws,Slit6, A	Signed Subtract from Accumulator	1	1

**Table 2-13.** FPU Instructions

Assembled Syntax	Description	Words	Cycles
ABS Fs, Fd	Absolute value of Fs	1	1
ADD Fb, Fs, Fd	Fd = Fb + Fs	1	2

**Table 2-13.** FPU Instructions (continued)

Assembled Syntax	Description	Words	Cycles
AND lit16, FSR	FSR = FSR AND lit16	1	1
AND lit16, FCR	FCR = FCR AND lit16	1	1
AND lit16, FEAR	FEAR = FEAR AND lit16	1	1
COS Fs , Fd	Fd = COS(Fs)	1	4
CPQ Fb, Fs	Compare Fb with Fs, Quiet Signaling	1	1
CPS Fb, Fs	Compare Fb with Fs, Signaling	1	1
DI2F Fs, Fd	Convert Double Word (64-bit) Integer to Floating-Point, Fs (integer) -->Fd (float)	1	2
DIV Fb, Fs, Fd	Signed Floating-Point Divide, Fd = Fb/Fs	1	11/32
FBRA EQ,Label	Floating Point Branch if Equal	1	1(2/3)
FBRA NE,Label	Floating Point Branch if Not Equal	1	1(2/3)
FBRA GT,Label	Floating Point Branch if Greater Than	1	1(2/3)
FBRA GE,Label	Floating Point Branch if Greater Than or Equal	1	1(2/3)
FBRA LT,Label	Floating Point Branch if Less Than	1	1(2/3)
FBRA LE,Label	Floating Point Branch if Less Than or Equal	1	1(2/3)
FBRA OR,Label	Floating Point Branch if Ordered	1	1(2/3)
FBRA UNE,Label	Floating Point Branch if Unordered or Not Equal	1	1(2/3)
FBRA UEQ,Label	Floating Point Branch if Unordered or Equal	1	1(2/3)
FBRA ULE,Label	Floating Point Branch if Unordered or Less Than or Equal	1	1(2/3)
FBRA ULT,Label	Floating Point Branch if Unordered or Less Than	1	1(2/3)
FBRA UGE,Label	Floating Point Branch if Unordered or Greater Than or Equal	1	1(2/3)
FBRA UGT,Label	Floating Point Branch if Unordered or Greater Than	1	1(2/3)
FBRA UN,Label	Floating Point Branch if Unordered	1	1(2/3)
FLIM Fb, Fs, Fd	Force Signed Data Limit, If Fd > Fs Then Fd = Fs If Fd < Fb then Fd = Fb	1	1
F2DI Fs, Fd	Convert Floating-Point Fs to Double Word (64-bit) Integer,Fs (float)--> Fd (integer)	1	1/2
F2LI Fs, Fd	Convert Floating-Point Fs to Long Word (32-bit) Integer,Fs (float)--> Fd (integer)	1	1/2
IOR lit16, FSR	Inclusive OR FSR, FSR = FSR .IOR. lit16	1	1

**Table 2-13.** FPU Instructions (continued)

Assembled Syntax	Description	Words	Cycles
IOR lit16, FCR	Inclusive OR FCR, FCR = FCR .IOR. Lit16	1	1
IOR lit16, FEAR	Inclusive OR FEAR, FEAR = FEAR .IOR. Lit16	1	1
LI2F Fs, Fd	Convert Long Word (32-bit) Integer to Floating-Point,Fs (integer)-->Fd (float)	1	1
MAC Fb, Fs, Fd	Floating-Point Signed Multiply and Accumulate,Fd = Fd +(Fb * Fs)	1	3/4
MAX Fb, Fs, Fd	Select the Signed Maximum of Fb and Fs {IEEE 754-2019 maximum(x,y)}, if Fs >= Fb then Fd = Fs Else Fd= Fb	1	1
MAXNM Fb, Fs, Fd	Select the Signed Maximum of Fb and Fs {IEEE 754-2019 maximumNumber(x,y)},, if Fs >= Fb then Fd = Fs Else Fd= Fb	1	1
MIN Fb, Fs, Fd	Select the Signed Minimum of Fb and Fs {IEEE 754-2019 minimum(x,y)}, if Fs <= Fb then Fd = Fs Else Fd= Fb	1	1
MINNM Fb, Fs, Fd	Select the Signed Minimum of Fb and Fs [minimumNumber()], if Fs <= Fb then Fd = Fs Else Fd= Fb	1	1
MOV.l Fs, Wd	Move coprocessor register to Wd	0.5/1	1
MOV.l Ws, Fd	Move Ws to coprocessor register	0.5/1	1
MOV.l lit32, Fd	Move 32-bit unsigned literal to coprocessor register	2	2
MOV Fs, [Wnd+Slit12]	Move Fs to [Wnd+Slit12]	1	1
MOV [Wns+Slit12], Fd	Move [Wns+Slit12] to Fd	1	1
MOV Fs, Fd	Move Fs to Fd	1	1
MOV index, Fd	Fd = Constant table (index) Fd	1	1
MUL Fb, Fs, Fd	Fd = Fb * Fs	1	3
NEG Fs, Fd	Fd = -Fs	1	1
SIN Fs, Fd	Fd = SIN(Fs)	1	4
SQRT Fs, Fd	Fd = $\sqrt{Fs}$	1	10/13
SUB Fb, Fs, Fd	Fd= Fb- Fs	1	2
TST Fs	Test Fs	1	1

## 3. Instruction Set Details

### 3.1. Data Addressing Modes

The PIC32A devices support three native addressing modes for accessing data memory. Data accesses may be performed using File Register Addressing, Register Direct or Indirect Addressing, and a few Immediate Addressing types, allowing a fixed value to be used by the instruction.

File Register Addressing provides the ability to operate on data stored up to 64 KB (if a WREG operand is required), and the MOV instruction provides access to all 1 MB of data space. Register Direct Addressing is used to access the 16 memory-mapped Working registers, W0:W15. Register Indirect Addressing is used to efficiently operate on data stored in the entire data space, using the contents of the Working registers as an Effective Address (EA). Immediate Addressing does not access data memory but provides the ability to use a constant value as an instruction operand. The address range of each mode is summarized in [Table 3-1](#).

**Table 3-1. PIC32A Addressing Modes**

Addressing Mode	Address Range
File Register	Upto 64 KB/1 MB
Register Direct	0x000F (Working register array, W0:W15)
Register Indirect	0x0000-0xFFFF
Immediate	N/A (constant value)

#### 3.1.1. File Register Addressing

File Register Addressing is used by instructions that use a predetermined data address as an operand for the instruction. The majority of instructions that support File Register Addressing provide access up to 64 KB (if a WREG operand is required). However, the MOV instruction provides access to all 1 MB of memory using File Register Addressing. This allows the loading of the data from any location in data memory to any Working register and storing the contents of any Working register to any location in data memory. It should be noted that File Register Addressing supports byte, extended byte, word and long word data sizes. Refer to [Example 3-1](#) of File Register Addressing modes.

Most instructions which support File Register Addressing perform an operation on the specified file register and the default Working register, WREG. If only one operand is supplied in the instruction, WREG is an implied operand and the operation results are stored back to the file register. In these cases, the instruction is effectively a Read-Modify-Write instruction. However, when both the file register and the WREG register are specified in the instruction, the operation results are stored in the WREG register and the contents of the file register are unchanged. Sample instructions that show the interaction between the file register and the WREG register are shown in [Example 3-2](#).

**Note:** Instructions which support File Register Addressing use 'f' as an operand in the instruction summary tables of [Instruction Set Overview](#)

**Example 3-1. File Register Addressing**

DEC	0x5000	; decrement data stored at 0x5000
-----	--------	-----------------------------------

Before Instruction:

Data Memory 0x5000 = 0x55555555
---------------------------------

After Instruction:

```
Data Memory 0x5000 = 0x55555554
MOV.L 0x000027FE, W0 ; move data stored at 0x000027FE to W0
```

Before Instruction:

```
W0 = 0x55555555
Data Memory 0x000027FE = 0x12345678
```

After Instruction:

```
W0 = 0x12345678
Data Memory 0x000027FE = 0x12345678
```

### **Example 3-2. File Register Addressing and WREG**

```
AND 0x00001000 ; AND 0x00001000 with WREG, store to 0x00001000
```

Before Instruction:

```
W0 (WREG) = 0x0000332C
Data Memory 0x00001000 = 0x55555555
```

After Instruction:

```
W0 (WREG) = 0x0000332C
Data Memory 0x00001000 = 0x00001104
```

```
AND 0x00001000, WREG ; AND 0x00001000 with WREG, store to WREG
```

Before Instruction:

```
W0 (WREG) = 0x0000332C
Data Memory 0x00001000 = 0x55555555
```

After Instruction:

```
W0 (WREG) = 0x00001104
Data Memory 0x00001000 = 0x55555555
```

### **3.1.2. Register Direct Addressing**

Register Direct Addressing is used to access the contents of the 16 Working registers (W0:W15). The Register Direct Addressing mode is fully orthogonal, which allows any Working register to be specified for any instruction that uses Register Direct Addressing, and it supports byte, word, and long word accesses. Instructions which employ Register Direct Addressing use the contents of the specified Working register as data to execute the instruction; therefore, this addressing mode is useful only when data already resides in the Working register core. Sample instructions which utilize Register Direct Addressing are shown in [Example 3-3](#).

Another feature of Register Direct Addressing is that it provides the ability for dynamic flow control. Since variants of the REPEAT instruction support Register Direct Addressing, flexible looping constructs may be generated using these instructions.

**Note:** Instructions that must use Register Direct Addressing use the symbols  $W_b$ ,  $W_n$ ,  $W_{ns}$  and  $W_{nd}$  in the summary tables of [Instruction Set Overview](#). Commonly, Register Direct Addressing may also be used when Register Indirect Addressing may be used. Instructions that use Register Indirect Addressing use the symbols  $W_d$  and  $W_s$  in the summary tables of [Instruction Set Overview](#).

#### Example 3-3. Register Direct Addressing

```
EXCH      W2, W3           ; Exchange W2 and W3
```

##### Before Instruction:

```
W2 = 0x000003499
W3 = 0x00000003D
```

##### After Instruction:

```
W2 = 0x00000003D
W3 = 0x000003499
```

```
IOR.L     #0x44, W0           ; Inclusive-OR 0x44 and W0
```

##### Before Instruction:

```
W0 = 0x12349C2E
```

##### After Instruction:

```
W0 = 0x12349C6E
```

```
SL       W6, W7, W8           ; Shift left W6 by W7, and store to W8
```

##### Before Instruction:

```
W6 = 0x0000000C
W7 = 0x00000008
W8 = 0x12345678
```

##### After Instruction:

```
W6 = 0x0000000C
W7 = 0x00000008
W8 = 0x00000C00
```

### 3.1.3. Register Indirect Addressing

Register Indirect Addressing is used to access any location in data memory by treating the contents of a Working register as an Effective Address (EA) to data memory. Essentially, the contents of the Working register become a pointer to the location in data memory that is to be accessed by the instruction.

This addressing mode is powerful because it also allows one to modify the contents of the Working register, either before or after the data access is made, by incrementing or decrementing the EA. By modifying the EA in the same cycle that an operation is being performed, Register Indirect Addressing allows for the efficient processing of data that is stored sequentially in memory. The modes of Indirect Addressing supported by PIC32A devices are shown in [Table 3-2](#).

**Table 3-2.** Indirect Addressing Modes

Indirect Mode	Syntax	Function (Byte Instruction)	Function (Word Instruction)	Function (Long word instruction)	Description
No Modification	[Wn]	EA = [Wn]	EA = [Wn]	EA = [Wn]	The contents of Wn form the EA.
Pre-Increment	[++Wn]	EA = [Wn + = 1]	EA = [Wn + = 2]	EA = [Wn + = 4]	Wn is pre-incremented to form the EA.
Pre-Decrement	[--Wn]	EA = [Wn - = 1]	EA = [Wn - = 2]	EA = [Wn - = 4]	Wn is pre-decremented to form the EA.
Post-Increment	[Wn++]	EA = [Wn] + = 1	EA = [Wn] + = 2	EA = [Wn] + = 4	The contents of Wn form the EA, then Wn is post-incremented.
Post-Decrement	[Wn--]	EA = [Wn] - = 1	EA = [Wn] - = 2	EA = [Wn] - = 4	The contents of Wn form the EA, then Wn is post-decremented.
Register Offset	[Wn+Wb]	EA = [Wn + Wb]	EA = [Wn + Wb]	EA = [Wn + Wb]	The sum of Wn and Wb forms the EA. Wn and Wb are not modified.

**Table 3-2** shows that four addressing modes modify the EA used in the instruction, and this allows the following updates to be made to the Working register: post-increment, post-decrement, pre-increment and pre-decrement.

**Table 3-2** also shows that the Register Offset mode addresses data which is offset from a base EA stored in a Working register. This mode uses the contents of a second Working register to form the EA by adding the two specified Working registers. Note that neither of the Working registers used to form the EA is modified. [Example 3-4](#) shows how Register Offset Indirect Addressing may be used to access data memory.

**Note:** The MOV with offset instructions provides a literal addressing offset ability to be used with Indirect Addressing. In these instructions, the EA is formed by adding the contents of a Working register to a signed literal. [Example 3-5](#) shows how these instructions may be used to move data to and from the Working register array.

#### Example 3-4. Indirect Addressing with Effective Address Update

```
MOV.B      [W0++], [W13--] ; byte move [W0] to [W13]
; post-inc W0, post-dec W13
```

Before Instruction:

```
W0 = 0x2300
W13 = 0x2708
Data Memory 0x2300 = 0x7783
Data Memory 0x2708 = 0x904E
```

After Instruction:

```
W0 = 0x2301
W13 = 0x2707
Data Memory 0x2300 = 0x7783
Data Memory 0x2708 = 0x9083
```

```
ADD      W1, [--W5], [++W8] ; pre-dec W5, pre-inc W8
; add W1 to [W5], store in [W8]
```

## Before Instruction:

```

W1 = 0x0800
W5 = 0x2200
W8 = 0x2400
Data Memory 0x21FE = 0x7783
Data Memory 0x2402 = 0xAACC

```

## After Instruction:

```

W1 = 0x0800
W5 = 0x21FE
W8 = 0x2402
Data Memory 0x21FE = 0x7783
Data Memory 0x2402 = 0x7F83

```

**Example 3-5.** Indirect Addressing with Register Offset

```

MOV.B      [W0+W1], [W7++] ; byte move
; [W0+W1] to W7, post-inc W7

```

## Before Instruction:

```

W0 = 0x2300
W1 = 0x01FE
W7 = 0x1000
Data Memory 0x24FE = 0x7783
Data Memory 0x1000 = 0x11DC

```

## After Instruction:

```

W0 = 0x2300
W1 = 0x01FE
W7 = 0x1001
Data Memory 0x24FE = 0x7783
Data Memory 0x1000 = 0x1183

```

```

LAC.l      [W0+W8], A ; load ACCA with [W0+W8]
; (sign-extend and zero-backfill)

```

## Before Instruction:

```

W0 = 0x2344
W8 = 0x0008
ACCA = 0x00 7877 9321 0000 0000
Data Memory 0x234C = 0xE290

```

## After Instruction:

```

W0 = 0x2344
W8 = 0x0008
ACCA = 0xFF E290 0000 0000 0000
Data Memory 0x234C = 0xE290

```

**Example 3-6.** Move with Literal Offset Instructions

```

MOV      [W0+0x20], W1 ; move [W0+0x20] to W1

```

## Before Instruction:

```

W0 = 0x1200
W1 = 0x01FE
Data Memory 0x1220 = 0xFD27

```

After Instruction:

```
W0 = 0x1200
W1 = 0xFD27
Data Memory 0x1220 = 0xFD27
```

```
MOV      W4, [W8-0x300] ; move W4 to [W8-0x300]
```

Before Instruction:

```
W4 = 0x3411
W8 = 0x2944
Data Memory 0x2644 = 0xCB98
```

After Instruction:

```
W4 = 0x3411
W8 = 0x2944
Data Memory 0x2644 = 0x3411
```

### 3.1.3.1. Register Indirect Addressing and the Instruction Set

The addressing modes presented in [Table 4-2](#) demonstrate the Indirect Addressing mode capability of the PIC32A devices. Due to operation encoding and functional considerations, not every instruction which supports Indirect Addressing supports all modes shown in [Table 4-2](#). The majority of instructions which use Indirect Addressing support the No Modify, Pre-Increment, Pre-Decrement, Post-Increment and Post-Decrement Addressing modes. The `MOV` instructions and several accumulator-based DSP instructions are also capable of using the Register Offset Addressing mode.

**Note:** Instructions that use Register Indirect Addressing use the operand symbols `Wd` and `Ws` in the summary tables of [Instruction Set Overview](#).

### 3.1.3.2. DSP MAC Indirect Addressing Modes

A special class of Indirect Addressing modes is utilized by the DSP MAC instructions. As is described later in [DSP MAC Instructions](#), the DSP MAC class of instructions is capable of performing two fetches from memory using Effective Addressing. Since DSP algorithms frequently demand a broader range of address updates, the addressing modes offered by the DSP MAC instructions provide greater range in the size of the Effective Address update which may be made. The following table shows that both X and Y prefetches support Post-Increment and Post-Decrement Addressing modes with updates of two, four and six bytes. Since DSP instructions only execute in Word/Long-word mode, no provisions are made for odd-sized EA updates.

**Table 3-3. DSP MAC Indirect Addressing Modes**

Addressing Mode	Word-Mode		Long-Word Mode	
	X Memory	Y Memory	X Memory	Y Memory
Indirect with No Modification	EA = [Wx]	EA = [Wy]	EA = [Wx]	EA = [Wy]
Indirect with Post-Increment by two	EA = [Wx] + = 2	EA = [Wy] + = 2	EA = [Wx] + = 4	EA = [Wy] + = 4
Indirect with Post-Increment by four	EA = [Wx] + = 4	EA = [Wy] + = 4	EA = [Wx] + = 8	EA = [Wy] + = 8
Indirect with Post-Increment by six	EA = [Wx] + = 6	EA = [Wy] + = 6	EA = [Wx] + = 12	EA = [Wy] + = 12
Indirect with Post-Decrement by two	EA = [Wx] - = 2	EA = [Wy] - = 2	EA = [Wx] - = 4	EA = [Wy] - = 4
Indirect with Post-Decrement by four	EA = [Wx] - = 4	EA = [Wy] - = 4	EA = [Wx] - = 8	EA = [Wy] - = 8
Indirect with Post-Decrement by six	EA = [Wx] - = 6	EA = [Wy] - = 6	EA = [Wx] - = 12	EA = [Wy] - = 12
Indirect with Register Offset	EA = [W9 + W12]	EA = [W11 + W12]	EA = [W9 + W12]	EA = [W11 + W12]

### 3.1.3.3. Modulo and Bit-Reversed Addressing Modes

The PIC32A architecture provides support for two special Register Indirect Addressing modes that are commonly used to implement DSP algorithms. Modulo (or circular) Addressing provides an automated means to support circular data buffers in X and/or Y memory. Modulo buffers remove the need for software to perform address boundary checks, which can improve the performance of certain algorithms. Similarly, Bit-Reversed Addressing allows one to access the elements of a buffer in a nonlinear fashion. This addressing mode simplifies data re-ordering for radix-2 FFT algorithms and provides a significant reduction in FFT processing time.

Both of these addressing modes are powerful features of the PIC32A architectures, which can be exploited by any instruction that uses Indirect Addressing.

### 3.1.4. Immediate Addressing

In Immediate Addressing, the instruction encoding contains a predefined constant operand that is used by the instruction. This addressing mode may be used independently, but it is more frequently combined with the File Register, Direct and Indirect Addressing modes. The size of the immediate operand which may be used varies with the instruction type. Constants of size 1-bit (#lit1), 5-bit (#bit5, #lit5 and #Slit5), 7-bit (#lit7), 6-bit (#Slit6), 8-bit (#lit8), 14-bit (#Slit14), 16-bit (#lit16), 20-bit(#lit20, #Slit20), and 32-bit(#lit32) may be used. Constants may be signed or unsigned, and the symbols #Slit6, #Slit8, and #Slit20 designate a signed constant. All other immediate constants are unsigned. [Table 3-4](#) shows the usage of each immediate operand in the instruction set.

**Table 3-4.** Immediate Operands in the Instruction Set

Operand	Instruction Usage
#lit1	PWRSAV
#lit3	CTXTSWP DISICTL
#bit5	BCLR, BSET, CLRWDT, BTG, BTST, BTSTS
#lit5	ASR, LSR, SL
#wid6	BFEXT, BFINS
#lit7	ADD, ADDC, AND, CP, CPB, IOR, SUB, SUBB, SUBBR, SUBR, XOR
#Slit6	SFTAC
#lit8	MOV.B, MULSU, MULUU, MULIUU, MULISU
#lit7	ADD, ADDC, AND, CP, CPB, IOR, RETLW, SUB, SUBB, XOR
#Slit6	LAC, LLAC, LUAC, SAC, SLAC, SUAC, SFTAC, SACR, SUB
#lit16	LNK, RETLW, MOV, SUB, SUBB, XOR, ADD, ADDC, AND, CP, IOR
#lit20	REPEAT
#lit16	MOV
#lit32	MOV, CALL, GOTO

The syntax for Immediate Addressing requires that the number sign (#) must immediately precede the constant operand value. The “#” symbol indicates to the assembler that the quantity is a constant. If an out-of-range constant is used with an instruction, the assembler will generate an error. Several examples of Immediate Addressing are shown in [Example 3-7](#).

#### Example 3-7. Immediate Addressing

```
PWRSAV      #1          ; Enter IDLE mode
ADD.B      #0x10, W0    ; Add 0x10 to W0 (byte mode)
```

Before Instruction:

```
W0 = 0x12A9
```

After Instruction:

```
W0 = 0x12B9
XOR      W0, #1, [W1++]
          ; Exclusive-OR W0 and 0x1
          ; Store the result to
[W1]           ; Post-increment W1
```

Before Instruction:

```
W0 = 0xFFFF
W1 = 0x0890
Data Memory 0x0890 = 0x0032
```

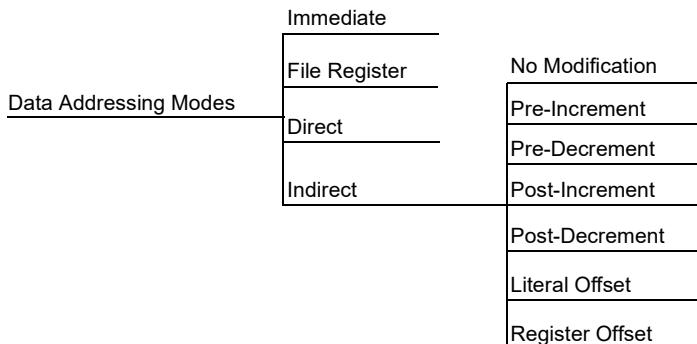
After Instruction:

```
W0 = 0xFFFF
W1 = 0x0892
Data Memory 0x0890 = 0xFFFFE
```

## 3.2. Data Addressing Mode Tree

The Data Addressing modes of the PIC32A family are summarized in [Figure 3-1](#).

**Figure 3-1.** Data Addressing Mode Tree



## 3.3. Program Addressing Modes

The PIC32A devices have a 24-bit Program Counter (PC). The PC addresses the 24-bit wide program memory to fetch instructions for execution and it may be loaded in several ways. Instructions are either 16-bit, 32-bit or 64-bit entities; therefore, the PC is incremented by 2, 4 or 8 during sequential 16-bit, 32-bit or 64-bit instruction execution, respectively.

Several methods may be used to modify the PC in a non-sequential manner and both absolute and relative changes may be made to the PC. The change to the PC may be from an immediate value encoded in the instruction or a dynamic value contained in a Working register. For exception handling, the PC is loaded with the address of the exception handler, which is stored in the Interrupt Vector Table (IVT). When required, the software stack is used to return scope to the foreground process from where the change in program flow occurred.

[Table 3-5](#) summarizes the instructions which modify the PC. When performing function calls, it is recommended that `RCALL` be used instead of `CALL`, since `RCALL` only consumes one word of program memory.

**Table 3-5.** Methods of Modifying Program Flow

Condition/Instruction	PC Modification	Software Stack Usage
Sequential Execution	$PC = (PC + 2) / (PC + 4) / (PC + 8)$ Depending on size of the instruction.	None
BRA Expr <sup>(1)</sup> (Branch Unconditionally)	$PC = PC = (PC + 4) + 2 * \text{slit20}$	None
BRA Condition, Expr <sup>(1)</sup> (Branch Conditionally)	$PC = PC + 4$ (condition false) $PC = (PC + 4) + 2 * \text{slit20}$ (condition true)	None
CALL Expr <sup>(1)</sup> (Call Subroutine)	$PC = \text{lit24}$	PC + 4 is PUSHed on the stack <sup>(2)</sup>
CALL Wn (Call Subroutine Indirect)	$PC = Wn$	PC + 4 is PUSHed on the stack <sup>(2)</sup>
GOTO Expr <sup>(1)</sup> (Unconditional Jump)	$PC = \text{lit24}$	None
GOTO Wn (Unconditional Indirect Jump)	$PC = Wn$	None
RCALL Expr <sup>(1)</sup> (Relative Call)	$PC = (PC + 4) + 2 * \text{slit20}$	PC + 4 is PUSHed on the stack <sup>(2)</sup>
RCALL Wn (Computed Relative Call)	$PC = (PC + 4) + 2 * Wn$	PC + 4 is PUSHed on the stack <sup>(2)</sup>
Exception Handling	$PC = \text{Address of the exception handler}$ (read from vector table)	PC + 4 is PUSHed on the stack <sup>(3)</sup>
PC = Target REPEAT instruction (REPEAT Looping)	PC not modified (if REPEAT active)	None
<b>Notes:</b>		
<ol style="list-style-type: none"> <li>For BRA, CALL and GOTO, the Expr may be a label, absolute address or expression, which is resolved by the linker to a 20-bit or 24-bit value (slit20 or lit24). When representing an address offset value, Expr can also be indicated by using a “.” and a sign, “+” or “-”. For example, the expression, “.+2”, means an address offset of +2 (i.e., the next instruction address relative to the current position of the Program Counter). See <a href="#">Instruction Descriptions</a> for details.</li> <li>After CALL or RCALL is executed, RETURN or RETLW will POP the Top-of-Stack (TOS) back into the PC.</li> <li>After an exception is processed, RETFIE will POP the Top-of-Stack (TOS) back into the PC.</li> </ol>		

## 3.4. Instruction Stalls

In order to maximize device throughput, all instructions are pipelined. A consequence of this pipelining is that address register data dependencies may arise between successive read and write operations using common registers.

‘Read-After-Write’ (RAW) dependencies occur across instruction boundaries and are detected by the hardware. An example of a RAW dependency would be a write operation that modifies W5, followed by a read operation that uses W5 as an Address Pointer. The contents of W5 will not be valid for the read operation until the earlier write completes. This problem is resolved by stalling the instruction execution for one instruction cycle, which allows the write to complete before the next read is started.

### 3.4.1. RAW Dependency Detection

During the Read-After-Write (RAW) instruction predecode, the core determines if any address register dependency is imminent across an instruction boundary. The stall detection logic compares the W register (if any) used for the destination EA of the instruction currently being executed with the W register to be used by the source EA (if any) of the prefetched instruction. When a match between the destination and source registers is identified, a set of rules is applied to decide whether or not to stall the instruction by one cycle. [Table 3-6](#) lists various RAW conditions that cause an instruction execution stall.

**Table 3-6.** Raw Dependency Rules (Detection By Hardware)

Destination Addressing Mode Using Wn	Source Addressing Mode Using Wn	Stall Required?	Examples <sup>(2)</sup> (Wn = W2)
Direct	Direct	No Stall	ADD.W W0, W1, W2 MOV.W W2, W3
Indirect	Direct	No Stall	ADD.W W0, W1, [W2] MOV.W W2, W3
Indirect	Indirect	No Stall	ADD.W W0, W1, [W2] MOV.W [W2], W3
Indirect	Indirect with Pre/Post-Modification	No Stall	ADD.W W0, W1, [W2] MOV.W [W2++], W3
Indirect with Pre/Post-Modification	Direct	No Stall	ADD.W W0, W1, [W2++] MOV.W W2, W3
Direct	Indirect	Stall <sup>(1)</sup>	ADD.W W0, W1, W2 MOV.W [W2], W3
Direct	Indirect with Pre/Post-Modification	Stall <sup>(1)</sup>	ADD.W W0, W1, W2 MOV.W [W2++], W3
Indirect	Indirect	Stall <sup>(1)</sup>	ADD.W W0, W1, [W2] (2) MOV.W [W2], W3 (2)
Indirect	Indirect with Pre/Post-Modification	Stall <sup>(1)</sup>	ADD.W W0, W1, [W2] (2) MOV.W [W2++], W3 (2)
Indirect with Pre/Post-Modification	Indirect	Stall <sup>(1)</sup>	ADD.W W0, W1, [W2++] MOV.W [W2], W3
Indirect with Pre/Post-Modification	Indirect with Pre/Post-Modification	Stall <sup>(1)</sup>	ADD.W W0, W1, [W2++] MOV.W [W2++], W3

When stalls are detected, one cycle is added to the instruction execution time.  
For these examples, the contents of W2 = the mapped address of W2 (0x0004).

**Note:** When Register Indirect with Offset Addressing is used to specify the destination for an instruction, and Ws is the same register as Wd, the old value of Ws is used for Wd (i.e., the address offset is ignored).

### 3.4.2. Instruction Stalls and Exceptions

In order to maintain deterministic operation, instruction stalls are allowed to happen, even if they occur immediately prior to exception processing.

### 3.4.3. Instruction Stalls and Instructions that Change Program Flow

CALL and RCALL write to the stack using W15; therefore, they may be subject to an instruction stall if the source read of the subsequent instruction uses W15.

GOTO, RETFIE and RETURN instructions are never subject to an instruction stall because they do not perform write operations to the Working registers.

### 3.4.4. Instruction Stalls and REPEAT Loops

Instructions operating in a REPEAT loop are subject to instruction stalls, just like any other instruction. Stalls may occur on loop entry, loop exit and also during loop processing.

## 3.5. Byte Operations

Since the data memory is byte-addressable, most of the base instructions may operate in either Byte mode or Word mode. When these instructions operate in Byte mode, the following rules apply:

- All direct Working register references use the Least Significant Byte (LSB) of the 32-bit Working register and leave the Most Significant Byte (MSB) unchanged

- All indirect Working register references use the data byte specified by the 32-bit address stored in the Working register
- All file register references use the data byte specified by the byte address
- The STATUS Register (SR) is updated to reflect the result of the byte operation

It should be noted that data addresses are always represented as byte addresses. Additionally, the native data format is little-endian, which means that words are stored with the LSB at the lower address and the MSB at the adjacent, higher address (as shown in [Figure 3-2](#)). [Example 3-8](#) shows sample byte move operations and [Example 3-9](#) shows sample byte math operations.

**Note:** The instructions that operate in Byte mode must use the ".b" or ".B" instruction extension to specify a byte instruction. For example, the following two instructions are valid forms of a byte clear operation: CLR.b W0 or CLR.B W0.

#### Example 3-8. Sample Byte Move Operations

```
MOV.B      #0x30, W0          ; move the literal byte 0x30 to W0
```

Before Instruction:

```
W0 = 0x5555
```

After Instruction:

```
W0 = 0x5530
```

```
MOV.B      0x1000, W0          ; move the byte at 0x1000 to W0
```

Before Instruction:

```
W0 = 0x5555
Data Memory 0x1000 = 0x1234
```

After Instruction:

```
W0 = 0x5534
Data Memory 0x1000 = 0x1234
```

```
MOV.B      W0, 0x1001        ; byte move W0 to address 0x1001
```

Before Instruction:

```
W0 = 0x1234
Data Memory 0x1000 = 0x5555
```

After Instruction:

```
W0 = 0x1234
Data Memory 0x1000 = 0x3455
```

```
MOV.B      W0, [W1++]        ; byte move W0 to [W1], then post-inc W1
```

Before Instruction:

```
W0 = 0x1234
W1 = 0x1001
Data Memory 0x1000 = 0x5555
```

After Instruction:

```
W0 = 0x1234
W1 = 0x1002
Data Memory 0x1000 = 0x3455
```

### Example 3-9. Sample Byte Math Operations

```
CLR.B      [W6--]      ; byte clear [W6], then post-dec W6
```

Before Instruction:

```
W6 = 0x1001
Data Memory 0x1000 = 0x5555
```

After Instruction:

```
W6 = 0x1000
Data Memory 0x1000 = 0x0055
```

```
SUB.B      W0, #0x10, W1      ; byte subtract literal 0x10 from W0
; and store to W1
```

Before Instruction:

```
W0 = 0x1234
W1 = 0xFFFF
```

After Instruction:

```
W0 = 0x1234
W1 = 0xFF24
```

```
ADD.B      W0, W1, [W2++]      ; byte add W0 and W1, store to [W2]
; and post-inc W2
```

Before Instruction:

```
W0 = 0x1234
W1 = 0x5678
W2 = 0x1000
Data Memory 0x1000 = 0x5555
```

After Instruction:

```
W0 = 0x1234
W1 = 0x5678
W2 = 0x1001
Data Memory 0x1000 = 0x55AC
```

## 3.6. Word Move Operations

Even though the data space is byte-addressable, all move operations made in Word mode must be word-aligned. This means that for all source and destination operands, the Least Significant Address bit must be '0'. Likewise, for long-word accesses, last 2 bits of addresses must be zero. If a word/long-word move is made to or from an unaligned address, an address error exception is generated. [Figure 3-2](#) shows how bytes and words may be aligned in data memory. [Example 3-10](#) contains several legal word move operations.

When an exception is generated due to a misaligned access, the exception is taken after the instruction executes. If the illegal access occurs from a data read, the operation will be allowed to complete, but the Least Significant bit (LSb) of the source address will be cleared to force word alignment. If the illegal access occurs during a data write, the write will be inhibited. [Example 3-11](#) contains several illegal word move operations.

**Figure 3-2.** Data Alignment in Memory

0x4001	b0	0x4000
0x4003	b1	0x4002
0x4005	b3	0x4004
0x4007	b5	0x4006
0x4009	b7	0x4008
0x400B	b8	0x400A

**Legend:**

- b0 – byte stored at 0x4000
- b1 – byte stored at 0x4003
- b5:b2 – long-word stored at 0x4007:4004 (b2 is LSB)
- b7:b6 – word stored at 0x4009:0x4008 (b6 is LSB)
- b8 – byte stored at 0x400A

**Note:** Instructions that operate in Word mode are not required to use an instruction extension. However, they may be specified with an optional ".w" or ".W" extension. For example, the following instructions are valid forms of a word clear operation: CLR W0 or CLR.w W0 or CLR.W W0.

#### Example 3-10. Legal Word Move Operations

```
MOV      #0x30, W0      ; move the literal word 0x30 to W0
```

Before Instruction:

```
W0 = 0x5555
```

After Instruction:

```
W0 = 0x0030
```

```
MOV      0x1000, W0      ; move the word at 0x1000 to W0
```

Before Instruction:

```
W0 = 0x5555
Data Memory 0x1000 = 0x1234
```

After Instruction:

```
W0 = 0x1234
Data Memory 0x1000 = 0x1234
```

```
MOV      [W0], [W1++]      ; word move [W0] to [W1],
; then post-inc W1
```

Before Instruction:

```
W0 = 0x1234
W1 = 0x1000
```

```
Data Memory 0x1000 = 0x5555
Data Memory 0x1234 = 0xAAAA
```

After Instruction:

```
W0 = 0x1234
W1 = 0x1004
Data Memory 0x1000 = 0xAAAA
Data Memory 0x1234 = 0xAAAA
```

### Example 3-11. Illegal Word Move Operations

```
MOV      0x1001, W0          ; move the word at 0x1001 to W0
```

Before Instruction:

```
W0 = 0x5555
Data Memory 0x1000 = 0x1234
Data Memory 0x1002 = 0x5678
```

After Instruction:

```
W0 = 0x1234
Data Memory 0x1000 = 0x1234
Data Memory 0x1002 = 0x5678
```

ADDRESS ERROR TRAP GENERATED

(source address is misaligned, so MOV is performed)

```
MOV      W0, 0x1001          ; move W0 to the word at 0x1001
```

Before Instruction:

```
W0 = 0x1234
Data Memory 0x1000 = 0x5555
Data Memory 0x1002 = 0x6666
```

After Instruction:

```
W0 = 0x1234
Data Memory 0x1000 = 0x5555
Data Memory 0x1002 = 0x6666
```

ADDRESS ERROR TRAP GENERATED

(destination address is misaligned, so MOV is not performed)

```
MOV      [W0], [W1++]
; then post-inc W1          ; word move [W0] to [W1],
```

Before Instruction:

```
W0 = 0x1235
W1 = 0x1000
Data Memory 0x1000 = 0x1234
Data Memory 0x1234 = 0xAAAA
Data Memory 0x1236 = 0xBBB
```

After Instruction:

```
W0 = 0x1235
W1 = 0x1002
Data Memory 0x1000 = 0xAAAA
```

```
Data Memory 0x1234 = 0xAAAA
Data Memory 0x1236 = 0xBBB
```

ADDRESS ERROR TRAP GENERATED

(source address is misaligned, so MOV is performed)

### 3.7. Using 16-Bit Literal Operands

Several instructions that support Byte and Word mode have 16-bit operands. For byte instructions, a 16-bit literal is too large to use. Therefore, when 16-bit literals are used in Byte mode, the range of the operand must be reduced to eight bits or the assembler will generate an error. Likewise, the literal is zero-extended to 32-bits for long-word instructions. Likewise, the literal is zero-extended to 32-bits for long-word instructions. [Table 3-7](#) shows that the range of a 16-bit literal is 0:65535 in Word/long-word mode and 0:255 in Byte mode.

Instructions that employ 16-bit literals in Byte and Word/Long-Word mode are ADD, ADDC, AND, IOR, RETLW, SUB, SUBB and XOR. [Example 3-12](#) shows how positive and negative literals are used in Byte mode for the ADD instruction.

**Table 3-7.** 16-Bit Literal Coding

Literal Value	Word/long-word Mode kk kkkk kkkk	Byte Mode kkkk kkkk
0	0000 0000 0000 0000	0000 0000
1	0000 0000 0000 0001	0000 0001
2	0000 0000 0000 0010	0000 0010
127	0000 0000 0111 1111	0111 1111
128	0000 0000 1000 0000	1000 0000
255	0000 0000 1111 1111	1111 1111
256	0000 0001 0000 0000	N/A
512	0000 0010 0000 0000	N/A
1023	0000 0011 1111 1111	N/A
65535	1111 1111 1111 1111	N/A

**Note:** Using a literal value greater than 127 in Byte mode is functionally identical to using the equivalent negative two's complement value, since the MSb of the byte is set. When operating in Byte mode, the assembler will accept either a positive or negative literal value (i.e., #-10).

**Example 3-12.** Using 16-Bit Literals for Byte Operands

```
ADD.B      #0x80, W0          ; add 128 (or -128) to W0
ADD.B      #0x380, W0          ; ERROR... Illegal syntax for byte mode
ADD.B      #0xFF, W0           ; add 255 (or -1) to W0
ADD.B      #0x3FF, W0          ; ERROR... Illegal syntax for byte mode
ADD.B      #0xF, W0            ; add 15 to W0
ADD.B      #0x7F, W0           ; add 127 to W0
ADD.B      #0x100, W0          ; ERROR... Illegal syntax for byte mode
```

### 3.8. Bit Field Insert/Extract Instructions

PIC32A provides a set of instructions that operate on bit fields within a target word.

#### 3.8.1. BFEXT

This instruction can extract multiple bits from a W register or data memory location into a destination W register.

### 3.8.2. BFINS

This instruction can insert multiple bits from a source W register or an 8-bit literal value into a W register or data memory location.

In both instructions, the location and width of the bit field within the target word are defined as literal values within the instruction.

## 3.9. Software Stack Pointer and Frame Pointer

### 3.9.1. Software Stack Pointer

The PIC32A devices feature a software stack which facilitates function calls and exception handling. W15 is the default Stack Pointer (SP) and after any reset, it is initialized to 0x4000. This ensures that the SP will point to valid RAM and permits stack availability for exceptions, which may occur before the SP is set by the user software. The user may reprogram the SP during initialization to any location within data space.

The SP always points to the first available free word at TOS (Top-of-Stack) and fills the software stack, working from lower addresses towards higher addresses. It pre-decrements for a stack `POP` (read) and post-increments for a stack `PUSH` (write).

The software stack is manipulated using the `PUSH` and `POP` instructions. The `PUSH` and `POP` instructions are the equivalent of a `MOV` instruction with W15 used as the destination pointer. For example, the contents of W0 can be PUSHed onto the TOS by:

```
PUSH.L W0
```

This syntax is equivalent to:

```
MOV.L W0, [W15++]
```

The contents of the TOS can be returned to W0 by:

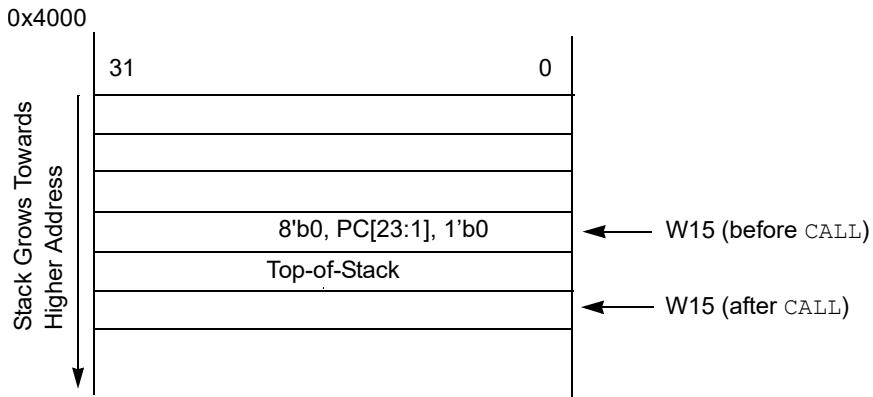
```
POP.L W0
```

This syntax is equivalent to:

```
MOV.L [--W15], W0
```

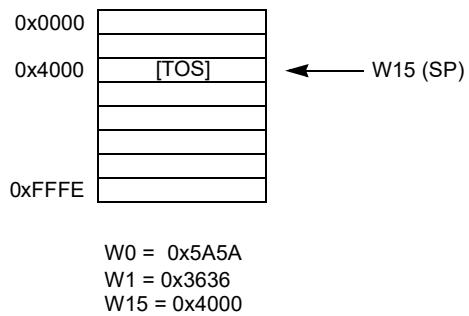
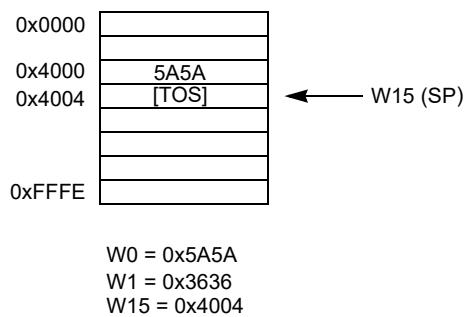
During any `CALL` instruction, the PC is PUSHed onto the stack, such that when the subroutine completes execution, program flow may resume from the correct location. When the PC is PUSHed onto the stack, the Most Significant Byte of PC is zero-extended, as shown in [Figure 3-3](#). During exception processing, the Most Significant seven bits of the PC are concatenated with the lower byte of the STATUS Register (SRL) and IPL[3] (CORCON[3]). This allows the primary STATUS Register contents and CPU Interrupt Priority Level to be automatically preserved during interrupts.

**Note:** In order to protect against misaligned stack accesses, W15[0] is always clear.

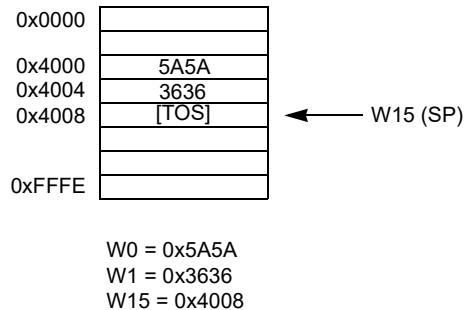
**Figure 3-3.** Stack Operation for CALL Instruction

### 3.9.1.1. Stack Pointer Example

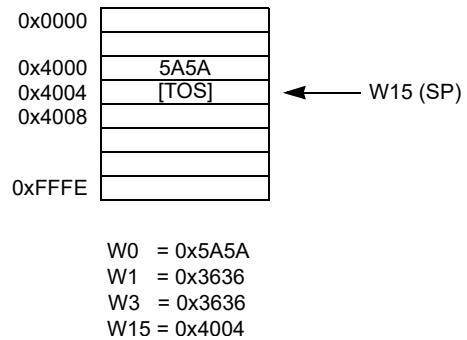
Figure 3-4 through Figure 3-7 show how the software stack is modified for the code snippet shown in [Example 3-13](#). Figure 3-4 shows the software stack before the first `PUSH` has executed. Note that the SP has the initialized value of 0x4000. Furthermore, the example loads 0x5A5A and 0x3636 to W0 and W1, respectively. The stack is PUSHed for the first time in [Figure 3-5](#) and the value contained in W0 is copied to TOS. W15 is automatically updated to point to the next available stack location and the new TOS is 0x4004. In [Figure 3-6](#), the contents of W1 are PUSHed onto the stack and the new TOS becomes 0x4008. In [Figure 3-7](#), the stack is POPped, which copies the last PUSHed value (W1) to W3. The SP is decremented during the `POP` operation and at the end of the example, the final TOS is 0x4004.

**Figure 3-4.** Stack Pointer Before the First `PUSH`**Figure 3-5.** Stack Pointer After “`PUSH.L W0`” Instruction

**Figure 3-6.** Stack Pointer After “PUSH.L W1” Instruction



**Figure 3-7.** Stack Pointer After “POP.L W3” Instruction



### **Example 3-13.** Stack Pointer Usage

```

MOV      #0x5A5A, W0          ; Load W0 with 0x5A5A
MOV      #0x3636, W1          ; Load W1 with 0x3636
PUSH.L   W0                  ; Push W0 to TOS
PUSH.L   W1                  ; Push W1 to TOS
POE.L   W3                  ; Pop TOS to W3

```

### 3.9.2. Software Stack Frame Pointer

A stack frame is a user-defined section of memory residing in the software stack. It is used to allocate memory for temporary variables, which a function uses, and one stack frame may be created for each function. W14 is the default Stack Frame Pointer (FP) and it is initialized to 0x0000 on any reset. If the Stack Frame Pointer is not used, W14 may be used like any other Working register.

The Link (`LNK`) and Unlink (`ULNK`) instructions provide stack frame functionality. The `LNK` instruction is used to create a stack frame. It is used during a call sequence to adjust the SP, such that the stack may be used to store temporary variables utilized by the called function. After the function completes execution, the `ULNK` instruction is used to remove the stack frame created by the `LNK` instruction. The `LNK` and `ULNK` instructions must always be used together to avoid stack overflow.

### 3.9.2.1. Stack Frame Pointer Example

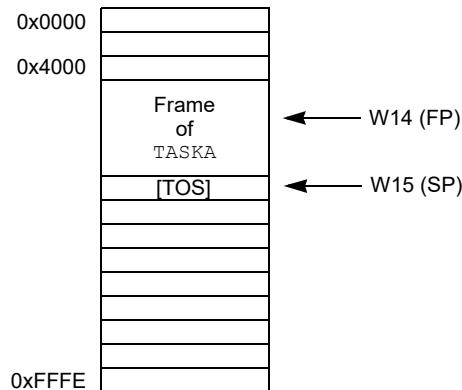
[Figure 3-8](#) through [Figure 3-10](#) show how a stack frame is created and removed for the code snippet shown in [Example 3-14](#). This example demonstrates how a stack frame operates and is not indicative of the code generated by the compiler. [Figure 3-8](#) shows the stack condition at the beginning of the example, before any registers are pushed to the stack. Here, W15 points to the first free stack location (TOS) and W14 points to a portion of stack memory allocated for the routine that is currently executing.

Before calling the function, "COMPUTE", the parameters of the function (W0, W1 and W2) are PUSHed on the stack. After the "CALL COMPUTE" instruction is executed, the PC changes to the address of "COMPUTE" and the return address of the function, "TASKA", is placed on the stack (Figure 3-9). Function "COMPUTE" then uses the "LNK #4" instruction to PUSH the calling routine's Frame Pointer value onto the stack and the new Frame Pointer will be set to point to the current Stack Pointer. Then, the literal 4 is added to the Stack Pointer address in W15, which reserves memory for two long-word of temporary data (Figure 3-10).

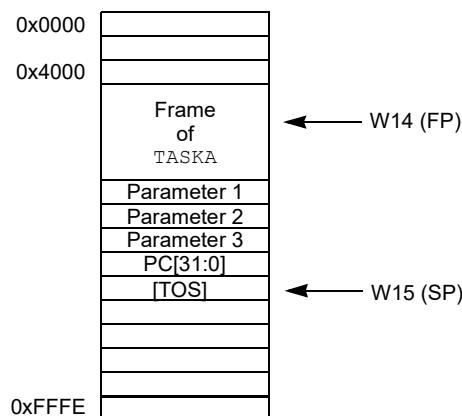
Inside the function, "COMPUTE", the FP is used to access the function parameters and temporary (local) variables. [W14 + n] will access the temporary variables used by the routine and [W14 - n] is used to access the parameters. At the end of the function, the ULNK instruction is used to copy the Frame Pointer address to the Stack Pointer and then POP the calling subroutine's Frame Pointer back to the W14 register. The ULNK instruction returns the stack back to the state shown in Figure 3-9.

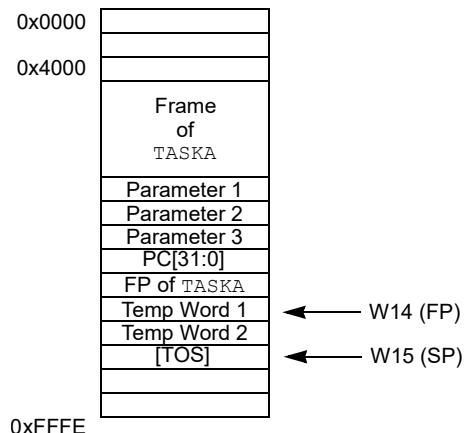
A RETURN instruction will return to the code that called the subroutine. The calling code is responsible for removing the parameters from the stack. The RETURN and POP instructions restore the stack to the state shown in Figure 3-8.

**Figure 3-8.** Stack at the Beginning of Frame Pointer Usage Example



**Figure 3-9.** Stack After "CALL COMPUTE" Executes



**Figure 3-10.** Stack After “LNK #4” Executes**Example 3-14.** Frame Pointer Usage

```

TASKA: ...
    PUSH W0      ; Push parameter 1
    PUSH W1      ; Push parameter 2
    PUSH W2      ; Push parameter 3
    CALL COMPUTE ; Call COMPUTE function
    POP W2       ; Pop parameter 3
    POP W1       ; Pop parameter 2
    POP W0       ; Pop parameter 1
    ...

COMPUTE:
    LNK #4      ; Stack FP, allocate 4 bytes for local variables
    ...
    ULINK       ; Free allocated memory, restore original FP
    RETURN      ; Return to TASKA

```

### 3.9.3. Stack Pointer Overflow

There is a Stack Limit register (SPLIM) associated with the Stack Pointer that is reset to 0x00000000. SPLIM is a 32-bit register, but SPLIM[1:0] is fixed to '00', because all stack operations must be long word-aligned. To match the Stack Pointer, the device address space is limited to 24-bits, so the upper 8-bits of SPLIM are always all 0.

The stack overflow check will not be enabled until a word write to SPLIM occurs, after which time it can only be disabled by a device reset. All Effective Addresses generated using W15 as a source or destination are compared against the value in SPLIM. Should the Effective Address be greater than the contents of SPLIM, then a stack error trap is generated.

If stack overflow checking has been enabled, a stack error trap will also occur if the W15 Effective Address calculation wraps over the end of data space (0x00FFFFFF).

### 3.9.4. Stack Pointer Underflow

The stack is initialized to 0x4000 during reset. A stack error trap will be initiated should the Stack Pointer address ever be less than 0x4000 .

## 3.10. Conditional Branch Instructions

Conditional branch instructions are used to direct program flow based on the contents of the STATUS Register. These instructions are generally used in conjunction with a compare class instruction, but they may be employed effectively after any operation that modifies the STATUS Register.

The compare instructions, CP, CP0 and CPB, perform a subtract operation (minuend – subtrahend), but do not actually store the result of the subtraction. Instead, compare instructions just update the flags in the STATUS Register, such that an ensuing conditional branch instruction may change program flow by testing the contents of the updated STATUS Register. If the result of the STATUS Register test is true, the branch is taken. If the result of the STATUS Register test is false, the branch is not taken.

The supported conditional branch instructions are shown in [Table 3-8](#). This table identifies the condition in the STATUS Register that must be true for the branch to be taken. In some cases, just a single bit is tested (as in BRA C), while in other cases, a complex logic operation is performed (as in BRA GT). Both signed and unsigned conditional tests are supported, and support is provided for DSP algorithms with the OA, OB, SA and SB condition mnemonics.

**Table 3-8.** Conditional Branch Instructions

Condition Mnemonic <sup>(1)</sup>	Description	Status Test
C	Carry (not Borrow)	C
GE	Signed Greater Than or Equal	(N&&OV)    (N&&OV)
GEU <sup>(2)</sup>	Unsigned Greater Than or Equal	C
GT	Signed Greater Than	(Z&&N&&OV)    (Z&&N&&OV)
GTU	Unsigned Greater Than	C&&Z
LE	Signed Less Than or Equal	Z    (N&&OV)    (N&&OV)
LEU	Unsigned Less Than or Equal	C    Z
LT	Signed Less Than	(N&&OV)    (N&&OV)
LTU <sup>(3)</sup>	Unsigned Less Than	C
N	Negative	N
NC	Not Carry (Borrow)	C
NN	Not Negative	N
NOV	Not Overflow	OV
NZ	Not Zero	Z
OA	Accumulator A Overflow	OA
OB	Accumulator B Overflow	OB
OV	Overflow	OV
SA	Accumulator A Saturate	SA
SB	Accumulator B Saturate	SB
Z	Zero	Z

**Notes:**

- Instructions are of the form: BRA mnemonic, Expr.
- GEU is identical to C and will reverse assemble to BRA C, Expr.
- LTU is identical to NC and will reverse assemble to BRA NC, Expr.

**Note:** The “Compare and Skip” instructions (CPBEQ, CPBGT, CPBLT, CPBNE, CPSEQ, CPSGT, CPSLT and CPSNE) do not modify the STATUS Register.

### 3.10.1. Floating Point Branch Instruction

Subsequent to floating point CPS/CPQ instructions setting one of the FSR ordering relations status bits, a subsequent floating point conditional branch (FBRA) instruction will (indirectly) examine these status bits, applying them to a logical predicate that represents the required condition. A list of the supported floating point branches and corresponding predicates is shown in [Table 3-9](#).

**Table 3-9.** FPU Conditional Branch Instruction

Condition Mnemonic <sup>(1)</sup>	Description	Status test
EQ	Equal	FSR.EQ
UNE	Unordered or Not Equal	(FSR.GT    FSR.LT    FSR.UN)
NE	Not Equal	(FSR.GT    FSR.LT)
UEQ	Unordered or Equal	(FSR.EQ    FSR.UN)
GT	Greater Than	FSR.GT
ULE	Unordered or Less Than or Equal	(FSR.LT    FSR.EQ    FSR.UN)
GE	Greater Than or Equal	(FSR.GT    FSR.EQ)
ULT	Unordered or Less Than	(FSR.LT    FSR.UN)
LT	Less Than	FSR.LT
UGE	Unordered or Greater Than or Equal	(FSR.GT    FSR.EQ    FSR.UN)
LE	Less Than or Equal	(FSR.LT    FSR.EQ)
UGT	Unordered or Greater Than	(FSR.GT    FSR.UN)
OR	Ordered	(FSR.GT    FSR.LT    FSR.EQ)
UN	Unordered	FSR.UN

**Note:**

- Instructions are of the form: FBRA mnemonic, Expr.

### 3.11. Z Status Bit

The Z Status bit is a special Zero Status bit that is useful for extended precision arithmetic. The Z bit functions like a normal Z flag for all instructions, except those that use the Carry/Borrow input (ADDC, CPB, SUBB and SUBBR). For the ADDC, CPB, SUBB and SUBBR instructions, the Z bit can only be cleared and never set. If the result of one of these instructions is non-zero, the Z bit will be cleared and will remain cleared, regardless of the result of subsequent ADDC, CPB, SUBB or SUBBR operations. This allows the Z bit to be used for performing a simple zero check on the result of a series of extended precision operations.

A sequence of instructions working on multiprecision data (starting with an instruction with no Carry/Borrow input) will automatically logically AND the successive results of the zero test. All results must be zero for the Z flag to remain set at the end of the sequence of operations. If the result of the ADDC, CPB, SUBB or SUBBR instruction is non-zero, the Z bit will be cleared and remain cleared for all subsequent ADDC, CPB, SUBB or SUBBR instructions.

### 3.12. DSP Data Formats

#### 3.12.1. Integer and Fractional Data

The DSP engine supports both integer and fractional data types. Integer data is inherently represented as a signed two's complement value, where the MSb is defined as a Sign bit. Generally speaking, the range of an N-bit two's complement integer is  $-2^{N-1}$  to  $2^{N-1} - 1$ . For a 16-bit integer, the data range is -32768 (0x8000) to 32767 (0x7FFF), including '0'. For a 32-bit integer, the data range is -2,147,483,648 (0x8000 0000) to 2,147,483,647 (0x7FFF FFFF).

Fractional data is represented as a two's complement number, where the MSb is defined as a Sign bit and the radix point is implied to lie just after the Sign bit. This format is commonly referred to as 1.15 (or Q15) format, where 1 is the number of bits used to represent the integer portion of the number and 15 is the number of bits used to represent the fractional portion. The range of an N-bit two's complement fraction with this implied radix point is -1.0 to  $(1 - 2^{1-N})$ . For a 16-bit fraction, the 1.15 data range is -1.0 (0x8000) to 0.999969482 (0x7FFF), including 0.0 and it has a precision of  $3.05176 \times 10^{-5}$ . In Normal Saturation mode, the 32-bit accumulators use a 1.31 format, which enhances the precision to  $4.6566 \times 10^{-10}$ . In Fractional mode, the 32x32-bit PIC32A multiplier generates a Q1.63 product, which has a precision of  $1.08420 \times 10^{-19}$ .

The dynamic range of the accumulators can be expanded by using the eight bits of the Upper Accumulator register (ACCxU) as Guard bits. Guard bits are used if the value stored in the accumulator overflows beyond the 62<sup>nd</sup> bit and they are useful for implementing DSP algorithms. This mode is enabled when the ACCSAT bit (CORCON[4]) is set to '1' and it expands the accumulators to 72 bits. The Guard bits are also used when the accumulator saturation is disabled. The accumulators then support an integer range of  $-2^{71}$  (0x80 0000 0000 0000 0000) to  $2^{71-1}$  (0x7F FFFF FFFF FFFF). In Fractional mode, the Guard bits of the accumulator do not modify the location of the radix point and the 72-bit accumulators use a 9.63 fractional format. Note that all fractional operation results are stored in the 72-bit accumulator, justified with a 1.63 radix point. As in Integer mode, the Guard bits merely increase the dynamic range of the accumulator. 9.63 fractions have a range of -256.0 (0x80 0000 0000 0000 0000) to 256.0  $-1.08420 \times 10^{-19}$  (0x7F FFFF FFFF FFFF).

**Table 3-10** identifies the range and precision of integers and fractions on 16-bit, 32-bit and 72-bit registers.

With the exception of DSP multiplies, the ALU operates identically on integer and fractional data. Namely, an addition of two integers will yield the same result (binary number) as the addition of two fractional numbers. The only difference is how the result is interpreted by the user. However, multiplies performed by DSP operations are different. In these instructions, data format selection is made by the IF bit (CORCON[0]) and it must be set accordingly ('0' for Fractional mode, '1' for Integer mode). This is required because of the implied radix point used by fractional numbers.

**Table 3-10.** PIC32A Data Ranges

Register Size	Integer Range	Fraction Range	Fraction Resolution
16-bit	-32768 to 32767	-1.0 to $(1.0 - 2^{-15})$	$3.052 \times 10^{-5}$
32-bit	-2,147,483,648 to 2,147,483,647	-1.0 to $(1.0 - 2^{-31})$	$4.657 \times 10^{-10}$
72-bit	$-2^{71}$ to $2^{71}$	-256.0 to 256.0 $-1.08420 \times 10^{-19}$	$1.08420 \times 10^{-19}$

### 3.12.2. Integer and Fractional Data Representation

Both integer and fractional data treat the MSb as a Sign bit, and the binary exponent decreases by one as the bit position advances toward the LSb. The binary exponent for an N-bit integer starts at (N-1) for the MSb and ends at '0' for the LSb. For an N-bit fraction, the binary exponent starts at '0' for the MSb and ends at (1-N) for the LSb (as shown in [Figure 3-11](#) for a positive value and in [Figure 3-12](#) for a negative value).

Conversion between integer and fractional representations can be performed using simple division and multiplication. To go from an N-bit integer to a fraction, divide the integer value by  $2^{N-1}$ . Similarly, to convert an N-bit fraction to an integer, multiply the fractional value by  $2^{N-1}$ .

**Figure 3-11.** Different Representations of 0x4001

Integer:

0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	1
$-2^{15}$	$2^{14}$	$2^{13}$	$2^{12}$	.....											$2^0$

$$0x4001 = 2^{14} + 2^0 = 16384 + 1 = \mathbf{16385}$$

1.15 Fractional:

0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	1
$-2^0$	$.2^{-1}$	$2^{-2}$	$2^{-3}$	.....											$2^{-15}$

↓  
Implied Radix Point

$$0x4001 = 2^{-1} + 2^{-15} = 0.5 + .000030518 = \mathbf{0.500030518}$$

**Figure 3-12.** Different Representations of 0xC002

Integer:

1	1	0	0	0	0	0	0	0	0	0	0	0	0	1	0
$-2^{15}$	$2^{14}$	$2^{13}$	$2^{12}$	.....											$2^0$

$$0xC002 = -2^{15} + 2^{14} + 2^1 = -32768 + 16384 + 2 = \mathbf{-16382}$$

1.15 Fractional:

1	1	0	0	0	0	0	0	0	0	0	0	0	0	1	0
$-2^0$	$.2^{-1}$	$2^{-2}$	$2^{-3}$	.....											$2^{-15}$

↓  
Implied Radix Point

$$0xC002 = -2^0 + 2^1 + 2^{14} = -1.0 + 0.5 + 0.000061035 = \mathbf{-0.499938965}$$

### 3.13. Accumulator Usage

Accumulators A and B are utilized by DSP instructions to perform mathematical and shifting operations. Since the accumulators are 72-bits wide and the X and Y data paths are only 32 bits, the method to load and store the accumulators must be understood.

Item A in [Figure 3-13](#) shows that each 72-bit accumulator (ACCA and ACCB) consists of an 8-bit upper register (ACCxU), a 32-bit high register (ACCxH) and a 32-bit low register (ACCxL). To address the bus alignment requirement and provide the ability for 1.31 math, ACCxH is used as a destination register for loading the accumulator (with the LAC instruction) and also as a source register for storing the accumulator (with the SACR instruction). This is represented by Item B in [Figure 3-13](#), where the upper and lower portions of the accumulator are shaded. In reality, during accumulator

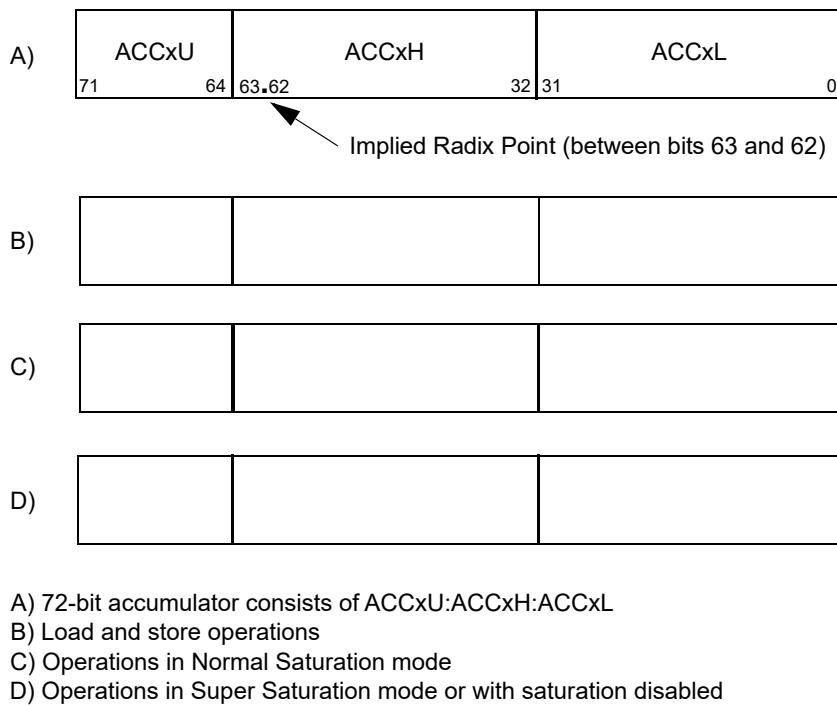
loads, ACCxL is zero backfilled and ACCxU is sign-extended to represent the sign of the value loaded in ACCxH.

**Note:** PIC32A devices provide SUAC, SAC and SLAC instructions to store ACCxU, ACCxH and ACCxL values respectively. Similarly, LUAC, LAC and LLAC instructions can be used to load values to ACCxU, ACCxH and ACCxL.

When normal (63-bit) saturation is enabled, DSP operations (such as ADD, MAC, MSC, etc.) solely utilize ACCxH:ACCxL (item C in [Figure 3-13](#)) and ACCxU is only used to maintain the sign of the value stored in ACCxH:ACCxL. For instance, when an MPY instruction is executed, the result is stored in ACCxH:ACCxL, and the sign of the result is extended through ACCxU.

When super saturation is enabled, or when saturation is disabled, all registers of the accumulator may be used (item D in [Figure 3-13](#)), and the results of DSP operations are stored in ACCxU:ACCxH:ACCxL. The benefit of ACCxU is that it increases the dynamic range of the accumulator, as described in [Integer and Fractional Data](#). Refer to [Table 3-10](#) to see the range of values which may be stored in the accumulator when in Normal and Super Saturation modes.

**Figure 3-13.** Accumulator Alignment and Usage



### 3.14. Accumulator Access

Accumulator can be accessed with File Register or Indirect Addressing, using any instruction which supports such addressing. However, it is recommended that the DSP instructions, LAC, SAC and SSCR, be used to load and store the accumulators, since they provide sign-extension, shifting and rounding capabilities. LAC, SAC, and SSCR instruction details are provided in [Instruction Descriptions](#).

**Notes:**

1. For convenience, ACCAU and ACCBU are sign-extended to 32 bits. This provides the flexibility to access these registers using either Byte or Word mode (when File Register or Indirect Addressing is used).
2. The OA, OB, SA or SB bit cannot be set by writing overflowed values to the memory-mapped accumulators using `MOV` instructions, as these Status bits are only affected by DSP operations.
3. PIC32A SUAC, SAC and SLAC instructions to store ACCxU, ACCxH and ACCxL values respectively. Similarly, LUAC, LAC and LLAC instructions can be used to load values to ACCxU, ACCxH and ACCxL.

## 3.15. DSP MAC Instructions

The DSP Multiply and Accumulate (`MAC`) operations are a special suite of instructions that provide the most efficient use of the PIC32A architectures. The DSP `MAC` instructions, shown in [Table 3-11](#), utilize both the X and Y data paths of the CPU core.

`MAC` instructions are also capable of performing an operation with one accumulator, while storing out the rounded contents of the alternate accumulator. This feature is called Accumulator Write-Back (WB) and it provides flexibility for the software developer. For instance, the Accumulator WB may be used to run two algorithms concurrently or efficiently process complex numbers, among other things.

**Table 3-11.** DSP MAC Instructions

Instruction	Description
<code>ED</code>	Euclidean Distance (No Accumulate)
<code>EDAC</code>	Euclidean Distance
<code>MAC</code>	Multiply and Accumulate
<code>SQRAC</code>	Square and Accumulate
<code>MPY</code>	Multiply to Accumulator
<code>MPY.N</code>	Negative Multiply to Accumulator
<code>MSC</code>	Multiply and Subtract from Accumulator
<code>SQRSC</code>	Square and Subtract from Accumulator
<code>SQR</code>	Square to Accumulator
<code>SQRN</code>	Negated Square to Accumulator

### 3.15.1. MAC Operations

The mathematical operations performed by the `MAC` class of DSP instructions center around multiplying or squaring the contents of two Working registers and either adding, subtracting or storing the result to either Accumulator A or Accumulator B. This is the operation of the `MAC`, `MPY`, `MPY.N`, `MSC`, `SQR`, `SQRAC`, `SQRSC` and `SQRN` instructions.

For the `ED` and `EDAC` instructions, the same multiplicand operand must be specified by the instruction because this is the definition of the Euclidean Distance operation.

### 3.15.2. Rounding Modes

The accumulator stores the instruction `SACR`, and all AWB write operations can operate with one of two different accumulator rounding modes. The rounding function, either conventional (biased) or convergent (unbiased), is applied during the accumulator write (store). The rounding mode is selected by `CORCON.RND`.

Conventional (biased) rounding examines bit 31 (for long word sized data op) or bit 47 (for word sized data op) of the 72-bit accumulator, zero extends it, and adds it to the MS portion of the accumulator. Consequently, if the LS word of the accumulator is between `32'h80000000` and `32'hFFFFFFFFF` (for long word sized data op) or `48'h8000000000000000` and `48'hFFFFFFFFFFFF` (for word

sized data op), the MS portion of the accumulator will be incremented. Alternatively, if the LS word of the accumulator is between 0x00000000 and 32'h7FFFFFFF (long word sized data op) or 48'h000000000000 and 48'h7FFFFFFFFF (word sized data op), the MS portion of the accumulator is left unchanged. A consequence of this algorithm is that over a succession of random rounding operations, the value will tend to be biased slightly positive.

Convergent (or unbiased) rounding operates in the same manner as conventional rounding, except when the LS word equals 32'h80000000 (for long word sized data op) or 48'h800000000000 (for word sized data op). If this is the case, the LS-bit of the MS portion of the accumulator (that is, bit 32 or bit 48) is examined. If it is 1, the MS-portion of the accumulator is incremented. If it is 0, the MS portion of the accumulator is not modified.

Assuming that the bit examined is effectively random in nature, this scheme removes any rounding bias that may accumulate.

The two rounding modes are defined in the following tables for each operand (and accumulator write) size. The rounding block will round based on the contents of AccX[31:0] (Long word round) or AccX[47:0] (Word round) depending upon the data size selected for the instruction.

**Table 3-12.** Rounding Operation (Long Word Operands and AWB Write)

ACCX[31:0] Value	Increment ACCX[71:32]?			
	Convergent		Biased	
	ACCX[32] = 0	ACCX[32] = 1	ACCX[32] = 0	ACCX[32] = 1
ACCX[31] = 0	No	No	No	No
ACCX[31:0] = 32'h80000000	No	Yes	Yes	Yes
ACCX[31] = 1	Yes	Yes	Yes	Yes

**Table 3-13.** Rounding Operation (Word Operands and AWB Write)

ACCX[31:0] Value	Increment ACCX[71:48]?			
	Convergent		Biased	
	ACCX[48] = 0	ACCX[48] = 1	ACCX[48] = 0	ACCX[48] = 1
ACCX[47] = 0	No	No	No	No
ACCX[47:0] = 48'h800000000000	No	Yes	Yes	Yes
ACCX[47] = 1	Yes	Yes	Yes	Yes

### 3.15.3. MAC Write-Back

The Write-Back ability of the MAC class of DSP instructions facilitates efficient processing of algorithms. This feature allows one mathematical operation to be performed with one accumulator and the rounded contents of the other accumulator to be stored in the same cycle.

The following addressing modes are supported:

- W0, W1, W2, W3, W13 or W14 register direct: the rounded contents of the non-target accumulator are written into the destination register as a 1.15 (Word mode) or 1.31 (Long Word mode) fraction. As is the case for any register direct word write, the value written is zero extended to 32-bits prior to being written into the destination WREG.
- [W13++] or [W15++] register indirect with post increment: the rounded contents of the non-target accumulator are written into the address pointed to by W13 or W15 as a 1.15 (Word mode) or 1.31 (Long Word mode) fraction. The destination EA W-reg is then incremented by two (for a word write) or four (for a long word write).

### 3.15.4. MAC Syntax

The syntax of the MAC class of instructions can have several formats, which depend on the instruction type and the operation it is performing with respect to prefetches and Accumulator

Write-Back. All MAC class instructions must specify a target accumulator along with two multiplicands, as shown in [Example 3-15](#).

If an Accumulator Write-Back is used in the instruction, it is specified last. By definition, the accumulator not used in the mathematical operation is stored, so the Write-Back accumulator is not specified in the instruction. Legal forms of Accumulator Write-Back (WB) are shown in [Example 3-16](#).

#### Example 3-15. Base MAC Syntax

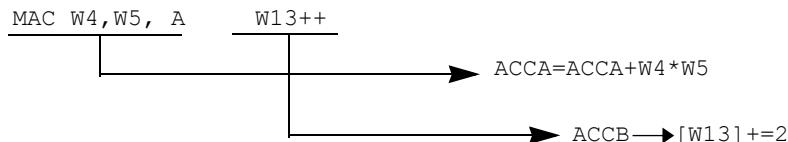
```
; MAC with no prefetch
MAC W4,W5, A
```

```
; MAC with no prefetch
MAC W7,W7, B
```



#### Example 3-16. MAC Accumulator WB Syntax

```
; MAC with indirect WB of ACCB
```



## 3.16. DSP Accumulator Instructions

The DSP accumulator instructions provide the ability to add, negate, shift, load and store the contents of either 72-bit accumulator. In addition, the ADD and SUB instructions allow the two accumulators to be added or subtracted from each other. DSP accumulator instructions are shown in [Table 3-14](#) and instruction details are provided in [Instruction Descriptions](#).

**Table 3-14.** DSP Accumulator Instructions

Instruction	Description	Accumulator WB?
ADD	Add Accumulators	No
ADD	16-Bit Signed Accumulator Add	No
LAC	Load Accumulator	No
LUAC	Load upper accumulator (ACCxU)	No
LLAC	Load lower accumulator (ACCxL)	No
NEG	Negate Accumulator	No
SAC	Store Accumulator	No
SUAC	Store upper accumulator (ACCxU)	No
SLAC	Store lower accumulator (ACCxL)	No
SACR	Store Rounded Accumulator	No
SFTAC	Arithmetic Shift Accumulator by Literal	No
SFTAC	Arithmetic Shift Accumulator by (Wn)	No
SUB	Subtract Accumulators	No

### 3.17. Scaling Data with the FBCL Instruction

To minimize quantization errors that are associated with data processing using DSP instructions, it is important to utilize the complete numerical result of the operations. This may require scaling data up to avoid underflow (i.e., when processing data from a 12-bit ADC) or scaling data down to avoid overflow (i.e., when sending data to a 10-bit DAC). The scaling, which must be performed to minimize quantization error, depends on the dynamic range of the input data which is operated on and the required dynamic range of the output data. At times, these conditions may be known beforehand and fixed scaling may be employed. In other cases, scaling conditions may not be fixed or known and then dynamic scaling must be used to process data.

The `FBCL` instruction (Find First Bit Change Left) can efficiently be used to perform dynamic scaling, because it determines the exponent of a value. A fixed point or integer value's exponent represents the amount that the value may be shifted before overflowing. This information is valuable because it may be used to bring the data value to "full scale", meaning that its numeric representation utilizes all the bits of the register it is stored in.

The `FBCL` instruction determines the exponent of a word by detecting the first bit change, starting from the value's Sign bit and working towards the LSB. Since the PIC32A device's barrel shifter uses negative values to specify a left shift, the `FBCL` instruction returns the negated exponent of a value. If the value is being scaled up, this allows the ensuing shift to be performed immediately with the value returned by `FBCL`. Additionally, since the `FBCL` instruction only operates on signed quantities, `FBCL` produces results in the range of -15:0 for word operation and -31:0 for long word operation. When the `FBCL` instruction returns 0, it indicates that the value is already at full scale. When the instruction returns -31, it indicates that the value cannot be scaled (as is the case with 0x0 and 0xFFFFFFFF). [Table 3-15](#) shows word data with various dynamic ranges, their exponents and the value after scaling the data to maximize the dynamic range. [Example 3-17](#) shows how the `FBCL` instruction may be used for block processing.

**Table 3-15.** Scaling Examples

Word Value	Exponent	Full-Scale Value (Word Value << Exponent)
0x0001	14	0x4000
0x0002	13	0x4000
0x0004	12	0x4000
0x0100	6	0x4000
0x01FF	6	0x7FC0
0x0806	3	0x4030
0x0007	1	0x400E
0x4800	0	0x4800
0x7000	0	0x7000
0x8000	0	0x8000
0x900A	0	0x900A
0xE001	2	0x8004
0xFF07	7	0x8380

**Note:** For the word values, 0x0000 and 0xFFFF, the `FBCL` instruction returns -15.

As a practical example, assume that block processing is performed on a sequence of data with very low dynamic range stored in 1.15 fractional format. To minimize quantization errors, the data may be scaled up to prevent any quantization loss that may occur as it is processed. The `FBCL` instruction can be executed on the sample with the largest magnitude to determine the optimal scaling value for processing the data. Note that scaling the data up is performed by left shifting the data. This is demonstrated with the code snippet below.

**Example 3-17.** Scaling with FBCL

```

; assume W0 contains the largest absolute value of the data block
; assume W4 points to the beginning of the data block
; assume the block of data contains BLOCK_SIZE words
; determine the exponent to use for scaling
FBCL      W0, W2          ; store exponent in W2
; scale the entire data block before processing
DO        # (BLOCK_SIZE-1), SCALE
LAC       [W4], A           ; move the next data sample to ACCA
SFTAC    A, W2             ; shift ACCA by W2 bits
SCALE:
SAC       A, [W4++]         ; store scaled input (overwrite
original)
; now process the data
; (processing block goes here)

```

**3.18. Data Range Limit Instructions**

The PIC32A family provides special instructions that automatically limit the data in a W register or an accumulator to lie within a user-specified numerical range. These include the FLIM, MAX, MIN and MINZ instructions.

**3.18.1. FLIM/FLIM.V**

The FLIM instruction simultaneously compares a maximum and a minimum data limit value with the specified W register (or data pointed to by the W register) and clamps the target data to the user-specified limit if the limit is exceeded. SR Status bits are set accordingly for subsequent signed branching. In the FLIM.V instruction, an additional W register is specified, in which the limit test result (known as "limit excess") value is loaded.

**3.18.2. MAX/MAX.V**

The MAX instruction compares a maximum data limit value (stored in a W register or the other accumulator) with the target accumulator and clamps the target accumulator to the user-specified limit if this upper limit is exceeded. SR Status bits are set accordingly for subsequent signed branching. In the MAX.V instruction, an additional W register (or W register in Indirect Addressing mode) is specified in which the limit excess value is loaded.

**3.18.3. MIN/MIN.V/MINZ**

The MIN instruction compares a minimum data limit value (stored in a W register or the other accumulator) with the target accumulator and clamps the target accumulator to the user-specified limit if the data is smaller than this minimum limit. SR Status bits are set accordingly for subsequent signed branching. In the MIN.V instruction, an additional W register (or W register in Indirect Addressing mode) is specified, in which the limit excess value is loaded. The MINZ instruction is simply a conditional version of the MIN instruction, which is executed only when Z = 1.

**3.19. Normalizing the Accumulator with the NORM Instruction**

The NORM instruction automatically normalizes the target accumulator to obtain the largest fractional value possible without overflow. The target accumulator must contain signed fractional data for the instruction result to be valid. It will shift the target accumulator right or left by as many bits as required to normalize the data, keeping the sign consistent. The shift value is stored in a destination address. The N Status bit reflects the direction of the accumulator shift.

If the accumulator cannot be normalized, the accumulator contents will not be affected.

## 4. Instruction Descriptions

### 4.1. Instruction Symbols

All the symbols used in [Instruction Descriptions](#) are listed in [Table 2](#).

### 4.2. Instruction Encoding Field Descriptors Introduction

All instruction encoding field descriptors used in [Instruction Descriptions](#) are shown in [Table 4-2](#) through [Table 4-10](#).

**Table 4-1.** Instruction Encoding Field Descriptors

Field	Description
A	Accumulator Selection bit: 0 = ACCA; 1 = ACCB
aaa	Accumulator Write Back mode (see <a href="#">Table 4-9</a> )
B	Byte Mode Selection bit: 0 = word operation; 1 = byte operation
bbb	3-bit bit position select: 000 = LSB; 111 = MSB
bbbb	5-bit bit position select: 00000 = LSB; 11111 = MSB
cccc	Bit field instructions LSb value
D	Destination Address bit: 0 = result stored in W0; 1 = result stored in file register
ddd	Wd destination register select: 0000 = W0; 1111 = W15
ddddd	Coprocessor destination register select (Fd for FPU where 00000 = F0; 11111 = F31)
E	MULxx Result size: 0 = 32-bit in Wnd; 1 = 64-bit in (Wnd+1, Wnd)
F	Selects between W15 (F = 0) and W14 (F = 1) registers
(ffff) ffff ffff ffff ffff	16-bit or 20-bit register file address (addressable space varies depending upon instruction class)
G	Bit test destination: 0 = Z flag bit; 1 = C flag bit
I	MULAxz Multiply mode: 0 = Fractional; 1 = Integer
IIii	X data fetch operation
JJJjjj	Y data fetch operation
k	1-bit literal field, constant data
kkk	3-bit literal field, constant data
k kkkk	5-bit literal field, constant data
kk kkkk	6-bit literal field, constant data
kkkk kkkk	8-bit literal field, constant data
kkkk kkkk kkkk kkkk	16-bit literal field, constant data
kkkk kkkk kkkk kkkk kkkk kkkk kkkk	32-bit literal field, constant data
L	Long (32-bit) Mode Selection bit: 0 = word or byte operation; 1 = long operation
mmmmm	Bit field instructions MSb value
nnnn nnnn nnnn nnnn nnnn n	21-bit signed instruction word offset field for relative branch/call instructions
nnnn nnnn nnnn nn00 nnnn nnnn	24-bit program address for goto/call instructions
ppp	Addressing mode for Ws source register (see <a href="#">Table 4-2</a> )
qqq	Addressing mode for Wd destination register (see <a href="#">Table 4-3</a> )
R	Selects between FPU coprocessor special registers or F-reg
rrr	Condition select for conditional move (MOVIF) instruction

**Table 4-1.** Instruction Encoding Field Descriptors (continued)

Field	Description
S	Opcode size select (16-bit: S = 0; 32-bit: S = 1)
ssss	Ws source or Wn source/destination register select: 0000 = W0; 1111 = W15
sssss	Coprocessor source register select (Fs for FPU where 00000 = F0; 11111 = F31)
T	Selects between Ws (T = 0) and SR (T = 1) target registers
U	Unused (don't care) Instruction bit. Assembler to assign '0'
V	FLIMW: Selects result format for Wnd (refer to instruction description) MULxx: Selects between unsigned Ws (V = 0) and signed Ws (V = 1).
W	Destination write control: 0 = Wd write not required; 1 = Wd write required
www	Source Wb base register select: 0000 = W0; 1111 = W15
zz	Coprocessor select

**Table 4-2.** Addressing Modes for Ws Source Register

ppp	Addressing Mode	Source Operand
000	Register Direct	Wns
001	Indirect	[Ws]
010	Indirect with Post-Decrement	[Ws--]
011	Indirect with Post-Increment	[Ws++]
100	Indirect with Pre-Decrement	[--Ws]
101	Indirect with Pre-Increment	[++Ws]
110	Status Register Direct	SR (Source)
111	Indirect with Register Offset	[Ws+Wb]

**Table 4-3.** Addressing Modes for Wd Destination Register

qqq	Addressing Mode	Destination Operand
000	Register Direct	Wnd
001	Indirect	[Wd]
010	Indirect with Post-Decrement	[Wd--]
011	Indirect with Post-Increment	[Wd++]
100	Indirect with Pre-Decrement	[--Wd]
101	Indirect with Pre-Increment	[++Wd]
110	Status Register (SR) Direct	SR (Destination)
111	Indirect with Register Offset	[Wd+Wb]

**Table 4-4.** Destination Addressing Modes for MCU Multiplications

dddd	Destination
0000	W1:W0
0001	W0
0010	W3:W2
0011	W2
0100	W5:W4
0101	W4
0110	W7:W6
0111	W6
1000	W9:W8
1001	W8
1010	W11:W10

**Table 4-4.** Destination Addressing Modes for MCU Multiplications (continued)

ddd	Destination
1011	W10
1100	W13:W12
1101	W12
1110	ACCA[39:0]
1111	ACCB[39:0]

**Table 4-5.** Offset Addressing Modes for Ws Source Register (with Register Offset)

ggg	Addressing Mode	Source Operand
000	Register Direct	Ws
001	Indirect	[Ws]
010	Indirect with Post-Decrement	[Ws--]
011	Indirect with Post-Increment	[Ws++]
100	Indirect with Pre-Decrement	[--Ws]
101	Indirect with Pre-Increment	[++Ws]
11x	Indirect with Register Offset	[Ws+Wb]

**Table 4-6.** Offset Addressing Modes for Wd Destination Register (with Register Offset)

hh	Addressing Mode	Source Operand
000	Register Direct	Wd
001	Indirect	[Wd]
010	Indirect with Post-Decrement	[Wd--]
011	Indirect with Post-Increment	[Wd++]
100	Indirect with Pre-Decrement	[--Wd]
101	Indirect with Pre-Increment	[++Wd]
11x	Indirect with Register Offset	[Wd+Wb]

**Table 4-7.** MAC or MPY Source Operands – Same Working Register

mm	Multiplicands
00	W4 * W4
01	W5 * W5
10	W6 * W6
11	W7 * W7

**Table 4-8.** MAC or MPY Source Operands – Different Working Register

mmm	Multiplicands
000	W4 * W5
001	W4 * W6
010	W4 * W7
011	Invalid
100	W5 * W6
101	W5 * W7
110	W6 * W7
111	Invalid

**Table 4-9.** MAC Accumulator Write-Back Selection

aaa	Write-Back Selection
000	W0 = Other Accumulator
001	W1 = Other Accumulator
010	W2 = Other Accumulator
011	W3 = Other Accumulator
100	W13 = Other Accumulator (Direct Addressing)
101	[W13++] = Other Accumulator (Indirect Addressing with Post-Increment)
110	[W15++] = Other Accumulator
111	No Accumulator Write Back

**Table 4-10.** Accumulator Selection

A	Target Accumulator
0	Accumulator A
1	Accumulator B

### 4.3. Instruction Description Example

The example description below is for the fictitious instruction, FOO. The following example instruction was created to demonstrate how the table fields (syntax, operands, operation, etc.) are used to describe the instructions presented in [Instruction Descriptions](#).

#### FOO

FOO	The Header field summarizes what the instruction does
Syntax:	The Syntax field consists of an optional label, the instruction mnemonic, any optional extensions which exist for the instruction and the operands for the instruction. Most instructions support more than one operand variant to support the various addressing modes. In these circumstances, all possible instruction operands are listed beneath each other and are enclosed in braces.
Operands:	The Operands field describes the set of values that each of the operands may take. Operands may be accumulator registers, file registers, literal constants (signed or unsigned) or Working registers.
Operation:	The Operation field summarizes the operation performed by the instruction.
Status Affected:	The Status Affected field describes which bits of the STATUS Register are affected by the instruction. Status bits are listed by bit position in descending order.
Encoding:	The Encoding field shows how the instruction is bit encoded. Individual bit fields are explained in the Description field and complete encoding details are provided in <a href="#">Table 5.2</a> .
Description:	The Description field describes in detail the operation performed by the instruction. A key for the encoding bits is also provided.
Words:	The Words field contains the number of program words that are used to store the instruction in memory.
Cycles:	The Cycles field contains the number of instruction cycles that are required to execute the instruction.
Examples:	The Examples field contains examples that demonstrate how the instruction operates. "Before" and "After" register snapshots are provided, which allow the user to clearly understand what operation the instruction performs.

## 4.4. Instruction Descriptions (A to BZ)

ADD		Add Wb and Ws					
Syntax:	{label:}	ADD.b	Wb,	Ws,	Wd		
		ADD.bz		[Ws],	[Wd]		
		ADD{.w}		[Ws++],	[Wd++]		
		ADD.l		[Ws--],	[Wd--]		
				[++Ws],	[++Wd]		
				[--Ws],	[--Wd]		
Operands:		$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$					
Operation:		$(Wb) + (Ws) \rightarrow Wd$					
Status Affected:		C, N, OV, Z					
Encoding:		S110	000L	dddd	ssss	pppq	qqww
Description:		Add the contents of the source register Ws and the contents of the base register Wb and place the result in the destination register Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.					
I-Words:		1 or 0.5					
Cycles:		1					

Example 1:	ADD.B	W5, W6, W7	;	Add W5 to W6, store result in W7 ;(Byte mode)
------------	-------	------------	---	--

Before Instruction		After Instruction	
W5	AB00	W5	AB00
W6	0030	W6	0030
W7	FFFF	W7	FF30
SR	0000	SR	0000

Example 2:	ADD	W5, W6, W7	;	Add W5 to W6, store result in W7 ;(Word mode)
------------	-----	------------	---	--

Before Instruction		After Instruction	
W5	AB00	W5	AB00
W6	0030	W6	0030
W7	FFFF	W7	AB30
SR	0000	SR	0008 (N = 1)

ADD					Add ACCA to ACCB			
Syntax:	{label:}	ADD	A	B				

Operands: none  
 Operation: ACCA + ACCB → ACC(A or B)  
 Status Affected: OA, SA or OB, SB  
 Encoding: 0111 001A UUUU 1000  
 Description: Add ACCA to ACCB and write results to selected accumulator.  
 The 'A' bits specify the destination accumulator.  
 I-Words: 0.5  
 Cycles: 1

Example:1	ADD	A	;ACCA = ACCA + ACCB	
		Before execution	After execution	
	ACCA	0x00_1022_2EE1_5 633_9078	ACCA	0x00_3066_9207_9 746_247B
	ACCB	0x00_2044_6326_4 112_9403	ACCB	0x00_2044_6326_4 112_9403

ADD	Signed Add to Accumulator						
Syntax:	{label:}	ADD{.w} Ws,	{ Slit6, }	A			
		ADD.I [Ws],		B			
		[Ws++]					
		[Ws--]					
		[~Ws],					
		[++Ws],					
		[Ws+Wb],					
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Slit6 ∈ [-32 ... +31]					
Operation:		(ACC) + Shift <sub>Slit6</sub> (Sign-extend(Ws)) → ACC					
Status Affected:		OA, SA or OB, SB					
Encoding:	1100 00AL	www ssss ppu	UUkk kkkk	1011			
Description:	<p>The operand contained at the Effective Address is assumed to be Q1.15 or Q1.31 fractional data for word and long data operations respectively. The operand is read, then automatically sign-extended and zero-backfilled to create a value the same size as the accumulator.</p> <p>The value is then optionally (arithmetically) shifted before being added to the target accumulator.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'A' bit specifies the destination accumulator.</p> <p>The 's' bits specify the source register Ws.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'w' bits specify the offset register Wb.</p> <p>The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator preshift; if the operand Slit6 is absent, the literal bit field is set to all 0's.</p>						
	<b>Note:</b>						
	<ol style="list-style-type: none"> <li>Positive values of operand Slit6 represent arithmetic shift right.</li> <li>Negative values of operand Slit6 represent shift left.</li> <li>This instruction operates in Long or Word mode only.</li> </ol>						
I-Words:	1						
Cycles:	1						
Example:1	ADD	w8, #1, B	; <b>ACCB</b> = (w8>>1):32'b0+ <b>ACCB</b> ;Shift right <b>ACCB</b> by 1				
	Before execution				After execution		
	w8	0x20446326	w8	0x20446326	ACCB	0x00_3066_94B9_4	
	ACCB	0x00_2044_6326_4 112_9403	ACCB	0x00_3066_94B9_4 112_9403			

<b>ADDC</b>		Add Wb and Ws with Carry						
Syntax:	{label:}	ADDC.b Wb,	Ws,	Wd				
		ADDC.bz	[Ws],	[Wd]				
		ADDC{.w}	[Ws++],	[Wd++]				
		ADDC.I	[Ws--],	[Wd--]				
			[++Ws],	[++Wd]				
			[--Ws],	[--Wd]				
Operands:		$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$						
Operation:		$(Wb) + (Ws) + (C) \rightarrow Wd$						
Status Affected:		C, N, OV, Z						
Encoding:		S110      001L      dddd      ssss      pppq      qqqw      wwUU      BU00						
Description:		Add the contents of the source register Ws and the contents of the base register Wb and the Carry bit and place the result in the destination register Wd. The Z bit is "sticky" (can only be cleared). The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.						
I-Words:		1 or 0.5						
Cycles:		1						

Example 1:	ADDC.B	W0,[W1++],[W2++]	; Add W0, [W1] and C bit (Byte mode)
			; Store the result in [W2]
			; Post-increment W1, W2

Before Instruction			After Instruction		
W0	CC20		W0	CC20	
W1	0800		W1	0801	
W2	1000		W2	1001	
Data 0800	AB25		Data 0800	AB25	
Data 1000	FFFF		Data 1000	FF46	
SR	0001 (C = 1)		SR	0000	

Example 2:	ADDC	W3,[W2++],[W1++]	; Add W3, [W2] and C bit (Word mode)
			; Store the result in [W1]
			; Post-increment W1, W2

Before Instruction			After Instruction		
W1	1000		W1	1002	
W2	2000		W2	2002	
W3	0180		W3	0180	
Data 1000	8000		Data 1000	2681	
Data 2000	2500		Data 2000	2500	
SR	0001 (C = 1)		SR	0000	

<b>ADDC</b>		Add Ws and Short Literal with Carry					
Syntax:	{label:}	ADDC.b Ws, ADDC.bz [Ws], ADDC.{w} [Ws++], ADDC.I [Ws--], [++Ws], [-Ws],	lit7, [Wd] [Wd++] [Wd--] [++Wd] [-Wd]				
Operands:		$Ws \in [W0 \dots W15]; lit7 \in [0 \dots 127]; Wd \in [W0 \dots W15]$					
Operation:		$(Ws) + lit7 + (C) \rightarrow Wd$					
		Note: The literal is zero-extended to the selected data size of the operation					
Status Affected:		C, N, OV, Z					
Encoding:		1110      001L      dddd      ssss      pppq      qqkk      kkkk      Bk10					
Description:		Add the contents of the source register Ws, the zero-extended unsigned literal operand and the Carry bit; place the result in the destination register Wd.					
		The Z bit is "sticky" (can only be cleared).					
		The 'L' and 'B' bits select operation data width.					
		The 'k' bits provide the signed literal operand.					
		The 's' bits select the source register.					
		The 'd' bits select the destination register.					
		The 'p' bits select the source addressing mode.					
		The 'q' bits select the destination addressing mode.					
		<b>Note:</b> Word (.w) and Zero-Extended Byte mode (.bz) mode with a register direct destination will 0 extend the result to 32-bits.					
I-Words:		1					
Cycles:		1					

Example 1:	ADDC.B	W0, #0x1F, [W7]	; Add W0, 31 and C bit (Byte mode)	
		; Store the result in [W7]		

Before Instruction			After Instruction		
W0	CC80		W0	CC80	
W7	12C0		W7	12C0	
Data 12C0	B000		Data 12C0	B09F	
SR	0000 (C = 0)		SR	0008 (N = 1)	

Example 2:	ADDC	W3, #0x6, [-W4]	; Add W3, 6 and C bit (Word mode)	
		; Store the result in [-W4]		

Before Instruction			After Instruction		
W3	6006		W3	6006	
W4	1000		W4	0FFE	
Data 0FFE	DDEE		Data 0FFE	600D	
Data 1000	DDEE		Data 1000	DDEE	
SR	0001 (C = 1)		SR	0000	

<b>ADDC</b>		Add Literal to Wn with Carry													
Syntax:	{label:}		ADDC.b lit16,		Wn										
			ADDC.bz												
			ADDC{.w}												
			ADDC.I												
Operands:	lit16 $\in$ [0 ... 65535]; Wn $\in$ [W0 ... W15]														
Operation:	(Wn) + lit16 + (C) $\rightarrow$ Wn		Note: The literal is zero-extended to the selected data size of the operation												
Status Affected:	C, N, OV, Z														
Encoding:	1100	001L	ssss	kkkk	kkkk	kkkk	kkkk	BU10							
Description:	Add the zero-extended literal operand to the contents of the Working register Wn and the Carry bit and place the result in the Working register Wn.		The Z bit is "sticky" (can only be cleared).												
	The 'L' and 'B' bits select operation data width.														
	The 's' bits select the Working register.														
	The 'k' bits specify the signed literal operand.														
I-Words:	1														
Cycles:	1														

<b>Example 1: ADDC.B</b>		#0xFF, W7	; Add -1 and C bit to W7 (Byte mode)			
Before Instruction		After Instruction				
W7	12C0	W7	12BF			
SR	0000	(C = 0)	SR	0009	(N, C = 1)	
<b>Example 2:</b>		<b>ADDC</b>	#0xFF, W1	; Add 255 and C bit to W1 (Word mode)		
Before Instruction			After Instruction			
W1	12C0		W1	13C0		
SR	0001	(C = 1)	SR	0000		

<b>ADDC</b>	<b>Add f and Carry bit and Wn</b>								
Syntax:	{label:}	ADDC.b	f,	Wn	{,WREG}				
		ADDC.bz							
		ADDC{.w}							
		ADDC.I							
Operands:		$f \in [0 \dots 64KB]$ ; $Wn \in [W0 \dots W15]$							
Operation:		$(f) + (Wn) + (C) \rightarrow$ destination designated by D							
Status Affected:		C, N, OV, Z							
Encoding:		1110	001L	ssss	ffff	ffff	ffff	BD01	
Description:		Add the contents of the Working register and the Carry Flag and the contents of the file register and place the result in the destination designated by D. If the optional Wn destination is specified, D=0 and store result in Wn; otherwise, D=1 and store result in the file register. The Z bit is "sticky" (can only be cleared). The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'f' bits select the address of the file register. The 's' bits select the Working register.							
I-words:		1							
Cycles:		1							

<b>Example 1:</b>	<b>ADDC.B</b>	<b>RAM100</b>	<b>; Add WREG and C bit to RAM100 ; (Byte mode)</b>
-------------------	---------------	---------------	---

Before Instruction			After Instruction		
WREG	CC60		WREG	CC60	
RAM100	8006		RAM100	8067	
SR	0001 (C = 1)		SR	0000	

<b>Example 2:</b>	<b>ADDC</b>	<b>RAM200, WREG</b>	<b>; Add RAM200 and C bit to the WREG ; (Word mode)</b>
-------------------	-------------	---------------------	---

Before Instruction			After Instruction		
WREG	5600		WREG	8A01	
RAM200	3400		RAM200	3400	
SR	0001 (C = 1)		SR	000C (N, OV = 1)	

<b>ADD</b>		Add Short Literal to Wn			
Syntax:	{label:}	ADD.I	lit5,	Wn	
Operands:		lit5 $\in [0 \dots 31]$	Wn $\in [W0 \dots W15]$		
Operation:		$(Wn) + \text{lit5} \rightarrow Wn$			
Note:	The literal is zero-extended to the selected data size of the operation				
Status Affected:		C, N, OV, Z			
Encoding:	0111	010k	kkkk	ssss	
Description:	Add the zero-extended literal operand to the contents of the Working register Wn and place the result in the Working register Wn. If literal >31 and/or word or byte operation is required, assemble as ADDLW instruction.				
	The 's' bits select the Working register.				
	The 'k' bits specify the signed literal operand.				
I-Words:	0.5				
Cycles:	1				

Example 1:	ADD.B	#0xFF, W7	; Add -1 to W7 (Byte mode)	
		Before Instruction		After Instruction
		W7 12C0		W7 12BF
		SR 0000		SR 0009 (N, C = 1)
Example 2:	ADD	#0xFF, W1	; Add 255 to W1 (Word mode)	
		Before Instruction		After Instruction
		W1 12C0		W1 13BF
		SR 0000		SR 0000

<b>ADD</b>		Add Ws and Short Literal					
Syntax:	{label:}	ADD.b	Ws,	lit7,	Wd		
		ADD.bz [Ws],			[Wd]		
		ADD.{w} [Ws++],			[Wd++]		
		ADD.l [Ws--],			[Wd--]		
		[++Ws],			[++Wd]		
		[~-Ws],			[~-Wd]		
Operands:		Ws ∈ [W0 ... W15]; lit7 ∈ [0 ... 127]; Wd ∈ [W0 ... W15]					
Operation:		(Ws) + lit7 → Wd					
		Note: The literal is zero-extended to the selected data size of the operation					
Status Affected:		C, N, OV, Z					
Encoding:		1110	000L	dddd	ssss	pppq	qqkk kkkk Bk10
Description:		Add the contents of the source register Ws and the zero-extended unsigned literal operand and place the result in the destination register Wd.					
		The 'L' and 'B' bits select operation data width.					
		The 'k' bits provide the literal operand.					
		The 's' bits select the source register.					
		The 'd' bits select the destination register.					
		The 'p' bits select the source addressing mode.					
		The 'q' bits select the destination addressing mode.					
I-Words:		1					
Cycles:		1					

Example 1:	ADD.B	W0, #0x1F, W7	<b>; Add W0 and 31 (Byte mode)</b>
			<b>; Store the result in W7</b>

Before Instruction			After Instruction		
W0	2290	W0		2290	
W7	12C0	W7		12AF	
SR	0000		SR	0008	(N = 1)

Example 2:	ADD	W3, #0x6, [-W4]	<b>; Add W3 and 6 (Word mode)</b>
			<b>; Store the result in [-W4]</b>

Before Instruction			After Instruction		
W3	6006		W3	6006	
W4	1000		W4	0FFE	
Data 0FFE	DDEE		Data 0FFE	600C	
Data 1000	DDEE		Data 1000	DDEE	
SR	0000		SR	0000	

ADD	Add Literal to Wn							
Syntax:	{label:}	ADD.b	lit16,	Wn				
		ADD.bz						
		ADD{.w}						
		ADD.l						
Operands:	$\text{lit16} \in [0 \dots 65535]; \text{Wn} \in [\text{W}0 \dots \text{W}15]$							
Operation:	$(\text{Wn}) + \text{lit16} \rightarrow \text{Wn}$							
	Note: The literal is zero-extended to the selected data size of the operation							
Status Affected:	C, N, OV, Z							
Encoding:	1100	000L	ssss	kkkk	kkkk	kkkk	kkkk	BU10
Description:	Add the zero-extended literal operand to the contents of the Working register Wn and place the result in the Working register Wn. The 'L' and 'B' bits select operation data width. The 's' bits select the Working register. The 'k' bits specify the signed literal operand.							
I-Words:	1							
Cycles:	1							

ADD	Add f and Wn							
Syntax:	{label:}	ADD.b	f	Wn	{,WREG}			
		ADD.bz						
		ADD{.w}						
		ADD.l						
Operands:	$f \in [0 \dots 64KB]$ ; $Wn \in [W0 \dots W15]$							
Operation:	$(f) + (Wn) \rightarrow \text{destination designated by D}$							
Status Affected:	C, N, OV, Z							
Encoding:	1110	000L	ssss	ffff	ffff	ffff	ffff	BD01
Description:	Add the contents of the Working register and the contents of the file register and place the result in the destination designated by D. If the optional Wn is specified, D = 0 and store result in Wn; otherwise, D = 1 and store result in the file register. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'f' bits select the address of the file register. The 's' bits select the Working register.							
I-Words:	1							
Cycles:	1							

Example 1:	ADD.B	RAM100	; Add WREG to RAM100 (Byte mode)
Before Instruction		After Instruction	
WREG CC80		WREG CC80	
RAM100 FFC0		RAM100 FF40	
SR 0000		SR 0005 (OV, C = 1)	
Example 2:	ADD	RAM200, WREG	; Add RAM200 to WREG (Word mode)
Before Instruction		After Instruction	
WREG CC80		WREG CC40	
RAM200 FFC0		RAM200 FFC0	
SR 0000		SR 0001 (C = 1)	

AND	AND Wb and Ws								
Syntax:	{label:}	AND.b	Wb,	Ws,	Wd				
		AND.bz		[Ws],	[Wd]				
		AND.{w}		[Ws++],	[Wd++]				
		AND.l		[Ws--],	[Wd--]				
				[++Ws],	[++Wd]				
				[--Ws],	[--Wd]				
				SR	SR				
Operands:		$Wb \in [W0 \dots W14]; Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$							
Operation:		$(Wb).AND.(Ws) \rightarrow Wd$							
Status Affected:		N, Z							
Encoding:		S110	100L	dddd	ssss	pppq	qqww	wwUU	BU00
Description:		Compute the AND of the contents of the source register Ws and the contents of the base register Wb and place the result in the destination register Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode. See and for modifier addressing information.							
		<b>Notes:</b>							
		1. When the SR is selected, SR data write will take priority over any SR update resulting from the AND operation. 2. .bz data size/mode is disallowed when writing to the SR.							
I-Words:		1 or 0.5							
Cycles:		1							

Example 1:	AND.B	W0, W1 [W2++]	; AND W0 and W1, and ; store to [W2] (Byte mode) ; Post-increment W2
------------	-------	---------------	--

Before Instruction		After Instruction	
W0	AA55	W0	AA55
W1	2211	W1	2211
W2	1001	W2	1002
Data 1000	FFFF	Data 1000	11FF
SR	0000	SR	0000

Example 2:	AND	W0, [W1++], W2	; AND W0 and [W1], and ; store to W2 (Word mode) ; Post-increment W1
------------	-----	----------------	--

Before Instruction		After Instruction	
W0	AA55	W0	AA55
W1	1000	W1	1002
W2	55AA	W2	2214
Data 1000	2634	Data 1000	2634
SR	0000	SR	0000

AND	AND Ws and 0's Extended Short Literal						
Syntax:	{label:}	AND.b	Ws,	lit7,	Wd		
		AND.bz [Ws],			[Wd]		
		AND{.w} [Ws++],			[Wd++]		
		AND.l [Ws--],			[Wd--]		
		[++Ws],			[++Wd]		
		[--Ws],			[--Wd]		
		SR			SR		
Operands:		Ws ∈ [W0 ... W15]; lit7 ∈ [0 ... 127]; Wd ∈ [W0 ... W15]					
Operation:		(Ws).AND.lit7 <sup>3</sup> → Wd					
Status Affected:		N, Z (see note 1)					
Encoding:	1110	100L	dddd	ssss	pppq	qqkk	kkkk
Description:	Compute the AND of the contents of the source register Ws and the zero-extended literal operand and place the result in the destination register Wd. The 'L' and 'B' bits select operation data width. The 'k' bits provide the unsigned literal operand. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.						
Notes:	<ol style="list-style-type: none"> <li>When the SR is selected, SR data write will take priority over any SR update resulting from the AND operation.</li> <li>.bz data size/mode is disallowed when writing to the SR.</li> <li>The literal is zero-extended to the selected data size of the operation.</li> </ol>						
I-Words:	1						
Cycles:	1						

<b>AND1</b>		<b>AND Ws and 1's Extended Short Literal</b>															
Syntax:	{label:}	AND1.b	Ws,	lit7,	Wd												
		AND1.bz [Ws],			[Wd]												
		AND1{.w} [Ws++],			[Wd++]												
		AND1.l [Ws--],			[Wd--]												
		[++Ws],			[++Wd]												
		[--Ws],			[--Wd]												
		SR			SR												
Operands:		Ws ∈ [W0 ... W15]; lit7 ∈ [0 ... 127]; Wd ∈ [W0 ... W15]															
Operation:		(Ws).AND.lit7 <sup>3</sup> → Wd															
Status Affected:		N, Z (see note 1)															
Encoding:		1110	110L	dddd	ssss	pppq	qqkk	kkkk	Bk10								
Description:		Compute the AND of the contents of the source register Ws and the ones-extended literal operand and place the result in the destination register Wd. The 'L' and 'B' bits select operation data width. The 'k' bits provide the unsigned literal operand. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.															
<b>Notes:</b>																	
1. When the SR is selected, SR data write will take priority over any SR update resulting from the AND operation.																	
2. .bz data size/mode is disallowed when writing to the SR.																	
3. The literal is ones-extended to the selected data size of the operation.																	
I-Words:		1															
Cycles:		1															
Example 1:	<b>AND1.I W5, #0x1F, W7</b>				<b>AND W5 and 1's extended 0x1F, and ; store to [W7] (long word mode) ; Post-increment W2</b>												

Before Instruction

W5 7FFFFFFF

W7 12345678

SR 0000

After Instruction

W5 7FFFFFFF

W7 7FFFFFF9F

SR 0000

AND	AND Literal and Wn							
Syntax:	{label:}	AND.b	lit16,	Wn				
		AND.bz		SR				
		AND{.w}						
		AND.l						
Operands:		lit16 ∈ [0 ... 65536]; Wn ∈ [W0 ... W15]; Status Register (SR)						
Operation:		(Wn) .AND. lit16 <sup>3</sup> → Wn or (SR) .AND. lit16 → SR						
Status Affected:		N, Z (see note 1)						
Encoding:	1100	100L	ssss	kkkk	kkkk	kkkk	kkkk	BT10
Description:	Compute the AND of the literal operand and the contents of the Working register Wn or SR and place the result in the Working register Wn or SR. The 'L' and 'B' bits select operation data width. The 's' bits select the Working register. The 'k' bits specify the unsigned literal operand. The 'T' bit selects between Ws (T = 0) and SR (T = 1) target registers.							
Notes:	<ol style="list-style-type: none"> <li>When the SR is selected, SR data write will take priority over any SR update resulting from the AND operation.</li> <li>.bz data size/mode is disallowed when writing to the SR.</li> <li>The literal is zero-extended to 32-bits for long word operations.</li> </ol>							
I-Words:	1							
Cycles:	1							

Example 1:	AND.B #0x83, W7	; AND 0x83 to W7 (Byte mode)
	Before Instruction	After Instruction
	W7 12C0	W7 1280
	SR 0000	SR 0008 (N = 1)
Example 2:	AND #0x333, W1	; AND 0x333 to W1 (Word mode)
	Before Instruction	After Instruction
	W1 12D0	W1 0210
	SR 0000	SR 0000

AND	And f and Wn							
Syntax:	{label:}	AND.b	f	,	Wn	,	WREG}	
		AND.bz						
		AND{.w}						
		AND.l						
Operands:	$f \in [0 \dots 64KB]$ ; $Wn \in [W0 \dots W14]$							
Operation:	$(f).AND.(Wn) \rightarrow \text{destination designated by } D$							
Status Affected:	N, Z							
Encoding:	1110	100L	ssss	ffff	ffff	ffff	ffff	BD01
Description:	<p>Compute the AND of the contents of the Working register and the contents of the file register and place the result in the destination designated by D. If the optional Wn is specified, D=0 and store result in Wn; otherwise, D=1 and store result in the file register.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>							
I-Words:	1							
Cycles:	1							

Example 1:	AND.B RAM100	;	AND WREG to RAM100 (Byte mode)
------------	--------------	---	--------------------------------

Before Instruction				After Instruction			
WREG	CC80			WREG	CC80		
RAM100	FFC0			RAM100	FF80		
SR	0000			SR	0008	(N = 1)	

Example 2:	AND RAM200, WREG	;	AND RAM200 to WREG (Word mode)
------------	------------------	---	--------------------------------

Before Instruction				After Instruction			
WREG	CC80			WREG	0080		
RAM200	12C0	RAM200			12C0		
SR	0000			SR	0000		

ASR                          Arithmetic Shift Right by 1							
Syntax:	{label:}	ASR.b	Ws,	Wd			
		ASR.bz	[Ws],	[Wd]			
		ASR.{w}	[Ws++],	[Wd++]			
		ASR.l	[Ws--],	[Wd--]			
			[++Ws],	[++Wd]			
			[--Ws],	[--Wd]			
Operands:	Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]						
Operation:	For long word operation: (Ws[31]) → Wd[31], (Ws[31:1]) → Wd[30:0], (Ws[0]) → C For word operation: (Ws[15]) → Wd[15], (Ws[15:1]) → Wd[14:0], (Ws[0]) → C For byte operation: (Ws[7]) → Wd[7], (Ws[7:1]) → Wd[6:0], (Ws[0]) → C						
Status Affected:	C, N,Z						
Encoding:	0010	100L	dddd	ssss	pppq	qqUU	UUUU
Description:	Arithmetic shift right the contents of the source register by one bit, placing the result in the destination register Wd. Destination register direct Extended Byte or Word mode will zero-extend the result to 32-bits, then write to Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.						
I-Words:	1 or 0.5						
Cycles:	1						

ASR                          Arithmetic Shift Right f							
Syntax:	{label:}	ASR.b	f	{,Wnd}	{,WREG}		
		ASR.bz					
		ASR{.w}					
		ASR.l					
Operands:	$f \in [0 \dots 64KB]$ ; Wnd $\in [W0 \dots W14]$						
Operation:	For long word operation: $(f[31]) \rightarrow \text{Dest}[31], (f[31]) \rightarrow \text{Dest}[30],$ $(f[30:1]) \rightarrow \text{Dest}[29:0], (f[0]) \rightarrow C$ For word operation: $(f[15]) \rightarrow \text{Dest}[15], (f[15]) \rightarrow \text{Dest}[14]$ $(f[14:1]) \rightarrow \text{Dest}[13:0], (f[0]) \rightarrow C$ For byte operation: $(f[7]) \rightarrow \text{Dest}[7:1], (f[7]) \rightarrow \text{Dest}[6],$ $(f[6:1]) \rightarrow \text{Dest}[5:0], (f[0]) \rightarrow C$						
Status Affected:	C, N, Z						
Encoding:	1010	000L	dddd	ffff	ffff	ffff	BD01
Description:	Shift the contents of the file register f one bit to the right through the carry flag, and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'f' bits select the address of the file register. The 'd' bits select the Working register.						
I-Words:	1						
Cycles:	1						

Example 1:	ASR.B RAM400, WREG	;	ASR RAM400 and store to WREG
		;	(Byte mode)

Before Instruction				After Instruction			
WREG	0600			WREG	0611		
RAM400	0823			RAM400	0823		
SR	0000			SR	0001	(C = 1)	

Example 2:	ASR RAM200	;	ASR RAM200 (Word mode)
Before Instruction		After Instruction	
RAM200	8009	RAM200	C004
SR	0000	SR	0009 (N, C = 1)

ASR Arithmetic Shift Right by Short Literal					
Syntax:	{label:}	ASR{.w}	Ws,	lit5,	Wd
		ASR.I	[Ws],		[Wd]
			[Ws++],		[Wd++]
			[Ws--],		[Wd--]
			[++Ws],		[++Wd]
			[~-Ws],		[~-Wd]
Operands:	Ws ∈ [W0 ... W15]; lit5 ∈ [0...31]; Wd ∈ [W0 ... W15]				
Operation:	lit5[4:0] → Shift_Val For long word operation: Ws[31] → Right shift input (arithmetic shift) Ws[31:0], 32'b0 → Shift_In[63:0] Arithmetic shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[63:32] → Wd For word operation: Ws[15] → Right shift input (arithmetic shift) {16{Ws[15]}}, Ws[15:0],32'b0 → Shift_In[63:0] Arithmetic shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[47:32] → Wd[15:0]				
Status Affected:	N,Z				
Encoding:	S010        000k        kkkk        dddd        pppq        qqss        ssUU        LU00				
Description:	Arithmetic shift right the contents of the source register Ws by lit5 bits (up to 31 positions), placing the result in the destination Wd.  This instruction will generate the correct result for a word operation with any shift value in lit5. If lit5 > 15 and Ws[15] = 0, then Wd[15:0]=0x0000. If lit5 > 15 and Ws[15] = 1, then Wd[15:0]=0xFFFF.  Register Direct Word mode will zero-extend the result to 32-bits, then write to Wd.  The 'S' bit selects instruction size. The 'L' bit selects word or long word operation. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode. The 'k' bits provide the literal operand. I-Words: 1 or 0.5 Cycles: 1				
	<b>Note:</b> 1. This instruction only operates in Word or Long Word mode.				

Example 1:	ASR W0, #0x4, W1	; ASR W0 by 4 and store to W1
	Before Instruction	After Instruction
	W0 060F	W0 060F
	W1 1234	W1 0060
	SR 0000	SR 0000

Example 2:	ASR W0, #0x6, W1	; ASR W0 by 6 and store to W1
	Before Instruction	After Instruction
	W0 80FF	W0 80FF
	W1 0060	W1 FE03
	SR 0000	SR 0008 (N = 1)

Example 3:	ASR W0, #0xF, W1	; ASR W0 by 15 and store to W1
	Before Instruction	After Instruction
	W0 70FF W1 CC26 SR 0000	W0 70FF W1 0000 SR 0002 (Z = 1)

<b>ASRM</b>		<b>Arithmetic Shift Right Multi-Precision by Short Literal</b>							
Syntax:	{label:}	ASRM{.l} Ws, [Ws], [Ws++], [Ws--], [++Ws], [--Ws],							
Operands:	Ws ∈ [W0 ... W15]; lit5 ∈ [0...31]; Wnd ∈ [W1 ... W14]								
Operation:	lit5[4:0] → Shift_Val Ws[31] → Right shift input (arithmetic shift) Ws[31:0], 32'b0 → Shift_In[63:0] Arithmetic shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[63:32] → Wnd Shift_Out[31:0]   Wnd-1 → Wnd-1								
Status Affected:	N,Z								
Encoding:	1110 000k kkkk dddd pppU UUss ssUU 0011								
Description:	Arithmetic shift right the contents of the source register Ws by lit5 bits (up to 31 positions), placing the result in the destination register Wnd. The register containing the next least significant data word will already contain an intermediate shift result. Bitwise OR this value with the data shifted out of Ws in order to create the final shift result, then update the corresponding destination register.  The Z bit is “sticky” (can only be cleared). The ‘s’ bits select the source register. The ‘d’ bits select the destination register. The ‘p’ bits select the source addressing mode. The ‘k’ bits provide the literal operand.								
I-Words:	1								
Cycles:	2								
1. <b>Note:</b> This instruction only operates in Long Word mode.									

<b>ASRM</b>		<b>Arithmetic Shift Right Multi-Precision by Wb</b>							
Syntax:	{label:}	ASRM{.l} Ws, [Ws], [Ws++], [Ws--], [++Ws], [--Ws],							
Operands:	Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Wnd ∈ [W1 ... W14]								
Operation:	Wb[15:0] → Shift_Val Ws[31] → Right shift input (arithmetic shift) Ws[31:0], 32'b0 → Shift_In[63:0] Arithmetic shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[63:32] → Wnd Shift_Out[31:0]   Wnd-1 → Wnd-1								
Status Affected:	N,Z								
Encoding:	1110 001U ssss dddd pppU UUww wwUU 0011								

**Instruction Descriptions (A to BZ) (continued)****ASRM****Arithmetic Shift Right Multi-Precision by Wb**

Description:

Arithmetic shift right the contents of the source register Ws by Wb bits, placing the result in the destination register Wnd. The register containing the next least significant data word will already contain an intermediate shift result. Bitwise OR this value with the data shifted out of Ws in order to create the final shift result, then update the corresponding destination register or memory address.

The arithmetic right shift may be by any amount between 0 and 32 bits. Should the shift value held in Wb[15:0] exceed 2'd32, the shift value will saturate to 2'd32 for consistency. Any data held in Wb[31:16] will have no effect.

The Z bit is "sticky" (can only be cleared).

The 'w' bits select the base (shift count) register.

The 's' bits select the source register.

The 'd' bits select the destination register.

The 'p' bits select the source addressing mode.

**Note:** This instruction only operates in Long Word mode.

I-Words:

1

Cycles:

2

ASR Arithmetic Shift Right by Wb					
Syntax:	{label:}	ASR.b	Ws,	Wb,	Wd
		ASR.bz	[Ws],		[Wd]
		ASR.{w}	[Ws++],		[Wd++]
		ASR.l	[Ws--],		[Wd--]
			[++Ws],		[++Wd]
			[~-Ws],		[~-Wd]
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]			
Operation:		Wb[15:0] → Shift_Val			
		For long word operation:			
		Ws[31] → Right shift input (arithmetic shift)			
		Ws[31:0], 32'b0 → Shift_In[63:0]			
		Arithmetic shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0]			
		Shift_Out[63:32] → Wd			
		For word operation:			
		Ws[15] → Right shift input (arithmetic shift)			
		{16{Ws[15]}}, Ws[15:0], 32'b0 → Shift_In[63:0]			
		Arithmetic shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0]			
		Shift_Out[47:32] → Wd[15:0]			
		For byte operation:			
		Ws<7> → Right shift input (arithmetic shift)			
		{24{Ws<7>}}, Ws[7:0], 32'b0 → Shift_In[63:0]			
		Arithmetic shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0]			
		Shift_Out[39:32] → Wd[7:0]			
Status Affected:		N,Z			
Encoding:		1010      100L      wwwwww      dddd      ppppq      qqss      ssUUU      B100			
Description:		Arithmetic shift right the contents of the source register by Wb bits, placing the result in the destination register Wd.			
		This instruction will generate the correct result for any shift value in Wb[15:0]:			
		<ul style="list-style-type: none"> <li>For a byte operation shift value &gt; 7: If Ws[7] = 0 then Wd[7:0]=0x00 else Wd[7:0]=0xFF</li> <li>For a word operation shift value &gt; 15: If Ws[15] = 0 then Wd[15:0]=0x0000 else Wd[15:0]=0xFFFF</li> <li>For a Long Word operation shift value &gt; 31: If Ws[31] = 0 then Wd[31:0]=0x00000000 else Wd[31:0]=0xFFFFFFFF Any data held in Wb[31:16] will have no effect.</li> </ul>			
		Destination register direct Extended Byte or Word mode will zero-extend the result to 32-bits, then write to Wd.			
		The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base (shift count) register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.			
I-Words:	1				
Cycles:	1				

BRA C	Branch if Carry/Unsigned Greater Than or Equal							
Syntax:	{label:}	BRA	C,	Label				
	{label:}	BRA	GEU,					
Operands:	Label is resolved by the linker to a signed word offset							
Operation:	<p>Condition = C</p> <p>If (condition) then {</p> <p>If slit20 = 1 then skip next (16-bit) instruction</p> <p>Else If (slit20 = 2 &amp;&amp; next_op[31] = 1) then skip next (32-bit) instruction</p> <p>Else (PC+4) + 2*slit20 → PC</p> <p>}</p> <p>Else no branch</p>							
Status Affected:	None							
Encoding:	1010	101U	nnnn	nnnn	nnnn	nnnn	nnnn	0010
Description:	<p>If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or branch to any size of instruction with a forward or backward range of up to 1MB.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals four, and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.</p> <p>Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.</p> <p>The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).</p> <p><b>Note:</b> Should the byte offset equal fourF and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.</p>							
I-Words:	1							
Cycles:	1 (2 or 3)							
	<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.							

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA C, _CHECK_FAIL	;Branch if (N&&OV)    (!N&&!OV), to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
		...	...	
		...	...	
		...	...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
		...	CP w6, w7	
		...	...	

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	Before execution PC 0x804000 w15 0x4500 SR[7:0] 8'b00000001		(C = 1;)	After execution PC 0x804030 w15 0x4500 SR[7:0] 8'b00000001 (C = 1;)

<b>BCLR</b>							
<b>Bit Clear in Ws</b>							
Syntax:	{label:}	BCLR.b	Ws,	bit5			
		BCLR{.w}	[Ws],				
		BCLR.I	[Ws++],				
			[Ws-],				
			[++Ws],				
			[--Ws],				
			SR				
Operands:	Ws ∈ [W0 ... W15]; Byte: bit5 ∈ [0 ... 7]; Word: bit5 ∈ [0 ... 15]; Long word: bit5 ∈ [0 ... 31]						
Operation:	0 → Ws<bit5>						
Status Affected:	None (see Note 1)						
Encoding:	S100      000b      bbbb      ssss      pppU      UUUU      UUUU      LB00						
Description:	Bit 'bit5' in register Ws is cleared. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'b' bits define value or bit5 of the bit position to be cleared (bit5[4:0] for the 16-bit opcode). The 's' bits select the source/destination register. The 'p' bits select the source addressing mode.						
<b>Note:</b>	1. When targeting the SR, a bit within the SR will be cleared as a result of the instruction operation.						
I-Words:	1 or 0.5						
Cycles:	1						

<b>Example 1:</b>	<b>BCLR.B W2, #0x2</b>	<b>; Clear bit 3 in W2</b>
	Before Instruction	After Instruction
	W2 F234	W2 F230
	SR 0000	SR 0000
<b>Example 2:</b>	<b>BCLR [W0++], #0x0</b>	<b>; Clear bit 0 in [W0] ; Post-increment W0</b>
	Before Instruction	After Instruction
	W0 2300	W0 2302
	Data 2300 5607	Data 2300 5606
	SR 0000	SR 0000

BCLR		Bit Clear in f														
Syntax:	{label:}	BCLR.b	f,		bit3											
Operands:		$\text{bit3} \in [0 \dots 7]; \text{f} \in [0 \dots 1\text{MB}]$														
Operation:		$0 \rightarrow \text{f}[bit3]$														
Status Affected:		None														
Encoding:	1100	000b	ffff	ffff	ffff	ffff	ffff	bb01								
Description:	<p>Bit 'bit3' in file register f is cleared.            The 'w' bits select value bit3 of the bit position to be cleared.            The 'f' bits select the address of the file register.</p>															
<b>Notes:</b>																
<ol style="list-style-type: none"> <li>1. This instruction operates in Byte mode only.</li> <li>2. The .b extension must be included with the opcode.</li> </ol>																
I-Words:	1															
Cycles:	1															

<b>Example 1:</b>	BCLR.B 0x800, #0x7	; Clear bit 7 in 0x800	
		Before Instruction	After Instruction
	Data 0800    66EF		Data 0800    666F
	SR    0000		SR    0000

BFEXT						Bit Field Extract from Ws into Wb								
Syntax:	{label:}		BFEXT{.w} bit5,	wid6,	Ws,	Wb								
			BFEXT.I		[Ws],			[Ws++],						
					[Ws-],			[++Ws],						
					[--Ws],			SR						
Operands:	Word: bit5 ∈ [0 .. 15]; wid6 ∈ [0 .. 16] Long: bit5 ∈ [0 .. 31]; wid6 ∈ [0 .. 32] Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W14]													
Operation:	See text													
Status Affected:	None													
Encoding:	1111	100m	mmmm	ssss	pppc	ccww	wwcc	L111						
Description:	A bit field is extracted (copied) from (Ws), and written into Wb. The bit field data loaded into Wb starts at Wb[0], and all MSbs within Wb that are beyond the defined bit field width will be cleared. The bit location within Ws of the Lsb of the bit field to be extracted is defined by operand bit5. The width of the bit field is defined by operand wid6 and may be between zero and up to 16-bits (word operation) or 32-bits (long word operation).  Word mode will zero-extend the result to 32-bits prior to the write to Wb. The 'L' bit selects word or long word operation. The 'w' bits select the bit field destination register. The 's' bits select the data source register. The 'p' bits select the source addressing mode. The 'cccc' bits define the bit field Lsb position within the target word. The 'mmmmm' bits define the bit field MSb position within the target word.													
<b>Note:</b>	1. Literal wid6 = 0 is a legal value for both long and word sized operations.													
I-Words:	1													
Cycles:	1													

BFEXT									Bit Field Extract from f into Wb								
Syntax:	{label:} BFEXT{.w} bit5, wid6, f, Wb			BFEXT.I													
Operands:	bit5 ∈ [0 .. 31]; wid6 ∈ [0 .. 32] Wb ∈ [W0 ... W14]; f ∈ [0 ... 1MB]																
Operation:	See text																
Status Affected:	None																
Encoding:																	
1st word	1100	110U	ffff	ffff	ffff	ffff	ffff	1001									
2nd word	1100	110L	UUUm	mmmm	UUUc	ccww	wwcc	1101									
Description:	A bit field is extracted (copied) from the file register address and written into Wb. The bit field data loaded into Wb starts at Wb[0], and all MSbs within Wb that are beyond the defined bit field width will be cleared. The bit location within the file register of the LSb of the bit field to be extracted is defined by operand bit5. The width of the bit field is defined by operand wid6 and may be between 0 and up to 16-bits (word operation) or 32-bits (long word operation). Word mode will zero-extended the result to 32-bits and then write to Wb. The 'L' bit selects word or long word operation. The 'w' bits select the bit field destination register. The 'f' bits select the address of the source file register. The 'cccc' bits define the bit field LSb position within the target word. The 'mmmmmm' bits define the bit field MSb position within the target word.																
Notes:																	
	1. Literal wid6 = 0 is a legal value for both long and word sized operations. 2. Accessible file address space is 1MB. LSb of 'f' bits is 1'b0 for word operation. LS 2-bits of 'f' bits is 2'b00 for long word operation.																
I-Words:	2																
Cycles:	2																

<b>BFINS</b>		Bit Field Insert from Wb into Ws						
Syntax:	{label:}	BFINS{.w} bit5, BFINS.I	wid6,	Wb,	Ws [Ws] [Ws++] [Ws--] [++Ws] [--Ws]	SR		
Operands:		Word: bit5 ∈ [0 .. 15]; wid6 ∈ [0 .. 16] Long: bit5 ∈ [0 .. 31]; wid6 ∈ [0 .. 32] Wb ∈ [W0 ... W14]; Ws ∈ [W0 ... W15]						
Operation:		See text						
Status Affected:		None						
Encoding:		1111      100m      mmmm      ssss      pppc      ccww      wwcc      L011						
Description:		A bit field is read from (Wb) and inserted (copied) into Ws. The bit field data sourced from Wb starts at Wb[0]. All MSbs within Wb that are beyond the defined bit field width are ignored and may be set to any value. The inserted bit field will overwrite the bits already in Ws or SR (i.e., the instruction does not shift any existing bits to accommodate the new bit field). The bit location within Ws of the LSb of the bit field to be inserted is defined by operand bit5. The width of the bit field is defined by operand wid6 and may be between zero and up to 16-bits (word operation) or 32-bits (long word operation). Word mode will zero-extend a Ws register direct result write to 32-bits. The 'L' bit selects word or long word operation. The 'w' bits select the bit field source register. The 's' bits select the data source/destination register. The 'p' bits select the source addressing mode. The 'cccc' bits define the bit field LSb position within the target word. The 'mmmm' bits define the bit field MSb position within the target word.						
<b>Note:</b>		1. Literal wid6 = 0 is a legal value for both long and word sized operations.						
I-Words:		1						
Cycles:		1						

<b>BFINS</b>		<b>Bit Field Insert from Wb into f</b>						
Syntax:	{label:}	BFINS{.w} bit5, BFINS.I	wid6,	Wb,	f			
Operands:	Word: bit5 ∈ [0 .. 15]; wid6 ∈ [0 .. 16] Long: bit5 ∈ [0 .. 31]; wid6 ∈ [0 .. 32] Wb ∈ [W0 ... W14]; f ∈ [0 ... 1MB]							
Operation:	See text							
Status Affected:	None							
Encoding:								
1st word	1100	110U	ffff	ffff	ffff	ffff	ffff	0001
2nd word	1100	110L	UUUm	mmmm	UUUc	ccww	wwcc	0101
Description:	<p>A bit field is read from (Wb) and inserted (copied) into the file register address. The bit field data sourced from Wb starts at Wb[0]. All MSbs within Wb that are beyond the defined bit field width are ignored and may be set to any value.</p> <p>The inserted bit field will overwrite the bits already in the file register (i.e., the instruction does not shift any existing bits to accommodate the new bit field).</p> <p>The bit location within the file register of the LSb of the bit field to be inserted is defined by operand bit5. The width of the bit field is defined by operand wid6 and may be between zero and up to 16-bits (word operation) or 32-bits (long word operation).</p> <p>The 'w' bits select the bit field source register.</p> <p>The 'f' bits select the source/destination file register.</p> <p>The 'cccc' bits define the bit field LSb position within the target word.</p> <p>The 'mmmmm' bits define the bit field MSb position within the target word.</p>							
Notes:	<ol style="list-style-type: none"> <li>1. Literal wid6 = 0 is a legal value for both long and word sized operations.</li> <li>2. Accessible file address space is 1MB. LSb of 'f' bits is 1'b0 for word operation. LS 2-bits of 'f' bits is 2'b00 for long word operation.</li> </ol>							
I-Words:	2							
Cycles:	2							

BFINS									Bit Field Insert Literal into Ws								
Syntax:	{label:}	BFINS{.w} bit5, BFINS.I	wid6,	lit16,	Ws	[Ws]	[Ws++]	[Ws--]	[++Ws]	[--Ws]							
Operands:		bit5 ∈ [0 .. 31]; wid6 ∈ [0 .. 32] lit16 ∈ [0 .. 65536]; Ws ∈ [W0 ... W15]															
Operation:		See text															
Status Affected:		None															
Encoding:																	
1st word	1100	101U	UUUU	kkkk	kkkk	kkkk	kkkk	kkkk	U001								
2nd word	1100	101m	mmmm	ssss	pppc	ccUU	UUcc	L101									
Description:		A bit field literal value is inserted (copied) into Ws. The bit field data sourced from the literal starts at the LSb of the literal. All MSbs within the literal value that are beyond the defined bit field width are ignored and may be set to any value. The inserted bit field will overwrite the bits already in Ws (i.e., the instruction does not shift any existing bits to accommodate the new bit field). The bit location within Ws of the LSb of the bit field to be inserted is defined by operand bit5. The width of the bit field is defined by operand wid6 and may be between zero and up to 16-bits (word operation) or 32-bits (long word operation). Word mode will zero-extend a Ws register direct result write to 32-bits. The 'L' bit selects word or long word operation. The 'K' bits contain the bit field source value. The 'S' bits select the source/destination register. The 'P' bits select the source addressing mode. The 'CCCC' bits define the bit field LSb position within the target word. The 'mmmmm' bits define the bit field MSb position within the target word.															
		<b>Note:</b>															
		1. Literal wid6 = 0 is a legal value for both long and word sized operations.															
I-Words:	2																
Cycles:	2																

BRA GE		Branch if Signed Greater Than or Equal					
Syntax:	{label:}	BRA	GE,	Label			
Operands:							
Operation:							
		Label is resolved by the linker to a signed word offset (slit20)					
		Condition = (N&&OV)   !(N&&!OV)					
		If (condition) {					
		If slit20 = 1 then skip next (16-bit) instruction					
		Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction					
		Else (PC+4) + 2*slit20 → PC					
		}					
		Else no branch					
Status Affected:		None					
Encoding:	1010	100U	nnnn	nnnn	nnnn	nnnn	0110
Description:		<p>If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.</p> <p>Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.</p> <p>The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).</p> <p><b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.</p>					
I-words:	1						
Cycles:	1 (2 or 3)						
		The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.					
Example:1	Address	Label	Instruction			Comment	
	0x804000	-	BRA GE, _CHECK_FAIL			;Branch if (N&&OV)   !(N&&!OV), to 0x804030(_CHECK_FAIL)	
	0x804004	_PASS_CONDITION:	SL.I	w5, w6			
		ON:					
		...		...			
		...		...			
		...		...			
	0x804030	_CHECK_FAIL:	MOV.I	SR, W6		; _CHECK_FAIL SubRoutine	
			CP w6, w7				
			...		...		
Before execution				After execution			

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	PC	0x804000	PC	0x804030
	w15	0x4500	w15	0x4500
	SR[7:0]	8'b0001100	(N = 1; OV = 1) SR[7:0]	8'b0001100 (N = 1; OV = 1)

BRA GT	Branch if Signed Greater Than						
Syntax:	{label:}	BRA	GT,	Label			
Operands:				Label is resolved by the linker to a signed word offset (slit20)			
Operation:				Condition = (!Z&&N&&OV)   !(IZ&&N&&!OV); If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch			
Status Affected:				None			
Encoding:	1010	100U	nnnn	nnnn	nnnn	nnnn	0010
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.  If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.  If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.  Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.  The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.			
I-Words:	1						
Cycles:	1 (2 or 3)						
				<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.			

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA GT, _CHECK_FAIL	;Branch if (IZ&&N&&OV)   !(IZ&&N&&!OV), to0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITION:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	;_CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution			After execution
PC	0x804000		PC	0x804030
w15	0x4500		w15	0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[7:0]	8'b0001100	(Z = 0;OV = 1; SR[7:0] 8'b0001100 N=1)	(Z = 0;OV = 1; N=1)

BRA GTU	Branch if Unsigned Greater Than						
Syntax:	{label:}	BRA	GTU,	Label			
Operands:				Label is resolved by the linker to a signed word offset (slit20)			
Operation:				Condition = (C&&!Z); If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch			
Status Affected:				None			
Encoding:	1010	111U	nnnn	nnnn	nnnn	nnnn	0110
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. Likewise, the literal is zero-extended to 32-bits for long-word instructions. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four, and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC, to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.			
I-Words:		1					
Cycles:		1 (2 or 3)					
				<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.			

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA GTU, _CHECK_FAIL	; Branch if (C&&!Z), to0x804030(_CH ECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	PC	0x804000	PC	0x804030
	w15	0x4500	w15	0x4500
	SR[7:0]	8'b00000001	(C = 1, Z = 0)	8'b00000001 (C = 1, Z = 0)

BRA LE	Branch if Signed Less Than or Equal						
Syntax:	{label:}	BRA	LE,	Label			
Operands:				Label is resolved by the linker to a signed word offset (slit20)			
Operation:				Condition = Z    (N&&!OV)    (!N&&OV); If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch			
Status Affected:				None			
Encoding:	1010	100U	nnnn	nnnn	nnnn	nnnn	1110
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four, and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.			
I-Words:	1						
Cycles:	1 (2 or 3)						
				<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.			

Example:1	Address	Label	Instruction	Comment		
	0x804000	-	BRA LE, _CHECK_FAIL	;Branch if Z    (N&&!OV)    (!N&&OV), to0x804030(_CHECK_FAIL)		
	0x804004	_PASS_CONDITION:	SL.I w5, w6			
	...		...			
	...		...			
	...		...			
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine		
	...		CP w6, w7			
	...		...			
	Before execution			After execution		
	PC	0x804000		PC	0x804030	
	w15	0x4500		w15	0x4500	
	SR[7:0]	8'b0000010	(Z = 1)	SR[7:0]	8'b0000010	(Z = 1)

<b>BRA LEU</b>		Branch if Unsigned Less Than or Equal						
Syntax:	{label:}	BRA	LEU,	Label				
Operands:				Label is resolved by the linker to a signed word offset (slit20)				
Operation:				Condition = !C   Z; If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 & next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch				
Status Affected:				None				
Encoding:	1010	111U	nnnn	nnnn	nnnn	nnnn	nnnn	0010
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. Likewise, the literal is zero-extended to 32-bits for long-word instructions. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four, and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.				
I-Words:		1						
Cycles:		1 (2 or 3)						
				<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.				

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA LT, _CHECK_FAIL	;Branch if !C Z, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	
PC	0x804000		PC	0x804030
w15	0x4500		w15	0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[7:0]	8'b00000000	(Z = 0;C = 0)	SR[7:0] 8'b00000000 (Z = 0;C = 0)

BRA LT	Branch if Signed Less Than							
Syntax:	{label:}	BRA	LT,	Label				
Operands:				Label is resolved by the linker to a signed word offset (slit20)				
Operation:				Condition = (N&&!OV)   (!N&&OV); If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch				
Status Affected:				None				
Encoding:	1010	100U	nnnn	nnnn	nnnn	nnnn	1010	
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four, and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.				
I-Words:		1						
Cycles:		1 (2 or 3)						
				<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.				

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA LT, _CHECK_FAIL	;Branch if IC Z, to0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	
PC	0x804000		PC	0x804030
w15	0x4500		w15	0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[7:0]	8'b00000000	(Z = 0;C = 0)	SR[7:0] 8'b00000000 (Z = 0;C = 0)

<b>BRA NC/LTU</b>		Branch if Not Carry / Unsigned Less Than									
Syntax:	{label:}	BRA	NC,	Label							
	{label:}	BRA	LTU,								
Operands:	Label is resolved by the linker to a signed word offset (slit20)										
Operation:	Condition = !C If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else If (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch										
Status Affected:	None										
Encoding:	1010	101U	nnnn	nnnn	nnnn	nnnn	nnnn	0110			
Description:	If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.										
I-Words:	1										
Cycles:	1 (2 or 3) <b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.										

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA LTU, _CHECK_FAIL	; Branch if SR.C == 0, to 0x804030(_CH_ECK_FAIL)
	0x804004	_PASS_CONDITION:	SL.I w5, w6	
		ON:		
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
Before execution			After execution	

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	PC 0x804000			PC 0x804030
	w15 0x4500			w15 0x4500
	SR[7:0] 8'b00000000		(C = 0)	SR[7:0] 8'b00000000 (C = 0)

BRA NN	Branch if Not Negative						
Syntax:	{label:}	BRA	NN,	Label			
Operands:				Label is resolved by the linker to a signed word offset (slit20)			
Operation:				Condition = !N If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch			
Status Affected:				None			
Encoding:	1010	011U	nnnn	nnnn	nnnn	nnnn	0110
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four, and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.			
I-Words:		1					
Cycles:		1 (2 or 3)					
				<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.			

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA NN, _CHECK_FAIL	; Branch if SR.N == 0, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	
PC	0x804000		PC	0x804004
w15	0x4500		w15	0x4500
				(No Branch)

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[7:0]	8'b00001000	(N = 1)	SR[7:0] 8'b00001000 (N = 1)
Example:2	Address	Label	Instruction	Comment
	0x804000	-	BRA N, _CHECK_FAIL	; Branch if SR.N == 1, to0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	
	PC 0x804000		PC 0x804030	
	w15 0x4500		w15 0x4500	
	SR[7:0] 8'b00001000		SR[7:0] 8'b00001000	
	(N = 1)		(N = 1)	

<b>BRA NOV</b>		Branch if Not Overflow						
Syntax:	{label:}	BRA	NOV,	Label				
Operands:				Label is resolved by the linker to a signed word offset (slit20)				
Operation:				Condition = !OV				
				If (condition) then {				
				If slit20 = 1 then skip next (16-bit) instruction				
				Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction				
				Else (PC+4) + 2*slit20 → PC				
				}				
				Else no branch				
Status Affected:				None				
Encoding:	1010	101U	nnnn	nnnn	nnnn	nnnn	nnnn	1110
Description:	<p>If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.</p> <p>Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.</p> <p>The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).</p> <p><b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.</p>							
I-Words:	1							
Cycles:	1 (2 or 3)							
	<p><b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.</p>							

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA NOV, _CHECK_FAIL	; Branch if SR.OV == 0, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITION:	SLI w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
Before execution			After execution	
PC 0x804000			PC 0x804004 (No Branch)	
w15 0x4500			w15 0x4500	

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[7:0] 8'b000000100		(OV = 1)	SR[7:0] 8'b000000100 (OV = 1)

BRA NZ	Branch if Not Zero						
Syntax:	{label:}	BRA	NZ,		Label		
Operands:							
Operation:							
	Label is resolved by the linker to a signed word offset (slit20)						
	Condition = !Z						
	If (condition) then {						
	If slit20 = 1 then skip next (16-bit) instruction						
	Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction						
	Else (PC+4) + 2*slit20 → PC						
	}						
	Else no branch						
Status Affected:							
	None						
Encoding:	1010	011U	nnnn	nnnn	nnnn	nnnn	1110
Description:							
	If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB.						
	If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.						
	If the 2's complement byte offset value '2*slit20' (the PC offset) equals four, and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.						
	If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.						
	Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.						
	The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).						
	<b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.						
I-Words:	1						
Cycles:	1 (2 or 3)						
	<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.						

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA NZ, _CHECK_FAIL	; Branch if SR.Z == 0, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
Before execution			After execution	
	PC	0x804000	PC	0x804030
	w15	0x4500	w15	0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[15:8]	8'b00000000	(Z = 0)	SR[15:8] 8'b00000000 (Z = 0)

BRA OA	Branch if Overflow Accumulator A						
Syntax:	{label:}	BRA	OA,	Label			
Operands:	Label is resolved by the linker to a signed word offset (slit20)						
Operation:	Condition = OA If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch						
Status Affected:	None						
Encoding:	1010	110U	nnnn	nnnn	nnnn	nnnn	0010
Description:	<p>If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.</p> <p>Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.</p> <p>The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).</p> <p><b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.</p>						
I-Words:	1						
Cycles:	1 (2 or 3)						
	<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.						

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA OA, _CHECK_FAIL	; Branch if SR.OA == 1, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
Before execution				After execution
PC    0x804000				PC    0x804030
w15    0x4500				w15    0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[15:8]	8'b10001000	(OA, OAB = 1)	SR[15:8] 8'b10001000 (OA, OAB = 1)

<b>BRA OB</b>		Branch if Overflow Accumulator B						
Syntax:	{label:}	BRA	OB,	Label				
Operands:				Label is resolved by the linker to a signed word offset (slit20)				
Operation:				Condition = OB				
				If (condition) then {				
				If slit20 = 1 then skip next (16-bit) instruction				
				Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction				
				Else (PC+4) + 2*slit20 → PC				
				}				
				Else no branch				
Status Affected:				None				
Encoding:	1010	110U	nnnn	nnnn	nnnn	nnnn	nnnn	0110
Description:	<p>If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.</p> <p>Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.</p> <p>The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).</p> <p><b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.</p>							
I-Words:	1							
Cycles:	1 (2 or 3)							
	<p><b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.</p>							

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA OB, _CHECK_FAIL	; Branch if SR.OA == 1, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	
PC	0x804000		PC	0x804004
w15	0x4500		w15	0x4500
				(No Branch)

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[15:8]	8'b00000000	(OB, OAB = 0)	SR[15:8] 8'b00000000 (OB, OAB = 0)



<b>BRA OV</b>		Branch if Overflow						
Syntax:	{label:}	BRA	OV,		Label			
Operands:								
Operation:								
		Condition = OV						
		If (condition) then {						
		If slit20 = 1 then skip next (16-bit) instruction						
		Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction						
		Else (PC+4) + 2*slit20 → PC						
		}						
		Else no branch						
Status Affected:		None						
Encoding:	1010	101U	nnnn	nnnn	nnnn	nnnn	nnnn	1010
Description:								
		If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB.						
		If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.						
		If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.						
		If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.						
		Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.						
		The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).						
		<b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.						
I-Words:		1						
Cycles:		1 (2 or 3)						
		<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.						

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA OV, _CHECK_FAIL	; Branch if SR.OV == 1, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
Before execution				After execution
	PC	0x804000		PC 0x804030
	w15	0x4500		w15 0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[7:0]	8'b000000100	(OV = 1)	SR[7:0] 8'b000000100 (OV = 1)

<b>BRA</b>	<b>Branch Unconditionally</b>						
------------	-------------------------------	--	--	--	--	--	--

Syntax: {label:} BRA Label

Operands: Label is resolved by the linker to a signed word offset (slit20)

Operation:  $(PC+4) + 2*slit20 \rightarrow PC$

Status Affected: None

Encoding: 1010 111U nnnn nnnn nnnn nnnn nnnn 1010

Description: The program will branch unconditionally with a forward or backward range of 1MB. The 2's complement byte offset value '2\*slit20' (the PC offset) is added to the PC. Since the PC will have incremented to fetch the next instruction, the new address will be  $(PC+4) + 2*slit20$ . The 'n' bits are a signed literal that specifies the number of PS words to branch from  $(PC+4)$ .

I-Words: 1

Cycles: 1

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA _CHECK_FAIL	; Branch to label _CEHCK_FAIL
	0x804004	_PASS_CONDITION:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
Before execution				After execution
PC	0x804000		PC	0x804030
w15	0x4500		w15	0x4500

Example:2	Address	Label	Instruction	Comment
	0x804000	-	BRA _CHECK_FAIL + 4	; Branch to label _CEHCK_FAIL
	0x804004	_PASS_CONDITION:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
Before execution				After execution
PC	0x804000		PC	0x804034
w15	0x4500		w15	0x4500

BRA	Computed Branch							
Syntax:	{label:}	BRA	Wns					
Operands:	Wns ∈ [W0 ... W14]							
Operation:	(PC + 4) + 2*Wns[19:0] → PC, NOP → Instruction Register.							
Status Affected:	None							
Encoding:	1101 011U UUUU ssss UUUU UUUU UUUU 0010							
Description:	Computed branch with a forward or backward branch address range of 1MB. The signed value in Wns[19:0] represents the PS (16-bit) word offset from the current PC. This value is multiplied by two to create a byte address that is then added to the contents of the PC to form the target address. Therefore, Wn must contain a signed value that specifies the number of PS words to offset from (PC+4) for the branch. The 's' bits select the source register. <b>Note:</b> If Wns[31:19] is not all 0's or all 1's, an address error trap will be initiated.							
I-Words:	1							
Cycles:	2							
	<b>Note:</b> The branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.							

Example:1	Address	Label	Instruction	Comment	
	0x803FFC		MOV.I #0x16, W7		
	0x804000	-	BRA W7	; Branch to offset contained in W7 (0x804004 + 2*0x16)	
	0x804004	_PASS_CONDITON:	SL.I w5, w6		
	...		...		
	...		...		
	...		...		
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine	
	...		CP w6, w7		
	...		...		
	Before execution		After execution		
	PC	0x804000	PC	0x804030	
	w15	0x4500	w15	0x4500	

BREAK	Break
Syntax:	{label:} BREAK
Operands:	none
Operation:	<p>Application (User) Mode: Execute as NOP</p> <p>Debugger Mode: Mission Mode: Stop user code execution. Debug Mode: Execute as NOP</p>
Status Affected:	None
Encoding:	S111 001U UUUU 0010 UUUU UUUU UUUU UU00
Description:	<p>BREAK will only execute as such when the device is operating in Mission mode. If encountered in any other mode, it will be executed as an NOP.</p> <p>BREAK will stop user code execution and switch from Mission into Debug mode where it will execute the resident Debug Executive (DE). The User PC is not modified by this instruction. All exceptions (including traps) are blocked while executing BREAK and while in Debug mode.</p> <p>In order to avoid possible hazards between Debug and Mission mode code, BREAK will execute two FNOPs prior to starting Debug mode execution. The total instruction cycle count includes these FNOPs.</p> <p>The 'S' bit selects instruction size.</p> <p><b>Note:</b> 16-bit encoding (<b>bold</b>).</p>
I-Words:	1 or 0.5
Cycles:	3

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA SA, _CHECK_FAIL	; Branch if SR.SB == 1, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITION:	SL.I w5, w6	
			...	...
			...	...
			...	...
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
			CP w6, w7	
			...	...
	Before execution			After execution
PC	0x804000		PC	0x804004 (No Branch)
w15	0x4500		w15	0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[15:8]	8'b00000000	(SA, SAB = 0)	SR[15:8] 8'b00000000 (SA, SAB = 0)

BRA SB	Branch if ACCB Saturation						
Syntax:	{label:}	BRA	SB,	Label			
Operands:				Label is resolved by the linker to a signed word offset (slit20)			
Operation:				Condition = SB If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch			
Status Affected:				None			
Encoding:	1010	110U	nnnn	nnnn	nnnn	nnnn	1110
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.			
I-Words:		1					
Cycles:		1 (2 or 3)					
				<b>Note:</b> The branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.			

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA SB, _CHECK_FAIL	; Branch if SR.SB == 1, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITI ON:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
		Before execution		After execution
PC	0x804000		PC	0x804030
w15	0x4500		w15	0x4500

**Instruction Descriptions (A to BZ) (continued)**

Example:1	Address	Label	Instruction	Comment
	SR[15:8]	8'b000010100	(SB, SAB = 1)	SR[15:8] 8'b000000010 (SB, SAB = 1)

BSET							
Bit Set in Ws							
Syntax:	{label:}	BSET.b	Ws,	bit5			
		BSET{.w}	[Ws],				
		BSET.I	[Ws++],				
			[Ws-],				
			[++Ws],				
			[--Ws],				
			SR				
Operands:	Ws ∈ [W0 ... W15]; Byte: bit5 ∈ [0 ... 7]; Word: bit5 ∈ [0 ... 15]; Long word: bit5 ∈ [0 ... 31]						
Operation:	1 → Ws<bit5>						
Status Affected:	None (see note 2)						
Encoding:	S100      001b      bbbb      ssss      pppU      UUUU      UUUU      LB00						
Description:	Bit 'bit5' in register Ws is set. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'b' bits define value or bit5 of the bit position to be cleared (bit5[4:0] for the 16-bit opcode). The 's' bits select the source/destination register. The 'p' bits select the source addressing mode. See for modifier addressing information.						
<b>Note:</b>							
1.	When targeting the SR, a bit within the SR will be set as a result of the instruction operation.						
I-Words:	1 or 0.5						
Cycles:	1						

Example 1: BSET.B W3, #0x7		; Set bit 7 in W3	
Before Instruction		After Instruction	
W3	0026	W3	00A6
SR	0000	SR	0000
<b>Example 2: BSET [W4++], #0x0</b>		<b>; Set bit 0 in [W4] ; Post-increment W4</b>	
Before Instruction		After Instruction	
W4	6700	W4	6702
Data 6700	1734	Data 6700	1735
SR	0000	SR	0000

BSET	Bit Set f							
Syntax:	{label:}	BSET.b	f,	bit3				
Operands:			bit3 ∈ [0 ... 7]; f ∈ [0 ... 1MB]					
Operation:			1 → f<bit3>					
Status Affected:			None					
Encoding:	1100	001b	ffff	ffff	ffff	ffff	ffff	bb01
Description:	Bit 'bit3' in file register f is set. The 'w' bits select value bit3 of the bit position to be set. The 'f' bits select the address of the file register.							
<b>Notes:</b>								
1.	This instruction operates in Byte mode only.							
2.	The .b extension must be included with the opcode.							
3.	Accessible file address space is 1MB.							
I-Words:	1							
Cycles:	1							

Example 1:	BSET.B 0x601, #0x3	; Set bit 3 in 0x601	
		Before Instruction	After Instruction
	Data 0600 F234		Data 0600 FA34
	SR 0000		SR 0000
Example 2:	BSET 0x444, #0xF	; Set bit 15 in 0x444	
		Before Instruction	After Instruction
	Data 0444 5604		Data 0444 D604
	SR 0000		SR 0000

BSW	Bit Write in Ws								
Syntax:	{label:}	BSW.bC	Ws,	Wb					
		BSW.bZ	[Ws],						
		BSW.{w}C	[++Ws],						
		BSW.{w}Z	[--Ws],						
		BSW.IC	[Ws++],						
		BSW.IZ	[Ws--],						
Operands:		Wb ∈ [W0 ... W14]; Ws ∈ [W0 ... W15]							
Operation:		If ".Z" option, then Z → Ws<(Wb)> If ".C" option, then C → Ws<(Wb)>							
Status Affected:		None							
Encoding:	S100	101L	wwww	ssss	pppU	UUUU	UUUG	BU00	
Description:	<p>The bit number defined in Wb is written in Ws with the value of the C or Z bit. For byte, word and long word operations, the target bit number is defined by Wb[2:0], Wb[3:0] and Wb[4:0], respectively. Any bit within Wb, beyond those bits required to select the target bit, will be ignored and may be set to any value.</p> <p>The 'S' bit selects instruction size.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'w' bits select the bit select register.</p> <p>The 'G' bit selects the Z or C flag bit as source (G = 0 selects Z flag).</p> <p>The 's' bits select the source register.</p> <p>The 'p' bits select the source addressing mode.</p>								
I-Words:	1 or 0.5								
Cycles:	1								

Example 1:	BSW.C W2, W3	; Set bit W3 in W2 to the value ; of the C bit	
------------	--------------	---	--

Before Instruction		After Instruction	
W2	F234	W2	7234
W3	111F	W3	111F
SR	0002 (Z = 1, C = 0)	SR	0002 (Z = 1, C = 0)

Example 2:	BSW.Z W2, W3	; Set bit W3 in W2 to the complement ; of the Z bit	
------------	--------------	---	--

Before Instruction		After Instruction	
W2	E235	W2	E234
W3	0550	W3	0550
SR	0002 (Z = 1, C = 0)	SR	0002 (Z = 1, C = 0)

Example 3:	BSW.C [++W0], W6	; Set bit W6 in [W0++] to the value ; of the C bit	
------------	------------------	--	--

Before Instruction		After Instruction	
W0	1000	W0	1002
W6	34A3	W6	34A3
Data 1002	2380	Data 1002	2388
SR	0001 (Z = 0, C = 1)	SR	0001 (Z = 0, C = 1)

Example 4:      BSW.Z [W1-], W5 ; Set bit W5 in [W1] to the  
; complement of the Z bit  
; Post-decrement W1

Before Instruction		After Instruction	
W1	1000	W1	0FFE
W5	888B	W5	888B
Data 1000	C4DD	Data 1000	CCDD
SR	0001 (C = 1)	SR	0001 (C = 1)

BTG	Bit Toggle in Ws						
Syntax:	{label:}	BTG.b	Ws,	bit5			
		BTG{.w}	[Ws],				
		BTG.l	[Ws++],				
			[Ws-],				
			[++Ws],				
			[--Ws],				
			SR				
Operands:	Ws ∈ [W0 ... W15]; Byte: bit5 ∈ [0 ... 7]; Word: bit5 ∈ [0 ... 15]; Long word: bit5 ∈ [0 ... 31]						
Operation:	(Ws)[bit5] → Ws[bit5]						
Status Affected:	None (see note 2)						
Encoding:	S100      010b      bbbb      ssss      pppU      UUUU      UUUU      LB00						
Description:	Bit 'bit5' in register Ws is toggled. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'b' bits define value or bit5 of the bit position to be cleared (bit5[4:0] for the 16-bit opcode) The 's' bits select the source/destination register. The 'p' bits select the source addressing mode.						
<b>Note:</b>	<ol style="list-style-type: none"> <li>When targeting the SR, a bit within the SR will be toggled as a result of the instruction operation.</li> </ol>						
I-Words:	1 or 0.5						
Cycles:	1						

BTG	Bit Toggle f							
Syntax:	{label:}	BTG.b	f,	bit3				
Operands:	bit3 ∈ [0 ... 7]; f ∈ [0 ... 1MB]							
Operation:	(f)[bit3] → (f)[bit3]							
Status Affected:	None							
Encoding:	1100	010b	ffff	ffff	ffff	ffff	ffff	bb01
Description:	Bit 'bit3' in file register f is toggled. The 'w' bits select value bit3 of the bit position to be cleared. The 'f' bits select the address of the file register.							
<b>Notes:</b>								
1. This instruction operates in Byte mode only. 2. The .b extension must be included with the opcode. 3. Accessible file address space is 1MB.								
I-Words:	1							
Cycles:	1							

<b>Example 1:</b>	<b>BTG W2, #0x0</b>	<b>; Toggle bit 0 in W2</b>
	Before Instruction	After Instruction
	W2 F234	W2 F235
	SR 0000	SR 0000
<b>Example 2:</b>	<b>BTG [W0++], #0x0</b>	<b>; Toggle bit 0 in [W0] ; Post-increment W0</b>
	Before Instruction	After Instruction
	W0 2300	W0 2302
	Data 2300 5606	Data 2300 5607
	SR 0000	SR 0000
<b>Example 3:</b>	<b>BTG.B 0x1001, #0x4</b>	<b>; Toggle bit 4 in 0x1001</b>
	Before Instruction	After Instruction
	Data 1000 F234	Data 1000 E234
	SR 0000	SR 0000
<b>Example 4:</b>	<b>BTG 0x1660, #0x8</b>	<b>; Toggle bit 8 in RAM660</b>
	Before Instruction	After Instruction
	Data 1660 5606	Data 1660 5706
	SR 0000	SR 0000

BTST Bit Test in Ws						
Syntax:	{label:}	BTST.bC	Ws,	bit5		
		BTST.bZ	[Ws],			
		BTST.{w}C	[Ws++],			
		BTST.{w}Z	[Ws--],			
		BTST.IC	[++Ws],			
		BTST.IZ	[--Ws],			
Operands:	Ws ∈ [W0 ... W15];					
	Byte: bit5 ∈ [0 ... 7];					
	Word: bit5 ∈ [0 ... 15];					
	Long word: bit5 ∈ [0 ... 31]					
Operation:	If ".Z" option, (Ws)[bit5] → Z If ".C" option, (Ws)[bit5] → C					
Status Affected:	C or Z					
Encoding:	S100      011b      bbbb      ssss      pppU      UUUU      UUUG      LB00					
Description:	Bit 'bit5' in register Ws is tested. The Zero Flag contains the inversion of the bit or the Carry Flag contains the bit. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'G' bit selects the Z or C flag bit as source (G = 0 selects Z flag). The 'b' bits define value or bit5 of the bit position to be cleared (bit5[4:0] for the 16-bit opcode). The 's' bits select the source/destination register. The 'p' bits select the source addressing mode.					
I-Words:	1 or 0.5					
Cycles:	1					

Example 1:	BTST.C [W0++], #0x3	; Set C = bit 3 in [W0] ; Post-increment W0
------------	---------------------	--

Before Instruction			After Instruction		
W0	1200		W0	1202	
Data 1200	FFF7		Data 1200	FFF7	
SR	0001 (C = 1)		SR	0000	

Example 2:	BTST.Z W0, #0x7	; Set Z = complement of bit 7 in W0
------------	-----------------	-------------------------------------

Before Instruction			After Instruction		
W0	F234		W0	F234	
SR	0000		SR	0002 (Z = 1)	

BTST	Bit Test f			
------	------------	--	--	--

Syntax:                    {label:}      BTST.b      f,      bit3

Operands:                bit3 ∈ [0 ... 7]; f ∈ [0 ... 1MB]

Operation:                (f)[bit3] → Z

Status Affected:        Z

Encoding:                1100      011b      ffff      ffff      ffff      ffff      bb01

Description:               Bit 'bit3' in file register f is tested, the Zero Flag bit is set if it is zero and cleared otherwise. The file register contents are unchanged.

The 'b' bits select value bit3 of the bit position to be cleared.

The 'f' bits select the address of the file register.

**Notes:**

1. This instruction operates in Byte mode only.
2. The .b extension must be included with the opcode.
3. Accessible file address space is 1MB.

I-Words:                1

Cycles:                 1

Example 1:	BTST.B 0x1201, #0x3	; Set Z = complement of ; bit 3 in 0x1201
------------	---------------------	--

Before Instruction	After Instruction
Data 1200      F7FF SR      0000	Data 1200      F7FF SR      0002      (Z = 1)

Example 2:	BTST 0x1302, #0x7	; Set Z = complement of ; bit 7 in 0x1302
------------	-------------------	--

Before Instruction	After Instruction
Data 1302      F7FF SR      0002      (Z = 1)	Data 1302      F7FF SR      0000

BTST	Bit Test/Set in Ws						
Syntax:	{label:}	BTST.bC	Ws,	bit5			
		BTST.bZ	[Ws],				
		BTST.{w}C	[Ws++],				
		BTST..{w}Z	[Ws--],				
		BTST.IC	[++Ws],				
		BTST..IZ	[--Ws],				
Operands:	Ws ∈ [W0 ... W15]; Byte: bit5 ∈ [0 ... 7]; Word: bit5 ∈ [0 ... 15]; Long word: bit5 ∈ [0 ... 31]						
Operation:	if ".Z" option, first (Ws)[bit5] → Z, then 1 → Ws[bit5] if ".C" option, first (Ws)[bit5] → C, then 1 → Ws[bit5]						
Status Affected:	C or Z						
Encoding:	S100      100b      bbbb      ssss      pppU      UUUU      UUUG      LB00						
Description:	Bit 'bit5' in register Ws is tested and then set. For the ".Z" option, the Z Flag is set to the complement of the 'bit5' value prior to being set, and the C Flag is not modified. For the ".C" option, the C Flag is set to the 'bit5' value prior to being set, and the Z Flag is not modified. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'G' bit selects the Z or C Flag bit as source (G = 0 selects Z Flag). The 'b' bits define value or bit5 of the bit position to be cleared. The 's' bits select the source/destination register. The 'p' bits select the source addressing mode.						
I-Words:	1 or 0.5						
Cycles:	1						

Example 1:	BTSTS.C [W0++], #0x3	; Set C = bit 3 in [W0] ; Set bit 3 in [W0] = 1 ; Post-increment W0
------------	----------------------	---

Before Instruction				After Instruction			
W0	1200			W0	1202		
Data 1200	FFF7			Data 1200	FFFF		
SR	0001	(C = 1)		SR	0000		

Example 2:	BTSTS.Z W0, #0x7	; Set Z = complement of bit 7 ; in W0, and set bit 7 in W0 = 1
------------	------------------	---

Before Instruction				After Instruction			
W0	F234			W0	F2BC		
SR	0000			SR	0002	(Z = 1)	

<b>BTSTS</b>		<b>Bit Test/Set f</b>							
Syntax:	{label:}	BTSTS.b      f,      bit3							
Operands:	bit3 ∈ [0 ... 7]; f ∈ [0 ... 1MB]								
Operation:	First (f)[bit3] → Z, then 1 → (f)[bit3]								
Status Affected:	Z								
Encoding:	1100      100b      ffff      ffff      ffff      ffff      ffff      bb01								
Description:	Bit 'bit3' in file register f is tested and then set. The Z Flag is set to the complement of the 'bit3' value prior to being set. The 'w' bits select value bit3 of the bit position to be test/set. The 'f' bits select the address of the file register.								
Notes:	<ol style="list-style-type: none"> <li>1. This instruction operates in Byte mode only.</li> <li>2. The .b extension must be included with the opcode.</li> <li>3. Accessible file address space is 1MB.</li> </ol>								
I-Words:	1								
Cycles:	1								

Example 1:	BTSTS.B 0x1201, #0x3	; Set Z = complement of bit 3 in 0x1201, ; then set bit 3 of 0x1201 = 1
------------	----------------------	--

Before Instruction				After Instruction			
Data 1200 F7FF				Data 1200 FFFF			
SR 0000				SR 0002 (Z = 1)			

Example 2:	BTSTS 0x808, #15	; Set Z = complement of bit 15 in 0x808, ; then set bit 15 of 0x808 = 1
------------	------------------	--

Before Instruction				After Instruction			
RAM300 8050				RAM300 8050			
SR 0002 (Z = 1)				SR 0000			

BTST	Bit Test in Ws								
Syntax:	{label:}	BTST.bC	Ws,	Wb					
		BTST.bZ	[Ws],						
		BTST.{w}C	[Ws++]						
		BTST.{w}Z	[Ws--],						
		BTST.IC	[++Ws],						
		BTST.IZ	[--Ws],						
Operands:		Wb ∈ [W0 ... W14]; Ws ∈ [W0 ... W15]							
Operation:		if ".Z" option, (Ws) < (Wb) → Z if ".C" option, (Ws) < (Wb) → C							
Status Affected:		C or Z							
Encoding:	S100	110L	wwww	ssss	pppU	UUUU	UUUG	BU00	
Description:	<p>The bit number defined in Wb is tested in source register Ws. For byte, word and long word operations, the target bit number is defined by Wb[2:0], Wb[3:0] and Wb[4:0], respectively. Any bit within Wb, beyond those bits required to select the target bit, will be ignored and may be set to any value.</p> <p>For the ".Z" option, the Z Flag is set to the complement of the bit value tested, and the C Flag is not modified. For the ".C" option, the C Flag is set to the bit value tested, and the Z Flag is not modified. Wb[31:5] are ignored and may be set to any value.</p> <p>The 'S' bit selects instruction size.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'w' bits select the bit select register.</p> <p>The 'G' bit selects the Z or C Flag bit as source (G = 0 selects Z Flag).</p> <p>The 's' bits select the source register.</p> <p>The 'p' bits select the source addressing mode.</p>								
I-words:	1 or 0.5								
Cycles:	1								

Example 1:	BTST.C W2, W3	; Set C = bit W3 of W2	
		Before Instruction	After Instruction
		W2 F234	W2 F234
		W3 2368	W3 2368
		SR 0001 (C = 1)	SR 0000
Example 2:	BTST.Z [W0++], W1	; Set Z = complement of ; bit W1 in [W0], ; Post-increment W0	
		Before Instruction	After Instruction
		W0 1200	W0 1202
		W1 CCC0	W1 CCC0
		Data 1200 6243	Data 1200 6243
		SR 0002 (Z = 1)	SR 0000

BRA Z	Branch if Zero						
Syntax:	{label:}	BRA	Z,	Label			
Operands:				Label is resolved by the linker to a signed word offset (slit20)			
Operation:				Condition = Z If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch			
Status Affected:				None			
Encoding:	1010	011U	nnnn	nnnn	nnnn	nnnn	1010
Description:				If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). <b>Note:</b> Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.			
I-Words:		1					
Cycles:		1 (2 or 3)					
				<b>Note:</b> The branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.			

Example:1	Address	Label	Instruction	Comment
	0x804000	-	BRA Z, _CHECK_FAIL	; Branch if Z == 1, to 0x804030(_CHECK_FAIL)
	0x804004	_PASS_CONDITION:	SL.I w5, w6	
	...		...	
	...		...	
	...		...	
	0x804030	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
		Before execution		After execution
		PC 0x804000		PC 0x804030
		w15 0x4500		w15 0x4500
		SR[7:0] 8'b00000010	(Z = 1)	SR[7:0] 8'b00000010
				(Z = 1)

## 4.5. Instruction Descriptions (C to DTB)

CALL		Call Subroutine								
Syntax:	{label:}	CALL	lit24							
Operands:	lit24 $\in [0 \dots 16\text{MB}]$									
Operation:	$(\text{PC}) + 4 \rightarrow \text{PC}$ , $(8'b0, \text{PC}[23:2]) \rightarrow \text{TOS}[31:2]$ ; $2'b00 \rightarrow \text{TOS}[1:0]$ , $(W15) + 4 \rightarrow W15$ , $\text{lit24} \rightarrow \text{PC}[23:0]$ ; NOP $\rightarrow$ Instruction Register									
Status	None									
Affected:										
Encoding:	1101	111n	nnnn	nnnn	nnnn	nnnn	nnnn	nnn10		
Description:	Subroutine call to any address within all of executable memory address space. A call to either a 32-bit or 16-bit instruction is permitted. The long word aligned return address (PC+4) is pushed onto the system stack. The 24-bit value 'lit24' is then loaded into the PC (the opcode does not store the LSb which is always 1'b0). The 'n' bits form the target address. <b>Note:</b> The (byte) PC address is always either word or long word aligned. The opcode does not store the LSb because it is always 1'b0.									
I-Words:	1									
Cycles:	1									

Example:1	Address	Label	Instruction	Comment
	0x804000	-	CALL _CHECK_FAIL	; Calls function in address 0x805000(_CHECK_FAIL)
	...			
	0x805000	_CHECK_FAIL	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	
	PC	0x804000	PC	0x805000
	w15	0x4500	w15	0x4504
	Data @4500	0x0000	Data @4500	0x804004
	SR	0x00000000	SR	0x00000000

<b>CALL</b>		Call Indirect Subroutine												
Syntax:	{label:}	CALL	Wns											
Operands:	Wns $\in [W0 \dots W14]$													
Operation:	$(PC) + 2 \rightarrow PC,$ $(8'b0, PC[23:2]) \rightarrow TOS[31:2]; 2'b00 \rightarrow TOS[1:0],$ $(W15)+4 \rightarrow W15,$ $Wns[23:0] \rightarrow PC[23:0];$ NOP $\rightarrow$ Instruction Register.													
Status Affected:	None													
Encoding:	1101	011U	UUUU	ssss	UUUU	UUUU	UUUU	1110						
Description:	Indirect subroutine call of any instruction address (16-bit or 32-bit) within program memory using a computed call PS (word) address. The long word aligned return address (PC+4) is pushed onto the system stack. The 24-bit value (Wns[23:0]) is then loaded into the PC[23:0]. Wns must therefore contain a PS byte address. The value of Wns[0] is ignored, and PC[0] is always set to 1'b0. The 's' bits specify the source register.													
<b>Note:</b>	1. If Wns[31:24] != 8'h00, an address error trap will be initiated.													
I-Words:	1													
Cycles:	2													
<b>Note:</b> The call target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.														

Example:1	Address	Label	Instruction	Comment
	0x803FFC	-	MOV.I #_CHECK_FAIL, W0	
	0x804000	-	CALL W0	; Calls function in address 0x805000(_CHECK_FAIL)
	...			
	0x805000	_CHECK_FAIL:	MOV.I SR, W6	; _CHECK_FAIL SubRoutine
	...		CP w6, w7	
	...		...	
	Before execution		After execution	
	PC	0x804000	PC	0x805000
	W0	0x805000	PC	0x805000
	w15	0x4500	w15	0x4504
	Data @4500	0x0000	Data @4500	0x804004
	SR	0x00000000	SR	0x00000000

CLR	Clear Accumulator			
Syntax: {label}	CLR	A		
		B		
Operands:	None			
Operation:	0 → ACC(A or B)			
Status Affected:	OA,SA or OB,SB			
Encoding:	0111      001A      UUUU      1100			
Description:	Clear the specified accumulator (A or B). The 'A' bit selects the accumulator to clear.			
I-Words:	0.5			
Cycles:	1			

Example 1:	CLR A; Clear ACCA	
	Before Instruction	After Instruction
	ACCA 0x00_0012 _3412_452 4_1234	ACCA 0x00_0000 _0000_000 0_0000

CLR	Clear f							
Syntax:	{label:}	CLR.b	f					
		CLR{.w}						
		CLR.I						
Operands:	$f \in [0 \dots 1\text{MB}]$							
Operation:	0x00000000 → file register for long operation 0x0000 → file register for word operation 0x00 → file register for byte operation							
Status Affected:	None							
Encoding:	1010	011L	ffff	ffff	ffff	ffff	ffff	B101
Description:	Clear the file register. The 'L' and 'B' bits select operation data width. The 'f' bits select the address of the file register.							
<b>Notes:</b>	<ol style="list-style-type: none"> <li>Same flow as any other file operation with D-bit (opcode[2]) set to 1.</li> <li>Accessible file address space is 1MB.</li> </ol>							
I-Words:	1							
Cycles:	1							

Example 1:	CLR.B RAM200 ; Clear RAM200 (Byte mode)							
Before Instruction				After Instruction				
RAM200 8009				RAM200 8000				
SR 0000				SR 0000				

Example 2:	CLR WREG ; Clear WREG (Word mode)							
Before Instruction				After Instruction				
WREG 0600				WREG 0000				
SR 0000				SR 0000				

CLRWDT	Clear Watchdog Timer
Syntax:	{label:} CLRWDT
Operands:	None
Operation:	0 → WDT Reg
Status Affected:	None
Encoding:	0111 001U UUUU 0101
Description:	Clear the WatchDog Timer register.
I-Words:	0.5
Cycles:	1

COM	Complement Ws				
Syntax:	{label:}	COM.b	Ws,	Wd	
		COM.bz	[Ws],	[Wd]	
		COM{.w}	[Ws++],	[Wd++]	
		COM.l	[Ws-],	[Wd--]	
			[++Ws],	[++Wd]	
			[~-Ws],	[~-Wd]	
Operands:	$Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$				
Operation:	$(Ws) \rightarrow Wd$				
Status Affected:	Z, N				
Encoding:	S111	100L	dddd	ssss	pppq
Description:	Compute the 1's complement of the contents of the source register Ws and place the result in the destination register Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.				
I-words:	1 or 0.5				
Cycles:	1				

Example 1:	<b>COM.B [W0++], [W1++]; COM [W0] and store to [W1] (Byte mode)</b> <b>; Post-increment W0, W1</b>					
	Before Instruction					
	W0 2301					
	W1 2400					
	Data 2300 5607					
	Data 2400 ABCD					
	SR 0000					
	After Instruction					
	W0 2302					
	W1 2401					
	Data 2300 5607					
	Data 2400 ABA9					
	SR 0008 (N = 1)					

Example 2:	<b>COM W0, [W1++]; COM W0 and store to [W1] (Word mode)</b> <b>; Post-increment W1</b>					
	Before Instruction					
	W0 D004					
	W1 1000					
	Data 1000 ABA9					
	SR 0000					
	After Instruction					
	W0 D004					
	W1 1002					
	Data 1000 2FFB					
	SR 0000					

COM	Complement f						
Syntax:	{label:}	COM.b	f	{,Wnd}	{,WREG}		
		COM.bz					
		COM{.w}					
		COM.I					
Operands:		f ∈ [0 ... 64KB]; Wnd ∈ [W0 ... W14]					
Operation:		(f) → destination designated by D					
Status Affected:		Z, N					
Encoding:	1101	100L	dddd	ffff	ffff	ffff	BD01
Description:							
		Compute the 1's complement of the contents of the file register and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register.					
		The 'L' and 'B' bits select operation data width.					
		The 'D' bit selects the destination.					
		The 'f' bits select the address of the file register.					
		The 'd' bits select the Working register.					
I-Words:	1						
Cycles:	1						

**Example 1: COM.b RAM200 ; COM RAM200 (Byte mode)**

Before Instruction				After Instruction			
RAM200	80FF			RAM200	8000		
SR	0000			SR	0002	(Z)	

**Example 2: COM RAM400, WREG ; COM RAM400 and store to WREG ; (Word mode)**

Before Instruction				After Instruction			
WREG	1211			WREG	F7DC		
RAM400	0823			RAM400	0823		
SR	0000			SR	0008	(N = 1)	

CP	Compare Wb with Ws, Set Status Flags						
Syntax:	{label:}	CP.b	Wb,	Ws			
		CP{.w}		[Ws]			
		CP.l		[Ws++]			
				[Ws--]			
				[++Ws]			
				[--Ws]			
Operands:		$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]$					
Operation:		$(Wb) - (Ws)$					
Status Affected:		C, N, OV, Z					
Encoding:	S011	110L	wwww	ssss	pppU	UUUU	UUUU BU00
Description:	Compute $(Wb) - (Ws)$ , then set flags but do not store result. Equivalent to SUB instruction (without a destination result write). The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'w' bits select the Wb source register. The 's' bits select the Ws source register. The 'p' bits select the source addressing mode.						
I-Words:	1 or 0.5						
Cycles:	1						

Example 1:	<b>CP.B W0, [W1++] ; Compare [W1] with W0 (Byte mode)</b> <b>; Post-increment W1</b>						
	Before Instruction						
	W0 ABA9						
	W1 2000						
	Data 2000 D004						
	SR 0000						
	After Instruction						
	W0 ABA9						
	W1 2001						
	Data 2000 D004						
	SR 0009 (N, C = 1)						

Example 2:	<b>CP W5, W6 ; Compare W6 with W5 (Word mode)</b>						
	Before Instruction						
	W5 2334						
	W6 8001						
	SR 0000						
	After Instruction						
	W5 2334						
	W6 8001						
	SR 000C (N, OV = 1)						

<b>CPB</b>		Compare Wb with Ws with Borrow, Set Status Flags						
Syntax:	{label:}	CPB.b	Wb,	Ws				
		CPB{.w}		[Ws]				
		CPB.l		[Ws++]				
				[Ws--]				
				[++Ws]				
				[--Ws]				
Operands:		$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]$						
Operation:		$(Wb) - (Ws) - (C)$						
Status Affected:		C, N, OV, Z						
Encoding:		S011	111L	wwww	ssss	pppU	UUUU	UUUU BU00
Description:		Compute $(Wb) - (Ws) - (C)$ , then set flags but do not store result. Equivalent to SUBB instruction without a destination result write. The Z bit is "sticky" (can only be cleared). The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'w' bits select the Wb source register. The 's' bits select the Ws source register. The 'p' bits select the source addressing mode.						
I-Words:		1 or 0.5						
Cycles:		1						

CPB							Compare Ws with lit13 with Borrow, Set Status Flags						
Syntax:	{label:}	CPB.b	Ws,	lit13									
		CPB{.w}	[Ws],										
		CPB.l	[Ws++],										
			[Ws-],										
			[++Ws],										
			[--Ws],										
Operands:		Ws ∈ [W0 ... W15];											
		lit13 ∈ [0 ... 8191] (32-bit opcode); lit4 ∈ [0 ... 15] (16-bit opcode) <sup>1</sup>											
Operation:		(Ws) - lit13 - (C)											
		or											
		(Ws) - lit4 - (C) <sup>1</sup>											
		<b>Note:</b> The literal is zero-extended to the selected data size of the operation											
Status Affected:		C, N, OV, Z											
Encoding:		S011      101L      kkkk      ssss      pppk      kkkk      kkkk      BU00											
Description:		Zero-extend literal then compute (Ws) - literal - (C). Set flags but do not store result.											
		The Z bit is sticky (can only be cleared).											
		The 'L' and 'B' bits select operation data width.											
		The 'p' bits select the source addressing mode.											
		The 's' bits select the source register.											
		The 'k' bits provide the literal operand.											
		lit13[12:0] = Opcode[12:4], Opcode[23:20]											
		lit4[3:0] = Opcode[23:20]											
I-Words:		1 or 0.5											
Cycles:		1											

CP	Compare Ws with lit13, Set Status Flags						
Syntax:	{label:}	CP.b	Ws,	lit13			
		CP{.w}	[Ws],				
		CP.l	[Ws++],				
			[Ws-],				
			[++Ws],				
			[-Ws],				
Operands:	Ws ∈ [W0 ... W15]; lit13 ∈ [0 ... 8191] (32-bit opcode); lit4 ∈ [0 ... 15] (16-bit opcode) <sup>1</sup>						
Operation:	(Ws) - lit13 (32-bit opcode) or (Ws) - lit4 (16-bit opcode) <sup>1</sup>						
	Note: The literal is zero-extended to the selected data size of the operation						
Status Affected:	C, N, OV, Z						
Encoding:	S011      100L      kkkk      ssss      pppk      kkkk      kkkk      BU00						
Description:	Zero-extend literal as necessary, then compute (Ws) - literal. Set flags but do not store result. The Z bit is sticky (can only be cleared). The 'L' and 'B' bits select operation data width. The 'p' bits select the source addressing mode. The 's' bits select the source register. The 'k' bits provide the literal operand. lit13[12:0] = Opcode[12:4], Opcode[23:20] lit4[3:0] = Opcode[23:20]						
I-Words:	1 or 0.5						
Cycles:	1						

<b>CPB</b>		<b>Compare Wb with lit16 with Borrow, Set Status Flags</b>												
Syntax:		{label:} CPB.b		Wb,	lit16									
		CPB{.w}												
		CPB.l												
Operands:		Wb ∈ [W0 ... W15]; lit16 ∈ [0 ... 65535]												
Operation:		(Wb) - lit16 - (C)												
Status Affected:		C, N, OV, Z												
Encoding:		1100	011L	www	kkkk	kkkk	kkkk	kkkk	B110					
Description:		Compute (Wb) - lit16 - (C), set flags but do not store result. The Z bit is “sticky” (can only be cleared). The ‘L’ and ‘B’ bits select operation data width. The ‘w’ bits select the source register. The ‘k’ bits provide the literal operand.												
I-Words:		1												
Cycles:		1												

CP	Compare f with Ws, Set Status Flags							
Syntax:	{label:}	CP.b	f,	Ws				
		CP{.w}						
		CP.l						
Operands:		$f \in [0 \dots 64KB]$ ; $Ws \in [W0 \dots W15]$						
Operation:		$(f) - (Ws)$						
Status Affected:		C, N, OV, Z						
Encoding:	1010	101L	ssss	ffff	ffff	ffff	ffff	BU01
Description:	<p>Compute <math>(f) - (Ws)</math>, set flags but do not store result. Equivalent to SUBWF instruction with a stack sink destination ([W15]).</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>							
I-Words:	1							
Cycles:	1							

<b>CP0</b>		<b>Compare f with Zero, Set Status Flags</b>								
Syntax:	{label:}	CP0.b	f							
		CP0{.w}								
		CP0.l								
Operands:	f ∈ [0 ... 1MB]									
Operation:	Long: (f) - 0x00000000									
	Word: (f) - 0x0000									
	Byte: (f) - 0x00									
Status Affected:	C, N, OV, Z									
Encoding:	1010	100L	ffff	ffff	ffff	ffff	ffff	ffff	BU01	
Description:	Compute (f) minus zero value for selected data size, set flags but do not store result. The C bit is always set, and OV is always cleared by this operation.									
	The 'L' and 'B' bits select operation data width.									
	The 'f' bits select the address of the file register.									
	Accessible file address space is 1 MB.									
I-Words:	1									
Cycles:	1									

<b>Example 1: CP0.B RAM100 ; Compare RAM100 with 0x0 (Byte mode)</b>	
Before Instruction	After Instruction
RAM100    44C3	RAM100    44C3
SR    0000	SR    0009    (N, C = 1)

<b>Example 2: CP0 0x1FFE ; Compare (0x1FFE) with 0x0 (Word mode)</b>	
Before Instruction	After Instruction
Data 1FFE    0001	Data 1FFE    0001
SR    0000	SR    0001    (C = 1)

<b>CPB</b>		Compare f with Ws with Borrow, Set Status Flags							
Syntax:		{label:}	CPB.b	f,	Ws				
			CPB{.w}						
			CPB.l						
Operands:		$f \in [0 \dots 64KB]$ ; $Ws \in [W0 \dots W15]$							
Operation:		$(f) - (Ws) - (C)$							
Status Affected:		C, N, OV, Z							
Encoding:		1010	110L	ssss	ffff	ffff	ffff	ffff	BU01
Description:		Compute $(f) - (Ws) - (C)$ , set flags but do not store result. Equivalent to SUBBWF instruction with a stack sink destination write ([W15]). The Z bit is “sticky” (can only be cleared). The ‘L’ and ‘B’ bits select operation data width. The ‘F’ bits select the address of the file register. The ‘s’ bits select the Working register.							
I-Words:	1								
Cycles:	1								

CP	Compare Wb with lit16, Set Status Flags								
Syntax:	{label:}	CP.b	Wb,	lit16					
		CP{.w}							
		CP.l							
Operands:		$Wb \in [W0 \dots W15]; lit16 \in [0 \dots 65535]$							
Operation:		$(Wb) - lit16$							
Status Affected:		C, N, OV, Z							
Encoding:	1100	010L	www	kkkk	kkkk	kkkk	kkkk	B110	
Description:	<p>Compute <math>(Wb) - lit16</math>, set flags but do not store result.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'w' bits select the source register.</p> <p>The 'k' bits provide the literal operand.</p>								
I-Words:	1								
Cycles:	1								

CTXTSWP	CPU Register Context Swap Literal			
Syntax:	{label:} CTXTSWP lit3			
Operands:	lit3 ∈ [0 ... 7]			
Operation:	Switch CPU register context to context defined by lit3 lit3 → SR.CTX[2:0]			
Status Affected:	None			
Encoding:	0111	001U	Ukkk	0110
Description:	<p>This instruction will force a CPU register context switch from the current context to the target context defined by value lit3. If supported, coprocessor contexts will also be switched accordingly. A context switch will update the Current Context Identifier (SR.CTX[2:0]) to reflect the new active CPU register context.</p> <p>The 'k' bits select the target register context.</p>			
I-Words:	0.5			
Cycles:	2			

CTXTSWP	CPU Register Context Swap Ws							
Syntax:	{label:} CTXTSWP Wns							
Operands:	Wns ∈ [W0 ... W14]							
Operation:	Switch CPU register context to context defined in Wns[2:0] <sup>1</sup> Wns[2:0] → SR.CTX[2:0]							
Status Affected:	None							
Encoding:	1101	100U	UUUU	ssss	UUUU	UUUU	UUUU	U110
Description:	<p>This instruction will force a CPU register context switch (W0 through W7) from the current context to the target context defined by the contents of Wns[2:0]. Supported register contexts in any instantiated coprocessors will also be switched.</p> <p>A context switch will update the Current Context (SR.CTX[2:0]) to reflect the new active CPU register context.</p> <p>The 's' bits select the source register.</p> <p><b>Note:</b></p> <ol style="list-style-type: none"> <li>1. The contents of Wns[31:3] are ignored.</li> </ol>							
I-Words:	1							
Cycles:	2							

DEC	Decrement f							
Syntax:	{label:}	DEC.b	f	{,Wnd}	{,WREG}			
		DEC.bz						
		DEC{.w}						
		DEC.l						
Operands:		f ∈ [0 ... 64KB]; Wnd ∈ [W0 ... W14]						
Operation:		(f) - 1 → destination designated by D						
Status Affected:		C, N, OV, Z						
Encoding:	1101	011L	dddd	ffff	ffff	ffff	BD01	
Description:		Subtract one from the contents of the file register and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register.						
		The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'f' bits select the address of the file register. The 'd' bits select the Working register.						
I-Words:	1							
Cycles:	1							

**Example 1: DEC.B 0x200 ; Decrement (0x200) (Byte mode)**

Before Instruction				After Instruction			
Data 200	80FF			Data 200	80FE		
SR	0000			SR	0009	(N, C = 1)	

**Example 2: DEC RAM400, WREG ; Decrement RAM400 and store to WREG ; (Word mode)**

Before Instruction				After Instruction			
WREG	1211			WREG	0822		
RAM400	0823			RAM400	0823		
SR	0000			SR	0000		

DEC2	Decrement f by 2						
Syntax:	{label:}	DEC2.b	f	{,Wnd}	{,WREG}		
		DEC2.bz					
		DEC2{.w}					
		DEC2.l					
Operands:		$f \in [0 \dots 64KB]$ ; Wnd $\in [W0 \dots W14]$					
Operation:		(f) - 2 → destination designated by D					
Status Affected:		C, N, OV, Z					
Encoding:	1101	111L	dddd	ffff	ffff	ffff	BD01
Description:	<p>Subtract two from the contents of the file register and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 'd' bits select the Working register.</p>						
I-Words:	1						
Cycles:	1						

DISICTL								Disable Interrupts Control Literal							
Syntax:	{label:}	DISICTL	lit3		{ ,Wd}										
					{ ,[Wd]}										
					{ ,[Wd++]}										
					{ ,[Wd-]}										
					{ ,[++Wd]}										
					{ ,[--Wd]}										
Operands:	lit3 ∈ [0 ... 7]; Wd ∈ [W0 ... W15]														
Operation:	Disable interrupts at or below IPL threshold (IPLT[2:0]) defined by lit3 Setting lit3 = 0 will enable all interrupt levels up to and including those at SR.IPL[2:0]. If Wd declared, zero extended prior IPLT[2:0] → Wd[31:0]														
Status Affected:	None														
Encoding:	1101	110U	dddd	UUUU	UUUq	qqUU	Ukkk	0W10							
Description:	This instruction disables interrupts at or below the IPL threshold (IPLT[2:0]) defined by lit3, effective starting from the subsequent instruction. Traps cannot be inhibited by this instruction. The current DISICTL 3-bit IPLT is memory mapped (DISIIP1.IPLT[2:0]) and can be read by the user at any time. A write to DISIIP1 will have no effect. In addition, the IPLT established prior to DISICTL execution may be optionally written to destination Wd if a destination is declared (W = 1). The prior IPLT[2:0] is zero extended to 32-bits prior to being written. If a destination operand is not declared, nothing will be written (W = 0). This facilitates nesting of DISICTL and/or DISICTLW instructions.														
	The 'W' bit determines if a destination write is required. The 'K' bits are an unsigned literal that specifies the DISICTL IPL threshold. The 'd' bits select the destination register. The 'q' bits select the destination addressing mode. <b>Note:</b> SR.IPL[2:0] is not modified by this instruction.														
I-Words:	1														
Cycles:	1														

DISICTL Disable Interrupts Control Wns							
Syntax:	{label:}	DISICTL	Wns	{ ,Wd}			
				{ ,[Wd]}			
				{ ,[Wd++]}			
				{ ,[Wd-]}			
				{ ,[++Wd]}			
				{ ,[--Wd]}			
Operands:	Wns ∈ [W0 ... W14]; Wd ∈ [W0 ... W15]						
Operation:	Disable interrupts at or below IPL threshold (IPLT[2:0]) defined by Wns[2:0] Wns[2:0] = 0 will enable all interrupt levels If Wd declared, zero extended prior IPLT[2:0] → Wd[31:0]						
Status Affected:	None						
Encoding:	1101	110U	dddd	ssss	UUUq	qqUU	UUUU
Description:	This instruction disables interrupts at or below the IPL threshold (IPLT[2:0]) defined by the LS 3-bits of source Wns (remaining MSbs of Wns are ignored), effective starting from the subsequent instruction. Traps cannot be inhibited by this instruction. The current DISICTLW 3-bit IPLT is memory mapped (DISIIPPL.IPLT[2:0]) and can be read by the user at any time. A write to DISIIPPL will have no effect. In addition, the IPLT established prior to DISICTL execution may be optionally written to destination Wd if a destination is declared (W = 1). The prior IPLT[2:0] is zero extended to 32-bits prior to being written. If a destination operand is not declared, nothing will be written (W = 0). This facilitates nesting of DISICTL(W) instructions. This instruction can be used (optionally in conjunction with DISICTRLR) to control the impact of system interrupts. The 'W' bit determines if a destination write is required. The 's' bits select the source register. The 'd' bits select the destination register. The 'q' bits select the destination addressing mode. <b>Note:</b> SR.IPL[2:0] is not modified by this instruction.						
I-Words:	1						
Cycles:	1						

DIVF Signed Fractional 16/16 and 32/16 Divide																											
Syntax:	{label:}	DIVF.w	Wm	Wn																							
		DIVF.I																									
Operands:		$Wn \in [W0 \dots W14]; Wm \in [W0 \dots W13]$																									
Operation:		DIVF.w (16/16): $Wm[15:0] = \text{Dividend}, Wn[15:0] = \text{Divisor}$ : $Wm << 16; 16'b0 \rightarrow Wm[15:0]$ : $Wm / Wn \rightarrow 16'b0, Wm[15:0]$ : Remainder $\rightarrow 16'b0, W(m+1)[15:0]$																									
		DIVF.I (32/16): $Wm[31:0] = \text{Dividend}, Wn[15:0] = \text{Divisor}$ : $Wm / Wn \rightarrow 16'b0, Wm[15:0]$ : Remainder $\rightarrow 16'b0, W(m+1)[15:0]$																									
Status Affected:		C, N, OV, Z																									
Encoding:		1110	100L	wwww	ssss	UUUU	UUUU 1011																				
Description:		Iterative, signed fractional 32-bit (or 16-bit) by 16-bit divide to a 16-bit quotient and a 16-bit remainder, both of which are zero-extended prior to being written to Wm and W(m+1), respectively. The sign of the remainder will be the same as that of the dividend. The 16-bit by 16-bit divide scales the dividend to become a 32-bit fractional value prior to executing the divide.																									
This instruction must be executed six times to generate the correct quotient and remainder. This may only be achieved by executing a REPEAT with an iteration count of five (i.e. 5+1 iterations in all) and the DIVF.x instruction as its target. The REPEAT loop may be interrupted at any iteration boundary.																											
C is modified as per the divide algorithm. Z is set if the remainder is clear. Z is cleared otherwise. N is set if the remainder is negative. N is cleared otherwise. OV is set if the divide will result in an overflow. The quotient and remainder will be deterministic but meaningless.																											
The 'w' bits select the source (dividend) register. The 's' bits select the source (divisor) register.																											
<b>Notes:</b>																											
1. The divisor is tested for zero during the first iteration. An attempt to divide by zero will initiate an arithmetic error trap during the first cycle of DIVF.x execution.																											
2. Wn cannot share the same W-reg with Wm or W(m+1).																											
I-Words:	1																										
Cycles:	6																										
<b>Example:1</b>																											
<table border="1"> <thead> <tr> <th>Address</th><th>Label</th><th>Instruction</th><th>Comment</th></tr> </thead> <tbody> <tr> <td>0x804000</td><td>-</td><td>REPEAT #5</td><td>; Execute DIV 5 times.</td></tr> <tr> <td>0x804004</td><td>-</td><td>DIVF.I w4, w8</td><td>;Fractional Divide contents of W4 bycontents of W5</td></tr> </tbody> </table>								Address	Label	Instruction	Comment	0x804000	-	REPEAT #5	; Execute DIV 5 times.	0x804004	-	DIVF.I w4, w8	;Fractional Divide contents of W4 bycontents of W5								
Address	Label	Instruction	Comment																								
0x804000	-	REPEAT #5	; Execute DIV 5 times.																								
0x804004	-	DIVF.I w4, w8	;Fractional Divide contents of W4 bycontents of W5																								
<table border="1"> <thead> <tr> <th colspan="2">Before execution</th><th colspan="2">After execution</th></tr> </thead> <tbody> <tr> <td>w4</td><td>0x1001</td><td>w4</td><td>0x2000</td></tr> <tr> <td>w8</td><td>0x4000</td><td>w8</td><td>0x4000</td></tr> <tr> <td>w5</td><td>0x0000</td><td>w5</td><td>0x0001</td></tr> <tr> <td>SR</td><td>0x00000000</td><td>SR</td><td>0x00000000</td></tr> </tbody> </table>								Before execution		After execution		w4	0x1001	w4	0x2000	w8	0x4000	w8	0x4000	w5	0x0000	w5	0x0001	SR	0x00000000	SR	0x00000000
Before execution		After execution																									
w4	0x1001	w4	0x2000																								
w8	0x4000	w8	0x4000																								
w5	0x0000	w5	0x0001																								
SR	0x00000000	SR	0x00000000																								

DIVFL	Signed Fractional 32/32 Divide													
Syntax:	{label:}	DIVFL	Wm	Wn										
Operands:	$Wn \in [W0 \dots W14]; Wm \in [W0 \dots W13]$													
Operation:	32/32; Wm = Dividend, Wn = Divisor: $Wm / Wn \rightarrow Wm[31:0]$ ; Remainder $\rightarrow W(m+1)[31:0]$													
Status Affected:	C, N, OV, Z													
Encoding:	1110	101U	wwww	ssss	UUUU	UUUU	UUUU							
Description:	Iterative, signed fractional 32-bit by 32-bit divide to a 32-bit quotient and a 32-bit remainder. The sign of the remainder will be the same as that of the dividend.													
This instruction must be executed ten times to generate the correct quotient and remainder. This may only be achieved by executing a REPEAT with an iteration count of nine (i.e. 9+1 iterations in all) and the DIVFL instruction as its target. The REPEAT loop may be interrupted at any iteration boundary.														
C is modified as per the divide algorithm. Z is set if the remainder is clear. Z is cleared otherwise. N is set if the remainder is negative. N is cleared otherwise. OV is set if the divide will result in an overflow. The quotient and remainder will be deterministic but meaningless.														
The 'w' bits select the source (dividend) register. The 's' bits select the source (divisor) register.														
<b>Notes:</b>														
1. The divisor is tested for zero during the first iteration. An attempt to divide by zero will initiate an arithmetic error trap during the first cycle of DIVFL execution.														
2. Wn cannot share the same W-reg with Wm or W(m+1).														
I-Words:	1													
Cycles:	10													

DIVS	Signed Integer 16/16 and 32/16 Divide						
Syntax:	{label:}	DIVS.w	Wm,	Wn			
		DIVS.I					
Operands:		Wn ∈ [W0 ... W14], Wm ∈ [W0 ... W13]					
Operation:		DIVS.w (16/16): Wm[15:0] = Dividend, Wn[15:0] = Divisor: {16{Wm[15]}}, Wm[15:0] → Wm[31:0]; Wm / Wn → 16'b0, Wm[15:0]; Remainder → 16'b0, W(m+1)[15:0]					
		DIVS.I (32/16): Wm[31:0] = Dividend, Wn[15:0] = Divisor: Wm / Wn → 16'b0, Wm[15:0]; Remainder → 16'b0, W(m+1)[15:0]					
Status Affected:	C, N, OV, Z						
Encoding:	1110      100L      wwww      ssss      UUUU      UUUU      UUUU      0011						
Description:	Iterative, signed integer 32-bit (or 16-bit) by 16-bit divide to a 16-bit quotient and a 16-bit remainder, both of which are zero-extended to 32-bits prior to being written to Wm and W(m+1), respectively. The sign of the remainder will be the same as that of the dividend. The 16-bit by 16-bit divide sign-extends the dividend prior to executing the divide. This instruction must be executed six times to generate the correct quotient and remainder. This may only be achieved by executing a REPEAT with an iteration count of five (i.e. 5+1 iterations in all) and the DIVS instruction as its target. The REPEAT loop may be interrupted at any iteration boundary.						
	C is modified as per the divide algorithm.						
	Z is set if the remainder is clear. Z is cleared otherwise. N is set if the remainder is negative. N is cleared otherwise. OV is set if the divide will result in an overflow. OV is cleared otherwise. The quotient and remainder will be deterministic but meaningless.						
	The 'L' bit selects 32-bit or 16-bit dividend size. The 'w' bits select the source (dividend) register. The 's' bits select the source (divisor) register.						
	<b>Notes:</b>						
	1. The divisor is tested for zero during the first iteration. An attempt to divide by zero will initiate an arithmetic error trap during the first cycle of DIVS.x 2. Wn cannot share the same W-reg with Wm or W(m+1).						
I-Words:	1						
Cycles:	6						

Example:1	Address	Label	Instruction	Comment
	0x804000	-	REPEAT #5	; Execute DIV 5 times.
	0x804004	-	DIVS.I w4, w8	;Divide contents of W4 by contents of W5
Before execution			After execution	
w4	0x3000		w4	0x013B
w8	0x27		w8	0x27
w5	0x0000		w5	0x0003
SR	0x00000000		SR	0x00000000

DIVSL	Signed Integer 32/32 Divide						
Syntax:	{label:}	DIVSL	Wm,	Wn			
Operands:	$Wn \in [W0 \dots W14], Wm \in [W0 \dots W13]$						
Operation:	32/32; Wm = Dividend, Wn = Divisor: $Wm / Wn \rightarrow Wm$ ; Remainder $\rightarrow W(m+1)$						
Status Affected:	C, N, OV, Z						
Encoding:	1110	101U	wwww	ssss	UUUU	UUUU	UUUU 0011
Description:	<p>Iterative, signed integer 32-bit by 32-bit divide to a 32-bit quotient and a 32-bit remainder. The sign of the remainder will be the same as that of the dividend.</p> <p>This instruction must be executed ten times to generate the correct quotient and remainder. This may only be achieved by executing a REPEAT with an iteration count of nine (i.e. 9+1 iterations in all) and the DIVSL instruction as its target. The REPEAT loop may be interrupted at any iteration boundary.</p>						
	<p>C is modified as per the divide algorithm.</p> <p>Z is set if the remainder is clear. Z is cleared otherwise.</p> <p>N is set if the remainder is negative. N is cleared otherwise.</p> <p>OV is set if the divide will result in an overflow. OV is cleared otherwise. The quotient and remainder will be deterministic but meaningless.</p>						
	<p>The 'w' bits select the source (dividend) register.</p> <p>The 's' bits select the source (divisor) register.</p>						
	<b>Notes:</b>						
	<ol style="list-style-type: none"> <li>1. The divisor is tested for zero during the first iteration. An attempt to divide by zero will initiate an arithmetic error trap during the first cycle of DIVSL.</li> <li>2. Wn cannot share the same W-reg with Wm or W(m+1).</li> </ol>						
I-Words:	1						
Cycles:	10						

DIVU	Unsigned Integer 16/16 and 32/16 Divide						
Syntax:	{label:}	DIVU.w	Wm,	Wn			
		DIVU.I					
Operands:		$Wn \in [W0 \dots W14], Wm \in [W0 \dots W13]$					
Operation:		DIVU.w (16/16); $Wm[15:0] = \text{Dividend}, Wn[15:0] = \text{Divisor}$ : $16'b0 \rightarrow Wm[31:16]$ ; $Wm / Wn \rightarrow 16'b0, Wm[15:0]; \text{Remainder} \rightarrow 16'b0, W(m+1)[15:0]$					
		DIVU.I (32/16); $Wm[31:0] = \text{Dividend}, Wn[15:0] = \text{Divisor}$ : $Wm / Wn \rightarrow 16'b0, Wm[15:0]; \text{Remainder} \rightarrow 16'b0, W(m+1)[15:0]$					
Status Affected:		C, N, OV, Z					
Encoding:	1110	100L	wwww	ssss	UUUU	UUUU	0111
Description:	Iterative, unsigned integer 32-bit (or 16-bit) by 16-bit divide to a 16-bit quotient and a 16-bit remainder, both of which are zero-extended to 32-bits prior to being written to Wm and W(m+1), respectively. The 16-bit by 16-bit divide also zero extends the dividend prior to executing the divide. This instruction must be executed six times to generate the correct quotient and remainder. This may only be achieved by executing a REPEAT with an iteration count of five (i.e. 5+1 iterations in all) and the DIVU instruction as its target. The REPEAT loop may be interrupted at any iteration boundary.						
	C is modified as per the divide algorithm. Z is set if the remainder is clear. Z is cleared otherwise. N is always cleared. OV (DIVU.w) is always cleared because an overflow is not possible. OV (DIVU.I) is set if the divide will result in an overflow. OV is cleared otherwise.						
	The 'L' bit selects 32-bit or 16-bit dividend size. The 'w' bits select the source (dividend) register. The 's' bits select the source (divisor) register.						
<b>Notes:</b>	<ol style="list-style-type: none"> <li>The divisor is tested for zero during the first iteration. An attempt to divide by zero will initiate an arithmetic error trap during the first cycle of DIVU.x</li> <li>Wn cannot share the same W-reg with Wm or W(m+1).</li> </ol>						
I-Words:	1						
Cycles:	6						

DIVUL	Unsigned Integer 32/32 Divide						
Syntax:	{label:}	DIVUL	Wm,	Wn			
Operands:	$Wn \in [W0 \dots W14], Wm \in [W0 \dots W13]$						
Operation:	32/32; Wm = Dividend, Wn = Divisor: $Wm / Wn \rightarrow Wm$ ; Remainder $\rightarrow W(m+1)$						
Status Affected:	C, N, OV, Z						
Encoding:	1110	101U	wwww	ssss	UUUU	UUUU	0111
Description:	Iterative, unsigned integer 32-bit by 32-bit divide to a 32-bit quotient and a 32-bit remainder. This instruction must be executed ten times to generate the correct quotient and remainder. This may only be achieved by executing a REPEAT with an iteration count of nine (i.e. 9+1 iterations in all) and the DIVUL instruction as its target. The REPEAT loop may be interrupted at any iteration boundary.						
	C is modified as per the divide algorithm. Z is set if the remainder is clear. Z is cleared otherwise. N is always cleared. OV is always cleared because an overflow is not possible.						
	The 'w' bits select the source (dividend) register. The 's' bits select the source (divisor) register.						
	<b>Notes:</b>						
	<ol style="list-style-type: none"> <li>The divisor is tested for zero during the first iteration. An attempt to divide by zero will initiate an arithmetic error trap during the first cycle of DIVUL.</li> <li>Wn cannot share the same W-reg with Wm or W(m+1).</li> </ol>						
I-Words:	1						
Cycles:	10						

DTB Decrement, Test and Branch							
Syntax:	{label:}	DTB	Wn,	Label			
Operands:		Wn ∈ [W0 ... W14]; Label is resolved by the linker to a signed word offset (slit16)					
Operation:		Wn = Wn - 1; If (Wn != 0 [see text]) then { If slit16 = 1 then skip next (16-bit) instruction Else if (slit16 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit16 → PC } Else No branch					
Status Affected:		None					
Encoding:	1011	100U	ssss	nnnn	nnnn	nnnn	UU01
Description:		Decrement Wn[31:0] and write result back to Wn (see note 2). Test Wn[31:0] after decrement and branch to target address if Wn[31:0] != 0. Do not branch when Wn[31:0] = 0. When utilized as a code block loop counter where DTB is located at the end of the loop, DTB will iterate the loop (i.e., branch) Wn times and will exit the loop with Wn = 0x0000_0000 (see note 2). If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 64KB. If the 2's complement byte offset value '2*slit16' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2*slit16' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit16' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit16. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.					
		The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).					
		<b>Notes:</b>					
		1. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met. 2. Execution of DTB when Wn = 0 will result in a loop count of $2^{31}$ .					
I-Words:		1					
Cycles:		1 (2 or 3)					
		<b>Note:</b> The taken branch target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.					

Example:1	Address	Label	Instruction	Comment
	0x804000		MOV.I #8, w6	;Update loop counter
	0x804004	_START_BUTTERFLY	MOV.I #1, W4	;Start of loop
		:		
	0x804008	-	SL.I w4, w6, w4	
	0x80400C	-	...	
	...	-	...	

## Instruction Descriptions (C to DTB) (continued)

Example:1	Address	Label	Instruction	Comment
	0x804040	-	DTB W6, _START_BUTTERFLY	;Branch to 0x804004 if w6 != 0

## 4.6. Instruction Descriptions (E to MULUU)

ED Subtract and Square to Accumulator (Partial Euclidean Distance)					
Syntax: {label}	ED{.w}	Wx,	Wy,	A	{,AWB}
	ED.I	[Wx], [Wx]+=kx [Wx]-=kx [Wx+=kx] [Wx-=kx] [Wx+W12]	[Wy], , [Wy]+=ky, , [Wy]-=ky, , [Wy+=ky], , [Wy-=ky], , [Wy+W12],	B	
Operands:		$Wx \in \{W0 \dots W14\}$ ; $Wy \in \{W0 \dots W14\}$			
		Word mode: $kx \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ; $ky \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ;			
		Long Word mode: $kx \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ; $ky \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ;			
		$AWB \in \{W0, W1, W2, W3, W13, [W13++], [W15++]^4\}$			
Operation:		$((Wx) - (Wy))^2 \rightarrow ACC(A \text{ or } B);$ (ACC(B or A)) rounded $\rightarrow AWB$			
		When Indirect Pre/Post Modified Addressing: $(Wx)+kx \rightarrow Wx$ or $(Wx)-kx \rightarrow Wx$ ; $(Wy)+ky \rightarrow Wy$ or $(Wy)-ky \rightarrow Wy$ ;			
Status Affected:		OA,SA or OB,SB			
Encoding:	1101	10AL	wwww	ssss	IIIi iJJJ jjaa a011

**Instruction Descriptions (E to MULUU) (continued)****ED                   Subtract and Square to Accumulator  
(Partial Euclidean Distance)**

Description: Instruction to compute  $(A-B)^2$  functions. Signed or unsigned (defined by CORCON.US) subtract then square of data concurrently read from Wx and Wy or fetched from X and Y address space (see note 3). The result is sign-extended or zero-extended to 72-bits then written to the specified accumulator. Fractional results are also scaled prior to the accumulator update to align the operand and accumulator (msw) fractional points.

When indirect addressing is selected for either or both X and Y address space, the Wx and Wy registers provide the corresponding indirect addresses. The address modifier values are kx and ky, respectively, and represent the number of data bytes by which to modify the Effective Address.

The optional AWB specifies the direct or indirect (see note 4) store of the (32-bit) rounded fractional contents of the accumulator not targeted by the ED operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size (see note 5). AWB is not intended for use when the DSP engine is operating in Integer mode.

Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly.

The 'L' bit selects word or long word operation.

The 'A' bit selects the accumulator for the result.

The 'I' bits select the Operation.X-Space Addressing mode.

The 'i' bits select the kx modification value.

The 'J' bits select the Operation.Y-Space Addressing mode.

The 'j' bits select the ky modification value.

The 's' bits select the Wx register

The 'w' bits select the Wy register.

The 'a' bits select the accumulator write-back destination and addressing mode.

**Notes:**

1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.
2. Operands are always regarded as signed or unsigned based on the state of CORCON.US.
3. Use of the same W-reg for both indirect source (X or Y) and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.
4. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.

I-Words: 1

Cycles: 2

EDAC                   Subtract, Square and Accumulate (Euclidean Distance)				
Syntax: {label}	EDAC{.w} Wx,	Wy,	A	{,AWB}
EDAC.I	[Wx],	[Wy],	B	
	[Wx]+=kx	, [Wy]+=ky,		
	[Wx]=-kx	, [Wy]=-ky,		
	[Wx+=kx]	, [Wy+=ky],		
	[Wx-=kx]	, [Wy-=ky],		
	[Wx+W12]	, [Wy+W12],		

Operands:  $Wx \in \{W0 \dots W14\}$ ;  $Wy \in \{W0 \dots W14\}$

Word mode:  $kx \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ;

$ky \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ;

Long Word mode:  $kx \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ;  
 $ky \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ;

AWB  $\in \{W0, W1, W2, W3, W13, [W13++], [W15++]\}^4$

**Instruction Descriptions (E to MULUU) (continued)**

<b>EDAC                      Subtract, Square and Accumulate (Euclidean Distance)</b>							
Operation:	$ACC(A \text{ or } B) + ((Wx) - (Wy))^2 \rightarrow ACC(A \text{ or } B);$ (ACC(B or A)) rounded $\rightarrow AWB$						
When Indirect Pre/Post Modification Addressing:	$(Wx)+kx \rightarrow Wx$ or $(Wx)-kx \rightarrow Wx;$ $(Wy)+ky \rightarrow Wy$ or $(Wy)-ky \rightarrow Wy;$						
Status	OA,SA or OB,SB						
Affected:							
Encoding:	1101	10AL	wwww	ssss	IIII	iJJJ	jjaa a111
Description:	Instruction to compute $(A-B)^2$ functions. Signed or unsigned (defined by CORCON.US) subtract then square of data read from Wx and Wy or concurrently fetched from the X and Y address space. The result is sign-extended or zero-extended to 72-bits and then added to the specified accumulator. Fractional or integer operation (defined by CORCON.IF) will determine if the result is scaled or not prior to the accumulator update. Fractional operation will scale the result to align the operand and accumulator (msw) fractional points (see note 3). Integer operation will align the LSb of the result with the LSb of the accumulator. When indirect addressing is selected for either or both X and Y address space, the Wx and Wy registers provide the corresponding indirect addresses. The address modifier values are kx and ky, respectively, and represent the number of data words by which to modify the Effective Address. The optional AWB specifies the direct or indirect (see note 4) store of the (32-bit) rounded fractional contents of the accumulator not targeted by the EDAC operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size (see note 5). AWB is not intended for use when the DSP engine is operating in Integer mode.						
	Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly. The 'L' bit selects word or long word operation. The 'A' bit selects the accumulator for the result. The 'I' bits select the Operation.X-Space Addressing mode. The 'i' bits select the kx modification value. The 'j' bits select the ky modification value. The 's' bits select the Wx register. The 'w' bits select the Wy register. The 'a' bits select the accumulator write-back destination and addressing mode.						
<b>Notes:</b>							
1.	Operates in Fractional or Integer Data mode as defined by CORCON.IF.						
2.	Operands are always regarded as signed or unsigned based on the state of CORCON.US.						
3.	The LS portion of ACCx is unaffected when operating in Fractional mode with word sized data. Lower significance data that may be present from prior (32-bit data) operations is therefore preserved. Users not requiring this should clear ACCx during initialization.						
4.	Use of the same W-reg for both indirect source (X or Y) and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.						
5.	Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.						
I-Words:	1						
Cycles:	2						

<b>EXCH                      Exchange Ws and Wd</b>							
Syntax:	{label:} EXCH Wns, Wnd						
Operands:	Wns $\in$ [W0 ... W15]; Wnd $\in$ [W0 ... W15]						
Operation:	(Wns) $\leftrightarrow$ (Wnd)						
Status Affected:	None						
Encoding:	1000	001U	dddd	ssss	UUUU	UUUU	0100

**Instruction Descriptions (E to MULUU) (continued)****EXCH Exchange Ws and Wd**

Description: This instruction exchanges the contents of two Working registers and is a two cycle operation. In cycle one, Wnd is read and moved to Wns. In cycle two, Wns is read and moved to Wnd. The 's' bits select the Wns register.

The 'd' bits select the Wnd register.

**Notes:**

1. Opcode shared with MOV (32-bit variant) using sub-opcode.
2. Long word operation is assumed.
3. Although operand order has no effect on final outcome, it could influence the detection of hazards because operands are read in different cycles.

I-Words: 1

Cycles: 2

**FBCL Find First Bit Change from Left**

Syntax: {label:} FBCL{.w} Ws, Wnd  
                   FBCL.I [Ws],  
                   [Ws++],  
                   [Ws-],  
                   [+Ws],  
                   [-Ws],

Operands: Ws ∈ [W0 ... W15]; Wnd ∈ [W0 ... W14]

Operation: See description

Status Affected: C

Encoding: 1110 011L dddd ssss pppU UUUU UUUU 1011

Description: Finds the first occurrence of a one (for a positive signed value) or zero (for a negative signed value) starting from the next MSb after the Sign bit working towards the Lsb of the word operand. The bit number is sign-extended to 32-bits and is written to the destination register.

The next MSb after the Sign bit is assigned number zero. For a word operation, the Lsb number is assigned -14 and a result of -15 (C=1) indicates that the bit was not found. For a long word operation, the Lsb number is assigned -30 and a result of -31 (C=1) indicates that the bit was not found.

C is cleared for all non-zero results.

The 'L' bit selects word or long word operation.

The 's' bits select the source register.

The 'd' bits select the destination register.

The 'p' bits select the source addressing mode.

See and for modifier addressing information.

**Note:** This instruction only operates in Word and Long Word mode.

I-Words: 1

Cycles: 1

**Example 1:** FBCL W1, W9 ; Find 1st bit change from left in W1; and store result to W9

## Before Instruction

W1	55FF
----	------

W9	FFFF
----	------

SR	0000
----	------

## After Instruction

W1	55FF
----	------

W9	0000
----	------

SR	0000
----	------

<b>Example 2:</b>	<b>FBCL W1, W9</b>	<b>; Find 1st bit change from left in W1; and store result to W9</b>
	Before Instruction	After Instruction
	W1 FFFF	W1 FFFF
	W9 BBBB	W9 FFF1
	SR 0000	SR 0001 (C = 1)

<b>Example 3:</b>	<b>FBCL [W1++], W9</b>	<b>; Find 1st bit change from left in [W1] ; and store result to W9; Post-increment W1</b>
	Before Instruction	After Instruction
	W1 2000	W1 2002
	W9 BBBB	W9 FFF9
	Data 2000 FF0A	Data 2000 FF0A
	SR 0000	SR 0000

<b>FBRA EQ</b>		<b>Coprocessor Branch 0 (CP0 FPU Branch if Equal)</b>							
Syntax:	{label:}	FBRA	EQ,	Expr					
Operands:		Label is resolved by the linker to a signed word offset (slit20)							
Operation:		Condition = FSR.EQ; If (condition) then { If slit20 = 1 then skip next (16-bit) instruction Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction Else (PC+4) + 2*slit20 → PC } Else no branch							
Status Affected:		None							
Encoding:	1011	0000	nnnn	nnnn	nnnn	nnnn	nnnn	zz10	

**Instruction Descriptions (E to MULUU) (continued)****FBRA EQ****Coprocessor Branch 0  
(CP0 FPU Branch if Equal)****Description:**

If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20.

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

**I-Words:**

1

**Cycles:**

1 (2 or 3)

**FBRA GE****Coprocessor Branch 6  
(CP0 FPU Branch if Greater Than or Equal)****Syntax:**

{label:} FBRA GE, Expr

**Operands:**

Expr is resolved by the linker to a signed word offset (slit20)

**Operation:**

```
Condition = (FSR.GT || FSR.EQ);
If (condition) then {
    If slit20 = 1 then skip next (16-bit) instruction
    Else If (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction
    Else (PC+4) + 2*slit20 → PC
}
Else no branch
```

**Status Affected:**

None

**Encoding:**

1011 0110 nnnn nnnn nnnn nnnn nnnn zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA GE****Coprocessor Branch 6  
(CP0 FPU Branch if Greater Than or Equal)**

**Description:** If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

**FBRA GT****Coprocessor Branch 4  
(CP0 FPU Branch if Greater Than)**

Syntax: {label:} FBRA GT, Expr

Operands: Expr is resolved by the linker to a signed word offset (slit20)

Operation:

```
Condition = FSR.GT;
If (condition) then {
    If slit20 = 1 then skip next (16-bit) instruction
    Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction
    Else (PC+4) + 2*slit20 → PC
}
Else no branch
```

Status Affected:

None

Encoding:

1011 0100 nnnn nnnn nnnn nnnn nnnn zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA GT****Coprocessor Branch 4  
(CP0 FPU Branch if Greater Than)**

**Description:** If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20.

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

**FBRA LE****Coprocessor Branch 10  
(CP0 FPU Branch if Less Than or Equal)**

Syntax: {label:} FBRA LE, Expr

Operands: Expr is resolved by the linker to a signed word offset (slit20)

Operation: Condition = (FSR.LT || FSR.EQ);  
 If (condition) then {  
   If slit20 = 1 then skip next (16-bit) instruction  
   Else if (slit20 = 2 && next\_op[31] = 1) then skip next (32-bit) instruction  
   Else (PC+4) + 2\*slit20 → PC  
   }  
 Else no branch

Status Affected: None

Encoding: 1011      1010      nnnn      nnnn      nnnn      nnnn      nnnn      zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA LE****Coprocessor Branch 10  
(CP0 FPU Branch if Less Than or Equal)****Description:**

If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20.

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

**I-Words:**

1

**Cycles:**

1 (2 or 3)

**FBRA LT****Coprocessor Branch 8  
(CP0 FPU Branch if Less Than)****Syntax:**

{label:} FBRA LT, Expr

**Operands:**

Expr is resolved by the linker to a signed word offset (slit20)

**Operation:**

```
Condition = FSR.LT;
If (condition) then {
    If slit20 = 1 then skip next (16-bit) instruction
    Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction
    Else (PC+4) + 2*slit20 → PC
}
Else no branch
```

**Status Affected:**

None

**Encoding:**

1011 1000 nnnn nnnn nnnn nnnn nnnn zz10

## Instruction Descriptions (E to MULUU) (continued)

**FBRA LT****Coprocessor Branch 8  
(CP0 FPU Branch if Less Than)**

## Description:

If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20.

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

## I-Words:

1

## Cycles:

1 (2 or 3)

**FBRA NE****Coprocessor Branch 2  
(CP0 FPU Branch if Not Equal)**

## Syntax:

{label:} FBRA NE, Expr

## Operands:

Expr is resolved by the linker to a signed word offset (slit20)

## Operation:

```
Condition = (FSR.GT || FSR.LT);
If (condition) then {
    If slit20 = 1 then skip next (16-bit) instruction
    Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction
    Else (PC+4) + 2*slit20 → PC
}
Else no branch
```

## Status Affected:

None

## Encoding:

1011 0010 nnnn nnnn nnnn nnnn nnnn zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA NE****Coprocessor Branch 2  
(CP0 FPU Branch if Not Equal)**

Description:

If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20.

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

**FBRA OR****Coprocessor Branch 12  
(CP0 FPU Branch if Ordered)**

Syntax: {label:} FBRA OR, Expr

Operands: Expr is resolved by the linker to a signed word offset (slit20)

Operation: Condition = (FSR.GT || FSR.LT || FSR.EQ)

If (condition) then {

If slit20 = 1 then skip next (16-bit) instruction

Else if (slit20 = 2 &amp;&amp; next\_op[31] = 1) then skip next (32-bit) instruction

Else (PC+4) + 2\*slit20 → PC

}

Else no branch

Status Affected: None

Encoding: 1011 1100 nnnn nnnn nnnn nnnn nnnn zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA OR****Coprocessor Branch 12  
(CP0 FPU Branch if Ordered)****Description:**

If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20.

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

**I-Words:**

1

**Cycles:**

1 (2 or 3)

**FBRA UGE****Coprocessor Branch 9  
(CP0 FPU Branch if Unordered or Greater Than or Equal)****Syntax:**

{label:} FBRA UGE Expr

**Operands:**

Expr is resolved by the linker to a signed word offset (slit20)

**Operation:**

Condition = (FSR.GT || FSR.EQ || FSR.UN);

If (condition) then {

If slit20 = 1 then skip next (16-bit) instruction

Else if (slit20 = 2 &amp;&amp; next\_op[31] = 1) then skip next (32-bit) instruction

Else (PC+4) + 2\*slit20 → PC

}

Else no branch

**Status Affected:**

None

**Encoding:**

1011 1001 nnnn nnnn nnnn nnnn nnnn zz10

## Instruction Descriptions (E to MULUU) (continued)

FBRA UGE	Coprocessor Branch 9 (CP0 FPU Branch if Unordered or Greater Than or Equal)							
Description:	<p>If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.</p> <p>Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.</p> <p>The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).</p> <p>The 'z' bits select the target coprocessor.</p>							
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> <li>2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.</li> <li>3. Branch conditions are evaluated within the coprocessor.</li> </ol>							
I-Words:	1							
Cycles:	1 (2 or 3)							
FBRA UGT	Coprocessor Branch 11 (CP0 FPU Branch if Unordered or Greater Than)							
Syntax:	{label:}	FBRA	UGT,	Expr				
Operands:	Expr is resolved by the linker to a signed word offset (slit20)							
Operation:	<p>Condition = (FSR.GT    FSR.UN);</p> <p>If (condition) then {</p> <p>If slit20 = 1 then skip next (16-bit) instruction</p> <p>Else if (slit20 = 2 &amp;&amp; next_op[31] = 1) then skip next (32-bit) instruction</p> <p>Else (PC+4) + 2*slit20 → PC</p> <p>}</p> <p>Else no branch</p>							
Status Affected:	None							
Encoding:	1011	1011	nnnn	nnnn	nnnn	nnnn	nnnn	zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA UGT****Coprocessor Branch 11  
(CP0 FPU Branch if Unordered or Greater Than)**

**Description:** If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

**FBRA UEQ****Coprocessor Branch 3  
(CP0 FPU Branch if Unordered or Equal)**

Syntax: {label:} FBRA UEQ Expr

Operands: Expr is resolved by the linker to a signed word offset (slit20)

Operation: Condition = (FSR.EQ || FSR.UN);  
If (condition) then {  
  If slit20 = 1 then skip next (16-bit) instruction  
  Else if (slit20 = 2 && next\_op[31] = 1) then skip next (32-bit) instruction  
  Else (PC+4) + 2\*slit20 → PC  
}  
Else no branch

Status Affected: None

Encoding: 1011 0011 nnnn nnnn nnnn nnnn nnnn zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA UEQ****Coprocessor Branch 3  
(CP0 FPU Branch if Unordered or Equal)**

**Description:** If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20.

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

**FBRA ULE****Coprocessor Branch 5  
(CP0 FPU Branch if Unordered or Less Than or Equal)**

Syntax: {label:} FBRA ULE, Expr

Operands: Expr is resolved by the linker to a signed word offset (slit20)

Operation: Condition = (FSR.LT || FSR.EQ || FSR.UN);  
 If (condition) then {  
   If slit20 = 1 then skip next (16-bit) instruction  
   Else if (slit20 = 2 && next\_op[31] = 1) then skip next (32-bit) instruction  
   Else (PC+4) + 2\*slit20 → PC  
   }  
   Else no branch

Status Affected: None

Encoding: 1011      0101      nnnn      nnnn      nnnn      nnnn      nnnn      zz10

**Instruction Descriptions (E to MULUU) (continued)**

FBRA ULE	Coprocessor Branch 5 (CP0 FPU Branch if Unordered or Less Than or Equal)
Description:	<p>If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.</p> <p>If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.</p> <p>Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.</p> <p>The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).</p> <p>The 'z' bits select the target coprocessor.</p>
	<b>Notes:</b>
	<ol style="list-style-type: none"> <li>1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> <li>2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.</li> <li>3. Branch conditions are evaluated within the coprocessor.</li> </ol>
I-Words:	1
Cycles:	1 (2 or 3)

FBRA ULT	Coprocessor Branch 7 (CP0 FPU Branch if Unordered or Less Than)
Syntax:	{label:} FBRA ULT, Expr
Operands:	Expr is resolved by the linker to a signed word offset (slit20)
Operation:	<p>Condition = (FSR.LT    FSR.UN);</p> <p>If (condition) then {</p> <p>If slit20 = 1 then skip next (16-bit) instruction</p> <p>Else if (slit20 = 2 &amp;&amp; next_op[31] = 1) then skip next (32-bit) instruction</p> <p>Else (PC+4) + 2*slit20 → PC</p> <p>}</p> <p>Else no branch</p>
Status Affected:	None
Encoding:	1011 0111 nnnn nnnn nnnn nnnn nnnn zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA ULT****Coprocessor Branch 7  
(CP0 FPU Branch if Unordered or Less Than)**

Description:

If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be  $(PC+4) + 2*slit20$ . Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from  $(PC+4)$ . The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

for details.

I-Words:

1

Cycles:

1 (2 or 3)

**FBRA UN****Coprocessor Branch 13  
(CP0 FPU Branch if Unordered)**

Syntax:      {label:}      FBRA      UN,      Expr

Operands:      Expr is resolved by the linker to a signed word offset (slit20)

Operation:      Condition = FSR.UN;

If (condition) then {

If slit20 = 1 then skip next (16-bit) instruction

Else if (slit20 = 2 &amp;&amp; next\_op[31] = 1) then skip next (32-bit) instruction

Else  $(PC+4) + 2*slit20 \rightarrow PC$ 

}

Else no branch

Status Affected:      None

Encoding:      1011      1101      nnnn      nnnn      nnnn      nnnn      zzzz

**Instruction Descriptions (E to MULUU) (continued)****FBRA UN****Coprocessor Branch 13  
(CP0 FPU Branch if Unordered)**

**Description:** If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined. If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2\*slit20. Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path. The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4). The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

**FBRA UNE****Coprocessor Branch 1  
(CP0 FPU Branch if Unordered or Not Equal)**

Syntax: {label:} FBRA UNE Expr

Operands: Label is resolved by the linker to a signed word offset (slit20)

Operation: Condition = (FSR.GT || FSR.LT || FSR.UN);  
 If (condition) then {  
   If slit20 = 1 then skip next (16-bit) instruction  
   Else if (slit20 = 2 && next\_op[31] = 1) then skip next (32-bit) instruction  
   Else (PC+4) + 2\*slit20 → PC  
   }  
 Else no branch

Status Affected: None

Encoding: 1011 0001 nnnn nnnn nnnn nnnn nnnn zz10

**Instruction Descriptions (E to MULUU) (continued)****FBRA UNE**

**Coprocessor Branch 1  
(CP0 FPU Branch if Unordered or Not Equal)**

Description:

If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB. If the 2's complement byte offset value '2\*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If the 2's complement byte offset value '2\*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.

If requirements for a conditional skip are not met, the 2's complement byte offset value '2\*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be  $(PC+4) + 2*slit20$ .

Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken. The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.

The 'n' bits are a signed literal that specifies the number of PS words offset from  $(PC+4)$ .

The 'z' bits select the target coprocessor.

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

FF1L Find First One from Left							
Syntax:	{label:}	FF1L{.w}	Ws,	Wnd			
		FF1L.I	[Ws],				
			[Ws++],				
			[Ws-],				
			[++Ws],				
			[--Ws],				
Operands:	Ws ∈ [W0 ... W15]; Wnd ∈ [W0 ... W14]						
Operation:	See description						
Status Affected:	C						
Encoding:	1110 011L dddd ssss pppp UUUU UUUU 0011						
Description:	Finds the first occurrence of a one starting from the MSb working towards the LSb of the word operand. The bit number result is zero-extended to 32-bits and is written to the destination register. The MSb is assigned number one. For a word operation, the LSb number is assigned 16, and a result of zero (C=1) indicates that the bit was not found. For a long word operation, the LSb number is assigned 32 and a result of zero (C=1) indicates that the bit was not found.						
	C is cleared for all non-zero results.						
	The 'L' bit selects word or long word operation.						
	The 's' bits select the source register.						
	The 'd' bits select the destination register.						
	The 'p' bits select the source addressing mode.						
	<b>Note:</b> This instruction only operates in word and long word mode.						
I-Words:	1						
Cycles:	1						

**Example 1:** FF1L W2, W5 ; Find the 1st one from the left in W2 ; and store result to W5

Before Instruction		After Instruction	
W2	000A	W2	000A
W5	BBBB	W5	000D
SR	0000	SR	0000

**Example 2:** FF1L [W2++], W5 ; Find the 1st one from the left in [W2] ; and store the result to W5 ; Post-increment W2

Before Instruction		After Instruction	
W2	2000	W2	2002
W5	BBBB	W5	0000
Data 2000	0000	Data 2000	0000
SR	0000	SR	0001 (C = 1)

FF1R	Find First One from Right							
Syntax:	{label:}	FF1R{.w}	Ws,	Wnd				
		FF1R.I	[Ws],					
			[Ws++],					
			[Ws-],					
			[++Ws],					
			[--Ws],					
Operands:	Ws	∈ [W0 ... W15]; Wnd	∈ [W0 ... W14]					
Operation:	See description							
Status Affected:	C							
Encoding:	1110	011L	dddd	ssss	pppU	UUUU	UUUU	0111
Description:	<p>Finds the first occurrence of a one starting from the LSb working towards the MSb of the word operand. The bit number result is zero-extended to 32-bits and is written to the destination register. The LSb is assigned number one. For a word operation, the MSb number is assigned 16, and a result of zero (C=1) indicates that the bit was not found. For a long word operation, the MSb number is assigned 32 and a result of zero (C=1) indicates that the bit was not found.</p> <p>C is cleared for all non-zero results.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 's' bits select the source register.</p> <p>The 'd' bits select the destination register.</p> <p>The 'p' bits select the source addressing mode.</p> <p><b>Note:</b> This instruction only operates in word and long word mode.</p>							
I-Words:	1							
Cycles:	1							

**Example 1:** FF1R W1, W9 ; Find the 1st one from the right in W1 ; and store the result to W9

Before Instruction		After Instruction	
W1	000A	W1	000A
W9	BBBB	W9	0002
SR	0000	SR	0000

**Example 2:** FF1R [W1++], W9 ; Find the 1st one from the right in [W1] ; and store the result to W9 ; Post-increment W1

Before Instruction		After Instruction	
W1	2000	W1	2002
W9	BBBB	W9	0010
Data 2000	8000	Data 2000	8000
SR	0000	SR	0000

<b>FLIM</b>		<b>Force (Signed) Data Range Limit</b>								
Syntax:	{label:}	FLIM{.w}	Wb,	Ws,						
		FLIM.I		[Ws],						
				[Ws++],						
				[Ws--],						
				[++Ws],						
				[~-Ws],						
Operands:	Ws ∈ [W0 ... W14]; Wb ∈ [W0 ... W13];									
Operation:	<pre> IF (Ws) &gt; (Wb) THEN { (Wb) → (Ws); 0→ Z; 0→ N; 0→ OV } ELSE IF (Ws) &lt; (Wb+1) THEN { (Wb+1) → Ws; 0→ Z; 1→ N; 0→ OV } ELSE { (Ws) → Ws<sup>4</sup>; 1→ Z; 0→ N; 0→ OV } </pre>									
Status Affected:	N, Z, OV									
Encoding:	1110	010L	UUUU	ssss	pppU	UUww	wwUU	0011		

**Instruction Descriptions (E to MULUU) (continued)**

FLIM	Force (Signed) Data Range Limit
Description:	<p>Simultaneously compare any size of signed data value in Ws to a maximum signed limit value held in Wb and a minimum signed limit value held in Wb+1.</p> <p>For long word data size, if Ws[31:0] is greater than Wb[31:0], set Ws[31:0] to the limit value held in Wb[31:0]. For word data sizes, if Ws[15:0] is greater than Wb[15:0], set Ws[15:0] to the limit value held in Wb[15:0] (Ws[31:16] is ignored). The Z, N and OV status bits are updated such that a subsequent BGT instruction will take a branch.</p> <p>For long word data size, if Ws[31:0] is less than Wb+1[31:0], set Ws[31:0] to the limit value held in Wb+1[31:0]. For word data sizes, if Ws[15:0] is less than Wb+1[15:0], set Ws[15:0] to the limit value held in Wb+1[15:0] (Ws[31:16] is ignored). The Z, N and OV status bits are updated such that a subsequent BLT instruction will take a branch.</p> <p>If Ws[31:0] is less than or equal to the maximum limit value in Wb[31:0], and greater than or equal to the minimum limit value in Wb+1[31:0] (for long word sized data), or Ws[15:0] is less than or equal to the maximum limit value in Wb[15:0] and greater than or equal to the minimum limit value in Wb+1[15:0] (for word sized data), neither data limit is applied such that the contents of Ws will not change (for the given data size). The Z status bit is set such that a subsequent BZ instruction will take a branch (OV and N status are both cleared).</p> <p>For word sized register direct operations, Ws is always 0 extended to 32-bits (even if neither of the data limits are applied)<sup>4</sup>.</p> <p>Note that the instruction always executes the maximum compare first. Should Ws be less than the maximum, it will then execute the minimum compare, and if Ws is found to be not less than the minimum, Ws will remain unchanged.</p> <p>Furthermore, should the (maximum) limit value Wb be inadvertently set to less than or equal to the (minimum) limit value Wb+1, the same compare sequence will still be executed. This condition is not detected.</p> <p>The OV status bit is always cleared by this instruction.</p> <p>The 's' bits select the source (data value) register.</p> <p>The 'w' bits select the base (data limit) register.</p> <p>The 'p' bits select the source addressing mode.</p> <p><b>Notes:</b></p> <ol style="list-style-type: none"> <li>1. Although the instruction assumes signed values for all operands, both upper and lower limit values may be of the same sign.</li> <li>2. The status bits are set based upon the data compare result.</li> <li>3. Word data FLIM{.w} writes to Ws are zero extended to 32-bits.</li> <li>4. Zero extension to 32-bits is consistent with any register direct word write. However, the potential for a RAW hazard from the Ws write therefore always exists, even if the limit is not exceeded.</li> </ol> <p>I-Words: 1 Cycles: 1</p>

FLIM Force (Signed) Data Range Limit with Limit Excess Result								
Syntax:	{label:}	FLIM{.w}{v} Wb,	Ws,	Wd				
		FLIM.l{v}	[Ws],	[Wd]				
			[Ws++],	[Wd++]				
			[Ws--],	[Wd--]				
			[++Ws],	[++Wd]				
			[~-Ws],	[~-Wd]				
Operands:	$Ws \in [W0 \dots W14]; Wb \in [W0 \dots W13]; Wnd \in [W0 \dots W14]$							
Operation:	<pre> IF (Ws) &gt; (Wb) THEN (   IF FLIM.v THEN (Ws) - (Wb) → Wnd   ELSE +1 → Wnd;   (Wb) → (Ws);   0 → Z; 0 → N; 0 → OV; ) ELSE IF (Ws) &lt; (Wb+1) THEN (   IF FLIM.v THEN (Ws) - (Wb+1) → Wnd   ELSE -1 → Wnd;   (Wb+1) → Ws;   0 → Z; 1 → N; 0 → OV; ) ELSE (   0 → Wnd;   (Ws) → Ws<sup>4</sup>;   1 → Z; 0 → N; 0 → OV; ) </pre>							
Status Affected:	N, Z, OV							
Encoding:	1110	010L	dddd	ssss	pppq	qqww	wwUV	0111

**Instruction Descriptions (E to MULUU) (continued)**

FLIM	Force (Signed) Data Range Limit with Limit Excess Result
Description:	<p>Simultaneously compare any size of signed data value in Ws to a maximum signed limit value held in Wb and a minimum signed limit value held in Wb+1. Write a limit excess value into Wnd.</p> <p>For long word data size, if Ws[31:0] is greater than Wb[31:0], set Ws[31:0] to the limit value held in Wb[31:0]. For word data sizes, if Ws[15:0] is greater than Wb[15:0], set Ws[15:0] to the limit value held in Wb[15:0] (Ws[31:16] is ignored).</p> <p>In all cases, write the (signed) value by which the limit is exceeded to Wnd[31:0] (FLIM.v, where opcode bit field V = 1) or set Wnd to +1 (FLIM, where opcode bit field V = 0). Whenever Ws is greater than Wb, Wnd will always be a positive value. The Z, N and OV status bits are updated such that a subsequent BGT instruction will take a branch.</p> <p>For long word data size, if Ws[31:0] is less than Wb+1[31:0], set Ws[31:0] to the limit value held in Wb+1[31:0]. For word data sizes, if Ws[15:0] is less than Wb+1[15:0], set Ws[15:0] to the limit value held in Wb+1[15:0] (Ws[31:16] is ignored). In all cases, write the (signed) value by which the limit is exceeded to Wnd[31:0] (FLIM.v, where opcode bit field V = 1) or set Wnd to +1 (FLIM, where opcode bit field V = 0). Whenever Ws is less than Wb+1, Wnd will always be a negative value. The Z, N and OV status bits are updated such that a subsequent BLT instruction will take a branch.</p> <p>If Ws[31:0] is less than or equal to the maximum limit value in Wb[31:0] and greater than or equal to the minimum limit value in Wb+1[31:0] (for long word sized data), or Ws[15:0] is less than or equal to the maximum limit value in Wb[15:0] and greater than or equal to the minimum limit value in Wb+1[15:0] (for word sized data), neither data limit is applied such that contents of Ws will not change (for the given data size). The Z status bit is set such that a subsequent BZ instruction will take a branch (OV and N status are both cleared).</p> <p>For word sized register direct operations, Ws is always zero extended to 32-bits (even if neither of the data limits are applied)<sup>4</sup>.</p> <p>Note that the instruction always executes the maximum compare first. Should Ws be less than the maximum, it will then execute the minimum compare, and if Ws is found to be not less than the minimum, Ws will remain unchanged.</p> <p>Furthermore, should the (maximum) limit value Wb be inadvertently set to less than or equal to the (minimum) limit value Wb+1, the same compare sequence will still be executed. This condition is not detected.</p> <p>The OV status bit is always cleared by this instruction.</p> <p>The 's' bits select the source (data value) register.</p> <p>The 'w' bits select the base (data limit) register.</p> <p>The 'd' bits select the destination (limit test result) register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'q' bits select the destination addressing mode.</p> <p>The 'V' bit selects the result format for Wnd.</p> <p>See for modifier addressing information.</p> <p><b>Notes:</b></p> <ol style="list-style-type: none"> <li>1. Although the instruction assumes signed values for all operands, both upper and lower limit values may be of the same sign.</li> <li>2. The status bits are set based upon the value loaded into Wnd.</li> <li>3. Word data FLIM{.w} writes to Ws (and FLIM{.w}v writes to Wnd) are zero extended to 32-bits.</li> <li>4. Zero extension to 32-bits is consistent with any register direct word write. However, the potential for a RAW hazard from the Ws write therefore always exists, even if the limit is not exceeded.</li> <li>5. Ws and Wd cannot be the same register.</li> </ol> <p>I-Words: 1</p> <p>Cycles: 2</p>

GOTO	Unconditional Jump							
Syntax:	{label:} GOTO lit24							
Operands:	lit24 ∈ [0 ... 16MB]							
Operation:	lit24 → PC[23:0], NOP → Instruction Register.							
Status Affected:	None							
Encoding:	1101	101n	nnnn	nnnn	nnnn	nnnn	nnnn	nn10
Description:	Unconditional jump to any address within (for latency) one cycle's memory space. The 24-bit value 'lit24' is loaded into the PC. A jump to either a 32-bit or 16-bit instruction is permitted. The 'n' bits form the target address.							
	<b>Note:</b> The (byte) PC address is always either word or long word aligned. The opcode does not store the LSb because it is always 1'b0.							
I-Words:	1							
Cycles:	1							
Example:1	Address	Label	Instruction				Comment	
	0x804000	-	GOTO _CHECK_FAIL				; Jumps to address 0x805000(_CHECK_ FAIL)	
	...		MOV.I SR, W6 CP w6, w7				; Random code	
	0x805000	_CHECK_FAIL:						

GOTO	Unconditional Indirect Jump							
Syntax:	{label:}	GOTO	Wns					
Operands:	Wns ∈ [W0 ... W14]							
Operation:	(Wns) → PC[23:0];							
	NOP → Instruction Register.							
Status Affected:	None							
Encoding:	1101	011U	UUUU	ssss	UUUU	UUUU	UUUU	1010
Description:	Unconditional indirect jump to any address within executable memory address space. The Wns[23:0] is loaded into PC[23:0]. Wns must therefore contain a PS byte address. The value of Wns[0] is ignored and PC[0] is always set to 1'b0. GOTO is a two-cycle instruction. The 's' bits select the source register. If Wns[31:24] != 8'h00, an address error trap will be initiated.							
I-Words:	1							
Cycles:	2							
<b>Note:</b> The goto target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.								

Example:1	Address	Label	Instruction	Comment
	0x804000	-	GOTO _CHECK_FAIL	; Jumps to address 0x805000(_CHECK_FAIL)
	...			
	0x805000	_CHECK_FAIL:	MOV.I SR, W6 CP w6, w7	; Random code

INC	Increment f						
Syntax:	{label:}	INC.b	f	{,Wnd}	{,WREG}		
		INC.bz					
		INC{.w}					
		INC.l					
Operands:		$f \in [0 \dots 64KB]$ ; Wnd $\in [W0 \dots W14]$					
Operation:		$(f) + 1 \rightarrow$ destination designated by D					
Status Affected:		C, N, OV, Z					
Encoding:	1101	010L	dddd	ffff	ffff	ffff	BD01
Description:	Add one to the contents of the file register and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register.						
	The 'L' and 'B' bits select operation data width.						
	The 'D' bit selects the destination.						
	The 'f' bits select the address of the file register.						
	The 'd' bits select the Working register.						
I-Words:	1						
Cycles:	1						

**Example 1:** INC.B 0x1000 ; Increment 0x1000 (Byte mode)

Before Instruction		After Instruction	
Data 1000	8FFF	Data 1000	8F00
SR	0000	SR	0101 (DC, C = 1)

**Example 2:** INC 0x1000, WREG ; Increment 0x1000 and store to WREG ; (Word mode)

Before Instruction		After Instruction	
WREG	ABCD	WREG	9000
Data 1000	8FFF	Data 1000	8FFF
SR	0000	SR	0108 (DC, N = 1)

<b>INC2</b>	<b>Increment f by 2</b>						
Syntax:	{label:}	INC2.b	f	{,Wnd}	{,WREG}		
		INC2.bz					
		INC2{.w}					
		INC2.l					
Operands:		f ∈ [0 ... 64KB]; Wnd ∈ [W0 ... W14]					
Operation:		(f) + 2 → destination designated by D					
Status Affected:		C, N, OV, Z					
Encoding:	1101	110L	dddd	fffff	fffff	fffff	BD01
Description:		Add two to the contents of the file register and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register.					
		The 'L' and 'B' bits select operation data width.					
		The 'D' bit selects the destination.					
		The 'f' bits select the address of the file register.					
		The 'd' bits select the Working register.					
I-Words:	1						
Cycles:	1						

<b>Example 1:</b>	<b>INC2.B 0x1000 ; Increment 0x1000 by 2 ; (Byte mode)</b>
	Before Instruction                          After Instruction
Data 1000	8FFF
SR	0000                                         Data 1000 8F01 SR 0101 (DC, C = 1)

<b>Example 2:</b>	<b>INC2 0x1000, WREG ; Increment 0x1000 by 2 and store to WREG ; (Word mode)</b>
	Before Instruction                          After Instruction
WREG	ABCD
Data 1000	8FFF
SR	0000                                         WREG 9001 SR 0108 (DC, N = 1) Data 1000 8FFF

IOR	Inclusive OR Wb and Ws								
Syntax:	{label:}	IOR.b	Wb,	Ws,	Wd				
		IOR.bz		[Ws],	[Wd]				
		IOR{.w}		[Ws++],	[Wd++]				
		IOR.I		[Ws--],	[Wd--]				
				[++Ws],	[++Wd]				
				[~-Ws],	[~-Wd]				
				SR	SR				
Operands:		Wb ∈ [W0 ... W14]; Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]							
Operation:		(Wb).IOR.(Ws) → Wd							
Status Affected:		N, Z							
Encoding:		S110	101L	dddd	ssss	pppq	qqww	wwUU	BU00
Description:		Compute the Inclusive OR of the contents of the source register Ws and the contents of the base register Wb and place the result in the destination register Wd.							
		The 'S' bit selects instruction size.							
		The 'L' and 'B' bits select operation data width.							
		The 's' bits select the source register.							
		The 'w' bits select the base register.							
		The 'd' bits select the destination register.							
		The 'p' bits select the source addressing mode.							
		The 'q' bits select the destination addressing mode.							
		<b>Notes:</b>							
		1. When the SR is selected, SR data write will take priority over any SR update resulting from the IOR operation.							
		2. .bz data size/mode is disallowed when writing to the SR.							
I-Words:		1 or 0.5							
Cycles:		1							

**Example 1:** IOR.B W1, [W5++], [W9++]; IOR W1 and [W5] (Byte mode); Store result to [W9]; Post-increment W5 and W9

Before Instruction		After Instruction	
W1	AAAA	W1	AAAA
W5	2000	W5	2001
W9	2400	W9	2401
Data 2000	1155	Data 2000	1155
Data 2400	0000	Data 2400	00FF
SR	0000	SR	0008 (N = 1)

**Example 2:** IOR W1, W5, W9; IOR W1 and W5 (Word mode); Store the result to W9

Before Instruction		After Instruction	
W1	AAAA	W1	AAAA
W5	5555	W5	5555
W9	A34D	W9	FFFF
SR	0000	SR	0008 (N = 1)

IOR		Inclusive OR Wb and Short Literal				
Syntax:	{label:}	IOR.b	Ws,	lit7,	Wd	
		IOR.bz	[Ws],		[Wd]	
		IOR.{.w}	[Ws++],		[Wd++]	
		IOR.I	[Ws-],		[Wd-]	
			[++Ws],		[++Wd]	
			[~-Ws],		[~-Wd]	
			SR		SR	
Operands:		Ws ∈ [W0 ... W15]; lit7 ∈ [0 ... 127]; Wd ∈ [W0 ... W15]				
Operation:		(Ws).IOR.lit7 → Wd				
		Note: The literal is zero-extended to the selected data size of the operation				
Status Affected:		N, Z				
Encoding:		1110      101L      dddd      ssss      pppq      qqkk      kkkk      Bk10				
Description:		Compute the Inclusive OR of the contents of the source register Ws and the zero-extended literal operand and place the result in the destination register Wd. The 'L' and 'B' bits select operation data width. The 'K' bits provide the literal operand (MSb in op[2]). The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode. See and for modifier addressing information.				
Notes:						
		1. When the SR is selected, SR data write will take priority over any SR update resulting from the AND operation. 2. .bz data size/mode is disallowed when writing to the SR.				
I-Words:		1				
Cycles:		1				

<b>Example 1:</b>	IOR.B W1, #0x5, [W9++] ; IOR W1 and 0x5 (Byte mode) ; Store to [W9] ; Post-increment W9		
	Before Instruction                          After Instruction		
W1	AAAA	W1	AAAA
W9	2000	W9	2001
Data	2000	Data	2000 00AF
SR	0000	SR	0008 (N = 1)

<b>Example 2:</b>	IOR W1, #0x0, W9 ; IOR W1 with 0x0 (Word mode) ; Store to W9		
	Before Instruction                          After Instruction		
W1	0000	W1	0000
W9	A34D	W9	0000
SR	0000	SR	0002 (Z = 1)

IOR	Inclusive OR Literal and Wn							
Syntax:	{label:}	IOR.b	lit16,	Wn				
		IOR.bz		SR				
		IOR{.w}						
		IOR.I						
Operands:		lit16 ∈ [0 ... 65535]; Wn ∈ [W0 ... W14]; Status Register (SR)						
Operation:		(Wn) .IOR. lit16 <sup>3</sup> → Wn or (SR) .IOR. lit16 → SR						
Status Affected:		N, Z						
Encoding:	1100	101L	ssss	kkkk	kkkk	kkkk	kkkk	BT10
Description:	Compute the Inclusive OR of the literal operand and the contents of the Working register Wn or SR and place the result in the Working register Wn or SR. The 'L' and 'B' bits select operation data width.							
	The 's' bits select the Working register.							
	The 'k' bits specify the literal operand.							
	The 'T' bit selects between Ws (T = 0) and SR (T = 1) target registers.							
Notes:								
	1. When the SR is selected, SR data write will take priority over any SR update resulting from the AND operation.							
	2. .bz data size/mode is disallowed when writing to the SR.							
	3. The literal is zero-extended to 32-bits for long word operations.							
I-Words:	1							
Cycles:	1							

Example 1:	IOR.B #0xAA, W9 ; IOR 0xAA to W9 ; (Byte mode)
	Before Instruction
W9	1234
SR	0000
	After Instruction
W9	12BE
SR	0008 (N = 1)

Example 2:	IOR #0x2AA, W4 ; IOR 0x2AA to W4 ; (Word mode)
	Before Instruction
W4	A34D
SR	0000
	After Instruction
W4	A3EF
SR	0008 (N = 1)

IOR	Inclusive OR f and Wn						
Syntax:	{label:}	IOR.b	f	,Wn	{,WREG}		
		IOR.bz					
		IOR{.w}					
		IOR.I					
Operands:		$f \in [0 \dots 64KB]$ ; $Wn \in [W0 \dots W14]$					
Operation:		$(f).IOR.(Wn) \rightarrow \text{destination designated by } D$					
Status Affected:		N, Z					
Encoding:	1110	101L	ssss	ffff	ffff	ffff	BD01
Description:	<p>Compute the Inclusive OR of the contents of the Working register and the contents of the file register and place the result in the destination designated by D. If the optional Wn is specified, D=0 and store result in Wn; otherwise, D=1 and store result in the file register.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>						
I-Words:	1						
Cycles:	1						

Load Accumulator							
Syntax:	{label:}	LAC{.w}	Ws,	{ Slit6, }	A		
		LAC.I	[Ws],		B		
			[Ws++]				
			[Ws-]				
			[~Ws],				
			[++Ws],				
			[Ws+Wb],				
Operands:	Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Slit6 ∈ [-32 ... +31]						
Operation:	<u>For Word operation:</u> 8{Ws[15]}, Ws[15:0], 48'h000000000000 → ACC[71:0] Shift <sub>Slit6</sub> (ACC) → ACC						
	<u>For Long Word operation:</u> 8{Ws[31]}, Ws[31:0], 32'h00000000 → ACC[71:0] Shift <sub>Slit6</sub> (ACC) → ACC						
Status Affected:	OA, SA or OB, SB						
Encoding:	1100	00AL	www	ssss	pppu	UUkk	kkkk
Description:	Read the contents of the Effective Address. Place the data read into the target accumulator, then optionally shift the entire accumulator. For a word operation (LAC{.w}), the value contained at the Effective Address is assumed to be Q1.15 signed fractional data and is written into ACCx[63:48]. ACCx[47:0] is set to all 0's, and ACCx[63:0] is automatically sign-extended (through bit 71). The accumulator is then arithmetically shifted by the value held in slit6. For a long word operation (LAC.I), the value contained at the Effective Address is assumed to be Q1.31 signed fractional data and is written into ACCx[63:32]. ACCx[31:0] is set to all 0's, and ACCx[63:0] is automatically sign-extended (through bit 71). The accumulator is then arithmetically shifted by the value held in slit6. The 'L' bit selects word or long word operation. The 'A' bits specify the destination accumulator. The 's' bits specify the source register Wns. The 'p' bits select the source addressing mode. The 'w' bits specify the offset register Wb. The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator shift; if the operand Slit6 is absent, the literal bit field is set to all 0's. <b>Note:</b> Positive values of operand Slit6 represent arithmetic shift right. Negative values of operand Slit6 represent shift left.						
I-Words:	1						
Cycles:	1						

Example 1:	LAC [W4++], #-3, B ; Load ACCB with [W4] << 3 ; Contents of [W4] do not change ; Post increment W4 ; Assume saturation disabled ; (SATB = 0)						
	Before Instruction						
	W4 2000						
	ACCB 00 1525 ABCD 2512 1922						
	Data 2000 1221						
	SR 0000						
	After Instruction						
	W4 2002						
	ACCB FF 9108 0000 0000 0000						
	Data 2000 1221						
	SR 4800 (OB, OAB = 1)						

Example 2:	LAC [--W2], #7, A ; Pre-decrement W2 ; Load ACCA with [W2] >> 7 ; Contents of [W2] do not change ; Assume saturation disabled ; (SATA = 0)						
	Before Instruction						
	W2 4002						
	After Instruction						
	W2 4000						

## Instruction Descriptions (E to MULUU) (continued)

**Example 2:**      **IAC [--W2], #7, A ; Pre-decrement W2 ; Load ACCA with [W2] >> 7 ; Contents of [W2] do not change ; Assume saturation disabled ; (SATA = 0)**

ACCA	00 1525 ABCD 2512	ACCA	FF FF22 1000 0000
	1922		0000
Data 4000	9108	Data 4000	9108
Data 4002	1221	Data 4002	1221
SR	0000	SR	0000

LLAC Load Lower Accumulator							
Syntax:	{label:}	LLAC.I	Ws, [Ws], [Ws++] [Ws-] [-Ws], [++Ws], [Ws+Wb],	{ Slit6, }	A B		
Operands:			Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Slit6 ∈ [-32 ... +31]				
Operation:			Ws[31:0] → ACC[31:0]				
			Shift <sub>Slit6</sub> (ACC) → ACC				
Status Affected:			OA, SA or OB, SB				
Encoding:	1100	00A1	wwww	ssss	pppu	UUkk	kkkk 0111
Description:			Read the contents of the Effective Address. Place the data read into the lower 32-bits of the target accumulator (ACCx[31:0]), then optionally arithmetically shift the entire accumulator.				
			The 'A' bits specify the destination accumulator.				
			The 's' bits specify the source register Wns.				
			The 'p' bits select the source addressing mode.				
			The 'w' bits specify the offset register Wb.				
			The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator shift; if the operand Slit6 is absent, the literal bit field is set to all 0's.				
			<b>Note:</b> Positive values of operand Slit6 represent arithmetic shift right. Negative values of operand Slit6 represent shift left.				
I-Words:	1						
Cycles:	1						
Example:1	LLAC	A, #4, w8					; Load ACCAL from w8 ;Right shift value of ACCA by 4
			Before execution				After execution
	ACCA		0x00_1022_2EE1_5 633_9078			ACCA	0xFF_F102_22EE_11 23_4567
	w8		0x12345678			w8	0x12345678

Allocate Stack Frame														
LNK														
Syntax:	{label:} LNK lit16													
Operands:	lit16 ∈ [0 ... 65535]													
Operation:	(W14) → (TOS) (W15) + 4 → W15 (W15) → W14; (W15) + lit16 → W15													
Status Affected:	None													
Encoding:	1101	010U	UUUU	kkkk	kkkk	kkkk	1010							
Description:	This instruction allocates a stack frame of size lit16 bytes (though always long word aligned) and adjusts the Stack Pointer and Frame Pointer. The 'k' bits specify the size of the stack frame in bytes.													
<b>Note:</b>														
1. This instruction operates with long word aligned operands only. The LS 2-bits of the literal encoding are therefore always set to 2'b00.														
I-Words:	1													
Cycles:	1													
Example:1	LNK	#8	;Allocate stack frame											
Before execution				After execution										
W14 0x1234				W14 0x5004										
W15 0x5000				W15 0x500C										
Data @0x5000 0x0000				Data @0x5000 0x1234										

Allocate Stack Frame (short)			
LNK			
Syntax:	{label:}	LNK	lit7
Operands:	lit7 ∈ [0 ... 127]		
Operation:	(W14) → (TOS) (W15) + 4 → W15 (W15) → W14; (W15) + lit5, 2'b00 → W15		
Status Affected:	None		
Encoding:	0111      001k      kkkk      1110		
Description:	This instruction allocates a stack frame of size lit7 bytes and adjusts the Stack Pointer and Frame Pointer. The 'k' bits specify the size of the stack frame (in long words).		
<b>Note:</b>			
1. This instruction operates with long word aligned operands only. The LS 2-bits of the frame size are therefore assumed to be always set to 2'b00 (i.e., such that only 5-bits of the literal need be encoded within the opcode).			
I-Words:	0.5		
Cycles:	1		

LSR		Logical Shift Right by 1						
Syntax:	{label:}	LSR.b	Ws,	Wd				
		LSR.bz	[Ws],	[Wd]				
		LSR.{w}	[Ws++],	[Wd++]				
		LSR.l	[Ws-],	[Wd--]				
			[++Ws],	[++Wd]				
			[~-Ws],	[~-Wd]				
Operands:		Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]						
Operation:		For long word operation: 0 → Wd[31], (Ws[31:1]) → Wd[30:0], (Ws[0]) → C						
		For word operation: 0 → Wd[15], (Ws[15:1]) → Wd[14:0], (Ws[0]) → C						
		For byte operation: 0 → Wd[7], (Ws[7:1]) → Wd[6:0], (Ws[0]) → C						
		0 C						
Status Affected:		C,N,Z						
Encoding:	S010	101L	dddd	ssss	pppq	qqUU	UUUU	B000
Description:	Logical shift right the contents of the source register Ws by one bit, placing the result in the destination register Wd. Destination register direct Extended Byte or Word mode will zero-extend the result to 32-bits, then write to Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.							
	<b>Note:</b> 1. The N flag can only be set if the MSb of Ws = 1 with a shift value of zero.							
I-Words:	1 or 0.5							
Cycles:	1							

Logical Shift Right f							
LSR	{label:}	LSR.b	f	{,Wnd}	{,WREG}		
Syntax:		LSR.b	f	{,Wnd}	{,WREG}		
		LSR.bz					
		LSR{.w}					
		LSR.I					
Operands:		$f \in [0 \dots 64KB]$		$Wnd \in [W0 \dots W14]$			
Operation:		<u>For byte operation:</u> 0 → Dest[7:1], (f[7:1]) → Dest[6:0], (f[0]) → C					
		<u>For word operation:</u> 0 → Dest[15], (f[15:1]) → Dest[14:0], (f[0]) → C					
		<u>For long word operation:</u> 0 → Dest[31], (f[31:1]) → Dest[30:0], (f[0]) → C					
Status Affected:		C, N, Z					
Encoding:	1010	001L	dddd	ffff	ffff	ffff	BD01
Description:	Shift the contents of the file register f one bit to the right and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register. The carry flag bit is set if the LSB of the file register is '1'. The N flag is always cleared. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'f' bits select the address of the file register. The 'd' bits select the Working register.						
I-Words:	1						
Cycles:	1						

**Example 1:** LSR.B 0x600 ; Logically shift right (0x600) by one ; (Byte mode)

Before Instruction				After Instruction			
Data 600	55FF			Data 600	557F		
SR	0000			SR	0001	(C = 1)	

**Example 2:** LSR 0x600, WREG ; Logically shift right (0x600) by one ; Store to WREG ; (Word mode)

Before Instruction				After Instruction			
Data 600	55FF			Data 600	55FF		
WREG	0000			WREG	2AFF		
SR	0000			SR	0001	(C = 1)	

Logical Shift Right by Short Literal						
Syntax:	{label:}	LSR{.w}	Ws,	lit5,	Wd	
		LSR.I	[Ws],		[Wd]	
			[Ws++],		[Wd++]	
			[Ws-],		[Wd-]	
			[++Ws],		[~Wd]	
			[~-Ws],		[++Wd]	
Operands:	Ws ∈ [W0 ... W15]; lit5 ∈ [0...32]; Wd ∈ [W0 ... W14]					
Operation:	lit5[4:0] → Shift_Val 1'b0 → Right shift input (logical shift) For long word operation: Ws[31:0], 32'b0 → Shift_In[63:0] Logical shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[63:32] → Wd For word operation: 16'b0,Ws[15:0],32'b0 → Shift_In[63:0] Logical shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[47:32] → Wd[15:0]					
Status Affected:	N,Z					
Encoding:	S010	001k	kkkk	dddd	pppq	qqss ssUU L000
Description:	Logical shift right the contents of the source register Ws by lit5 bits (up to 31 positions), placing the result in the destination Wd. The N flag is always cleared unless the shift value is zero (see note 1). This instruction will generate the correct result for any shift value in lit5 (a word operation shift value > 15, Wd[15:0]=0x0000). Register Direct Word mode will zero-extend the result to 32-bits, then write to Wd. The 'S' bit selects instruction size. The 'L' bit selects word or long word operation. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode. The 'k' bits provide the literal operand.					
Notes:	1. The N flag can only be set if the MSb of Ws = 1 with a shift value of zero. 2. This instruction only operates in Word or Long Word mode.					
I-Words:	1 or 0.5					
Cycles:	1					

**Example 1:** LSR W4, #14, W5 ; LSR W4 by 14 ; Store result to W5

Before Instruction		After Instruction	
W4	C800	W4	C800
W5	1200	W5	0003
SR	0000	SR	0000

**Example 2:** LSR W4, #1, W5 ; LSR W4 by 1 ; Store result to W5

Before Instruction		After Instruction	
W4	0505	W4	0505

**Instruction Descriptions (E to MULUU) (continued)**

<b>Example 2:</b> <code>LSR W4, #1, W5 ; LSR W4 by 1</code>				<code>; Store result to W5</code>
W5	F000	W5	0282	
SR	0000	SR	0000	

LSR	Logical Shift Right by Wb					
Syntax:	{label:}	LSR.b	Ws,	Wb,	Wd	
		LSR.bz	[Ws],		[Wd]	
		LSR.{w}	[Ws++],		[Wd++]	
		LSR.l	[Ws-],		[Wd-]	
			[++Ws],		[++Wd]	
			[--Ws],		[--Wd]	
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]				
Operation:		Wb[15:0] → Shift_Val				
		For long word operation:				
		Ws[31:0], 32'b0 → Shift_In[63:0]				
		Logical shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0]				
		Shift_Out[63:32] → Wd				
		For word operation:				
		16'b0, Ws[15:0], 32'b0 → Shift_In[63:0]				
		Logical shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0]				
		Shift_Out[47:32] → Wd[15:0]				
		For byte operation:				
		24'b0, Ws[7:0], 32'b0 → Shift_In[63:0]				
		Logical shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0]				
		Shift_Out[39:32] → Wd[7:0]				
Status Affected:		N,Z				
Encoding:	1010      101L      wwww      dddd      pppq      qqss      ssUU      B100					
Description:	<p>Logical shift right the contents of the source register Ws by Wb bits, placing the result in the destination register Wd.</p> <p>When used in isolation, this instruction will generate the correct result for any shift value in Wb[15:0]:</p> <ul style="list-style-type: none"> <li>• For a byte operation shift value &gt; 7, Wd[7:0]=0x00</li> <li>• For a word operation shift value &gt; 15, Wd[15:0]=0x0000</li> <li>• For a long word operation shift value &gt; 31, Wd=0x00000000</li> </ul> <p>When used in conjunction with ASRMW and/or LSRMW instructions for multi-precision multi-bit shift operations, this instruction will not generate a correct result for any shift value greater than 32.</p> <p>Any data held in Wb[31:16] will have no effect.</p> <p>The N flag is always cleared unless the shift value is 0 (see note 1).</p> <p>Destination register direct Extended Byte or Word mode will zero-extend the result to 32-bits, then write to Wd.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base (shift count) register.</p> <p>The 'd' bits select the destination register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'q' bits select the destination addressing mode.</p>					
	<b>Note:</b>	<ol style="list-style-type: none"> <li>1. The N flag can only be set if the MSb of Ws = 1 with a shift value of zero.</li> </ol>				
I-Words:	1					
Cycles:	1					

LSRM		Logical Shift Right Multi-Precision by Short Literal					
Syntax:	{label:}	LSRM{.l}	Ws, [Ws], [Ws++], [Ws-], [++Ws], [-Ws],	lit5,	Wnd		
Operands:			Ws ∈ [W0 ... W15]; lit5 ∈ [0...31]; Wnd ∈ [W1 ... W14]				
Operation:			lit5[4:0] → Shift_Val 0 → Right shift input (logical shift) Ws[31:0], 32'b0 → Shift_In[63:0]				
			Logical shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[63:32] → Wnd Shift_Out[31:0]   Wnd-1 → Wnd-1				
Status Affected:			N,Z				
Encoding:			1110      000k      kkkk      dddd      pppU      UUss      ssUU      0111				
Description:			Logical shift right the contents of the source register Ws by lit5 bits (up to 31 positions), placing the result in the destination register Wnd. The register containing the next least significant data word will already contain an intermediate shift result. Bitwise OR this value with the data shifted out of Ws in order to create the final shift result, then update the corresponding destination register. The Z bit is "sticky" (can only be cleared).  The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'k' bits provide the literal operand.				
Notes:			<ol style="list-style-type: none"> <li>The N flag can only be set if the MSb of Ws = 1 with a shift value of zero.</li> <li>This instruction operates in Long Word mode only.</li> </ol>				
I-Words:			1				
Cycles:			2				

Logical Shift Right Multi-Precision by Wb						
LSRM	{label:}	LSRM{.l}	Ws,	Wb,	Wnd	
Syntax:			[Ws], [Ws++], [Ws--], [++Ws], [--Ws],			
Operands:			Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Wnd ∈ [W1 ... W14]			
Operation:			Wb[15:0] → Shift_Val 0 → Right shift input (logical shift) Ws[31:0], 32'b0 → Shift_In[63:0] Logical shift right Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[63:32] → Wnd Shift_Out[31:0]   Wnd-1 → Wnd-1			
Status Affected:			N,Z			
Encoding:			1110      001U      ssss      dddd      pppU      UUww      wwUU      0111			
Description:			Logical shift right the contents of the source register Ws by Wb bits, placing the result in the destination register Wnd. The register containing the next least significant data word will already contain an intermediate shift result. Bitwise OR this value with the data shifted out of Ws in order to create the final shift result, then update the corresponding destination register. The right shift may be by any amount between zero and 32 bits. Should the shift value held in Wb[15:0] exceed 2'd32, the shift value will saturate to 2'd32 for consistency. Any data held in Wb[31:16] will have no effect. The Z bit is "sticky" (can only be cleared). The 'w' bits select the base (shift count) register. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode.			
			<b>Note:</b> This instruction operates in Long Word mode only.			
I-Words:			1			
Cycles:			2			

LUAC Load Upper Accumulator							
Syntax:	{label:}	LUAC.I	Ws, [Ws], [Ws++] [Ws-] [-Ws], [+Ws], [Ws+Wb],	{ Slit6, }	A B		
Operands:			Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Slit6 ∈ [-32 ... +31]				
Operation:			Ws[7:0] → ACC[71:64] Shift <sub>Slit6</sub> (ACC) → ACC				
Status Affected:			OA, SA or OB, SB				
Encoding:	1100	11A1	www	ssss	pppu	UUkk	kkkk 1011
Description:			Read the contents of the Effective Address. Place the LSB of the data read into the upper 8-bits of the target accumulator (ACCx[71:64]), then optionally arithmetically shift the entire accumulator. The 'A' bits specify the destination accumulator.				
			The 's' bits specify the source register Wns.				
			The 'p' bits select the source addressing mode.				
			The 'w' bits specify the offset register Wb.				
			The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator shift; if the operand Slit6 is absent, the literal bit field is set to all 0's.				
			<b>Note:</b> Positive values of operand Slit6 represent arithmetic shift right. Negative values of operand Slit6 represent shift left.				
I-Words:	1						
Cycles:	1						
Example:1	LUAC	A, #4, w8					; Load ACCAU from w8[7:0] ;Right shift value of ACCA by 4
			Before execution				After execution
			ACCA      0x00_1022_2EE1_5 633_9078				ACCA      0xFF_F102_22EE_15 63_3907
			w8        0x000000FF				w8        0x000000FF

<b>MAC</b>						
<b>Multiply and Accumulate</b>						
Syntax: {label:}	MAC{.w}	Wx	, Wy,	A	{,AWB}	
	MAC.I	[Wx]	, [Wy],	B		
		[Wx]+=kx	, [Wy]+=ky,			
		[Wx]=-kx	, [Wy]=-ky,			
		[Wx+=kx]	, [Wy+=ky],			
		[Wx-=kx]	, [Wy-=ky],			
		[Wx+W12]	, [Wy+W12],			
Operands:	$Wx \in \{W0 \dots W14\}; Wy \in \{W0 \dots W14\}$					
	Word mode: $kx \in \{-8, -6, -4, -2, 2, 4, 6, 8\};$					
	$ky \in \{-8, -6, -4, -2, 2, 4, 6, 8\};$					
	Long Word mode: $kx \in \{-16, -12, -8, -4, 4, 8, 12, 16\};$					
	$ky \in \{-16, -12, -8, -4, 4, 8, 12, 16\};$					
	$AWB \in \{W0, W1, W2, W3, W13, [W13++], [W15++]\}^3$					
Operation:	$(ACC(A \text{ or } B)) + ((Wx) * (Wy)) \rightarrow ACC(A \text{ or } B);$					
	$(ACC(B \text{ or } A)) \text{ rounded} \rightarrow AWB$					
	When Indirect Pre/Post Modified Addressing:					
	$(Wx)+kx \rightarrow Wx \text{ or } (Wx)-kx \rightarrow Wx;$					
	$(Wy)+ky \rightarrow Wy \text{ or } (Wy)-ky \rightarrow Wy;$					
Status Affected:	OA,SA or OB,SB					
Encoding:	1101	00AL	wwww	ssss	IIii	iJJJ jjaa a011

**Instruction Descriptions (E to MULUU) (continued)****MAC****Multiply and Accumulate****Description:**

Signed/unsigned (defined by CORCON.US) multiply of data read from Wx and Wy or concurrently fetched from the X and Y address space. The result is sign-extended (when at least one operand is considered signed) or zero-extended (when both operands are unsigned) to 72-bits and added to the specified accumulator. Fractional or integer operation (defined by CORCON.IF) will determine if the result is scaled or not prior to the accumulator update. Fractional operation will scale the result to align the operand and accumulator (msw) fractional points (see note 3). Integer operation will align the LSb of the result with the LSb of the accumulator.

When indirect addressing is selected for either or both the X and Y address space, Wx and Wy registers provide the corresponding indirect addresses. The address modifier values are kx and ky, respectively, and represent the number of data bytes by which to modify the Effective Address.

The optional AWB specifies the direct or indirect (see note 4) store of the (32-bit) rounded fractional contents of the accumulator not targeted by the MAC operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size. AWB is not intended for use when the DSP engine is operating in Integer mode.

Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly.

The 'L' bit selects word or long word operation.

The 'A' bit selects the accumulator for the result.

The 'I' bits select the Operation.X-Space Addressing mode.

The 'i' bits select the kx modification value.

The 'J' bits select the Operation.Y-Space Addressing mode.

The 'j' bits select the ky modification value.

The 's' bits select the Wx register

The 'w' bits select the Wy register.

The 'a' bits select the accumulator write-back destination and addressing mode.

**Notes:**

1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.
2. MAC register direct where Wx = Wy is equivalent to SQRAC Wx, ACCx
3. The LS portion of ACCx is unaffected when operating in Fractional mode with word sized data. Lower significance data that may be present from prior (32-bit data) operations is therefore preserved. Users not requiring this should clear ACCx during initialization.
4. Use of the same W-reg for both indirect source (X or Y) and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.
5. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.

**I-Words:**

1

**Cycles:**

1

<b>Example:1</b>	<b>MAC.I</b>	<b>[W2] += 4, [W5] -= 4, A, W13</b>	<b>;Multiply the contents of W2 and W5 and accumulate the results to ACCA ;Increment W2 by 4; Decrement W5 by 4; ;Write back ACCB to W13</b>
------------------	--------------	---	--

**Before execution**

W2 0x5000

W5 0x7010

W13 0x00000000

CORCON[7:0] 8'b10110000 (Super saturation enabled.)

ACCA 0x00\_40000000  
0\_00000000ACCB 0x00\_4021024  
2\_29220124

Data @0x5000 0x40000000

**After execution**

W2 0x5004

W5 0x700C

W13 0x40210242

CORCON[7:0] 8'b10110000 (Super saturation enabled.)

ACCA 0x00\_60000000  
0\_00000000ACCB 0x00\_4021024  
2\_29220124

Data @0x5000 0x40000000

## Instruction Descriptions (E to MULUU) (continued)

Example:1	MAC.I	[W2] += 4, [W5] -= 4, A, W13	;Multiply the contents of W2 and W5 and accumulate the results to ACCA ;Increment W2 by 4; Decrement W5 by 4; ;Write back ACCB to W13
		Data @0x7010 0x20000000	Data @0x7010 0x20000000
Example:2	MAC.I	[W2+=4], [W5+=4], B	;Increment W2 and W5 by 4; ;Multiply the contents of W2 and W5 and accumulate the results to ACCB
Before execution			After execution
W2 0x5000			W2 0x5004
W5 0x7010			W5 0x7014
W13 0x00000000			W13 0x40210242
CORCON[7:0] 8'b11110000 (Super saturation enabled.)			CORCON[7:0] 8'b11110000 (Super saturation enabled.)
ACCA 0x00_4000000 0_00000000			ACCA 0x00_4000000 0_00000000
ACCB 0xFF_E000000 0_00000000			ACCB 0x00_1FFFFFF F_00000000
Data @0x5000 0x40000000			Data @0x5000 0x40000000
Data @0x7010 0x20000000			Data @0x7010 0x20000000
Data @0x5004 0x7FFFFFFF			Data @0x5004 0x7FFFFFFF
Data @0x7014 0x40000000			Data @0x7014 0x40000000

MAX								Force (Signed) Maximum Data Range Limit							
Syntax:	{label:}	MAX{.w}	Wb,	Ws											
		MAX.I		[Ws]	[Ws++]	[Ws--]	[++Ws]								
Operands:		$Ws \in [W0 \dots W15]; Wb \in [W0 \dots W14]$													
Operation:		IF (Ws) > (Wb) THEN { (Wb) → (Ws); 0 → Z; 0 → N; 0 → OV } ELSE { 1 → Z; 0 → N; 0 → OV }													
Status Affected:		N, Z, OV													
Encoding:	1110	010L	UUUU	ssss	pppU	UUww	wwUU	1111							
Description:	Compare any size of signed data value in Ws to a maximum signed limit value held in Wb. For long word data size, if Ws[31:0] is greater than Wb[31:0], set Ws[31:0] to the limit value held in Wb[31:0]. For word data sizes, if Ws[15:0] is greater than Wb[15:0], set Ws[15:0] to the limit value held in Wb[15:0] (Ws[31:16] is ignored). The Z, N and OV status bits are set such that a subsequent BGT instruction will take a branch. If Ws[31:0] is less than or equal to the maximum limit value in Wb[31:0] (for long word data size), or if Ws[15:0] is less than or equal to the maximum limit value in Wb[15:0] (word data size), the data limit is not applied such that contents of Ws will not change (for the given data size). The Z, N and OV status bits are set such that a subsequent BZ or BLE instruction will take a branch. The OV status bit is always cleared by this instruction. The 's' bits select the source (data value) register. The 'w' bits select the base (data limit) register. The 'p' bits select the source addressing mode.														
	<b>Note:</b>														
	1. The status bits are set based upon the data compare result.														
I-Words:	1														
Cycles:	1														

MAX	Accumulator Force Maximum Data Range Limit			
Syntax:	{label:}	MAX	A	
		B		
Operands:	None			
Operation:	IF (MAX A) THEN (			
	IF ACCA - ACCB > 0 THEN (			
	ACCB → ACCA;			
	0 → Z; 0 → N; 0 → OV;			
	)			
	ELSE (			
	1 → Z; 0 → N; 0 → OV;			
	)			
	)			
	IF (MAX B) THEN (			
	IF ACCB - ACCA > 0 THEN (			
	ACCA → ACCB;			
	0 → Z; 0 → N; 0 → OV;			
	)			
	ELSE (			
	1 → Z; 0 → N; 0 → OV;			
	)			
	)			
Status Affected:	N, OV, Z			
Encoding:	0111          001A          UUUU          1101			
Description:	Clamp the target accumulator (defined in the instruction) to the maximum limit value previously loaded into the other accumulator. The compare examines the full 72-bit value of the target accumulator and will therefore clamp an overflowed accumulator (variable saturation). If the target accumulator is greater than the limit accumulator, load the target accumulator with the contents of the limit accumulator. The Z and N status bits are set such that a subsequent BGT instruction will take a branch. If the target accumulator is not greater than the limit accumulator, the target accumulator is unaffected. The Z status bit is set such that a subsequent BZ instruction will take a branch. The OV status bit is always cleared by this instruction. The 'A' bit specifies the destination accumulator.			
	<b>Note:</b>			
	1. OA and SA or OB and SB status bits are not modified by this instruction. Execute SFTAC <ACCx>, 0 after MAXAB operation to update DSP status to reflect contents of ACCx.			
I-Words:	0.5			
Cycles:	1			

<b>MAX</b>		Accumulator Force Maximum Data Range Limit with Limit Excess Result											
Syntax:	{label:}	MAX{.w}{.v} A, MAX.I{v} B,	Wd	[Wd]	[Wd++]	[Wd--]	[++Wd]						
							[--Wd]						
							[Wd+Wb]						
Operands:		Wd ∈ [W0 ... W15] (see note 4); Wb ∈ [W0 ... W15]											
Operation:	<pre> IF (MAX A) THEN (   IF ACCA - ACCB &gt; 0 THEN (     +1 → Wd OR ACCA - ACCB → Wd (see text);     ACCB → ACCA;     0 → Z; 0 → N; 0 → OV;   )   ELSE (     0 → Wd;     1 → Z; 0 → N; 0 → OV;   ) ) IF (MAX B) THEN (   IF ACCB - ACCA &gt; 0 THEN (     +1 → Wd OR ACCB - ACCA → Wd (see text);     ACCA → ACCB;     0 → Z; 0 → N; 0 → OV;   )   ELSE (     0 → Wd;     1 → Z; 0 → N; 0 → OV;   ) ) </pre>												
Status Affected:	N, OV, Z												
Encoding:	1100	10AL	wwww	dddd	UUUq	qqUU	UUUV	1011					

**Instruction Descriptions (E to MULUU) (continued)****MAX****Accumulator Force Maximum Data Range Limit  
with Limit Excess Result**

Description:	<p>Clamp the target accumulator (defined in the instruction) to the maximum limit value previously loaded into the other accumulator.</p> <p>The compare examines the full 72-bit value of the target accumulator and will therefore clamp an overflowed accumulator (variable saturation).</p> <p>If the target accumulator is greater than the limit accumulator, load the target accumulator with the contents of the limit accumulator. For MAX (instruction bit field V = 0, default when not declared), set Wd to +1. For MAX.v (instruction bit field V = 1), write the (signed) value by which the limit is exceeded to Wd:</p> <ul style="list-style-type: none"> <li>- Word operation: Consider only bits [47:32] of the compare result and write to Wd (destination is also word-sized, see note 3). If the limit is exceeded by a value greater than that which can be represented by a signed 16-bit number, saturate the Wd write to the maximum 16-bit positive value (i.e., set Wd[15:0] to 0x7FFF).</li> <li>- Long word operation: Consider only bits [31:0] of the compare result and write to Wd. If the limit is exceeded by a value greater than that which can be represented by a signed 32-bit number, saturate the Wd write to the maximum positive value (i.e., set Wd to 0xFFFFFFFF).</li> </ul> <p>The Z and N status bits are set such that a subsequent BGT instruction will take a branch if the limit is exceeded.</p> <p>If the target accumulator is not greater than the limit accumulator, the target accumulator is unaffected, and Wd is cleared. The Z status bit is set such that a subsequent BZ instruction will take a branch.</p> <p>The OV status bit is always cleared by this instruction.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'A' bit specifies the destination accumulator.</p> <p>The 'd' bits select the destination register.</p> <p>The 'q' bits select the destination addressing mode.</p> <p>The 'w' bits define the offset Wb.</p> <p>The 'V' bit defines the presence and result format for Wd.</p>
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. OA and SA or OB and SB status bits are not modified by this instruction. Execute SFTAC &lt;ACCx&gt;, 0 after MAXAB operation to update DSP status to reflect contents of ACCx.</li> <li>2. In keeping with all word sized register direct writes, Wd[15:0] will always be zero extended to 32-bits.</li> <li>3. Register direct destination W15 not permitted.</li> </ol>
I-Words:	1
Cycles:	2

<b>MIN</b>		<b>Force (Signed) Minimum Data Range Limit</b>						
Syntax:	{label:}	MIN{.w}	Wb,	Ws				
		MIN.I		[Ws]				
				[Ws++]				
				[Ws--]				
				[++Ws]				
				[~-Ws]				
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W14]						
Operation:		IF (Ws) < (Wb) THEN { (Wb) → (Ws); 0 → Z; 1 → N; 0 → OV } ELSE { 1 → Z; 0 → N; 0 → OV }						
Status Affected:		N, Z, OV						
Encoding:		1110	010L	UUUU	ssss	pppU	UUww	wwUU
Description:		Compare a word or long word signed data value in Ws to a minimum signed limit value held in Wb. For long word data size, if Ws[31:0] is less than Wb[31:0], set Ws[31:0] to the limit value held in Wb[31:0]. For word data sizes, if Ws[15:0] is less than Wb[15:0], set Ws[15:0] to the limit value held in Wb[15:0] (Ws[31:16] is ignored). The Z, N and OV status bits are set such that a subsequent BLT instruction will take a branch.						
		If Ws[31:0] is greater than or equal to the minimum limit value in Wb[31:0] (for long word data size), or if Ws[15:0] is greater than or equal to the minimum limit value in Wb[15:0] (word data size), the data limit is not applied such that contents of Ws will not change (for the given data size). The Z, N and OV status bits are set such that a subsequent BZ or BGE instruction will take a branch.						
		For word sized register direct operations, Ws is always zero extended to 32-bits (even if greater than or equal to the limit value in Wb) <sup>2</sup> .						
		The OV status bit is always cleared by this instruction.						
		The 's' bits select the source (data value) register.						
		The 'w' bits select the base (data limit) register.						
		The 'p' bits select the source addressing mode.						
		<b>Note:</b>						
		1. The status bits are set based upon the data compare result.						
I-Words:		1						
Cycles:		1						

MIN	Accumulator Force Minimum Data Range Limit (unconditional execution)			
Syntax:	{label:}	MIN	A	B
Operands:	None			
Operation:	IF (MIN A) THEN (			
	IF ACCA - ACCB < 0 THEN (			
	ACCB → ACCA;			
	0 → Z; 1 → N; 0 → OV;			
	)			
	ELSE (			
	1 → Z; 0 → N; 0 → OV;			
	)			
	)			
	IF (MIN B) THEN (			
	IF ACCB - ACCA < 0 THEN (			
	ACCA → ACCB;			
	0 → Z; 1 → N; 0 → OV;			
	)			
	ELSE (			
	1 → Z; 0 → N; 0 → OV;			
	)			
	)			
Status Affected:	N, OV, Z			
Encoding:	0111          001A          UUUU          1011			
Description:	Clamp the target accumulator (defined in the instruction) to the minimum limit value previously loaded into the other accumulator.			
	The compare examines the full 72-bit value of the target accumulator and will therefore clamp an overflowed accumulator (variable saturation).			
	If the target accumulator is less than the limit accumulator, load the target accumulator with the contents of the limit accumulator. The Z and N status bits are set such that a subsequent BLT instruction will take a branch.			
	If the target accumulator is not less than the limit accumulator, the target accumulator is unaffected. The Z status bit is set (Z = 1) such that a subsequent BLT instruction will take a branch.			
	The OV status bit is always cleared by this instruction.			
	The 'A' bit specifies the destination accumulator.			
<b>Note:</b>				
	1. OA and SA or OB and SB status bits are not modified by this instruction. Execute SFTAC <ACCx>, 0 after MAXAB execution to update DSP status to reflect contents of ACCx.			
I-Words:	0.5			
Cycles:	1			

<b>MIN</b>		Accumulator Force Minimum Data Range Limit with Limit Excess Result (unconditional execution)						
Syntax:	{label:}	MIN{.w}{.v} A, MIN.I{v} B,	Wd [Wd] [Wd++] [Wd--] [++Wd] [--Wd] [Wd+Wb]					
Operands:			Wd ∈ [W0 ... W15] (see note 4); Wb ∈ [W0 ... W15]					
Operation:			IF (MIN A) THEN ( IF ACCA - ACCB < 0 THEN ( -1 → Wd OR ACCA - ACCB → Wd (see text); ACCB → ACCA; 0 → Z; 1 → N; 0 → OV; ) ELSE ( 0 → Wd; 1 → Z; 0 → N; 0 → OV; ) IF (MIN B) THEN ( IF ACCB - ACCA < 0 THEN ( -1 → Wd OR ACCB - ACCA → Wd (see text); ACCA → ACCB; 0 → Z; 1 → N; 0 → OV; ) ELSE ( 0 → Wd; 1 → Z; 0 → N; 0 → OV; ) )					
Status Affected:		N, OV, Z						
Encoding:		1100      10AL      wwww      dddd      UUUq      qqUU      UUUU      0011						

**Instruction Descriptions (E to MULUU) (continued)****MIN****Accumulator Force Minimum Data Range Limit  
with Limit Excess Result (unconditional execution)**

Description:	<p>Clamp the target accumulator (defined in the instruction) to the minimum limit value previously loaded into the other accumulator.</p> <p>The compare examines the full 72-bit value of the target accumulator and will therefore clamp an overflowed accumulator (variable saturation).</p> <p>If the target accumulator is less than the limit accumulator, load the target accumulator with the contents of the limit accumulator. For MIN (instruction bit field V = 0, default when not declared), set Wd to -1. For MIN.v (instruction bit field V = 1), write the (signed) value by which the limit is exceeded to Wd:</p> <ul style="list-style-type: none"> <li>- Word operation: consider only bits [47:32] of the compare result and write to Wd (destination is also word sized, see note 3). If the limit is exceeded by a value greater than that which can be represented by a signed 16-bit number, saturate the Wd write to the maximum 16-bit negative value (i.e., set Wd[15:0] to 0x8000).</li> <li>- Long word operation: consider only bits [31:0] of the compare result and write to Wd. If the limit is exceeded by a value greater than that which can be represented by a signed 32-bit number, saturate the Wd write to the maximum negative value (i.e., set Wd to 0x80000000).</li> </ul> <p>The Z and N status bits are set such that a subsequent BLT instruction will take a branch if the limit is exceeded.</p> <p>If the target accumulator is not less than the limit accumulator, the target accumulator is unaffected, and Wd is cleared. The Z status bit is set such that a subsequent BZ instruction will take a branch.</p> <p>The OV status bit is always cleared by this instruction.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'A' bit specifies the destination accumulator.</p> <p>The 'd' bits select the destination register.</p> <p>The 'q' bits select the destination addressing mode.</p> <p>The 'w' bits define the offset Wb.</p> <p>The 'V' bit defines the result format for Wd.</p>
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. OA and SA or OB and SB status bits are not modified by this instruction. Execute SFTAC &lt;ACCx&gt;, 0 after MINABW execution to update DSP status to reflect contents of ACCx.</li> <li>2. In keeping with all word-sized register direct writes, Wd[15:0] will always be zero extended to 32-bits (irrespective of the sign of the result).</li> <li>3. Register direct destination W15 not permitted.</li> </ol>
I-Words:	1
Cycles:	2

<b>MOV</b>	<b>Move f to Wnd (Word or Long)</b>							
Syntax:	{label:}	MOV{.w}	f,	Wnd				
		MOV.I						
Operands:	$f \in [0 \dots 1MB]$ ; $Wnd \in [W0 \dots W15]$							
Operation:	$(f) \rightarrow Wnd$							
Status Affected:	None							
Encoding:	1001	00dd	ffff	ffff	ffff	ffff	ffffL	dd01
Description:	<p>Moves the contents of any file register to the specified W register. The file address is a word address. Word data will be zero-extended to 32-bits prior to destination write.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 'd' bits select the destination register.</p>							
<b>Notes:</b>	<ol style="list-style-type: none"> <li>Accessible file address space is 1MB.</li> <li>The file address is always either word or long word aligned. The opcode does not store the LSb because it is always 1'b0.</li> </ol>							
I-Words:	1							
Cycles:	1							

<b>MOV</b>	<b>Move f to Wnd (Byte)</b>							
Syntax:	{label:}	MOV.bz	f,		Wnd			
Operands:	$f \in [0 \dots 1\text{MB}]$ ; $\text{Wnd} \in [\text{W}0 \dots \text{W}14]$							
Operation:	$(f) \rightarrow \text{Wnd}$							
Status Affected:	None							
Encoding:	1001	10dd	ffff	ffff	ffff	ffff	ffff	dd01
Description:	<p>Moves contents of any file register to the specified W register. The file address is a byte address. The byte data will be zero-extended to 32-bits prior to destination write.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 'd' bits select the destination register.</p>							
I-Words:	1							
Cycles:	1							

<b>MOV</b>	<b>Move Wns to f (Byte)</b>						
Syntax:	{label:}	MOV.b	Wns	f			
Operands:	$f \in [0 \dots 1\text{MB}]$ ; $Wns \in [W0 \dots W15]$						
Operation:	$Wns \rightarrow (f)$						
Status Affected:	None						
Encoding:	1001	11ss	ffff	ffff	ffff	ffff	ss01
Description:	Moves contents of the specified W register to any file register. The file address is a byte address. The 'f' bits select the address of the file register. The 's' bits select the source register.						
I-Words:	1						
Cycles:	1						

<b>MOV</b>	<b>Move Wns to f (Word/Long)</b>							
Syntax:	{label:}	MOV{.w}	Wns	f				
		MOV.I						
Operands:	$f \in [0 \dots 1\text{MB}]$ ; $Wns \in [W0 \dots W15]$							
Operation:	$Wns \rightarrow (f)$							
Status Affected:	None							
Encoding:	1001	01ss	ffff	ffff	ffff	ffff	ffffL	ss01
Description:	<p>Moves contents of the specified W register to any file register. The file address is a word address. The 'L' bit selects word or long word operation.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the source register.</p> <p><b>Note:</b> The file address is always either word or long word aligned. The opcode does not store the LSb because it is always 1'b0.</p>							
I-Words:	1							
Cycles:	1							

<b>MOV</b>	<b>Move Wns with Signed Literal Offset to Wnd</b>							
Syntax:	{label:}	MOV.b	[Wns+Slit12], Wnd					
		MOV.bz	[Wns+Slit12], Wnd					
		MOV{.w}	[Wns+Slit13], Wnd					
		MOV.l	[Wns+Slit14], Wnd					
Operands:		Wns ∈ [W0 ... W15]						
		Wnd ∈ [W0 ... W15]						
Operation:		[Wns+Slit12/13/14] → Wnd						
Status Affected:		None						
Encoding:	1111	110L	dddd	ssss	kkkk	kkkk	kkkk	B011
Description:	<p>Moves contents of the source Effective Address to a specified W register. The contents of Wns are not modified by this operation.</p> <p>The instruction encoding includes space for a 12-bit literal which is scaled accordingly for byte, word and long word data moves in order to generate the corresponding byte offset value.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'd' bits select the destination register.</p> <p>The 's' bits select the source register.</p> <p>The 'k' bits specify the literal operand.</p>							
I-Words:	1							
Cycles:	1							

<b>MOV</b>							
<b>Move Ws to f (extended)</b>							
Syntax:	{label:}	MOV.I	Ws,	f			
		MOV{.w}	[Ws],				
		MOV{.b}	[Ws++],				
			[Ws-],				
			[++Ws],				
			[--Ws],				
Operands:	$f \in [1MB^1 \dots 4GB]$ ; $Ws \in [W0 \dots W15]$						
Operation:	$Ws \rightarrow (f)$						
Status Affected:	None						
Encoding:							
1st word	1111	001U	ffff	ffff	ffff	ffff	U011
2nd word	1111	001L	ssss	Ufff	pppf	ffff	ffff
Description:	<p>Moves contents of Ws to any file register. The file address is a byte address.</p> <p>The 's' bits select the Working register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'f' bits select the address of the file register:</p> <p>file[31:0] = Word1[18:13], Word2[9:4], Word1[23:4]</p>						
I-Words:	2						
Cycles:	2						

<b>MOV</b>		<b>Move Wns to Wnd with Signed Literal Offset</b>															
Syntax:	{label:}	MOV.b      Wns, [Wnd+Slit12] MOV{.w}    Wns, [Wnd+Slit13] MOV.l      Wns, [Wnd+Slit14]															
Operands:	Wns ∈ [W0 ... W15] Wns ∈ [W0 ... W15]																
Operation:	Wns → [Wnd+Slit12/13/14]																
Status Affected:	None																
Encoding:	1111	110L	dddd	ssss	kkkk	kkkk	kkkk	B111									
Description:	Moves contents of a specified W register to the destination Effective Address. The contents of Wnd are not modified by this operation. The instruction encoding includes space for a 12-bit literal which is scaled accordingly for byte, word and long word data moves in order to generate the corresponding byte offset value. The 'L' and 'B' bits select operation data width. The 'd' bits select the destination register. The 's' bits select the source register. The 'k' bits specify the literal operand.																
I-Words:	1																
Cycles:	1																

MOV	Move Ws to Wd						
Syntax:	{label:}	MOV.b	Ws,	Wd			
		MOV.bz	[Ws],	[Wd]			
		MOV{.w}	[Ws++]	[Wd++]			
		MOV.l	[Ws--]	[Wd--]			
			[~Ws],	[~Wd]			
			[++Ws],	[++Wd]			
			[Ws+Wb], [Wd+Wb]				
		SR	SR				
Operands:	Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]; Wb ∈ [W0 ... W14] (see Note 6)						
Operation:	(EAs) → (EAd)						
Status Affected:	None (see note 3)						
Encoding:	S000      001L      dddd      ssss      pppq      qqww      wwUU      B000						
Description:	Move the contents of the source register into the destination register. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode. The 'w' bits define the offset Wb.						
Notes:	1. When targeting the SR, the SR will be modified as a result of the instruction operation. 2. .bz data size/mode is disallowed when writing to the SR. 3. If register offset addressing is used for both source and destinations addressing modes, Wb cannot be set to W15.						
I-Words:	1 or 0.5						
Cycles:	1						

Example 1:	MOV.B [W0--], W4 ; Move [W0] to W4 (Byte mode) ; Post-decrement W0							
	Before Instruction				After Instruction			
W0	0A01				W0	0A00		
W4	2976				W4	2989		
Data 0A00	8988				Data 0A00	8988		
SR	0000				SR	0000		

<b>MOV</b>		<b>Move Literal to Wnd</b>						
Syntax:	{label:}	MOV.bz	lit8		Wnd			
		MOV{.w}	lit16					
		MOV.sl	lit24					
Operands:		lit8 ∈ [0 ... 255]; lit16 ∈ [0 ... 65535]; lit24 ∈ [0 ... 16MB]; Wnd ∈ [W0 ... W15]						
Operation:		8'h00, lit24 → Wnd						
Status Affected:		None						
Encoding:	10dd	ddkk	kkkk	kkkk	kkkk	kkkk	kkkk	kk11
Description:	<p>The literal 'k' is zero-extended to 32-bits then loaded into the Wnd register.            The 'd' bits select the Working register.</p> <p>The 'K' bits specify the value of the 8-bit, 16-bit or 24-bit literal. The 8-bit and 16-bit literals are zero extended by the assembler to generate the 24-bit value for the opcode.</p>							
I-Words:	1							
Cycles:	1							

<b>MOV</b>	<b>Move Long Literal to Wd</b>							
Syntax:	{label:}	MOV.I	lit32,	Wd				
				[Wd]				
				[Wd++]				
				[Wd--]				
				[++Wd]				
				[--Wd]				
Operands:	$\text{lit32} \in [0 \dots 4\text{GB}]$ ; $\text{Wd} \in [\text{W}0 \dots \text{W}15]$							
Operation:	$\text{lit32} \rightarrow \text{Wd}$							
Status Affected:	None							
Encoding:								
1st word	1000	000U	kkkk	kkkk	kkkk	kkkk	kkkk	U001
2nd word	1000	000U	dddd	Ukkk	kkkq	qqkk	kkkk	U101
Description:	<p>The literal 'k' is loaded into the Wd register.            The 'd' bits select the Working register.            The 'q' bits select the destination addressing mode.            The 'k' bits specify the value of the 32-bit literal:  <math>\text{lit32}[31:0] = \text{Word}2[18:13], \text{Word}2[9:4], \text{Word}1[23:4]</math></p>							
I-Words:	2							
Cycles:	2							

MOV	Move Short Literal to Wnd			
Syntax:	{label:}	MOV.I	lit5,	Wnd
Operands:	$\text{lit5} \in [0 \dots 31]$ ; $\text{Wnd} \in [\text{W0} \dots \text{W14}]$			
Operation:	27'h00000000, lit5 → Wnd			
Status Affected:	None			
Encoding:	0001	000k	kkkk	dddd
Description:	<p>The literal 'k' is zero-extended to 32-bits then loaded into the Wnd register. The 'd' bits select the Working register. The 'k' bits specify the value of the 5-bit literal.</p>			
I-Words:	0.5			
Cycles:	1			

MOV	Move Ws to Wd with Bit-Reversed Addressing							
Syntax:	{label:}	MOVR{.w} Ws,		[Wd++]				
		MOVR.I	[Ws],	[++Wd]				
			[Ws++],					
			[++Ws],					
Operands:		Ws ∈ [W0 ... W14]; Wd ∈ [W0 ... W14]						
Operation:		(EAs) → (EAd) with Bit-Reversed Addressing						
Status Affected:		None						
Encoding:	1000	001L	dddd	ssss	pppq	qqUU	UUUU	1100
Description:	<p>Move the contents of the source address to the destination address using Bit-Reversed Addressing to generate the destination EA. The destination Bit-Reversed Addressing modifier is sourced from XBREV.XB[14:0].</p> <p>The 'L' bit selects operation data width.</p> <p>The 's' bits select the source register.</p> <p>The 'd' bits select the destination register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'q' bits select the destination addressing mode.</p>							
<b>Note:</b>	<ol style="list-style-type: none"> <li>Ws = Wd only permitted when Source Addressing mode is register indirect without modification ([Ws]).</li> </ol>							
I-Words:	1							
Cycles:	1							

MOV	Move from Stack with Literal Offset to Wnd			
Syntax:	{label:}	MOV.I	[W15-lit7], Wnd	[W14+slit7], Wnd
Operands:	Wnd ∈ [W0 ... W14]; lit7 ∈ [4, 8 .... 124, 128] (see note 2) slit7 ∈ [-64, -60 .... +56, +60]			
Operation:	([W15 - lit7]) → Wnd or ([W14 + slit7]) → Wnd			
Status Affected:	None			
Encoding:	0000      11Fk      kkkk      dddd			
Description:	Move the contents of the system stack to the destination register using the SP (W15) or FP (W14) as source address with a (byte) offset. The stack source Effective Address is defined as [W15 - lit7] or [W14 + slit7], but is aligned to a 32-bit boundary (i.e., long word data moves only). The offset used by the CPU is a signed value generated from the 5-bit literal stored within the opcode. The Effective Address is long word aligned such that the LS 2-bits of the address are always 2'b00.			
	<ul style="list-style-type: none"> <li>When W15 is the source address register, the offset is always negative so the sign bit is implied (i.e., not stored within the opcode) as 1'b1. The assembler will generate the 2's complement of lit7, truncate it, and drop the sign bit to create the 5-bit opcode literal value. An offset of zero cannot be supported.</li> <li>When W14 is the source address register, the offset may be positive or negative, so the sign bit is stored within the 5-bit literal. The assembler will truncate Slit7 to create the 5-bit opcode literal value.</li> </ul> <p>The 'd' bits select the destination register. The 'F' bit selects between W15 (F = 0) and W14 (F = 1). The 'k' bits define a 5-bit signed literal.</p>			
	<b>Note:</b>			
	1. This instruction operates in Long Word mode only.			
I-Words:	0.5			
Cycles:	1			

MOV	Move Wns to Stack with Literal Offset			
Syntax:	{label:}	MOV.I	Wns,	[W15 - lit7] [W14 + slit7]
Operands:	Wns $\in$ [W0 ... W14]; lit8 $\in$ [4, 8 .... 124, 128] (see note 2) slit7 $\in$ [-64, -60 .... +56, +60]			
Operation:	Wns $\rightarrow$ ([W15 - lit7]) or Wns $\rightarrow$ ([W14 + slit7])			
Status Affected:	None			
Encoding:	0000	10Fk	kkkk	ssss
Description:	<p>Move the contents of the system stack to the destination register using the SP (W15) or FP (W14) as destination address with a (byte) offset. The stack destination Effective Address is defined as [W15 - lit7] or [W14 + slit7], but is aligned to a 32-bit boundary (i.e., long word data moves only). The offset used by the CPU is a signed value generated from the 5-bit literal stored within the opcode. The Effective Address is long word aligned such that the LS 2-bits of the address are always 2'b00.</p> <ul style="list-style-type: none"> <li>When W15 is the destination address register, the offset is always negative so the sign bit is implied (i.e., not stored within the opcode) as 1'b1. The assembler will generate the 2's complement of lit7, truncate it, and drop the sign bit to create the 5-bit opcode literal value. An offset of zero cannot be supported.</li> <li>When W14 is the destination address register, the offset may be positive or negative, so the sign bit is stored within the 5-bit literal. The assembler will truncate Slit7 to create the 5-bit opcode literal value.</li> </ul> <p>The 's' bits select the source register.  The 'F' bit selects between W15 (F = 0) and W14 (F = 1).  The 'k' bits specify the literal operand.</p> <p><b>Note:</b></p> <ol style="list-style-type: none"> <li>This instruction operates in Long Word mode only.</li> </ol>			
I-Words:	0.5			
Cycles:	1			

<b>MOVIF</b>						Conditionally Move Wb or Wns to Wd					
Syntax:	{label:}	MOVIF.b	CC,	Wb,	Wns,	Wd					
		MOVIF.bz				[Wd]					
		MOVIF{.w}				[Wd++]					
		MOVIF.I				[Wd--]					
						[-Wd]					
						[++Wd]					
Operands:		Wb ∈ [W0 ... W15]; Wns ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]									
		CC ∈ Z, N, C, OV, GT, LT, GTU									
Operation:		If CC = true then									
		Wb → (EAd)									
		Else									
		Wns → (EAd)									
Status Affected:		None									
Encoding:	1111	011L	dddd	ssss	rrrq	qqww	wwUU	B011			
Description:		Test the specified condition. If true, move the contents of the Wb source register into the destination, else move the contents of the Wns source register into the destination. The 'L' and 'B' bits select operation data width. The 'w' bits select the (condition true) Wb source register. The 's' bits select the (condition false) Ws source register. The 'd' bits select the destination register. The 'q' bits select the destination addressing mode. The 'r' bits select the move condition.									
I-Words:	1										
Cycles:	1										
<b>Example:1</b>	<b>MOVIF</b>	Z, W5, W6, W1				;Move W5 to W1 if Z == 1; else move W6 to W1					
		Before execution						After execution			
	W5	0x11111111				W5	0x11111111				
	W6	0x22222222				W6	0x22222222				
	W1	0x00000000				W1	0x11111111				
	SR[7:0]	8'bxxxxxxxxx1x	(Z = 1)			SR[7:0]	8'bxxxxxxxxx1x	(Z = 1)			
<b>Example:2</b>	<b>MOVIF</b>	C, W5, W6, W1				;Move W5 to W1 if C == 1; else move W6 to W1					
		Before execution						After execution			
	W5	0x11111111				W5	0x11111111				
	W6	0x22222222				W6	0x22222222				
	W1	0x00000000				W1	0x22222222				
	SR[7:0]	8'bxxxxxxxxx0	(C = 0)			SR[7:0]	8'bxxxxxxxxx0	(C = 0)			

<b>MOVS</b>		<b>Move Signed Literal to Wd</b>													
Syntax:	{label:}	MOVS.b	slit16,	Wd											
		MOVS{.w}		[Wd]											
		MOVS.l		[Wd++]											
				[Wd--]											
				[++Wd]											
				[--Wd]											
				SR											
Operands:		slit16 ∈ [-32768 ... 32767]; Wd ∈ [W0 ... W15]													
Operation:		Byte: slit16[7:0] → (EAd[7:0]) Word: slit16[15:0] → (EAd[15:0]) Long or Word Register Direct: {16{slit16[15]}}, slit16[15:0] → (EAd[31:0])													
Status Affected:		None (see note 4)													
Encoding:		1000	1kkk	ddd	kkkk	kkkq	qqkk	kkkk LB01							
Description:		Write a signed literal value to the destination. MOVS.b will move the LS 8-bits of the literal to the destination. For Register Direct Addressing, the remaining bits of the register will be unaffected. MOVS.w will move the 16-bit literal to the destination (see Note 1). For Register Direct Addressing, the value will be sign-extended to 32-bits. MOVS.l will sign extend the literal to 32-bits prior to moving the value to the destination. The 'L' and 'B' bits select operation data width. The 'd' bits select the Working register. The 'q' bits select the destination addressing mode. The 'k' bits specify the value of the literal. For .b mode, the 8-bit literal value is zero extended by the assembler to generate the 16-bit value for the opcode.													
<b>Notes:</b>															
1. For byte sized data, the literal can be considered to be a signed or unsigned value because only an 8-bit value is written. The declared slit16 literal value must be between 0x0000 and 0x00FF. 2. For word size data using any Indirect Addressing mode, the literal can be considered to be a signed or unsigned value because only a 16-bit value is written. 3. Long Word and Word Register Direct Addressing modes are equivalent operations because both sign-extend the literal to 32-bits. 4. When targeting the SR, the SR will be modified as a result of the instruction operation.															
I-Words:		1													
Cycles:		1													

MPY							Multiply to Accumulator						
Syntax: {label:}	MPY{.w}	Wx	,	Wy,	A	{,AWB}							
	MPY.I	[Wx]	,	[Wy],	B								
		[Wx]+=kx	,	[Wy]+=ky,									
		[Wx]=-kx	,	[Wy]=-ky,									
		[Wx+=kx]	,	[Wy+=ky],									
		[Wx-=kx]	,	[Wy-=ky],									
		[Wx+W12]	,	[Wy+W12],									
Operands:		$Wx \in \{W0 \dots W14\}; Wy \in \{W0 \dots W14\}$ Word mode: $kx \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ; $ky \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ; Long Word mode: $kx \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ; $ky \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ; $AWB \in \{W0, W1, W2, W3, W13, [W13++], [W15++]\}^3$											
Operation:		$((Wx) * (Wy)) \rightarrow ACC(A \text{ or } B);$ (ACC(B or A)) rounded $\rightarrow AWB$											
		When Indirect Pre/Post Modification Addressing: $(Wx)+kx \rightarrow Wx$ or $(Wx)-kx \rightarrow Wx$ ; $(Wy)+ky \rightarrow Wy$ or $(Wy)-ky \rightarrow Wy$ ;											
Status Affected:		OA,SA or OB,SB											
Encoding:		1101	01AL	www	sss	III	iJJJ	jjaa	a011				
Description:		Signed/unsigned (defined by CORCON.US) multiply of data read from Wx and Wy or concurrently fetched from the X and Y address space. The result is sign-extended (when at least one operand is considered signed) or zero-extended (when both operands are unsigned) to 72-bits then written to the specified accumulator. Fractional results are also scaled prior to the accumulator update to align the operand and accumulator (msw) fractional points. When indirect addressing is selected for either or both the X and Y address space, Wx and Wy registers provide the corresponding indirect addresses. The address modifier values are kx and ky, respectively, and represent the number of data bytes by which to modify the Effective Address. The optional AWB specifies the direct or indirect (see note 2) store of the (32-bit) rounded fractional contents of the accumulator not targeted by the MPY operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size (see note 3). AWB is not intended for use when the DSP engine is operating in Integer mode. Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly. The 'L' bit selects word or long word operation. The 'A' bit selects the accumulator for the result. The 'I' bits select the Operation.X-Space Addressing mode. The 'i' bits select the kx modification value. The 'J' bits select the Operation.Y-Space Addressing mode. The 'j' bits select the ky modification value. The 's' bits select the Wx register The 'w' bits select the Wy register. The 'a' bits select the accumulator write-back destination and addressing mode.											
		<b>Notes:</b>											
		1. Operates in Fractional or Integer Data mode as defined by CORCON.IF. 2. Use of the same W-reg for both indirect source (X or Y) and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address. 3. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.											
I-Words:		1											
Cycles:		1											

Example:1	MPY.L	[w4]+=4, [w5]+=4, A	; Multiply W4*W5 and store to ACCA; Post-increment W4 by 4; Post-increment W5 by 4; CORCON = 0x0000 (fractional multiply, no saturation)
-----------	-------	------------------------	--

Before Instruction		After Instruction	
W4	4014	W4	4018
W5	402C	W5	4030
ACCA	FF F780 2087 0000 0000	ACCA	00 3800 0000 0000 0000
Data	4014 C0000000	Data	4014 C0000000
Data	402C 90000000	Data	402C 90000000
CORCON	0000	CORCON	0000
SR	0000	SR	0000

Example:2	MPY.L	[w4]+=4, [w5+=4], B	; increment W5 by 4; Multiply W4*W5 and store to ACCB; Post-increment W4 by 4; CORCON = 0x0000 (fractional multiply, no saturation)
-----------	-------	------------------------	---

Before Instruction		After Instruction	
W4	4014	W4	4018
W5	4028	W5	402C
ACCB	00 9834 4500 0000 0055	ACCB	FF E954 3748 0000 0000
Data	4014 671F0000	Data	4014 671F0000
Data	402C E3DC0000	Data	402C E3DC0000
CORCON	0000	CORCON	0000
SR	0000	SR	0000

<b>MPYN</b>		<b>Negated Multiply to Accumulator</b>				
Syntax: {label:}	MPYN{.w} Wx	, Wy,	A	{,AWB}		
MPYN.I	[Wx]	, [Wy],	B			
	[Wx]+=kx	, [Wy]+=ky,				
	[Wx]=-kx	, [Wy]=-ky,				
	[Wx+=kx]	, [Wy+=ky],				
	[Wx-=kx]	, [Wy-=ky],				
	[Wx+W12]	, [Wy+W12],				
Operands:	$Wx \in \{W0 \dots W14\}; Wy \in \{W0 \dots W14\}$					
	Word mode: $kx \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ;					
	$ky \in \{-8, -6, -4, -2, 2, 4, 6, 8\}$ ;					
	Long Word mode: $kx \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ;					
	$ky \in \{-16, -12, -8, -4, 4, 8, 12, 16\}$ ;					
	$AWB \in \{W0, W1, W2, W3, W13, [W13++], [W15++]\^2\}$					
Operation:	<ul style="list-style-type: none"> <li>- <math>((Wx) * (Wy)) \rightarrow ACC(A \text{ or } B)</math>;</li> <li>(ACC(B or A)) rounded <math>\rightarrow AWB</math></li> </ul>					
	When Indirect Pre/Post Modification Addressing:					
	$(Wx)+kx \rightarrow Wx$ or $(Wx)-kx \rightarrow Wx$ ;					
	$(Wy)+ky \rightarrow Wy$ or $(Wy)-ky \rightarrow Wy$ ;					
Status Affected:	OA,SA or OB,SB					
Encoding:	1101	01AL	www	sss	III	iJJJ jjaa a111
Description:	<p>Signed/unsigned (defined by CORCON.US) multiply of data read from Wx and Wy or concurrently fetched from the X and Y address space. The result is sign-extended (when at least one operand is considered signed) or zero-extended (when both operands are unsigned) to 72-bits, negated and then written to the specified accumulator. Fractional results are also scaled prior to the accumulator update to align the operand and accumulator (msw) fractional points.</p> <p>When indirect addressing is selected for either or both the X and Y address space, Wx and Wy registers provide the corresponding indirect addresses. The address modifier values are kx and ky, respectively, and represent the number of data bytes by which to modify the Effective Address.</p> <p>The optional AWB specifies the direct or indirect store of the (32-bit) rounded fractional contents of the accumulator not targeted by the MPYN operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size (see note 3). AWB is not intended for use when the DSP engine is operating in Integer mode.</p> <p>Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'A' bit selects the accumulator for the result.</p> <p>The 'I' bits select the Operation.X-Space Addressing mode.</p> <p>The 'i' bits select the kx modification value.</p> <p>The 'J' bits select the Operation.Y-Space Addressing mode.</p> <p>The 'j' bits select the ky modification value.</p> <p>The 's' bits select the Wx register</p> <p>The 'w' bits select the Wy register.</p> <p>The 'a' bits select the accumulator write-back destination and addressing mode.</p>					
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.</li> <li>2. Use of the same W-reg for both indirect source (X or Y) and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.</li> <li>3. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.</li> </ol>					
I-Words:	1					
Cycles:	1					

<b>MSC</b>		<b>Multiply and Subtract from Accumulator</b>				
Syntax: {label:}	MSC{.w}	Wx	, Wy,	A	{,AWB}	
	MSC.I	[Wx]	, [Wy],	B		
		[Wx]+=kx	, [Wy]+=ky,			
		[Wx]=-kx	, [Wy]=-ky,			
		[Wx+=kx]	, [Wy+=ky],			
		[Wx-=kx]	, [Wy-=ky],			
		[Wx+W12]	, [Wy+W12],			
Operands:	$Wx \in \{W0 \dots W14\}; Wy \in \{W0 \dots W14\}$					
	Word mode: $kx \in \{-8, -6, -4, -2, 2, 4, 6, 8\};$					
	$ky \in \{-8, -6, -4, -2, 2, 4, 6, 8\};$					
	Long Word mode: $kx \in \{-16, -12, -8, -4, 4, 8, 12, 16\};$					
	$ky \in \{-16, -12, -8, -4, 4, 8, 12, 16\};$					
	$AWB \in \{W0, W1, W2, W3, W13, [W13++], [W15++]^3\}$					
Operation:	$(ACC(A \text{ or } B)) - ((Wx) * (Wy)) \rightarrow ACC(A \text{ or } B);$					
	$(ACC(B \text{ or } A)) \text{ rounded} \rightarrow AWB$					
	When Indirect Pre/Post Modified Addressing:					
	$(Wx)+kx \rightarrow Wx \text{ or } (Wx)-kx \rightarrow Wx;$					
	$(Wy)+ky \rightarrow Wy \text{ or } (Wy)-ky \rightarrow Wy;$					
Status Affected:	OA,SA or OB,SB					
Encoding:	1101	00AL	wwww	ssss	IIii	iJJJ jjaa a111

**Instruction Descriptions (E to MULUU) (continued)****MSC****Multiply and Subtract from Accumulator****Description:**

Signed/unsigned (defined by CORCON.US) multiply of data read from Wx and Wy or concurrently fetched from the X and Y address space. The result is sign-extended (when at least one operand is considered signed) or zero-extended (when both operands are unsigned) to 72-bits, and then subtracted from the specified accumulator. Fractional or integer operation (defined by CORCON.IF) will determine if the result is scaled or not prior to the accumulator update. Fractional operation will scale the result to align the operand and accumulator (rsw) fractional points. Integer operation will align the LSb of the result with the LSb of the accumulator.

For word sized operand operations, ACCx[31:0] (Integer mode) or ACCx[32:0] (Fractional mode) is preserved (see note 2).

When indirect addressing is selected for either or both the X and Y address space, Wx and Wy registers provide the corresponding indirect addresses. The address modifier values are kx and ky, respectively, and represent the number of data bytes by which to modify the Effective Address.

The optional AWB specifies the direct or indirect store of the (32-bit) rounded fractional contents of the accumulator not targeted by the MSC operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size (see note 4). AWB is not intended for use when the DSP engine is operating in Integer mode.

Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly.

The 'L' bit selects word or long word operation.

The 'A' bit selects the accumulator for the result.

The 'I' bits select the Operation.X-Space Addressing mode.

The 'i' bits select the kx modification value.

The 'J' bits select the Operation.Y-Space Addressing mode.

The 'j' bits select the ky modification value.

The 's' bits select the Wx register

The 'w' bits select the Wy register.

The 'a' bits select the accumulator write-back destination and addressing mode.

**Notes:**

1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.
2. The LS portion of ACCx is unaffected when operating in Fractional mode with word sized data. Lower significance data that may be present from prior (32-bit data) operations is therefore preserved. Users not requiring this should clear ACCx during initialization.
3. Use of the same W-reg for both indirect source (X or Y) and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.
4. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.

**I-Words:**

1

**Cycles:**

1

<b>Example:1</b>	<b>MSC.L</b>	<b>[W6]=-4, [W7]=-4, A, W7</b>	<b>; Multiply [W6]*[W7] and subtract theresult from ACCA ; W6 Post-decrement by 4 ;Post-decrement W7 by 4 ; CORCON = 0x0001 (integer multiply, no saturation)</b>
------------------	--------------	--------------------------------	---

Before Instruction		After Instruction	
W6	4C00 0000	W6	4BFC 0000
W7	4800 0000	W7	48FC 0000
W8	0C00 0000	W8	0BFC 0000
ACCA	00 0567 8000 0000 0000	ACCA	00 3738 5ED0 0000 0000
Data	4C00 9051 0000	Data	4BFC D309 0000
Data	4800 7230 0000	Data	48FC 100B 0000
CORCON	0001	CORCON	0001

## Instruction Descriptions (E to MULUU) (continued)

Example:1	MSC.L	[W6]=-4, [W7]-=4, A, W7	; Multiply [W6]*[W7] and subtract the result from ACCA ; W6 Post-decrement by 4 ; Post-decrement W7 by 4 ; CORCON = 0x0001 (integer multiply, no saturation)	
		SR 0000	SR	0000

Example:2	MSC.L	[W4+=12], [W5+w12], B, [W13]	; Multiply W4*W5 and subtract the result from ACCB ; Fetch [W11+W12] to W5 ; Write Back ACCA to W13 ; CORCON = 0x0000 (fractional multiply, no saturation)	
		Before Instruction	After Instruction	

W4	4000	W4	4000
W5	5000	W5	5000
W12	0800	W12	0800
W13	6233 1356	W13	3738 5ED0
ACCA	00 3738 5ED0 0000 0000	ACCA	00 3738 5ED0 0000 0000
ACCB	00 1000 0000 0000 0000	ACCB	00 0EC0 0000 0000 0000
Data	4800 0500 0000	Data	4800 0500 0000
Data	5800 2000 0000	Data	5800 2000 0000
CORCON	0000	CORCON	0000
SR	0000	SR	0000

<b>MUL</b>		Multiply f by Wn																	
Syntax:		{label:}		MUL.b	f,	Wns													
		MUL{.w}																	
		MUL.i																	
Operands:		$f \in [0 \dots 64KB]; Wn \in [W0 \dots W14]$																	
Operation:		If Byte mode: $(Wn)[7:0] * (f)[7:0] \rightarrow 16'b0, W2[15:0]$ . If Word mode: $(Wn)[15:0] * (f)[15:0] \rightarrow W2$ If Long Word mode: $(Wn) * (f) \rightarrow W2$																	
Status Affected:		None																	
Encoding:		1011	101L	ssss	ffff	ffff	ffff	ffff	BU01										
Description:		Unsigned integer multiply of Wn and the file register, then write the result to the default destination register W2. For long word operations, the LS 32-bits of the 64-bit result is written to W2. The 'L' and 'B' bits select operation data width. The 'f' bits select the address of the file register.																	
I-Words:		1																	
Cycles:		1																	

<b>MULISS/MULFSS</b>		<b>Signed-Signed Multiply to Accumulator</b>						
Syntax:	{label:}	MULISS.w Wb, MULFSS.w MULISS.I MULFSS.I	Ws, [Ws], [Ws++], [Ws--], [+Ws], [-Ws],	A	B			
Operands:		$Wb \in [W0 \dots W14]; Ws \in [W0 \dots W15]$						
Operation:		Signed (Wb) * signed (Ws) $\rightarrow$ ACC(A or B)[63:0] 8{ACC(A or B)[63]} $\rightarrow$ ACC(A or B)[71:64]						
Status Affected:		None						
Encoding:		1111      100L      wwww      ssss      pppA      UUUU      UUUU      I010						
Description:		<p>Performs word or long word integer (MULISS) or fractional (MULFSS) multiply of the signed contents of Wb and Ws. The source operands are interpreted as a two's-complement signed values.</p> <p>For MULISS.I, the 64-bit result will be sign-extended and written to ACCx[71:0].</p> <p>For MULFSS.I, the 64-bit result will be sign-extended, shifted left by one (to align the result and accumulator fractional points), and written to ACCx[71:0]. ACCx[0] will always be cleared.</p> <p>For MULISS.w, the 32-bit result will be sign-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>For MULFSS.w, the 32-bit result will be sign-extended, shifted left by 33 bits (to align the result and accumulator fractional points), and written to ACCx[71:33]. ACCx[32:0] will always be cleared.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'I' bit selects between integer and fractional operation.</p> <p>The 'A' bit selects the destination accumulator.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'p' bits select the source addressing mode.</p> <p><b>Note:</b> The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.</p>						
I-Words:		1						
Cycles:		1						

<b>MULISU/MULFSU</b>		<b>Signed-Unsigned Multiply to Accumulator</b>					
Syntax:	{label:}	MULISU.w Wb, MULFSU.w MULISU.I MULFSU.I	Ws, [Ws], [Ws++], [Ws--], [+Ws], [-Ws],	A	B		
Operands:		$Wb \in [W0 \dots W14]; Ws \in [W0 \dots W15]$					
Operation:		$\text{signed}(Wb) * \text{unsigned}(Ws) \rightarrow ACC(A \text{ or } B)[63:0]$ $8\{ACC(A \text{ or } B)[63]\} \rightarrow ACC(A \text{ or } B)[71:64]$					
Status Affected:		None					
Encoding:		1111	101L	wwww	ssss	pppA	UUUU UUUU I110
Description:		<p>Performs a word or long word integer (MULISU) or fractional (MULFSU) multiply of the signed contents of Wb and unsigned contents of Ws. Wb is interpreted as a two's-complement signed value. For MULISU.I, the 64-bit result will be sign-extended and written to ACCx[71:0].</p> <p>For MULFSU.I, the 64-bit result will be sign-extended, shifted left by one (to align the result and accumulator fractional points), and written to ACCx[71:0]. ACCx[0] will always be cleared.</p> <p>For MULISU.w, the 32-bit result will be sign-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>For MULFSU.w, the 32-bit result will be sign-extended, shifted left by 33 bits (to align the result and accumulator fractional points), and written to ACCx[71:33]. ACCx[32:0] will always be cleared.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'I' bit selects between integer and fractional operation.</p> <p>The 'A' bit selects the destination accumulator.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'p' bits select the source addressing mode.</p>					
Notes:		<ol style="list-style-type: none"> <li>1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.</li> <li>2. An unsigned fractional operand includes one integer bit (i.e., can be up to 1.999...). Consequently, the mixed sign fractional result includes one integer bit (i.e., can be up to 1.999...). Results equal to, or in excess of, the range +/-1.0 may require normalization before subsequent use.</li> </ol>					
I-Words:		1					
Cycles:		1					

MULISU/MULIUU	Signed/Unsigned - Unsigned Literal Multiply to Accumulator							
Syntax:	{label:}	MULISU.w Ws, MULIUU.w [Ws], MULISU.I [Ws++], MULIUU.I [Ws-], [+Ws], [-Ws],	lit8,		A	B		
Operands:		Ws ∈ [W0 ... W15]; lit8 ∈ [0 ... 255]						
Operation:		MULISU: signed (Ws) * unsigned lit8 → ACC(A or B)[63:0] 8{ACC(A or B)[63]} → ACC(A or B)[71:64] MULIUU: unsigned (Ws) * unsigned lit8 → ACC(A or B)[63:0] 8'b00 → ACC(A or B)[71:64]						
		Note: The literal is zero-extended to the selected data size of the operation						
Status Affected:		None						
Encoding:	1111      111L      UUUU      SSSS      pppA      kkkk      kkkk      1V10							
Description:	<p>Performs a word or long word integer multiply of the signed (MULISU) or unsigned (MULIUU) contents of Ws and unsigned lit8. The 'V' bit is set to 1'b1 to select a signed Ws value and set to 1'b0 to select an unsigned Ws value.</p> <p>For MULISU.I, the 64-bit result will be sign-extended and written to ACCx[71:0].</p> <p>For MULIUU.I, the 64-bit result will be zero-extended and written to ACCx[71:0].</p> <p>For MULISU.w, the 32-bit result will be sign-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>For MULIUU.w, the 32-bit result will be zero-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'V' bit selects between a signed or unsigned Ws value.</p> <p>The 'A' bit selects the destination accumulator.</p> <p>The 'S' bits select the source register.</p> <p>The 'P' bits select the source addressing mode.</p> <p>The 'K' bits determine the 8-bit literal value.</p>							
	<b>Notes:</b>							
	<ol style="list-style-type: none"> <li>1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.</li> <li>2. Fractional mode is not supported for these instructions.</li> </ol>							
I-Words:	1							
Cycles:	1							

<b>MULIUU/MULFUU</b>		<b>Unsigned Unsigned Multiply to Accumulator</b>						
Syntax:	{label:}	MULIUU.w Wb, MULFUU.w MULIUU.I MULFUU.I	Ws, [Ws], [Ws++], [Ws--], [+Ws], [-Ws],	A				
Operands:		$Wb \in [W0 \dots W14]; Ws \in [W0 \dots W15]$						
Operation:		unsigned (Wb) * unsigned (Ws) $\rightarrow$ ACC(A or B)[63:0] 8'h00 $\rightarrow$ ACC(A or B)[71:64]						
Status Affected:		None						
Encoding:		1111      101L      wwww      ssss      pppA      UUUU      UUUU      I010						
Description:		<p>Performs a word or long word integer (MULIUU) or fractional (MULFUU) multiply of the unsigned contents of Wb and unsigned contents of Ws (see Note 2).</p> <p>For MULIUU.I, the 64-bit result will be zero-extended and written to ACCx[71:0].</p> <p>For MULFUU.I, the 64-bit result will be zero-extended, shifted left by one (to align the result and accumulator fractional points), and written to ACCx[71:0]. ACCx[0] will always be cleared.</p> <p>For MULIUU.w, the 32-bit result will be zero-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>For MULFUU.w, the 32-bit result will be zero-extended, shifted left by 33 bits (to align the result and accumulator fractional points) and written to ACCx[71:33]. ACCx[32:0] will always be cleared.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'I' bit selects between integer and fractional operation.</p> <p>The 'A' bit selects the destination accumulator.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'p' bits select the source addressing mode.</p>						
Notes:		<ol style="list-style-type: none"> <li>1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.</li> <li>2. An unsigned fractional operand includes one integer bit (i.e., can be up to 1.999...). Consequently, the unsigned fractional result includes two integer bits (i.e., can be up to 3.999...). Results greater than or equal to 1.0 may require normalization before subsequent use.</li> </ol>						
I-Words:		1						
Cycles:		1						

<b>MULIUS/MULFUS</b>		<b>Unsigned-Signed Multiply to Accumulator</b>					
Syntax:	{label:}	MULIUS.w Wb, MULFUS.w MULIUS.I MULFUS.I	Ws, [Ws], [Ws++], [Ws--], [+Ws], [-Ws],	A			
Operands:		$Wb \in [W0 \dots W14]; Ws \in [W0 \dots W15]$					
Operation:		unsigned (Wb) * signed (Ws) $\rightarrow$ ACC(A or B)[63:0] 8{ACC(A or B)[63]} $\rightarrow$ ACC(A or B)[71:64]					
Status Affected:		None					
Encoding:		1111      100L      wwww      ssss      pppA      UUUU      UUUU      I110					
Description:		<p>Performs a word or long word integer (MULIUS) or fractional (MULFUS) multiply of the unsigned contents of Wb and signed contents of Ws (see Note 2).</p> <p>For MULIUS.I, the 64-bit result will be sign-extended and written to ACCx[71:0].</p> <p>For MULFUS.I, the 64-bit result will be sign-extended, shifted left by one (to align the result and accumulator fractional points), and written to ACCx[71:0]. ACCx[0] will always be cleared.</p> <p>For MULIUS.w, the 32-bit result will be sign-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>For MULFUS.w, the 32-bit result will be sign-extended, shifted left by 33 bits (to align the result and accumulator fractional points), and written to ACCx[71:33]. ACCx[32:0] will always be cleared.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'I' bit selects between integer and fractional operation.</p> <p>The 'A' bit selects the destination accumulator.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'p' bits select the source addressing mode.</p>					
Notes:		<ol style="list-style-type: none"> <li>1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.</li> <li>2. An unsigned fractional operand includes one integer bit (i.e., can be up to 1.999...). Consequently, the mixed sign fractional result includes one integer bit (i.e., can be up to 1.999...). Results equal to, or in excess of, the range +/-1.0 may require normalization before subsequent use.</li> </ol>					
I-Words:		1					
Cycles:		1					

<b>MULISS/MULIUS</b>		<b>Signed/Unsigned-Signed Literal Multiply to Accumulator</b>							
Syntax:	{label:}	MULISS.w Ws, MULIUS.w [Ws], MULISS.I [Ws++], MULIUS.I [Ws-], [++Ws], [-Ws],	slit8,	A	B				
Operands:	$Ws \in [W0 \dots W15]$ ; $slit8 \in [-128 \dots 127]$								
Operation:	MULxUS : unsigned (Ws) * signed slit8 $\rightarrow$ ACC(A or B)[63:0] MULxSS : signed (Ws) * signed slit8 $\rightarrow$ ACC(A or B)[63:0] $8\{ACC(A \text{ or } B)[63]\} \rightarrow ACC(A \text{ or } B)[71:64]$								
	<b>Note:</b> The literal is sign-extended to the selected data size of the operation								
Status Affected:	None								
Encoding:	1111	110L	UUUU	ssss	pppA	kkkk	kkkk	1V10	
Description:	<p>Performs a word or long word integer multiply of the signed (MULISS) or unsigned (MULIUS) contents of Ws and signed slit8. The 'V' bit is set to 1'b1 to select a signed Ws value and set to 1'b0 to select an unsigned Ws value.</p> <p>For MULISS.I, the 64-bit result will be sign-extended and written to ACCx[71:0].</p> <p>For MULIUS.I, the 64-bit result will be sign-extended and written to ACCx[71:0].</p> <p>For MULISS.w, the 32-bit result will be sign-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>For MULIUS.w, the 32-bit result will be sign-extended and written to ACCx[71:32]. ACCx[31:0] will hold the multiply result.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'V' bit selects between a signed or unsigned Ws value.</p> <p>The 'A' bit selects the destination accumulator.</p> <p>The 's' bits select the source register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'k' bits determine the 8-bit literal value.</p>								
	<b>Notes:</b>								
	<ol style="list-style-type: none"> <li>1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.</li> <li>2. Fractional mode is not supported for these instructions.</li> </ol>								
I-Words:	1								
Cycles:	1								

<b>MULSS/MULUS</b>		<b>Signed/Unsigned-Signed Literal Integer Multiply</b>																					
Syntax:		{label:} MULSS.d Ws, slit8, Wnd		MULSS.I [Ws],		MULSS.w [Ws++],		MULUS.d [Ws--],															
		MULUS.I [+Ws],		MULUS.w [-Ws],																			
Operands:		Ws ∈ [W0 ... W15]; slit8 ∈ [-128 ... 127]; Wnd ∈ [W0 ... W14]																					
Operation:		MULSS.d : signed (Ws) * signed slit8 → {Wnd+1[31:0], Wnd[31:0]} MULSS.I : signed (Ws) * signed slit8 → {Wnd[31:0]} MULSS.w : signed (Ws[15:0]) * signed slit8 → {Wnd[31:0]} MULUS.d : unsigned (Ws) * signed slit8 → {Wnd+1[31:0], Wnd[31:0]} MULUS.I : unsigned (Ws) * signed slit8 → {Wnd[31:0]} MULUS.w : unsigned (Ws[15:0]) * signed slit8 → {Wnd[31:0]}																					
Note: The literal is sign-extended to the selected data size of the operation																							
Status Affected:	None																						
Encoding:	1111	010L	dddd	ssss	ppPE	kkkk	kkkk	0V10															
Description:	<p>Performs a 32-bit x 32-bit or a 16-bit x 16-bit integer multiply of the signed or unsigned contents of Ws and sign-extended slit8. The 'V' bit is set to 1'b1 to select a signed Ws value and set to 1'b0 to select an unsigned Ws value.</p> <p>For MULSS.d and MULUS.d (L = 1, E = 1), a 32-bit x 32-bit multiply of Ws and sign-extended slit8 is executed. The MS 32-bits of the 64-bit result will be written to Wnd+1, and the LS 32-bits of the result will be written to Wnd.</p> <p>For MULSS.I and MULUS.I (L = 1, E = 0), a 32-bit x 32-bit multiply of Ws and sign-extended slit8 is executed. The LS 32-bits of the 64-bit result will be written to Wnd.</p> <p>For MULSS.w and MULUS.w (L = 0, E = 0), a 16-bit x 16-bit multiply of the lsw of Ws and sign-extended slit8 is executed. The result will be 32-bits and is written to Wnd.</p> <p>The 'L' bit selects the operand date size.</p> <p>The 'V' bit selects between a signed or unsigned Ws value.</p> <p>The 'E' bit selects between a 32-bit and 64-bit result write.</p> <p>The 's' bits select the source register.</p> <p>The 'd' bits select the destination register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'k' bits determine the 8-bit literal value.</p>																						
<b>Notes:</b>																							
1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.																							
2. For 32-bit result writes, user is responsible to ensure that the result does not overflow into the sign bit, Wnd[31].																							
I-Words:	1																						
Cycles:	1																						

<b>MULSS</b>		<b>Signed-Signed Integer Multiply</b>						
Syntax:	{label:}	MULSS.d	Wb,	Ws,	Wnd			
		MULSS.I		[Ws],				
		MULSS.w		[Ws++],				
				[Ws--],				
				[++Ws],				
				[--Ws],				
Operands:		Wb $\in$ [W0 ... W15]; Ws $\in$ [W0 ... W15]; Wnd $\in$ [W0 ... W14]						
Operation:		MULSS.d : signed (Wb) * signed (Ws) $\rightarrow$ {Wnd+1[31:0], Wnd[31:0]}						
		MULSS.I : signed (Wb) * signed (Ws) $\rightarrow$ {Wnd[31:0]}						
		MULSS.w : signed (Wb[15:0]) * signed (Ws[15:0]) $\rightarrow$ {Wnd[31:0]}						
Status Affected:		None						
Encoding:		S001	010L	dddd	ssss	ppPE	UUww	wwUU
Description:		Performs a 16-bit x 16-bit or 32-bit x 32-bit integer multiply of the signed contents of Wb and Ws. For MULSS.d (L=1, E=1), a 32-bit x 32-bit multiply of Wb and Ws is executed. The MS 32-bits of the 64-bit result will be written to Wnd+1, and the LS 32-bits of the result will be written to Wnd. For MULSS.I (L=1, E=0), a 32-bit x 32-bit multiply of Wb and Ws is executed. The LS 32-bits of the 64-bit result will be written to Wnd. For MULSS.w (L=0, E=0), a 16-bit x 16-bit multiply of the lsw of Wb and the lsw of Ws is executed. The result will be 32-bits and is written to Wnd.						
		Both source operands are interpreted as two's-complement integer signed values.						
		The 'S' bit selects instruction size.						
		The 'L' bit selects the operand data size.						
		The 'E' bit selects between a 32-bit and 64-bit result write.						
		The 's' bits select the source register.						
		The 'w' bits select the base register.						
		The 'd' bits select the destination register.						
		The 'p' bits select the source addressing mode.						
<b>Notes:</b>								
1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.								
2. For 32-bit result writes, user is responsible to ensure that the result does not overflow into the sign bit, Wnd[31].								
I-Words:		1 or 0.5						
Cycles:		1						

<b>MULSU</b>		Signed-Unsigned Integer Multiply					
Syntax:	{label:}	MULSU.d	Wb,	Ws,	Wnd		
		MULSU.I		[Ws],			
		MULSU.w		[Ws++],			
				[Ws--],			
				[++Ws],			
				[--Ws],			
Operands:		Wb $\in$ [W0 ... W15]; Ws $\in$ [W0 ... W15]; Wnd $\in$ [W0 ... W14]					
Operation:		MULSU.d : signed (Wb) * unsigned (Ws) $\rightarrow$ {Wnd+1[31:0], Wnd[31:0]}					
		MULSU.I : signed (Wb) * unsigned (Ws) $\rightarrow$ {Wnd[31:0]}					
		MULSU.w : signed (Wb[15:0]) * unsigned (Ws[15:0]) $\rightarrow$ {Wnd[31:0]}					
Status Affected:		None					
Encoding:	1001	011L	dddd	ssss	ppPE	UUww	wwUU
Description:	Performs a 16-bit x 16-bit or 32-bit x 32-bit integer multiply of the signed contents of Wb and unsigned contents of Ws. Wb is interpreted as a two's-complement signed value.						
	For MULSU.d (L = 1, E = 1), a 32-bit x 32-bit multiply of Wb and Ws is executed. The MS 32-bits of the 64-bit result will be written to Wnd+1, and the LS 32-bits of the result will be written to Wnd.						
	For MULSU.I (L=1, E = 0), a 32-bit x 32-bit multiply of Wb and Ws is executed. The LS 32-bits of the result will be written to Wnd.						
	For MULSU.w (L=0, E = 0), a 16-bit x 16-bit multiply of the lsw of Wb and the lsw of Ws is executed. The result will be 32-bits and is written to Wnd.						
	The 'L' bit selects the operand date size.						
	The 'E' bit selects between a 32-bit and 64-bit result write.						
	The 's' bits select the source register.						
	The 'w' bits select the base register.						
	The 'd' bits select the destination register.						
	The 'p' bits select the source addressing mode.						
	<b>Notes:</b>						
	1. The state of the Multiplier Mode bit (CORCON.US in CORCON) has no effect upon the operation of this instruction.						
	2. For 32-bit result writes, user is responsible to ensure that the result does not overflow into the sign bit, Wnd[31].						
I-Words:	1						
Cycles:	1						

<b>Example 1:</b>		MUL.SU W8, [W9], W0 ; Multiply W8*[W9] ; Store the result to W0:W1					
		Before Instruction					
		W0 68DC					
		W1 AA40					
		W8 F000					
		W9 178C					
		Data 178C F000					
		SR 0000					
		After Instruction					
		W0 0000					
		W1 F100					
		W8 F000					
		W9 178C					
		Data 178C F000					
		SR 0000					

<b>Example 2:</b>		MUL.SU W2, [++W3], W4 ; Pre-Increment W3 ; Multiply W2*[W3] ; Store the result to W4:W5					
		Before Instruction					
		W2 0040					
		W3 0280					
		After Instruction					
		W2 0040					
		W3 0282					

## Instruction Descriptions (E to MULUU) (continued)

**Example 2:**      **MUL.SU W2, [++W3], W4 ; Pre-Increment W3 ; Multiply W2\*[W3] ; Store the result to W4:W5**

W4	1819	W4	1A00
W5	2021	W5	0000
Data 0282	0068	Data 0282	0068
SR	0000	SR	0000

<b>MULSU/MULUU</b>		<b>Signed/Unsigned Unsigned Short Literal Integer Multiply</b>															
Syntax:		{label:} MULSU.d Ws, lit8, Wnd		MULSU.I [Ws],		MULSU.w [Ws++],		MULUU.d [Ws-],									
		MULUU.I [+Ws],		MULUU.w [-Ws],													
Operands:		Ws ∈ [W0 ... W15]; lit8 ∈ [0...255]; Wnd ∈ [W0 ... W14]															
Operation:		MULSU.d : signed (Ws) * unsigned lit8 → {Wnd+1[31:0], Wnd[31:0]}															
		MULSU.I : signed (Ws) * unsigned lit8 → {Wnd[31:0]}															
		MULSU.w : signed (Ws[15:0]) * unsigned lit8 → {Wnd[31:0]}															
		MULUU.d : unsigned (Ws) * unsigned lit8 → {Wnd+1[31:0], Wnd[31:0]}															
		MULUU.I : unsigned (Ws) * unsigned lit8 → {Wnd[31:0]}															
		MULUU.w : unsigned (Ws[15:0]) * unsigned lit8 → {Wnd[31:0]}															
Note: The literal is zero-extended to the selected data size of the operation																	
Status Affected:		None															
Encoding:		1111	011L	dddd	ssss	ppPE	kkkk	kkkk	0V10								
Description:		Performs a 16-bit x 16-bit or 32-bit x 32-bit integer multiply of the signed or unsigned contents of Ws and unsigned lit8. The 'V' bit is set to 1'b1 to select a signed Ws value and set to 1'b0 to select an unsigned Ws value.															
		For MULSU.d and MULUU.d (L = 1, E = 1), a 32-bit x 32-bit multiply of Ws and zero-extended lit8 is executed. The MS 32-bits of the 64-bit result will be written to Wnd+1, and the LS 32-bits of the result will be written to Wnd.															
		For MULSU.I and MULUU.I (L = 1, E = 0), a 32-bit x 32-bit multiply of Ws and zero-extended lit8 is executed. The LS 32-bits of the 64-bit result will be written to Wnd.															
		For MULSU.w and MULUU.w (L = 0, E = 0), a 16-bit x 16-bit multiply of the ls of Ws and zero-extended lit8 is executed. The result will be 32-bits and is written to Wnd.															
The 'L' bit selects the operand date size.																	
The 'V' bit selects between a signed or unsigned Ws value.																	
The 'E' bit selects between a 32-bit and 64-bit result write.																	
The 's' bits select the source register.																	
The 'd' bits select the destination register.																	
The 'p' bits select the source addressing mode.																	
The 'k' bits determine the 8-bit literal value.																	
<b>Notes:</b>																	
1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.																	
2. For signed 32-bit result writes, user is responsible to ensure that the result does not overflow into the sign bit, Wnd[31].																	
I-words:	1																
Cycles:	1																

<b>MULUS</b>		Unsigned-Signed Integer Multiply						
Syntax:	{label:}	MULUS.d	Wb,	Ws,	Wnd			
		MULUS.I		[Ws],				
		MULUS.w		[Ws++],				
				[Ws--],				
				[++Ws],				
				[--Ws],				
Operands:		Wb $\in$ [W0 ... W15]; Ws $\in$ [W0 ... W15]; Wnd $\in$ [W0 ... W14]						
Operation:		MULUS.d : unsigned (Wb) * signed (Ws) $\rightarrow$ {Wnd+1[31:0], Wnd[31:0]}						
		MULUS.I : unsigned (Wb) * signed (Ws) $\rightarrow$ {Wnd[31:0]}						
		MULUS.w : unsigned (Wb[15:0]) * signed (Ws[15:0]) $\rightarrow$ {Wnd[31:0]}						
Status Affected:		None						
Encoding:	1001	010L	dddd	ssss	ppPE	UUww	wwUU	0100
Description:	<p>Performs a 16-bit x 16-bit or 32-bit x 32-bit integer multiply of the unsigned contents of Wb and signed contents of Ws. Ws is interpreted as a two's-complement signed value.</p> <p>For MULUS.d (L = 1, E = 1), a 32-bit x 32-bit multiply of Wb and Ws is executed. The MS 32-bits of the 64-bit result will be written to Wnd+1, and the LS 32-bits of the result will be written to Wnd.</p> <p>For MULUS.I (L=1, E = 0), a 32-bit x 32-bit multiply of Wb and Ws is executed. The LS 32-bits of the result will be written to Wnd.</p> <p>For MULUS.w (L=0, E = 0), a 16-bit x 16-bit multiply of the lsw of Wb and the lsw of Ws is executed. The result will be 32-bits and is written to Wnd.</p> <p>The 'L' bit selects the operand date size.</p> <p>The 'E' bit selects between a 32-bit and 64-bit result write.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'd' bits select the destination register.</p> <p>The 'p' bits select the source addressing mode.</p>							
	<b>Notes:</b>							
	<ol style="list-style-type: none"> <li>1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.</li> <li>2. For 32-bit result writes, user is responsible to ensure that the result does not overflow into the sign bit, Wnd[31].</li> </ol>							
I-Words:	1							
Cycles:	1							

<b>MULUU</b>		<b>Unsigned Unsigned Integer Multiply</b>						
Syntax:	{label:}	MULUU.d	Wb,	Ws,	Wnd			
		MULUU.I		[Ws],				
		MULUU.w		[Ws++],				
				[Ws--],				
				[++Ws],				
				[--Ws],				
Operands:		Wb ∈ [W0 ... W15]; Ws ∈ [W0 ... W15]; Wnd ∈ [W0 ... W14]						
Operation:		MULUU.d : unsigned (Wb) * unsigned (Ws) → {Wnd+1[31:0], Wnd[31:0]}						
		MULUU.I : unsigned (Wb) * unsigned (Ws) → {Wnd[31:0]}						
		MULUU.w : unsigned (Wb[15:0]) * unsigned (Ws[15:0]) → {Wnd[31:0]}						
Status Affected:		None						
Encoding:		S001	011L	dddd	ssss	ppPE	UUww	wwUU
Description:		Performs a 16-bit x 16-bit or 32-bit x 32-bit integer multiply of the unsigned contents of Wb and Ws. For MULUU.d (L = 1, E = 1), a 32-bit x 32-bit multiply of Wb and Ws is executed. The MS 32-bits of the 64-bit result will be written to Wnd+1, and the LS 32-bits of the result will be written to Wnd.						
		For MULUU.I (L=1, E = 0), a 32-bit x 32-bit multiply of Wb and Ws is executed. The LS 32-bits of the result will be written to Wnd.						
		For MULUU.w (L=0, E = 0), a 16-bit x 16-bit multiply of the lsw of Wb and the lsw of Ws is executed. The result will be 32-bits and is written to Wnd.						
		The 'S' bit selects instruction size.						
		The 'L' bit selects the operand date size.						
		The 'E' bit selects between a 32-bit and 64-bit result write.						
		The 's' bits select the source register.						
		The 'w' bits select the base register.						
		The 'd' bits select the destination register.						
		The 'p' bits select the source addressing mode.						
		<b>Note:</b>						
		1. The state of the Multiplier Mode bit (CORCON.US) has no effect upon the operation of this instruction.						
I-Words:		1 or 0.5						
Cycles:		1						

## 4.7. Instruction Descriptions (N to XORWF)

NEG	Negate Accumulators		
Syntax:	{label:}	NEG	A B

Operands:	None
Operation:	if (NEGAB A) then -ACCA → ACCA if (NEGAB B) then -ACCB → ACCB
Status Affected:	OA, SA or OB, SB
Encoding:	0111 001A UUUU 1010
Description:	Compute the 2's complement of the contents of the specified accumulator and place the result back into the specified accumulator. The 'A' bit specifies the selected accumulator. OA and SA are affected if AccA is selected. OB and SB are affected if AccB is selected.
I-Words:	0.5
Cycles:	1

Example:1	NEG	A	;Negate the value of ACCA and store the results back in ACCA.
-----------	-----	---	---

Before execution		After execution	
ACCA	0x00_4000_0000_0 000_0000	ACCA	0xFF_C000_0000_0 000_0000

Example 1:	NEG A ; Negate ACCA ; Store result to ACCA ; CORCON = 0x0000 (no saturation)
	Before Instruction
ACCA	00 3290 59C8
CORCON	0000
SR	0000
	After Instruction
ACCA	FF CD6F A638
CORCON	0000
SR	0000

Example 2:	NEG B ; Negate ACCB ; Store result to ACCB ; CORCON = 0x00C0 (normal saturation)
	Before Instruction
ACCB	FF F230 10DC
CORCON	00C0
SR	0000
	After Instruction
ACCB	00 0DCF EF24
CORCON	00C0
SR	0000

NEG	Negate f						
Syntax:	{label:}	NEG.b	f	{,Wnd}	{,WREG}		
		NEG.bz					
		NEG{.w}					
		NEG.I					
Operands:	$f \in [0 \dots 64KB]$ ; $Wnd \in [W0 \dots W14]$						
Operation:	$(f) + 1 \rightarrow$ destination designated by D						
Status Affected:	C, N, OV, Z						
Encoding:	1101	101L	dddd	ffff	ffff	ffff	BD01
Description:	<p>Compute the 2's complement of the contents of the file register and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 'd' bits select the Working register.</p>						
I-Words:	1						
Cycles:	1						

**Example 1:** NEG.B 0x880, WREG ; Negate (0x880) (Byte mode) ; Store result to WREG

Before Instruction		After Instruction	
WREG (W0)	9080	WREG (W0)	90AB
Data 0880	2355	Data 0880	2355
SR	0000	SR	0008 (N = 1)

**Example 2:** NEG 0x1200 ; Negate (0x1200) (Word mode)

Before Instruction		After Instruction	
Data 1200	8923	Data 1200	76DD
SR	0000	SR	0000

NEOP	None Executable No Operation
Syntax:	{label:} NEOP
Operands:	None
Operation:	No Operation
Status Affected:	None
Encoding:	0000 0000 UUUU UUUU
Description:	No Operation is performed. This instruction will also consume no execution time. It can be regarded as a means to pad a 16-bit instruction (to use a 32-bit word) when the instruction cannot be paired with another 16-bit instruction (to maintain alignment).
I-Words:	0.5
Cycles:	0

<b>NOP</b>		<b>No Operation</b>			
Syntax:	{label:} NOP				
Operands:	None				
Operation:	No Operation				
Status Affected:	None				
Encoding:	0000	0001	UUUU	UUUU	
Description:	No Operation is performed.				
I-Words:	1				
Cycles:	1				

<b>Example 1:</b>		<b>NOP ; execute no operation</b>			
		Before Instruction			
PC	00 1092	PC	00 1094		
SR	0000	SR	0000		

<b>Example 2:</b>		<b>NOP ; execute no operation</b>			
		Before Instruction			
PC	00 08AE	PC	00 08B0		
SR	0000	SR	0000		

<b>NOPR</b>	<b>No Operation (32-bit)</b>							
-------------	------------------------------	--	--	--	--	--	--	--

Syntax:                    {label:} NOPR

Operands:                None

Operation:              No Operation

Status Affected:       None

Encoding:                1111        111U        UUUU        UUUU        UUUU        UUUU        UU11

Description:              No Operation is performed.

I-Words:                1

Cycles:                  1

<b>Example 1:</b>	<b>NOPR ; execute no operation</b>							
-------------------	------------------------------------	--	--	--	--	--	--	--

Before Instruction				After Instruction			
PC        00 2430				PC        00 2432			
SR        0000				SR        0000			

<b>Example 2:</b>	<b>NOPR ; execute no operation</b>							
-------------------	------------------------------------	--	--	--	--	--	--	--

Before Instruction				After Instruction			
PC        00 1466				PC        00 1468			
SR        0000				SR        0000			

NORM	Normalize Accumulator							
Syntax:	{label:}	NORM{.w} A, NORM.I B,	Wd [Wd] [Wd++] [Wd--] [++Wd] [--Wd] [Wd+Wb]					
Operands:			Wd ∈ [W0 ... W15], Wb ∈ [W0 ... W15]					
Operation:			refer to text					
Status Affected:			OA or OB, N, Z					
Encoding:	1100	10AL	www	ddd	UUUq	qqUU	UUUU	1111
Description:			Normalize the contents of the target accumulator. If the accumulator contains an overflowed value, the contents of the accumulator are shifted right by the minimum number of bits required to remove the overflow. If the accumulator does not contain an overflowed value, the contents of the accumulator are shifted left by the minimum number of bits required to produce the largest fractional data value without an overflow.					
			If it is not possible to normalize the target accumulator (i.e., it is already normalized, or it is all 0's or all 1's) Wd is cleared and the Z bit is set (and the N bit is cleared). The target accumulator is unaffected.					
			If it is possible to normalize the target accumulator, the exponent (shift value) required to normalize the target accumulator is written into Wd. A positive result indicates that a right shift of the accumulator was required for normalization. A negative result indicates that a left shift of the accumulator was required for normalization. The N bit is set to reflect the sign of the result and the Z bit is cleared.					
			The 'L' bit selects word or long word Wd destination.					
			The 'A' bit specifies the destination accumulator.					
			The 'd' bits select the destination register.					
			The 'q' bits select the destination addressing mode.					
			The 'w' bits define the offset Wb.					
			<b>Note:</b>					
			1.					
			OA or OB status bits are set based on the content of the target accumulator. Consequently, because NORM removes any overflow, OA or OB will always be cleared.					
			SA/SB are 'sticky' so will remain set if set prior to execution of the NORM, but they can never be affected by this instruction.					
I-Words:	1							
Cycles:	1							

POP	Pop Top of Return Stack															
Syntax:	{label:}	POP.I	f													
Operands:	$f \in [0 \dots 1\text{MB}]$															
Operation:	$(W15) - 4 \rightarrow W15;$ $(TOS) \rightarrow (f, 2'b00)$															
Status Affected:	None															
Encoding:	1011	1101	ffff	ffff	ffff	ffff	ffff	UU01								
Description:	The Stack Pointer (W15) is pre-decremented and Top-of-Stack (TOS) value is pulled off the stack and written to the file register. The 'f' bits select the address of the file register.															
<b>Notes:</b>																
1. This instruction operates in Long Word mode only. The LS 2-bits of the file address will therefore always be 2'b00.																
2. Accessible file address space is 1MB.																
I-Words:	1															
Cycles:	1															

POP		Pop Working (W) Register		
Syntax:	{label:}	POP.I		
Operands:				
Operation:				
Status Affected:				
Encoding:				
Description:				
I-Words:				
Cycles:				

## Notes:

1. This instruction operates in Long Word mode only. The LS 2-bits of the file address will therefore always be 2'b00.
  2. Accessible file address space is 1MB.

PUSH	Push Top of Return Stack (TOS)							
Syntax:	{label:}	PUSH.I	f					
Operands:	f ∈ [0 ... 1MB]							
Operation:	(f) → (TOS);							
	(W15)+4 → W15							
Status Affected:	None							
Encoding:	1011	1111	ffff	ffff	ffff	ffff	ffff	UU01
Description:	The file register contents are written to the Top of Stack (TOS) location. The Stack Pointer (W15) is then incremented. The 'f' bits select the address of the file register.							
Notes:	<ol style="list-style-type: none"> <li>This instruction operates in Long Word mode only. The LS 2-bits of the file address will therefore always be 2'b00.</li> <li>Accessible file address space is 1MB.</li> </ol>							
I-Words:	1							
Cycles:	1							

Example 1:	PUSH 0x2004 ; Push (0x2004) to TOS							
Before Instruction								After Instruction
W15 0B00								W15 0B02
Data 0B00 791C								Data 0B00 D400
Data 2004 D400								Data 2004 D400
SR 0000								SR 0000

Example 2:	PUSH 0xC0E ; Push (0xC0E) to TOS							
Before Instruction								After Instruction
W15 0920								W15 0922
Data 0920 0000								Data 0920 67AA
Data 0C0E 67AA								Data 2004 67AA
SR 0000								SR 0000

PUSH	Push Working (W) Register			
Syntax:	{label:}	PUSH.I	Wns	
Operands:	Wns ∈ [W0 ... W15];			
Operation:	PUSH.I 1:2: Wns → (TOS); (W15)+4 → W15			
Status Affected:	None			
Encoding:	0001      111L      dddd      ssss			
Description:	Move the contents of the source Working register, Wns, to the top of a LIFO stack data structure (or data memory). Use of the system Stack Pointer ([W15++]) is implied for PUSH.I. The 'L' bit selects operation data width. The 's' bits select the Working Source register. The 'd' bits select the Working Destination register.			
I-Words:	0.5			
Cycles:	1			

**Notes:**

1. PUSH.I Wns is functionally equivalent to MOV.I Wns, [W15++].
2. In order to preserve system stack Long Word alignment, the PUSH mnemonic only supports Long Word mode, and even the MOV.w mnemonic does not permit use of W15 as destination.

<b>PWRSAV</b> Enter Power Saving Mode			
Syntax:	{label:}	PWRSAV	lit1
Operands:	lit1 ∈ [0 ... 1]		
Operation:	0 → Sleep-mode WDT 1 → WDTO 0 → SLEEP 0 → IDLE		
	Enter either IDLE or SLEEP mode		
Status Affected:	SLEEP, IDLE		
Encoding:	0111	001U	UUUK      0111
Description:	If lit1 = 0, device is placed in SLEEP mode. If lit1 = 1, device is placed in IDLE mode.		
	The Sleep-mode WDT is reset. The SLEEP and IDLE status bits (located within the RCON register) are cleared. If SLEEP mode is selected, the device is shut down. The oscillator source is stopped. If IDLE is selected, the CPU is shut down, but the peripherals continue to operate.		
	The processor will exit from SLEEP or IDLE mode through an interrupt or a reset or a Watchdog Time-Out.		
	If exiting from IDLE mode, the clock source is reapplied to CPU. If exiting from SLEEP mode, the clock source is restarted.		
	If waking from SLEEP mode : 1 → SLEEP (in RCON register)		
	If waking from IDLE mode : 1 → IDLE (in RCON register)		
	If awakened by a WDT timeout: 1 → WDTO		
	The 'k' bit select the Power Saving mode.		
I-Words:	0.5		
Cycles:	2		

**Example 1:** PWRSAV #0 ; Enter SLEEP mode

Before Instruction		After Instruction	
SR	0040      (IPL = 2)	SR	0040      (IPL = 2)

**Example 2:** PWRSAV #1 ; Enter IDLE mode

Before Instruction		After Instruction	
SR	0020      (IPL = 1)	SR	0020      (IPL = 1)

<b>RCALL</b>		Relative Call							
Syntax:	{label:}	RCALL Label							
Operands:	Label is resolved by the linker to a signed word offset (slit20)								
Operation:	$(PC) + 4 \rightarrow PC,$ $(8'b0, PC[23:2]) \rightarrow TOS[31:2]; 2'b00 \rightarrow TOS[1:0],$ $(W15) + 4 \rightarrow W15,$ $PC + 2*slit20 \rightarrow PC;$ NOP → Instruction Register								
Status Affected:	None								
Encoding:	1101	010U	nnnn	nnnn	nnnn	nnnn	nnnn	0010	
Description:	Subroutine call with a forward or backward branch address range of 1MB to either a 32-bit or 16-bit instruction. The long word aligned return address (PC+4) is pushed onto the system stack. The 2's complement byte offset value '2*slit20' (the PC offset) is added to the PC. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20. The 'n' bits are a signed literal that specifies the number of PS words to branch from (PC+4).								
I-Words:	1								
Cycles:	1								

**Example 1:**    012004 RCALL \_Task1 ; Call \_Task1; \_Task1  
                   012006 ADD W0, W1, W2 subroutine  
                   . . . . .  
                   012458 \_Task1: SUB W0, W2, W3  
                   01245A ...

Before Instruction		After Instruction	
PC	01 2004	PC	01 2458
W15	0810	W15	0814
Data 0810	FFFF	Data 0810	2006
Data 0812	FFFF	Data 0812	0001
SR	0000	SR	0000

**Example 2:**    00620E RCALL \_Init006210 MOV W0, [W4+ ; Call \_Init; \_Init subroutine  
                   +] . . . . .007000 \_Init: CLR W2007002 ...

Before Instruction		After Instruction	
PC	00 620E	PC	00 7000
W15	0C50	W15	0C54
Data 0C50	FFFF	Data 0C50	6210
Data 0C52	FFFF	Data 0C52	0000
SR	0000	SR	0000

RCALL	Computed Call							
Syntax:	{label:}	RCALL	Wns					
Operands:	Wns ∈ [W0 ... W15]							
Operation:	(PC) + 4 → PC, (8'b0, PC[23:0]) → TOS[31:0], (W15) + 4 → W15, (PC) + 2*Wns[19:0] → PC; NOP → Instruction Register							
Status Affected:	None							
Encoding:	1101      011U      UUUU      ssss      UUUU      UUUU      UUUU      0110							
Description:	Computed subroutine call with a forward or backward branch address range of 1MB to any size of instruction. The long word aligned return address (PC+4) is pushed onto the system stack. The signed value in Wns[19:0] represents the PS (16-bit) word offset from the current PC. This value is multiplied by 2 to create a byte address which is then added to the contents of the PC to form the target address. Wn must therefore contain a signed value that specifies the number of PS words to offset from (PC+4) for the call. RCALLW is a two-cycle instruction. The 's' bits specify the source register.							
	<b>Note:</b> If Wns[31:19] is not all 0's or all 1's, an address error trap will be initiated.							
I-Words:	1							
Cycles:	2							
	<b>Note:</b> The call target instruction fetch is not executed if an exception is pending, making effective instruction execution time one cycle.							

REPEAT	Repeat next instruction, literal count							
Syntax:	{label:}	REPEAT	lit20					
Operands:	lit20	∈ [0 ... 1048575]						
Operation:	(12'h000, lit20) → RCOUNT[31:0] (Loop Count Register)							
	(PC) + 4 → PC							
	1 → RA (Enable Code Looping)							
Status Affected:	RA (if lit20 > 0)							
Encoding:	1101	010U	kkkk	kkkk	kkkk	kkkk	kkkk	0110
Description:	The instruction immediately following the REPEAT instruction is executed (lit20+1) times. The repeated instruction is held in the instruction register for all iterations and so is fetched only once (during the REPEAT instruction, as would be expected). The last iteration of the repeated instruction fetches the next instruction. The repeat count is decremented during each iteration. When it equals zero (during the penultimate loop), PC increment and instruction fetch is re-enabled (RA=0) allowing the final iteration to fetch the next instruction and normal execution continues. The repeated instruction can be interrupted before any iteration by any interrupt. Note that the user must save and restore RCOUNT to support nested repeats (e.g., from within the interrupt service routine). The 'k' bits are an unsigned literal that specifies the loop count.							
	<b>Note:</b>							
	1. If lit20 = 0, RA is not set and the subsequent instruction executes as normal (equivalent to a loop count of one).							
I-Words:	1							
Cycles:	1							

REPEAT	Repeat next instruction, short literal count			
Syntax:	{label:}	REPEAT	lit5	
Operands:	lit5 ∈ [0 ... 31]			
Operation:	(27'h000000, lit5) → RCOUNT[31:0] (Loop Count Register) (PC) + 2 → PC 1 → RA (Enable Code Looping)			
Status Affected:	RA (if lit5 > 0)			
Encoding:	0111      001k      kkkk      1111			
Description:	The instruction immediately following the REPEAT instruction is executed (lit5+1) times. If lit5 > 31, assemble as REPEAT instruction. The repeated instruction is held in the instruction register for all iterations and so is fetched only once (during the REPEAT instruction, as would be expected). The last iteration of the repeated instruction fetches the next instruction. The repeat count is decremented during each iteration. When it equals zero (during the penultimate loop), PC increment and instruction fetch is re-enabled (RA=0) allowing the final iteration to fetch the next instruction and normal execution continues. The repeated instruction can be interrupted before any iteration by any interrupt. Note that the user must save and restore RCOUNT to support nested repeats (e.g., from within the interrupt service routine). The 'k' bits are an unsigned literal that specifies the loop count. <b>Note:</b> If lit5 = 0, RA is not set and the subsequent instruction executes as normal (equivalent to a loop count of one)			
I-Words:	0.5			
Cycles:	1			

REPEAT	Repeat next instruction, variable count							
Syntax:	{label:}	REPEAT	Wn					
Operands:	Wn ∈ [W0 ... W15]							
Operation:	(Wn[31:0]) → RCOUNT[31:0] (Loop Count Register) (PC) + 2 → PC 1 → RA (Enable Code Looping)							
Status Affected:	RA (if (Wn) > 0)							
Encoding:	1101	100U	UUUU	ssss	UUUU	UUUU	UUUU	U010
Description:	The instruction immediately following the REPEAT instruction is executed (Wn)+1 times. The repeated instruction is held in the instruction register for all iterations and so is fetched only once (during the REPEAT instruction, as would be expected). The last iteration of the repeated instruction fetches the next instruction. The repeat count is decremented during each iteration. When it equals zero (during the penultimate loop), PC increment and instruction fetch is re-enabled (RA=0) allowing the final iteration to fetch the next instruction and normal execution continues. The repeated instruction can be interrupted before any iteration by any interrupt. Note that the user must save and restore RCOUNT to support nested repeats (e.g., from within the interrupt service routine). The 's' bits specify the Wn register that contains the loop count <b>Note:</b> If Wn = 0, RA is not set and the subsequent instruction executes as normal (equivalent to a loop count of one).							
I-Words:	1							
Cycles:	1							

RESET	Reset							
Syntax:	{label:} RESET							
Operands:	None							
Operation:	Force all registers and flag bits that are affected by an MCLR reset to their Reset condition.							
Status Affected:	None							
Encoding:	1111	101U	UUUU	UUUU	UUUU	UUUU	UUUU	0011
Description:	This instruction provides a means to execute a software reset.							
I-Words:	1							
Cycles:	1							

**Example 1: 00202A RESET ; Execute software RESET on PIC32A**

Before Instruction				After Instruction			
PC	00 202A			PC	00 0000		
W0	8901			W0	0000		
W1	08BB			W1	0000		
W2	B87A			W2	0000		
W3	872F			W3	0000		
W4	C98A			W4	0000		
W5	AAD4			W5	0000		
W6	981E			W6	0000		
W7	1809			W7	0000		
W8	C341			W8	0000		
W9	90F4			W9	0000		
W10	F409			W10	0000		
W11	1700			W11	0000		
W12	1008			W12	0000		
W13	6556			W13	0000		
W14	231D			W14	0000		
W15	1704			W15	0800		
SPLIM	1800			SPLIM	0000		
TBLPAG	007F			TBLPAG	0000		
PSVPAG	0001			PSVPAG	0000		
CORCON	00F0			CORCON	0020	(SATDW = 1)	
RCON	0000			RCON	0040	(SWR = 1)	
SR	0021	(IPL, C = 1)		SR	0000		

RETFIE      Return from Interrupt			
Syntax:	{label:}	RETFIE	
Operands:	None		
Operation:	(W15) - 4 → W15 TOS[23:1] → (PC[23:1]), (W15) - 4 → W15 TOS[15:0] → (SR[31:0]), 1'b0 → PC[0] NOP → Instruction Register		
Status Affected:	OA, OB, SA, SB, (OAB), (SAB), IPL[3:0], RA, N, OV, Z, C		
Encoding:	0111      001U      UUUU      0000		
Description:	Return from interrupt service routine to either a 32-bit or 16-bit instruction. The stack is popped and the Top of Stack (TOS) is loaded into PC[23:1] (PC[0] is always clear). The stack is popped again and the Top of Stack (TOS) is loaded into the SR. In addition, this instruction will also manage hardware context switching (IPL-based).		
<b>Note:</b>			
1. OAB and SAB will reflect the returned state of OA/OB and SA/SB respectively.			
I-Words:	0.5		
Cycles:	4		
<b>Note:</b>			
1. Return PC instruction fetch is not executed if an exception is pending, making effective instruction execution time (for latency) one cycle.			

**Example 1: 000A26 RETFIE ; Return from ISR**

Before Instruction		After Instruction	
PC	00 0A26	PC	01 0230
W15	0834	W15	0830
Data 0830	0230	Data 0830	0230
Data 0832	8101	Data 0832	8101
CORCON	0001	CORCON	0001
SR	0000	SR	0081      (IPL = 4, C = 1)

**Example 2: 008050 RETFIE ; Return from ISR**

Before Instruction		After Instruction	
PC	00 8050	PC	00 7008
W15	0926	W15	0922
Data 0922	7008	Data 0922	7008
Data 0924	0300	Data 0924	0300
CORCON	0000	CORCON	0000
SR	0000	SR	0003      (Z, C = 1)

RETLW	Return with Literal in Wd							
Syntax:	{label:}	RETLW.b	lit16,	Wnd				
		RETLW.bz						
		RETLW{.w}						
		RETLW.I						
Operands:		Wnd ∈ [W0 ... W14]; lit16 ∈ [0 ... 65535]						
Operation:	(W15) - 4 → W15							
	TOS[23:1] → (PC[23:1]),							
	1'b0 → PC[0]							
	Byte mode: lit16[7:0] → Wnd[7:0]							
	Extended Byte mode: 24'b0, lit16[7:0] → Wnd							
	Word or Long Word mode: 16'b0, lit16 → Wnd							
	NOP → Instruction Register							
Status Affected:	None							
Encoding:	1100	110L	dddd	kkkk	kkkk	kkkk	kkkk	BU10
Description:	Return with a literal value in Wnd. When operating in Long Word, Word or Extended Byte mode, the literal is zero-extended to 32-bits before being written into Wnd. When operating in Byte mode, the literal is loaded into the LSb of Wnd, leaving the MSb of Wnd unchanged. The 'L' and 'B' bits select operation data width. The 'd' bits select the destination register. The 'k' bits define the literal.							
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. W15 is excluded from being a valid destination for the return literal value.</li> <li>2. Word and Long Word operating modes are equivalent.</li> <li>3. Return PC instruction fetch is not executed if an exception is pending, making effective instruction execution time (for latency) one cycle.</li> </ol>							
I-Words:	1							
Cycles:	3 (refer to note 3)							

RETURN	Return			
Syntax:	{label:} RETURN			
Operands:	None			
Operation:	(W15) - 4 → W15 TOS[23:1] → (PC[23:1]) 1'b0 → PC[0] NOP → Instruction Register			
Status Affected:	None			
Encoding:	0111	001U	UUUU	0001
Description:	Return from subroutine to either a 32-bit or 16-bit instruction. The stack is popped and the Top-of-Stack (TOS) is loaded into PC[23:0]. PC[0] is always clear.			
I-Words:	0.5			
Cycles:	3			
<b>Note:</b>	<ol style="list-style-type: none"> <li>1. Return PC instruction fetch is not executed if an exception is pending, making effective instruction execution time (for latency) 1 cycle.</li> </ol>			

<b>Example 1: 001A06 RETURN ; Return from subroutine</b>				
Before Instruction			After Instruction	
PC	00 1A06		PC	01 0004
W15	1248		W15	1244
Data 1244	0004		Data 1244	0004
Data 1246	0001		Data 1246	0001
SR	0000		SR	0000

<b>Example 2: 005404 RETURN ; Return from subroutine</b>				
Before Instruction			After Instruction	
PC	00 5404		PC	00 0966
W15	090A		W15	0906
Data 0906	0966		Data 0906	0966
Data 0908	0000		Data 0908	0000
SR	0000		SR	0000

RLC	Rotate Left Ws through Carry				
Syntax:	{label:}	RLC.b	Ws,	Wd	
		RLC.bz	[Ws],	[Wd]	
		RLC.{w}	[Ws++],	[Wd++]	
		RLC.l	[Ws-],	[Wd-]	
			[++Ws],	[++Wd]	
			[--Ws],	[--Wd]	
Operands:	$Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$				
Operation:	<u>For Byte operation:</u> (C) $\rightarrow$ Wd[0], (Ws[6:0]) $\rightarrow$ Wd[7:1], (Ws<7>) $\rightarrow$ C <u>For Word operation:</u> (C) $\rightarrow$ Wd[0], (Ws[14:0]) $\rightarrow$ Wd[15:1], (Ws[15]) $\rightarrow$ C <u>For Long Word operation:</u> (C) $\rightarrow$ Wd[0], (Ws[30:0]) $\rightarrow$ Wd[31:1], (Ws[31]) $\rightarrow$ C				
Status Affected:	C, N, Z				
Encoding:	0011	000L	dddd	ssss	pppq qqUU UUUU BU00
Description:	Rotate the contents of the source register Ws one bit to the left through the carry flag and place the result in the destination register Wd. In all cases, N and Z are set based on an evaluation of the result using the data size of the operation. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.				
I-Words:	1 or 0.5				
Cycles:	1				

Example 1	RLC.B W0, W3	;Rotate Left w/C (W0) (Byte mode)
		;Store the result in W3

Before Instruction		After Instruction	
W0	9976	W0	9976
W3	5879	W3	58ED
SR	0001 (C = 1)	SR	0009 (N = 1)

Example 2	[W2++],[W8]	;Rotate Left w/C [W2] (Word mode)
		;Post-increment W2

Before Instruction		After Instruction	
W2	2008	W2	200A
W8	094E	W8	094E
Data 094E	3689	Data 094E	8082 (N = 1)
Data 2008	C041	Data 2008	C041
SR	0001 (C = 1)	SR	0009 (N, C = 1)

<b>RLC</b> <b>Rotate Left f through Carry</b>							
Syntax:	{label:}	RLC.b	f	{,Wnd}	{,WREG}		
		RLC.bz					
		RLC{.w}					
		RLC.l					
Operands:		$f \in [0 \dots 64KB]$ ; Wnd $\in [W0 \dots W14]$					
Operation:		<u>For Byte operation:</u> $(C) \rightarrow \text{Dest}[0], (f[6:0]) \rightarrow \text{Dest}[7:1], (f<7>) \rightarrow C$					
		<u>For Word operation:</u> $(C) \rightarrow \text{Dest}[0], (f[14:0]) \rightarrow \text{Dest}[15:1], (f[15]) \rightarrow C$					
		<u>For Long Word operation:</u> $(C) \rightarrow \text{Dest}[0], (f[30:0]) \rightarrow \text{Dest}[31:1], (f[31]) \rightarrow C$					
Status Affected:		C, N, Z					
Encoding:	1011	000L	dddd	ffff	ffff	ffff	BD01
Description:	Rotate the contents of the file register f one bit to the left through the carry flag and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'f' bits select the address of the file register. The 'd' bits select the Working register.						
I-Words:	1						
Cycles:	1						

<b>Example 1</b>	<b>RLC.B 0x1233</b>	<b>; Rotate Left w/ C (0x1233) (Byte mode)</b>
------------------	---------------------	--

Before Instruction				After Instruction			
Data 1232	E807			Data 1232	D007		
SR	0000			SR	0009	(N, C = 1)	

<b>Example 2</b>	<b>RLC 0x820, WREG</b>	<b>; Rotate Left w/ C (0x820) (Word mode)</b>
		<b>; Store result in WREG</b>

Before Instruction				After Instruction			
WREG (W0)	5601			WREG (W0)	42DD		
Data 0820	216E			Data 0820	216E		
SR	0001	(C = 1)		SR	0000	(C = 0)	

RLNC	Rotate Left Ws (No Carry)					
Syntax:	{label:}	RLNC.b	Ws,	Wd		
		RLNC.bz	[Ws],	[Wd]		
		RLNC{.w}	[Ws++],	[Wd++]		
		RLNC.I	[Ws-],	[Wd--]		
			[++Ws],	[++Wd]		
			[--Ws],	[--Wd]		
Operands:	$Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$					
Operation:	<u>For Byte operation:</u> $(Ws[6:0]) \rightarrow Wd[7:1], (Ws<7>) \rightarrow Wd[0]$					
	<u>For Word operation:</u> $(Ws[14:0]) \rightarrow Wd[15:1], (Ws[15]) \rightarrow Wd[0]$					
	<u>For Long Word operation:</u> $(Ws[30:0]) \rightarrow Wd[31:1], (Ws[31]) \rightarrow Wd[0]$					
Status Affected:	N, Z					
Encoding:	S011	001L	dddd	ssss	pppq	qqUU
Description:	Rotate the contents of the source register Ws one bit to the left and place the result in the destination register Wd. The Carry Flag bit is not affected. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.					
I-Words:	1 or 0.5					
Cycles:	1					

Example 1	RLNC.B W0, W3	; Rotate Left (W0) (Byte mode)
		; Store the result in W3

Before Instruction		After Instruction	
W0	9976	W0	9976
W3	5879	W3	58EC
SR	0001	SR	0009 (N, C = 1)

Example 2	RLNC [W2++], [W8]	; Rotate Left [W2] (Word mode)
		; Post-increment W2
		; Store result in [W8]

Before Instruction		After Instruction	
W2	2008	W2	200A
W8	094E	W8	094E
Data 094E	3689	Data 094E	8083
Data 2008	C041	Data 2008	C041
SR	0001 (C = 1)	SR	0009 (N, C = 1)

RLNC	Rotate Left f (No Carry)						
Syntax:	{label:}	RLNC.b	f	{,Wnd}	{,WREG}		
		RLNC.bz					
		RLNC{.w}					
		RLNC.I					
Operands:		$f \in [0 \dots 64KB]$ ; Wnd $\in [W0 \dots W14]$					
Operation:		<u>For Byte operation:</u> $(f[6:0]) \rightarrow \text{Dest}[7:1], (f<7>) \rightarrow \text{Dest}[0]$					
		<u>For Word operation:</u> $(f[14:0]) \rightarrow \text{Dest}[15:1], (f[15]) \rightarrow \text{Dest}[0]$					
		<u>For Long Word operation:</u> $(f[30:0]) \rightarrow \text{Dest}[31:1], (f[31]) \rightarrow \text{Dest}[0]$					
Status Affected:		N, Z					
Encoding:	1011	001L	dddd	ffff	ffff	ffff	BD01
Description:	Rotate the contents of the file register f one bit to the left and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register. The Carry Flag bit is not affected. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'F' bits select the address of the file register. The 'd' bits select the Working register.						
I-Words:	1						
Cycles:	1						

Example 1	RLNC.B 0x1233	;	Rotate Left (0x1233) (Byte mode)
Before Instruction		After Instruction	
Data 1232	E807	Data 1233	D107
SR	0000	SR	0008 (N = 1)

Example 2	RLNC 0x820, WREG	;	Rotate Left (0x820) (Word mode)
		;	Store result in WREG
Before Instruction		After Instruction	
WREG (W0) 5601		WREG (W0) 42DC	
Data 0820 216E		Data 0820 216E	
SR	0001 (C = 1)	SR	0000 (C = 0)

<b>RRC</b>		<b>Rotate Right Ws through Carry</b>					
Syntax:	{label:}	RRRC.b	Ws,	Wd			
		RRRC.bz	[Ws],	[Wd]			
		RRRC.{w}	[Ws++],	[Wd++]			
		RRRC.l	[Ws-],	[Wd-]			
			[++Ws],	[++Wd]			
			[--Ws],	[--Wd]			
Operands:		Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]					
Operation:		<u>For Byte operation:</u> (C) → Wd<7>, (Ws[7:1]) → Wd[6:0], (Ws[0]) → C					
		<u>For Word operation:</u> (C) → Wd[15], (Ws[15:1]) → Wd[14:0], (Ws[0]) → C					
		<u>For Long Word operation:</u> (C) → Wd[31], (Ws[31:1]) → Wd[30:0], (Ws[0]) → C					
Status Affected:		C, N, Z					
Encoding:		S011	010L	dddd	ssss	pppq	qqUU
Description:		Rotate the contents of the source register Ws one bit to the right through the carry flag and place the result in the destination register Wd.					
		The 'S' bit selects instruction size.					
		The 'L' and 'B' bits select operation data width.					
		The 's' bits select the source register.					
		The 'd' bits select the destination register.					
		The 'p' bits select the source addressing mode.					
		The 'q' bits select the destination addressing mode.					
I-Words:		1 or 0.5					
Cycles:		1					

<b>Example 1:</b>	RRRC.B W0, W3	<b>; Rotate Right w/ C (W0) (Byte mode)</b>	<b>; Store the result in W3</b>
-------------------	---------------	---	---------------------------------

Before Instruction		After Instruction	
W0	9976	W0	9976
W3	5879	W3	58BB
SR	0001 (C = 1)	SR	0008 (N = 1)

<b>Example 2:</b>	RRRC [W2++], [W8]	<b>; Rotate Right w/ C [W2] (Word mode)</b>	<b>; Post-increment W2</b>	<b>; Store result in [W8]</b>
-------------------	-------------------	---	--------------------------------	-----------------------------------

Before Instruction		After Instruction	
W2	2008	W2	200A
W8	094E	W8	094E
Data	094E 3689	Data	094E E020
Data	2008 C041	Data	2008 C041
SR	0001 (C = 1)	SR	0009 (N, C = 1)

<b>RRC</b>		<b>Rotate Right f through Carry</b>					
Syntax:	{label:}	RRC.b	f	{,Wnd}	{,WREG}		
		RRC.bz					
		RRC{.w}					
		RRC.l					
Operands:		$f \in [0 \dots 64KB]$ ; $Wnd \in [W0 \dots W14]$					
Operation:		<u>For Byte operation:</u> $(C) \rightarrow Dest[7], (f[7:1]) \rightarrow Dest[6:0], (f[0]) \rightarrow C$					
		<u>For Word operation:</u> $(C) \rightarrow Dest[15], (f[15:1]) \rightarrow Dest[14:0], (f[0]) \rightarrow C$					
		<u>For Long Word operation:</u> $(C) \rightarrow Dest[31], (f[31:1]) \rightarrow Dest[30:0], (f[0]) \rightarrow C$					
Status Affected:		C, N, Z					
Encoding:		1011	010L	dddd	fffff	fffff	fffff BD01
Description:		Rotate the contents of the file register f one bit to the left through the carry flag and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'f' bits select the address of the file register. The 'd' bits select the Working register.					
I-Words:		1					
Cycles:		1					

<b>Example 1</b>	<b>RRC.B 0x1233</b>	<b>;Rotate Right w/C (0x1233) (Byte mode)</b>
------------------	---------------------	---

Before Instruction		After Instruction	
Data 1232	E807	Data 1232	7407
SR	0000	SR	0000

<b>Example 2:</b>	<b>RRC 0x820, WREG</b>	<b>; Rotate Right w/C (0x820) (Word mode) ; Store result in WREG</b>
-------------------	------------------------	--

Before Instruction		After Instruction	
WREG (W0)	5601	WREG (W0)	90B7
Data 0820	216E	Data 0820	216E
SR	0001	(C = 1)	SR 0008 (N = 1)

RRNC	Rotate Right Ws (No Carry)						
Syntax:	{label:}	RRNC.b	Ws,	Wd			
		RRNC.bz	[Ws],	[Wd]			
		RRNC{.w} [Ws++],		[Wd++]			
		RRNC.l	[Ws-],	[Wd--]			
			[++Ws],	[++Wd]			
			[--Ws],	[--Wd]			
Operands:	Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]						
Operation:	<u>For Byte operation:</u> (Ws[7:1] → Wd[6:0], (Ws[0]) → Wd<7> <u>For Word operation:</u> (Ws[15:1] → Wd[14:0], (Ws[0]) → Wd[15]) <u>For Long Word operation:</u> (Ws[31:1] → Wd[30:0], (Ws[0]) → Wd[31])						
Status Affected:	N, Z						
Encoding:	S011	011L	dddd	ssss	pppq	qqUU	UUUU BU00
Description:	Rotate the contents of the source register Ws one bit to the right and place the result in the destination register Wd. The Carry Flag bit is not affected. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.						
I-Words:	1 or 0.5						
Cycles:	1						

Example 1	RRNC.B W0, W3	; Rotate Right (W0) (Byte mode)
		; Store the result in W3

Before Instruction			After Instruction		
W0	9976		W0	9976	
W3	5879		W3	583B	
SR	0001 (C = 1)		SR	0001 (C = 1)	

Example 2:	RRNC [W2++], [W8]	; Rotate Right [W2] (Word mode)
		; Post-increment W2
		; Store result in [W8]

Before Instruction			After Instruction		
W2	2008		W2	200A	
W8	094E		W8	094E	
Data 094E	3689		Data 094E	E020	
Data 2008	C041		Data 2008	C041	
SR	0000 (C = 1)		SR	0008 (N = 1)	

RRNC	Rotate Right f (No Carry)					
Syntax:	{label:}	RRNC.b	f	{,Wnd}	{,WREG}	
		RRNC.bz				
		RRNC{.w}				
		RRNC.l				
Operands:		$f \in [0 \dots 64KB]$ ; Wnd $\in [W0 \dots W14]$				
Operation:		<u>For Byte operation:</u> $(f[7:1]) \rightarrow \text{Dest}[6:0], (f[0]) \rightarrow \text{Dest}[7]$				
		<u>For Word operation:</u> $(f[15:1]) \rightarrow \text{Dest}[14:0], (f[0]) \rightarrow \text{Dest}[15]$				
		<u>For Long Word operation:</u> $(f[31:1]) \rightarrow \text{Dest}[30:0], (f[0]) \rightarrow \text{Dest}[31]$				
Status Affected:		N, Z				
Encoding:	1011	011L	dddd	ffff	ffff	ffff BD01
Description:	Rotate the contents of the file register f one bit to the right, and place the result in the destination designated by D. If the optional Wnd is specified, D=0 and store result in Wnd; otherwise, D=1 and store result in the file register. The Carry Flag bit is not affected. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'F' bits select the address of the file register. The 's' bits select the Working register.					
I-Words:	1					
Cycles:	1					

Example 1	RRNC.B 0x1233	; Rotate Right (0x1233) (Byte mode)	
		Before Instruction	After Instruction
		Data 1232 E807 SR 0000	Data 1232 7407 SR 0000

Example 2:	RRNC 0x820, WREG	; Rotate Right (0x820) (Word mode) ; Store result in WREG	
		Before Instruction	After Instruction
		WREG (W0) 5601 Data 0820 216E SR 0001 (C = 1)	WREG (W0) 10B7 Data 0820 216E SR 0001 (C = 1)

SAC	Store Accumulator							
Syntax:	{label:}	SAC{.w}	A,	{ Slit6, }	Wd			
		SAC.I	B,		[Wd]			
					[Wd++]			
					[Wd-]			
					[~-Wd]			
					[++Wd]			
					[Wd+Wb]			
Operands:	$\text{Slit6} \in [-32 \dots +31]; \text{Wd} \in [W0 \dots W15]$ (see note 4); $\text{Wb} \in [W0 \dots W15]$							
Operation:	<u>For Long Word operation:</u> $\text{Shift}_{\text{Slit6}}(\text{ACC})$ $(\text{ACC}[63:32]) \rightarrow \text{Wd}[31:0]$							
	<u>For Word operation:</u> $\text{Shift}_{\text{Slit6}}(\text{ACC})$ $(\text{ACC}[63:48]) \rightarrow \text{Wd}[15:0]$ (see note 3)							
Status Affected:	None							
Encoding:	1100	01AL	wwww	dddd	UUUq	qqkk	kkkk	0011
Description:	Read then optionally shift accumulator value, then store the truncated result to the destination Effective Address. Any shift will apply only to data read from the accumulator, so will not modify the accumulator contents. The shift value is sourced from a signed literal value. After the shift, a word operation (SAC{.w}) assumes that the value is a Q1.15 signed fraction. Accordingly, the post-shift result [63:48] is then written to the Effective Address. After the shift, a long word operation (SAC.I) assumes that the value is a Q1.31 signed fraction. Accordingly, the post-shift result [63:32] is then written to the Effective Address. The 'L' bit selects word or long word operation. The 'A' bit specifies the source accumulator. The 'd' bits specify the destination register Wd. The 'q' bits select the destination addressing mode. The 'w' bits specify the offset register Wb. The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator shift. If operand Slit6 is absent, set literal to all 0's.							
Notes:	<ol style="list-style-type: none"> <li>Positive values of operand Slit6 represent arithmetic shift right. Negative values of operand Slit6 represent shift left.</li> <li>When the destination is register direct (W-reg), the result is either zero extended (CORCON.US = 1) or sign extended (CORCON.US = 0) to 32-bits.</li> <li>Register direct destination W15 not permitted.</li> </ol>							
I-Words:	1							
Cycles:	1							

<b>Example 1:</b>	<b>SAC A, #4, W5</b>	<b>; Right shift ACCA by 4</b>
		<b>; Store result to W5</b>
		<b>; CORCON = 0x0010 (SATDW = 1)</b>

Before Instruction		After Instruction	
W5	B900	W5	0120
ACCA	00 120F FF00	ACCA	00 120F FF00
CORCON	0010	CORCON	0010
SR	0000	SR	0000

<b>Example 2:</b>	SAC B, #-4, [W5++]	; Left shift ACCB by 4 ; Store result to [W5], Post-increment W5 ; CORCON = 0x0010 (SATDW = 1)
-------------------	--------------------	--

	Before Instruction		After Instruction
W5	2000	W5	2002
ACCB	FF C891 8F4C	ACCB	FF C891 1F4C
Data 2000	5BBE	Data 2000	8000
CORCON	0010	CORCON	0010
SR	0000	SR	0000

SE	Sign Extend Ws							
Syntax:	{label:}	SE.b	Ws,	Wnd				
		SE.w	[Ws],					
			[Ws++],					
			[Ws-],					
			[++Ws],					
			[--Ws],					
			[Ws+Wb],					
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Wnd ∈ [W0 ... W14]						
Operation:		<u>If Byte Mode:</u> {24{Ws[7]}}, Ws[7:0] → Wnd[31:0]						
		<u>If Word Mode:</u> {16{Ws[15]}}, Ws[15:0] → Wnd[31:0]						
Status Affected:		C,N,Z						
Encoding:	S111	110B	dddd	ssss	pppU	UUww	wwUU	UU00
Description:		Sign-extend an 8-bit or 16-bit signed value in Ws to a 32-bit signed value, then write the result back to Wnd.						
		C is set to the complement of N.						
		The 'S' bit selects instruction size.						
		The 'B' bit selects byte or word operation.						
		The 's' bits select the source register.						
		The 'd' bits select the destination register.						
		The 'p' bits select the source addressing mode.						
		The 'w' bits define the offset Wb.						
		<b>Notes:</b>						
		1. This operation always writes a long word.						
		2. Unlike all other Word mode instructions, word data does not always zero extend.						
I-Words:		1 or 0.5						
Cycles:		1						

Example:1	SE.b W0, W3	;Sign extend a byte in W0, store it in W3			
		Before execution		After execution	
	W0	0x00000089		W0	0x00000089
	W3	0x00000000		W3	0xFFFFFFF89
	SR[7:0]	8'b00000000		SR[7:0]	8'b00000000

Example:2	SE W0, W3	;Sign extend a word in W0, store it in W3			
		Before execution		After execution	
	W0	0x12347879		W0	0x12347879
	W3	0x00000000		W3	0x00007879
	SR[7:0]	8'b00000000		SR[7:0]	8'b00001000 (N = 1)

<b>SETM</b>	<b>Set f</b>							
Syntax:	{label:}	SETM.b	f					
		SETM{.w}						
		SETM.I						
Operands:	$f \in [0 \dots 1\text{MB}]$							
Operation:	0xFFFFFFFF → file register for long operation 0xFFFF → file register for word operation 0xFF → file register for byte operation							
Status Affected:	None							
Encoding:	1010	111L	ffff	ffff	ffff	ffff	ffff	B101
Description:	Set the file register. The 'L' and 'B' bits select operation data width. The 'f' bits select the address of the file register.							
<b>Note:</b>	<ol style="list-style-type: none"> <li>Accessible file address space is 1MB.</li> </ol>							
I-Words:	1							
Cycles:	1							

SFTAC	Arithmetic Shift Accumulator							
Syntax:	{label:}	SFTAC{.w} A, SFTAC.I B,	Ws [Ws] [Ws++] [Ws--] [++Ws] [--Ws] [Ws+Wb]					
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]						
Operation:		Shift <sub>(Ws)</sub> (ACC)						
Status Affected:		OA, SA or OB, SB						
Encoding:	1100	11AL	wwww	ssss	pppU	UUUU	UUUU	0011
Description:	<p>Arithmetic shift of accumulator.</p> <p>The shift value is sourced from the word or long word signed contents of Ws[5:0]. If Ws is positive or zero, the shift will be a right shift of between 1 and 31 bits maximum (or no shift if Ws is 0). If Ws is negative, the shift will be a left shift of between 1 and 32 bits maximum.</p> <p>If Ws is out of range (i.e., when Ws[15:5] for word, or Ws[31:5] for long word are not equal to all 1's or all 0's), a math error trap will be requested during execution. The instruction will continue to execute but the result will not be written to the destination accumulator, and the SR will not be modified.</p> <p>The 'L' bit selects word or long word Ws (shift count).</p> <p>The 'A' bit selects the accumulator for the result.</p> <p>The 's' bits select the Ws register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'w' bits define the offset Wb.</p>							
I-Words:	1							
Cycles:	1							

<b>SFTAC                          Arithmetic Shift Accumulator</b>							
Syntax:	{label:}	SFTAC	A,	Slit6			
			B,				
Operands:		Slit6 $\in [-32 \dots 31]$					
Operation:		Shift <sub>k</sub> (ACC)					
Status Affected:		OA, SA or OB, SB					
Encoding:	1100	11AU	UUUU	UUUU	UUkk	kkkk	0111
Description:	Arithmetic shift of accumulator. The Slit6 is used as the shift amount. If Slit6 is positive or zero, the shift will be a right shift of between 0 and 31 bits maximum. If Slit6 is negative, the shift will be a left shift of between 1 and 32 bits maximum. The 'A' bit selects the accumulator for the result. The 'k' bits determine the number of bits to be shifted.						
I-Words:	1						
Cycles:	1						

SL	Shift Left by 1				
Syntax:	{label:}	SL.b	Ws,	Wd	
		SL.bz	[Ws],	[Wd]	
		SL{.w}	[Ws++],	[Wd++]	
		SL.I	[Ws-],	[Wd--]	
			[++Ws],	[++Wd]	
			[~-Ws],	[~-Wd]	
Operands:	$Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$				
Operation:	For long word operation: $(Ws[31]) \rightarrow C, (Ws[30:0]) \rightarrow Wd[31:1], 0 \rightarrow Wd[0]$ For word operation: $(Ws[15]) \rightarrow C, (Ws[14:0]) \rightarrow Wd[15:1], 0 \rightarrow Wd[0]$ For byte operation: $(Ws[7]) \rightarrow C, (Ws[6:0]) \rightarrow Wd[7:1], 0 \rightarrow Wd[0]$				
Status Affected:	C,N,Z				
Encoding:	S010	110L	dddd	ssss	pppq
Description:	Shift left the contents of the source register Ws by 1 bit, placing the result in the destination register Wd. Destination register direct Extended Byte or Word mode will zero-extend the result to 32-bits, then write to Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.				
I-Words:	1 or 0.5				
Cycles:	1				

SL	Shift Left f					
Syntax:	{label:}	SL.b	f	{,Wnd}	{,WREG}	
		SL.bz				
		SL{.w}				
		SL.l				
Operands:		$f \in [0 \dots 64KB]$ ; Wnd $\in [W0 \dots W14]$				
Operation:		For byte operation: $(f[7]) \rightarrow (C), (f[6:0]) \rightarrow \text{Dest}[7:1], 0 \rightarrow \text{Dest}[0]$				
		For word operation: $(f[15]) \rightarrow (C), (f[14:0]) \rightarrow \text{Dest}[15:1], 0 \rightarrow \text{Dest}[0]$				
		For long word operation: $(f[31]) \rightarrow (C), (f[30:0]) \rightarrow \text{Dest}[31:1], 0 \rightarrow \text{Dest}[0]$				
Status Affected:		C, N, Z				
Encoding:	1010	010L	dddd	ffff	ffff	ffff BD01
Description:	Shift the contents of the file register f one bit to the left with a '0' fill. The carry flag is set if the MSB of f is '1'. Place the result in the destination designated by D. If the optional Wnd is specified, D = 0 and store result in Wnd; otherwise, D = 1 and store result in the file register. The 'L' and 'B' bits select operation data width. The 'D' bit selects the destination. The 'F' bits select the address of the file register. The 'd' bits select the Working register.					
I-Words:	1					
Cycles:	1					

Example 1:	SL.B 0x909	; Shift left (0x909) (Byte mode)		
		Before Instruction		
	Data 0908	9439		
	SR	0000		
			Data 0908	0839
			SR	0001 (C = 1)

Example 2:	SL 0x1650, WREG	; Shift left (0x1650) (Word mode) ; Store result in WREG		
		Before Instruction		
	WREG (W0)	0900		
	Data 1650	4065		
	SR	0000		
			WREG (W0)	80CA
			Data 1650	4065
			SR	0008 (N = 1)

Shift Left by Short Literal						
SL						
Syntax:	{label:}	SL{.w}	Ws,	lit5,	Wd	
		SL.I	[Ws],		[Wd]	
			[Ws++],		[Wd++]	
			[Ws-],		[Wd-]	
			[++Ws],		[~Wd]	
			[~-Ws],		[++Wd]	
Operands:	Ws ∈ [W0 ... W14]; lit5 ∈ [0...31]; Wd ∈ [W0 ... W14]					
Operation:	lit5[4:0] → Shift_Val 1'b0 → Left shift input <u>For long word operation:</u> 32'b0,Ws[31:0] → Shift_In[63:0] Shift left Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[31:0] → Wnd <u>For word operation:</u> 48'b0, Ws[15:0], → Shift_In[63:0] Shift left Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[15:0] → Wnd[15:0]					
Status Affected:	N,Z					
Encoding:	S010	010k	kkkk	dddd	pppq	qqss ssUU LU00
Description:	Shift left the contents of the source register Ws by lit5 bits (up to 31 positions), placing the result in the destination Wd. This instruction will generate the correct result for any shift value in lit5 (a word operation shift value > 15, Wnd[15:0]=0x0000). When used in conjunction with the SLMK instruction for multi-precision multi-bit shift operations, this instruction will not generate a correct result for any shift value greater than 32. Register Direct Word mode will zero-extend the result to 32-bits, then write to Wd. The 'S' bit selects instruction size. The 'L' bit selects word or long word operation. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode. The 'k' bits provide the literal operand.					
	<b>Note:</b> 1. This instruction only operates in Word or Long Word mode.					
I-Words:	1 or 0.5					
Cycles:	1					

<b>Example 1:</b>	SL W2, #4, W2	; Shift left W2 by 4
		; Store result to W2

Before Instruction		After Instruction	
W2	78A9	W2	8A90
SR	0000	SR	0008 (N = 1)

<b>Example 2:</b>	SL W3, #12, W8	; Shift left W3 by 12
		; Store result to W8

Before Instruction		After Instruction	
W3	0912	W3	0912

## Instruction Descriptions (N to XORWF) (continued)

Example 2:      SL W3, #12, W8 ; Shift left W3 by 12  
                      ; Store result to W8

W8 1002

W8 2000

SR 0000

SR 0000

SL	Shift Left by Wb				
Syntax:	{label:}	SL.b	Ws,	Wb,	Wd
		SL.bz	[Ws],		[Wd]
		SL{.w}	[Ws++],		[Wd++]
		SL.I	[Ws-],		[Wd-]
			[++Ws],		[++Wd]
			[~-Ws],		[~-Wd]
Operands:	$Ws \in [W0 \dots W15]; Wb \in [W0 \dots W15]; Wd \in [W0 \dots W15]$				
Operation:	Wb[15:0] → Shift_Val <u>For long word operation:</u> 32'b0, Ws[31:0] → Shift_In[63:0] Shift left Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[31:0] → Wd				
	<u>For word operation:</u> 48'b0, Ws[15:0] → Shift_In[63:0] Shift left Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[15:0] → Wd				
	<u>For byte operation:</u> 56'b0, Ws[7:0] → Shift_In[63:0] Shift left Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[7:0] → Wd				
Status Affected:	N,Z				
Encoding:	1010	110L	www	ddd	pppq
Description:	Shift left the contents of the source register Ws by Wb bits, placing the result in the destination register Wd. When used in isolation, this instruction will generate the correct result for any shift value in Wb[15:0]: <ul style="list-style-type: none"><li>• For a byte operation shift value &gt; 7, Wd[7:0]=0x00</li><li>• For a word operation shift value &gt; 15, Wd[15:0]=0x0000</li><li>• For a long word operation shift value &gt; 31, Wd=0x00000000</li></ul> When used in conjunction with the SLMW instruction for multi-precision multi-bit shift operations, this instruction will not generate a correct result for any shift value greater than 32. Any data held in Wb[31:16] will have no effect. Destination register direct Extended Byte or Word mode will zero-extend the result to 32-bits, then write to Wd. The 'L' and 'B' bits select operation data width.				
	The 's' bits select the source register. The 'w' bits select the base (shift count) register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.				
I-Words:	1				
Cycles:	1				

SLAC Store Lower Accumulator							
Syntax:	{label:}	SLAC{.l}	A, B,	{ Slit6 , }	Wd [Wd] [Wd++] [Wd-] [-Wd] [++Wd] [Wd+Wb]		
Operands:				Slit6 ∈ [-32 ... +31]; Wd ∈ [W0 ... W15] (see note 2); Wb ∈ [W0 ... W15]			
Operation:				Shift <sub>Slit6</sub> (ACC) (ACC[31:0]) → Wd[31:0]			
Status Affected:				None			
Encoding:	1100	01A1	wwww	dddd	UUUq	qqkk	kkkk
Description:							0111
				Read then optionally shift accumulator value, then store post-shift ACC[31:0] to the destination Effective Address. Any shift will apply only to data read from the accumulator so will not modify the accumulator contents. The shift value is sourced from a signed literal value.			
				The 'A' bit specifies the source accumulator.			
				The 'd' bits specify the destination register Wd.			
				The 'q' bits select the destination addressing mode.			
				The 'w' bits specify the offset register Wb.			
				The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator shift. If operand Slit6 is absent, set literal to all 0's.			
				<b>Notes:</b>			
				1. Positive values of operand Slit6 represent arithmetic shift right. Negative values of operand Slit6 represent shift left.			
				2. Register direct destination W15 not permitted.			
I-Words:	1						
Cycles:	1						

Example:1	SLAC A, #4, w8	;Right shift value of ACCAL by 4; store the results in w8 ; with CORCON.US = 0
-----------	----------------	--

Before execution				After execution			
ACCA	0xFF_9022_2EE1		_5633_9078	ACCA	0xFF_9022_2EE1		_5633_9078
w8	0x00000000			w8	0x05633907		

Shift Left Multi-Precision by Short Literal						
SLM	{label:}	SLM{.l}	Ws,	lit5,	Wnd	
Syntax:			[Ws], [Ws], [Ws++], [Ws-], [++Ws], [-Ws],			
Operands:			Ws ∈ [W0 ... W15]; lit5 ∈ [0...31]; Wnd ∈ [W0 ... W13]			
Operation:			lit5[4:0] → Shift_Val 0 → Left shift input 32'b0,Ws[31:0] → Shift_In[63:0] Shift left Shift_In[63:0] by Shift_Val → Shift_Out[63:0] Shift_Out[31:0] → Wnd Shift_Out[63:32]   Wnd+1 → Wnd+1			
Status Affected:			Z			
Encoding:			1110      000k      kkkk      dddd      pppU      UUss      ssUU      1011			
Description:			Shift left the contents of the source register Ws by lit5 bits (up to 31 positions), placing the result in the destination register Wnd. The register containing the next most significant data word will already contain an intermediate shift result. Bitwise OR this value with the data shifted out of Ws in order to create the final shift result, then update the corresponding destination register. This instruction is intended to be used in conjunction with the SLK instruction to support multi-precision multi-bit shift operations. The Z bit is “sticky” (can only be cleared). The ‘s’ bits select the source register. The ‘d’ bits select the destination register. The ‘p’ bits select the source addressing mode. The ‘k’ bits provide the literal operand.			
			<b>Note:</b> This instruction operates in Long Word mode only.			
I-Words:			1			
Cycles:			2			

Example:1	SLAC	A, #4, w8	;Right shift value of ACCAL by 4; store the re
	Before execution ACCA w8	0xFF_9022_2EE1_5633 _9078 0x00000000	

SLM	Shift Left Multi-Precision by Wb							
Syntax:	{label:}	SLM{.I}	Ws,	Wb,	Wnd			
			[Ws],					
			[Ws++],					
			[Ws-],					
			[++Ws],					
			[--Ws],					
Operands:	Ws	$\in [W0 \dots W15]$	Wb	$\in [W0 \dots W15]$	Wnd	$\in [W0 \dots W13]$		
Operation:	Wb[15:0] → Shift_Val		0 → Left shift input					
	32'b0,Ws[31:0] → Shift_In[63:0]							
	Shift left Shift_In[63:0] by Shift_Val → Shift_Out[63:0]							
	Shift_Out[31:0] → Wnd							
	Shift_Out[63:32]   Wnd+1 → Wnd+1							
Status Affected:	Z							
Encoding:	1110	001U	ssss	dddd	pppU	UUww	wwUU	1011
Description:	<p>Shift left the contents of the source register Ws by Wb bits, placing the result in the destination register Wnd. The register containing the next most significant data word will already contain an intermediate shift result. Bitwise OR this value with the data shifted out of Ws in order to create the final shift result, then update the corresponding destination register.</p> <p>The left shift may be by any amount between 0 and 32 bits. Should the shift value held in Wb[15:0] exceed 2'd32, the shift value will saturate to 2'd32 for consistency. Any data held in Wb[31:16] will have no effect.</p> <p>This instruction is intended to be used in conjunction with the SLW instruction to support multi-precision multi-bit shift operations.</p> <p>The Z bit is “sticky” (can only be cleared).</p> <p>The ‘w’ bits select the base (shift count) register.</p> <p>The ‘s’ bits select the source register.</p> <p>The ‘d’ bits select the destination register.</p> <p>The ‘p’ bits select the source addressing mode.</p> <p><b>Note:</b> This instruction operates in Long Word mode only.</p>							
I-Words:	1							
Cycles:	2							

<b>SQR</b>													
<b>Square to Accumulator</b>													
Syntax: {label}	SQR{.w}	Wx,	A	{,AWB}									
	SQR.I	[Wx],	B										
		[Wx]+=kx,											
		[Wx]=-kx,											
		[Wx+=kx],											
		[Wx-=kx],											
		[Wx+W12],											
Operands:		Wx ∈ {W0 ... W15};											
		Word mode: kx ∈ {-8, -6, -4, -2, 2, 4, 6, 8};											
		Long Word mode: kx ∈ {-16, -12, -8, -4, 4, 8, 12, 16};											
		AWB ∈ {W0, W1, W2, W3, W13, [W13++], [W15++]} <sup>5</sup>											
Operation:		(Wx) <sup>2</sup> → ACC(A or B);											
		(ACC(B or A)) rounded → AWB											
		When Indirect Pre/Post Modification Addressing:											
		(Wx)+kx→Wx or (Wx)-kx→Wx;											
Status Affected:		OA,SA or OB,SB											
Encoding:		1101      01AL      wwww      ssss      IIIi      i110      UUaa      a011											
Description:		Instruction to compute (A) <sup>2</sup> functions. Signed or unsigned (defined by CORCON.US) square of data read from Wx or fetched from the X address space. Note that, because only one operand is required, DS is not split into the X and Y address space for concurrent read access, so X-space in this instruction represents all of available DS (which will include Y-space).											
		The result is sign-extended or zero-extended to 72-bits then written to the specified accumulator. Fractional results are also scaled prior to the accumulator update to align the operand and accumulator (msw) fractional points.											
		When Indirect Pre/Post Modified Addressing is selected, the address modifier value is kx, and represents the number of data bytes by which to modify the Effective Address.											
		The optional AWB specifies the direct or indirect store of the (32-bit) rounded fractional contents of the accumulator not targeted by the SQR operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size (see note 4). AWB is not intended for use when the DSP engine is operating in Integer mode.											
		Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly.											
		The 'L' bit selects word or long word operation.											
		The 'A' bit selects the accumulator for the result.											
		The 'I' bits select the Operation.Addressing mode.											
		The 'i' bits select the kx modification value.											
		The 's' bits select Wx, data register or X-space source address register.											
		The 'w' bits are to be a copy of the 's' bits.											
		The 'a' bits select the accumulator write-back destination and addressing mode.											
<b>Notes:</b>													
1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.													
2. Operands are always regarded as signed or unsigned based on the state of CORCON.US.													
3. Use of the same W-reg for both indirect source and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.													
4. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.													
I-Words:		1											
Cycles:		1											
Example:1	SQR.I	[w9]+=4, A		;Square the value stored in address contained by w9; ;store the results in ACCA and increment w9 by 4									

Before execution

After execution

**Instruction Descriptions (N to XORWF) (continued)**

Example:1	SQR.I	[w9]+=4, A	;Square the value stored in address contained by w9; ;store the results in ACCA and increment w9 by 4		
		ACCA	0x00_7022_2EE1_5 633_9078	ACCA	0x00_1000_0000_0 000_0000
		w9	0x5000	w9	0x5000
		Data @5000	0x20000000	Data @5000	0x20000000

SQRAC								Square and Accumulate							
Syntax: {label}	SQRAC{.w} Wx, A {,AWB}														
	SQRAC.I [Wx], B														
	[Wx]+=kx, [Wx]=-kx, [Wx+=kx], [Wx-=kx], [Wx+W12],														
Operands:	Wx ∈ {W0 ... W15}; Word mode: kx ∈ {-8, -6, -4, -2, 2, 4, 6, 8}; Long Word mode: kx ∈ {-16, -12, -8, -4, 4, 8, 12, 16}; AWB ∈ {W0, W1, W2, W3, W13, [W13++], [W15++]} <sup>5</sup>														
Operation:	ACC(A or B) + (Wx) <sup>2</sup> → ACC(A or B); (ACC(B or A)) rounded → AWB When Indirect Pre/Post Modification Addressing: (Wx)+kx→Wx or (Wx)-kx→Wx;														
Status Affected:	OA,SA or OB,SB														
Encoding:	1101	00AL	www	sss	IIIi	i110	UUaa	a011							
Description:	Instruction to compute Accumulator + (A) <sup>2</sup> functions. Signed or unsigned (defined by CORCON.US) square of data read from Wx or fetched from the X address space. Note that, because only one operand is required, DS is not split into X and Y address space for concurrent read access, so X-space in this instruction represents all of the available DS (which will include Y-space). The result is sign-extended or zero-extended to 72-bits then added to the specified accumulator. Fractional or integer operation (defined by CORCON.IF) will determine if the result is scaled or not prior to the accumulator update. Fractional operation will scale the result to align the operand and accumulator (msw) fractional points (see note 3). Integer operation will align the LSb of the result with the LSb of the accumulator. When indirect pre/post modified addressing is selected, the address modifier value is kx and represents the number of data bytes by which to modify the Effective Address. The optional AWB specifies the direct or indirect (see note 4) store of the (32-bit) rounded fractional contents of the accumulator not targeted by the SQRAC operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size. AWB is not intended for use when the DSP engine is operating in Integer mode. Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly. The 'L' bit selects word or long word operation. The 'A' bit selects the accumulator for the result. The 'I' bits select the Operation.Addressing mode. The 'i' bits select the kx modification value. The 's' bits select Wx, data register or X-space source address register. The 'w' bits are to be a copy of the 's' bits. The 'a' bits select the accumulator write-back destination and addressing mode.														
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.</li> <li>2. Operands are always regarded as signed or unsigned based on the state of CORCON.US.</li> <li>3. The LS portion of ACCx is unaffected when operating in Fractional mode with word sized data. Lower significance data that may be present from prior (32-bit data) operations is therefore preserved. Users not requiring this should clear ACCx during initialization.</li> <li>4. Use of the same W-reg for both indirect source and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.</li> <li>5. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.</li> </ol>														
I-Words:	1														
Cycles:	1														

Example:1	SQRAC.I [w9]+=4, A	;Square the value stored in address contained by w9; ;accumulate the results to ACCA and increment w9 by 4	
Before execution			After execution
ACCA	0x00_4000_0000_0 000_0000	ACCA	0x00_5000_0000_0 000_0000
w9	0x5000	w9	0x5000
Data @5000	0x20000000	Data @5000	0x20000000

SQRN Negated Square to Accumulator							
Syntax: {label}	SQRN{.w} Wx,	A	{,AWB}				
	SQRN.I [Wx],	B					
	[Wx]+=kx,						
	[Wx]=-kx,						
	[Wx+=kx],						
	[Wx-=kx],						
	[Wx+W12],						
Operands:	$Wx \in \{W0 \dots W15\};$ Word mode: $kx \in \{-8, -6, -4, -2, 2, 4, 6, 8\};$ Long Word mode: $kx \in \{-16, -12, -8, -4, 4, 8, 12, 16\};$ $AWB \in \{W0, W1, W2, W3, W13, [W13++], [W15++]^5\}$						
Operation:	$-(Wx)^2 \rightarrow ACC(A \text{ or } B);$ (ACC(B or A)) rounded $\rightarrow AWB$						
	When Indirect Pre/Post Modification Addressing: $(Wx)+kx \rightarrow Wx$ or $(Wx)-kx \rightarrow Wx;$						
Status Affected:	OA,SA or OB,SB						
Encoding:	1101	01AL	wwww	ssss	IIIi	i110	UUaa a111
Description:	<p>Instruction to compute <math>-(A)^2</math> functions. Signed or unsigned (defined by CORCON.US) square of data read from <math>Wx</math> or fetched from the X address space. Note that, because only one operand is required, DS is not split into X and Y address space for concurrent read access, so X-space in this instruction represents all of the available DS (which will include Y-space).</p> <p>The result is sign-extended or zero-extended to 72-bits, negated and then written to the specified accumulator. Fractional results are also scaled prior to the accumulator update to align the operand and accumulator (msw) fractional points. For word sized operand operations, when Indirect Pre/Post Modified Addressing is selected, the address modifier value is <math>kx</math> and represents the number of data bytes by which to modify the Effective Address.</p> <p>The optional AWB specifies the direct or indirect (see note 3) store of the (32-bit) rounded fractional contents of the accumulator not targeted by the SQRN operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size (see note 4). AWB is not intended for use when the DSP engine is operating in Integer mode.</p> <p>Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'A' bit selects the accumulator for the result.</p> <p>The 'I' bits select the Operation.Addressing mode.</p> <p>The 'i' bits select the <math>kx</math> modification value.</p> <p>The 's' bits select <math>Wx</math>, data register or X-space source address register.</p> <p>The 'w' bits are to be a copy of the 's' bits.</p> <p>The 'a' bits select the accumulator write-back destination and addressing mode.</p>						
Notes:	<ol style="list-style-type: none"> <li>1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.</li> <li>2. Use of the same W-reg for both indirect source and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.</li> <li>3. Use of the same W-reg for both indirect source and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.</li> <li>4. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.</li> </ol>						
I-Words:	1						
Cycles:	1						
Example:1	SQRN.I	[w9]+=4, A	;Square the value stored in address contained by w9; ;store the results in ACCA and increment w9 by 4				
	Before execution			After execution			

**Instruction Descriptions (N to XORWF) (continued)**

Example:1	SQRN.I	[w9]+=4, A	;Square the value stored in address contained by w9; ;store the results in ACCA and increment w9 by 4	
	ACCA	0x00_4000_0000_000_0000	ACCA	0xFF_F000_0000_000_0000
	w9	0x5000	w9	0x5000
	Data @5000	0x20000000	Data @5000	0x20000000

SQRSC								Square and Subtract from Accumulator							
Syntax: {label}	SQRSC{.w} Wx, A {,AWB}														
	SQRSC.I [Wx], B														
	[Wx]+=kx, [Wx]=kx, [Wx+=kx], [Wx-=kx], [Wx+W12],														
Operands:	Wx ∈ {W0 ... W15}; Word mode: kx ∈ {-8, -6, -4, -2, 2, 4, 6, 8}; Long Word mode: kx ∈ {-16, -12, -8, -4, 4, 8, 12, 16}; AWB ∈ {W0, W1, W2, W3, W13, [W13++], [W15++]} <sup>5</sup>														
Operation:	ACC(A or B) - (Wx) <sup>2</sup> → ACC(A or B); (ACC(B or A)) rounded → AWB When Indirect Pre/Post Modification Addressing: (Wx)+kx→Wx or (Wx)-kx→Wx;														
Status Affected:	OA,SA or OB,SB														
Encoding:	1101	00AL	www	sss	IIIi	i110	UUaa	a111							
Description:	Instruction to compute Accumulator - (A) <sup>2</sup> functions. Signed or unsigned (defined by CORCON.US) square of data read from Wx or fetched from the X address space. Note that, because only one operand is required, DS is not split into X and Y address space for concurrent read access, so X-space in this instruction represents all of the available DS (which will include Y-space). The result is sign-extended or zero-extended to 72-bits then subtracted from the specified accumulator. Fractional or integer operation (defined by CORCON.IF) will determine if the result is scaled or not prior to the accumulator update. Fractional operation will scale the result to align the operand and accumulator (msw) fractional points (see note 4). Integer operation will align the LSb of the result with the LSb of the accumulator. When Indirect Pre/Post Modified Addressing is selected, the address modifier value is kx and represents the number of data bytes by which to modify the Effective Address. The optional AWB specifies the direct or indirect (see note 4) store of the (32-bit) rounded fractional contents of the accumulator not targeted by the SQRSC operation. Rounding mode is defined by CORCON.RND. Write data width is determined by selected instruction data size. AWB is not intended for use when the DSP engine is operating in Integer mode. Data read may be 16-bit or 32-bit values. All indirect address modification is scaled accordingly. The 'L' bit selects word or long word operation. The 'A' bit selects the accumulator for the result. The 'I' bits select the Operation.Addressing mode. The 'i' bits select the kx modification value. The 's' bits select Wx, data register or X-space source address register. The 'w' bits are to be a copy of the 's' bits. The 'a' bits select the accumulator write-back destination and addressing mode.														
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. Operates in Fractional or Integer Data mode as defined by CORCON.IF.</li> <li>2. Operands are always regarded as signed or unsigned based on the state of CORCON.US.</li> <li>3. Use of the same W-reg for both indirect source and AWB indirect destination is not permitted if the source is a pre- or post-modified Effective Address.</li> <li>4. The LS portion of ACCx is unaffected when operating in Fractional mode with word sized data. Lower significance data that may be present from prior (32-bit data) operations is therefore preserved. Users not requiring this should clear ACCx during initialization.</li> <li>5. Stack must remain long word aligned. Consequently, [W15++] AWB is only permitted for use with long word MAC-class instructions.</li> </ol>														
I-Words:	1														
Cycles:	1														

Example:1	SQRSC.I	[w9]+=4, A	;Square the value stored in address contained by w9; ;subtract the results from ACCA and increment w9 by 4
-----------	---------	------------	--

Before execution

ACCA 0x00\_4000\_0000\_0  
000\_0000

w9 0x5000

Data @5000 0x20000000

After execution

ACCA 0x00\_3000\_0000\_0  
000\_0000

w9 0x5000

Data @5000 0x20000000

SACR	Store Rounded Accumulator, Literal Shift							
Syntax:	{label:}	SACR{.w} A, SACR.I B,	{ Slit6, }	Wd [Wd] [Wd++] [Wd-] [-Wd] [++Wd] [Wd+Wb]				
Operands:			Slit6 $\in$ [-32 ... +31]; Wd $\in$ [W0 ... W15]; Wb $\in$ [W0 ... W15];					
Operation:			<u>For long word operation:</u> Long Word mode: Post-shift(Round{ACC[63:32]}) $\rightarrow$ Wd[31:0]					
			<u>For word operation:</u> Word mode: Post-shift(Round{ACC[63:48]}) $\rightarrow$ Wd[15:0] <sup>1</sup>					
Status Affected:		None						
Encoding:	1100	01AL	www	ddd	UUUq	qqkkk	kkkk	1011
Description:		Perform an optional, signed 6-bit shift of the specified accumulator, then store the rounded contents of Acc[63:32] (in case of long-word operation) or Acc[63:48] (in case of word operation) to Wd. The shift range is -32:31, where a negative operand indicates an arithmetic left shift and a positive operand indicates an arithmetic right shift. The Rounding mode (Conventional or Convergent) is set by the RND bit (CORCON<1>). Either Register Direct or Indirect Addressing may be used for Wd.						
		The 'L' bit selects word or long word operation.						
		The 'A' bit specifies the source accumulator.						
		The 'd' bits specify the destination register Wd.						
		The 'q' bits select the destination addressing mode.						
		The 'w' bits specify the offset register Wb.						
		The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator shift. If operand Slit6 is absent, set literal to all 0's.						
		<b>Notes:</b>						
		1. When the destination is register direct (W-reg), the result is either zero extended (CORCON.US = 1) or sign extended (CORCON.US = 0) to 32 bits.						
		2. Register direct destination W15 not permitted.						
I-Words:		1						
Cycles:		1						
Example:1	SACR	B, #1, w8	;word mode;Right shift value of ACCB by 1; store rounded result of shifted value in w8					

		Before execution		After execution	
		ACCB	0x00_1022_2EE1_5 633_9078	ACCB	0x00_1022_2EE1_5 633_9078
		w8	0x00000000	w8	0x00007022

		Before execution		After execution	
		ACCA	0x00_7022_2EE1_5 633_9078	ACCA	0x00_7022_2EE1_5 633_9078
		w8	0x00000000	w8	0x7022EE1

SACR	Store Rounded Accumulator after Shift							
Syntax:	{label:}	SACR{.w} A, SACR.I      B,	Ws,	Wd [Wd] [Wd++] [Wd-] [-Wd] [++Wd] [Wd+Wb]				
Operands:		Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15] (see note 2); Wb ∈ [W0 ... W15];						
Operation:		<u>For long word operation:</u> Long Word mode: Post-shift(Round{ACC[63:32]}) → Wd[31:0]						
		<u>For word operation:</u> Word mode: Post-shift(Round{ACC[63:48]}) → Wd[15:0] <sup>1</sup>						
Status Affected:		None						
Encoding:	1100	10AL	www	ddd	UUUq	qqss	ssUU	0111
Description:	<p>Perform an optional, signed 6-bit shift of the specified accumulator, then store the rounded contents of Acc[63:32] (in case of long-word operation) or Acc[63:48] (in case of word operation) to Wd. The shift range is sourced from the singed contents of Ws[5:0], values in the range of -32:31, where a negative value indicates an arithmetic left shift and a positive value indicates an arithmetic right shift. The Rounding mode (Conventional or Convergent) is set by the RND bit (CORCON&lt;1&gt;. Either Register Direct or Indirect Addressing may be used for Wd.</p> <p>If Ws is out of range (that is., Ws[31:5] not equal to all 1's or all 0's), a math error trap will be requested during execution. The instruction will continue to execute based on the value held in Ws[5:0].</p> <p>After the shift, a word operation (SACR.w) assumes that the value is a Q1.15 signed fraction. Accordingly, a round into shift result bit 48 is executed and shift result [63:48] is then written to the Effective Address.</p> <p>After the shift, a long word operation (SACR.I) assumes that the value is a Q1.31 signed fraction. Accordingly, a round into shift result bit 32 is executed and shift result [63:32] is then written to the Effective Address.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'A' bit specifies the source accumulator.</p> <p>The 'd' bits specify the destination register Wd.</p> <p>The 's' bits specify the shift value source register Ws.</p> <p>The 'q' bits select the destination addressing mode.</p> <p>The 'w' bits specify the offset register Wb.</p> <p>See for modifier addressing information.</p>							
		<b>Notes:</b>						
		1. When the destination is register direct (W-reg), the result is either zero extended (CORCON.US = 1) or sign extended (CORCON.US = 0) to 32-bits.						
		2. Register direct destination W15 not permitted.						
I-Words:		1						
Cycles:		1						
Example:1	SACR.I	A, W9, w8	;Right shift value of ACCA by content in w9; store rounded result of shifted value in w8					
			Before execution		After execution			
			ACCA	0x00_7022_2EE1_5 633_9078		ACCA	0x00_7022_2EE1_5 633_9078	
			w8	0x00000000		w8	0x70222EE1	
			w9	0x1		w9	0x1	

SUAC	Store Upper Accumulator							
Syntax:	{label:}	SUAC{.l}	A,	{ Slit6 , }	Wd			
			B,		[Wd]			
					[Wd++]			
					[Wd-]			
					[~-Wd]			
					[++Wd]			
					[Wd+Wb]			
Operands:	$\text{Slit6} \in [-32 \dots +31]; \text{Wd} \in [W0 \dots W15]$ (see note 2); $\text{Wb} \in [W0 \dots W15]$							
Operation:	$\text{Shift}_{\text{Slit6}}(\text{ACC})$ $(24\{\text{ACC}[71]\}, \text{ACC}[71:64]) \rightarrow \text{Wd}[31:0]$							
Status Affected:	None							
Encoding:	1100	01A1	wwww	dddd	UUUq	qqkk	kkkk	1111
Description:	<p>Read then optionally shift accumulator value, then store the post-shifted upper byte (ACC[71:64]) to the destination Effective Address. Any shift will apply only to data read from the accumulator so will therefore not modify the accumulator contents. The shift value is sourced from a signed literal value.</p> <p>When the destination is register direct (W-reg), the byte result is either zero extended (CORCON.US = 1) or sign extended (CORCON.US = 0) to 32 bits prior to the write.</p> <p>The 'A' bit specifies the source accumulator.</p> <p>The 'd' bits specify the destination register Wd.</p> <p>The 'q' bits select the destination addressing mode.</p> <p>The 'w' bits specify the offset register Wb.</p> <p>The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator shift. If operand Slit6 is absent, set literal to all 0's.</p>							
Notes:	<ol style="list-style-type: none"> <li>Positive values of operand Slit6 represent arithmetic shift right. Negative values of operand Slit6 represent shift left.</li> <li>Register direct destination W15 not permitted.</li> </ol>							
I-Words:	1							
Cycles:	1							
Example:1	SACR.I	A, W9, w8	;Right shift value of ACCA by content in w9 ; Store ACCAU to w8					

Before execution

ACCA 0xC1\_1022\_2EE1\_5  
633\_9078  
w8 0x12345678

After execution

ACCA 0xC1\_1022\_2EE1\_5  
633\_9078  
w8 0xFFFFFFF

SUB	Subtract Ws from Wb					
Syntax:	{label:}	SUB.b	Wb,	Ws,	Wd	
		SUB.bz		[Ws],	[Wd]	
		SUB.{w}		[Ws++],	[Wd++]	
		SUB.I		[Ws--],	[Wd-]	
				[++Ws],	[++Wd]	
				[--Ws],	[--Wd]	
Operands:		$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$				
Operation:		$(Wb) - (Ws) \rightarrow Wd$				
Status Affected:		C, N, OV, Z				
Encoding:	S110	010L	dddd	ssss	pppq	qqww
Description:	Subtract the contents of the source register Ws from the contents of the base register Wb and place the result in the destination register Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.					
I-Words:	1 or 0.5					
Cycles:	1					

Example 1:	SUB.B W0, W1, W0	; Sub. W1 from W0 (Byte mode)
		; Store result to W0
	Before Instruction	After Instruction
	W0 1732	W0 17EE
	W1 7844	W1 7844
	SR 0000	SR 0108 (DC, N = 1)

Example 2:	SUB W7, [W8++], [W9++]	; Sub. [W8] from W7 (Word mode)
		; Store result to [W9]
		; Post-increment W8
		; Post-increment W9
	Before Instruction	After Instruction
	W7 2450	W7 2450
	W8 1808	W8 180A
	W9 2020	W9 2022
	Data 1808 92E4	Data 1808 92E4
	Data 2020 A557	Data 2020 916C
	SR 0000	SR 010C (DC, N, OV = 1)

SUB	Subtract Accumulators			
Syntax:	{label:}	SUB	A	B

Operands: none  
 Operation: if (SUB A) then ACCA - ACCB → ACCA  
              if (SUB B) then ACCB - ACCA → ACCB  
 Status Affected: OA, SA or OB, SB  
 Encoding: 0111      001A      UUUU      1001  
 Description: Subtract accumulators and write results to selected accumulator.  
 The 'A' bit specifies the destination accumulator.  
 I-Words: 0.5  
 Cycles: 1

Example:1	SUB	A	;ACCA = ACCA - ACCB ; CORCON = 0x0000 (no saturation)
-----------	-----	---	--

Before execution		After execution	
ACCA	0x76_120F_098A _1212_4593	ACCA	0x52_CC7B_E05 5_7FDD_B56A
ACCB	0x23_4593_2934 _9234_9029	ACCB	0x23_4593_2934 _9234_9029
CORCON	0x00000000	CORCON	0x00000000
SR[15:8]	8'b00000000	SR[15:8]	8'b11000000 (OA/OB = 1)

Example:2	SUB	B	;ACCB = ACCB - ACCA ; CORCON = 0x0000 (no saturation)
-----------	-----	---	--

Before execution		After execution	
ACCA	0xFF_9022_2EE1 _5633_9078	ACCA	0xFF_9022_2EE1 _5633_9078
ACCB	0x00_2456_8F4C _0922_1245	ACCB	0x00_7FFF_FFFF_ FFFF_FFFF
CORCON	0x00000040	CORCON	0x00000040 (SATB)
SR[15:8]	8'b00000000	SR[15:8]	8'b00010100 (SB/SAB = 1)

Example 1:	SUB	w8, #1, B	;ACCB = ACCB - (w8>>1):32'b0
		Before Execution	After Execution
	w8	0x20446326	w8      0x20446326
	ACCB	0x00_2044_6326_4112_94 03	ACCB 0x00_1022_3193_4112_94 03

<b>SUB</b>							
<b>Signed Subtract from Accumulator</b>							
Syntax:	{label:}	SUB{.w} Ws,	{ Slit6, }	A			
		SUB.I [Ws],		B			
		[Ws++]					
		[Ws--]					
		[~Ws],					
		[++Ws],					
		[Ws+Wb],					
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Slit6 ∈ [-32 ... +31]					
Operation:		(ACC) - Shift <sub>Slit6</sub> (Sign-extend(Ws)) → ACC					
Status Affected:		OA, SA or OB, SB					
Encoding:	1100 00AL	www ssss ppU	UUkk	kkkk	1111		
Description:	<p>The operand contained at the Effective Address is assumed to be Q1.15 or Q1.31 fractional data for word and long data operations, respectively. The operand is read, then automatically sign-extended and zero-backfilled to create a value the same size as the accumulator.</p> <p>The value is then optionally (arithmetically) shifted before being subtracted from the target accumulator.</p> <p>The 'L' bit selects word or long word operation.</p> <p>The 'A' bit specifies the destination accumulator.</p> <p>The 's' bits specify the source register Ws.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'w' bits specify the offset register Wb.</p> <p>The 'k' bits encode the optional operand Slit6 which determines the amount of the accumulator preshift; if the operand Slit6 is absent, the literal bit field is set to all 0's.</p>						
Notes:	<ol style="list-style-type: none"> <li>Positive values of operand Slit6 represent arithmetic shift right.</li> <li>Negative values of operand Slit6 represent shift left.</li> <li>This instruction operates in Long or Word mode only.</li> </ol>						
I-Words:	1						
Cycles:	1						

Example:1	SUB	w8, #1, B	;ACCB = ACCB - (w8>>1):32'b0
Before execution		After execution	
w8	0x20446326	w8	0x20446326
ACCB	0x00_2044_6326_4 112_9403	ACCB	0x00_1022_3193_4 112_9403

Example:2	SUB	A	;ACCA = ACCA - ACCB ;CORCON = 0x0000 (no saturation)
Before execution		After execution	
ACCA	0x76_120F_098A _1212_4593	ACCA	0x52_CC7B_E05 5_7FDD_B56A
ACCB	0x23_4593_2934 _9234_9029	ACCB	0x23_4593_2934 _9234_9029
CORCON	0x00000000	CORCON	0x00000000
SR[15:8]	8'b00000000	SR[15:8]	8'b11000000 (OA/OB = 1)

Example:3	SUB	B	;ACCB = ACCB - ACCA ;CORCON = 0x0000 (no saturation)
-----------	-----	---	---

## Instruction Descriptions (N to XORWF) (continued)

Example:3	SUB	B	$\text{ACCB} = \text{ACCB} - \text{ACCA}$ ; CORCON = 0x0000 (no saturation)
Before execution		After execution	
ACCA 0xFF_9022_2EE1 _5633_9078		ACCA 0xFF_9022_2EE1 _5633_9078	
ACCB 0x00_2456_8F4C _0922_1245		ACCB 0x00_7FFF_FFFF_ FFFF_FFFF	
CORCON 0x00000040		CORCON 0x00000040 (SATB)	
SR[15:8] 8'b00000000		SR[15:8] 8'b00010100 (SB/SAB = 1)	

SUBB	Subtract Ws from Wb with Borrow					
Syntax:	{label:}	SUBB.b	Wb,	Ws,	Wd	
		SUBB.bz		[Ws],	[Wd]	
		SUBB{.w}		[Ws++],	[Wd++]	
		SUBB.l		[Ws--],	[Wd-]	
				[++Ws],	[++Wd]	
				[--Ws],	[--Wd]	
Operands:		$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$				
Operation:		$(Wb) - (Ws) - (C) \rightarrow Wd$				
Status Affected:		C, N, OV, Z				
Encoding:	S110      011L      dddd      ssss      pppq      qqww      wwUU      BU00					
Description:	Subtract the contents of the source register Ws and the Carry flag from the contents of the base register Wb and place the result in the destination register Wd. The Z bit is "sticky" (can only be cleared). The 'S' bit selects instruction size . The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.					
I-Words:	1 or 0.5					
Cycles:	1					

**Example 1:** SUBB.B W0, W1, W0 ; Sub. W1 and C from W0  
 (Byte mode)  
 ; Store result to W0

Before Instruction	After Instruction
W0 1732	W0 17ED
W1 7844	W1 7844
SR 0000	SR 0108 (DC, N = 1)

**Example 2:** SUBB W7, [W8++], [W9++] ; Sub. [W8] and C from W7 (Word mode)  
 ; Store result to [W9]  
 ; Post-increment W8  
 ; Post-increment W9

Before Instruction	After Instruction
W7 2450	W7 2450
W8 1808	W8 180A
W9 2022	W9 2024
Data 1808 92E4	Data 1808 92E4
Data 2022 A557	Data 2022 916B
SR 0000	SR 010C (DC, N, OV = 1)

<b>SUBB</b>		Subtract Short Literal from Ws with Borrow				
Syntax:	{label:}	SUBB.b	Ws,	lit7,	Wd	
		SUBB.bz	[Ws],		[Wd]	
		SUBB{.w}	[Ws++],		[Wd++]	
		SUBB.l	[Ws-],		[Wd-]	
			[++Ws],		[++Wd]	
			[--Ws],		[--Wd]	
Operands:		Ws ∈ [W0 ... W15]; lit7 ∈ [0 ... 127]; Wd ∈ [W0 ... W15]				
Operation:		(Ws) - lit7 - (C) → Wd				
		Note: The literal is zero-extended to the selected data size of the operation				
Status Affected:		C, N, OV, Z				
Encoding:	1110      011L      dddd      ssss      pppq      qqkk      kkkk      Bk10					
Description:	Subtract the zero-extended unsigned literal operand and the Carry bit from the contents of the source register Ws, and place the result in the destination register Wd. The Z bit is "sticky" (can only be cleared).					
	The 'L' and 'B' bits select operation data width.					
	The 'k' bits provide the literal operand (MSb in op[2]).					
	The 's' bits select the source register.					
	The 'd' bits select the destination register.					
	The 'p' bits select the source addressing mode.					
	The 'q' bits select the destination addressing mode.					
I-Words:	1					
Cycles:	1					

SUBB	Subtract Literal from Wn with Borrow							
Syntax:	{label:}	SUBB.b	lit16,	Wn				
		SUBB.bz						
		SUBB{.w}						
		SUBB.I						
Operands:		lit16 $\in [0 \dots 65535]$	;	Wn $\in [W0 \dots W15]$				
Operation:		$(Wn) - \text{lit16} - (C) \rightarrow Wn$						
		<b>Note:</b> The literal is zero-extended to the selected data size of the operation						
Status Affected:		C, N, OV, Z						
Encoding:		1100	011L	ssss	kkkk	kkkk	kkkk	B010
Description:		Subtract the zero-extended literal operand and the Carry bit from the contents of the Working register Wn, and place the result in the Working register Wn. In Byte mode, only the LSb of Wn is written. In Word mode, the word result is zero-extended to 32-bits and written to Wn.						
		The Z bit is "sticky" (can only be cleared). The 'L' and 'B' bits select operation data width. The 's' bits select the Working register. The 'K' bits specify the literal operand.						
I-Words:		1						
Cycles:		1						

SUBBR	Subtract f and Carry bit from Wn							
Syntax:	{label:}	SUBBR.b f	,Wn	{,WREG}				
		SUBBR.bz						
		SUBBR{.w}						
		SUBBR.I						
Operands:	$f \in [0 \dots 64KB]$ ; $Wn \in [W0 \dots W15]$							
Operation:	$(Wn) - (f) - (C) \rightarrow \text{destination designated by } D$							
Status Affected:	C, N, OV, Z							
Encoding:	1110	011L	ssss	ffff	ffff	ffff	ffff	BD01
Description:	<p>Subtract the contents of the file register and the Carry bit from the contents of the Working register and place the result in the destination designated by D. If the optional Wn is specified, D=0 and store result in Wn; otherwise, D=1 and store result in the file register.</p> <p>The Z bit is "sticky" (can only be cleared).</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>							
I-Words:	1							
Cycles:	1							

SUBBR	Subtract Wb from Ws with Borrow						
Syntax:	{label:}	SUBBR.b	Wb,	Ws,	Wd		
		SUBBR.bz		[Ws],	[Wd]		
		SUBBR{.w}		[Ws++],	[Wd++]		
		SUBBR.I		[Ws--],	[Wd-]		
				[++Ws],	[++Wd]		
				[--Ws],	[--Wd]		
Operands:		$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$					
Operation:		$(Ws) - (Wb) - (C) \rightarrow Wd$					
Status Affected:		C, N, OV, Z					
Encoding:	S110	111L	dddd	ssss	pppq	qqww	wwUU BU00
Description:	Subtract the contents of the base register Wb and the Carry flag from the contents of the source register Ws, and place the result in the destination register Wd. The Z bit is "sticky" (can only be cleared). The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.						
I-Words:	1 or 0.5						
Cycles:	1						

Example 1:	SUBBR.B W0, W1, W0	; Sub. W0 and C from W1 (Byte mode)
		; Store result to W0

Before Instruction			After Instruction		
W0	1732		W0	1711	
W1	7844		W1	7844	
SR	0000		SR	0001	(C = 1)

Example 2:	SUBBR W7,[W8++],[W9++]	; Sub. W7 and C from [W8] (Word mode)
		; Store result to [W9]
		; Post-increment W8
		; Post-increment W9

Before Instruction			After Instruction		
W7	2450		W7	2450	
W8	1808		W8	180A	
W9	2022		W9	2024	
Data 1808	92E4		Data 1808	92E4	
Data 2022	A557		Data 2022	6E93	
SR	0000		SR	0005	(OV, C = 1)

SUBBR	Subtract Ws from Short Literal with Borrow				
Syntax:	{label:}	SUBBR.b	Ws,	lit7	Wd
		SUBBR.bz	[Ws],		[Wd]
		SUBBR{.w}	[Ws++],		[Wd++]
		SUBBR.I	[Ws-],		[Wd-]
			[++Ws],		[++Wd]
			[--Ws],		[--Wd]
Operands:	$Ws \in [W0 \dots W15]; lit7 \in [0 \dots 127]; Wd \in [W0 \dots W15]$				
	Note: The literal is zero-extended to the selected data size of the operation				
Operation:	$lit7 - (Ws) - (C) \rightarrow Wd$				
Status Affected:	C, N, OV, Z				
Encoding:	1110	111L	dddd	ssss	pppq
Description:	Subtract the contents of the source register Ws and the Carry flag from the zero-extended unsigned literal, and place the result in the destination register Wd. The Z bit is "sticky" (can only be cleared). The 'L' and 'B' bits select operation data width. The 'k' bits provide the literal operand (MSb in op[2]). The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.				
I-Words:	1				
Cycles:	1				

<b>SUBB</b>		Subtract Wn and Carry bit from f											
Syntax:	{label:}	SUBB.b	f	,Wn	{,WREG}								
		SUBB.bz											
		SUBB{.w}											
		SUBB.I											
Operands:	$f \in [0 \dots 64KB]; Wn \in [W0 \dots W15]$												
Operation:	$(f) - (Wn) - (C) \rightarrow \text{destination designated by D}$												
Status Affected:	C, N, OV, Z												
Encoding:	1110	111L	ssss	ffff	ffff	ffff	ffff	BD01					
Description:	<p>Subtract the contents of the Working register and the Carry bit from the contents of the file register, and place the result in the destination designated by D. If the optional Wn is specified, D=0 and store result in Wn; otherwise, D=1 and store result in the file register.</p> <p>The Z bit is "sticky" (can only be cleared).</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>												
I-Words:	1												
Cycles:	1												

SUB	Subtract f from Wn							
Syntax:	{label:}	SUBR.b	f	,	Wn	{,WREG}		
		SUBR.bz						
		SUBR{.w}						
		SUBR.I						
Operands:	$f \in [0 \dots 64KB]$ ; $Wn \in [W0 \dots W15]$							
Operation:	$(Wn) - (f) \rightarrow \text{destination designated by D}$							
Status Affected:	C, N, OV, Z							
Encoding:	1110	010L	ssss	ffff	ffff	ffff	ffff	BD01
Description:	<p>Subtract the contents of the file register from the contents of the Working register, and place the result in the destination designated by D. If the optional Wn is specified, D=0 and store result in Wn; otherwise, D=1 and store result in the file register.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>							
I-Words:	1							
Cycles:	1							

SUB	Subtract Short Literal from Wn			
Syntax:	{label:}	SUB.I	lit5,	Wn
Operands:	lit5 $\in [0 \dots 31]$	Wn $\in [W0 \dots W15]$		
Operation:	$(Wn) - lit5 \rightarrow Wn$			
Status Affected:	C, N, OV, Z			
Encoding:	0111	011k	kkkk	ssss
Description:	<p>Subtract the zero-extended literal operand from the contents of the Working register and place the result in the Working register Wn. If literal &gt;31 and/or word or byte operation is required, assemble as SUBLW instruction.</p> <p>The 's' bits select the Working register.</p> <p>The 'k' bits specify the literal operand.</p>			
I-Words:	0.5			
Cycles:	1			

SUB	Subtract Short Literal from Ws				
Syntax:	{label:}	SUB.b	Ws,	lit7,	Wd
		SUB.bz	[Ws],		[Wd]
		SUB{.w}	[Ws++],		[Wd++]
		SUB.l	[Ws-],		[Wd-]
			[++Ws],		[++Wd]
			[--Ws],		[--Wd]
Operands:	$Ws \in [W0 \dots W15]; lit7 \in [0 \dots 127]; Wd \in [W0 \dots W15]$				
Operation:	$(Ws) - lit7 \rightarrow Wd$				
	<b>Note:</b> The literal is zero-extended to the selected data size of the operation				
Status Affected:	C, N, OV, Z				
Encoding:	1110	010L	dddd	ssss	pppq
Description:	<p>Subtract the zero-extended unsigned literal operand from the contents of the source register Ws, and place the result in the destination register Wd.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'k' bits provide the literal operand (MSb in op[2]).</p> <p>The 's' bits select the source register.</p> <p>The 'd' bits select the destination register.</p> <p>The 'p' bits select the source addressing mode.</p> <p>The 'q' bits select the destination addressing mode.</p>				
I-Words:	1				
Cycles:	1				

<b>SUB</b>		Subtract Literal from Wn												
Syntax:		{label:}	SUB.b	lit16,	Wn									
			SUB.bz											
			SUB{.w}											
			SUB.I											
Operands:		lit16 $\in$ [0 ... 65535]; Wn $\in$ [W0 ... W15]												
Operation:		(Wn) - lit16 $\rightarrow$ Wn												
Status Affected:		C, N, OV, Z												
Encoding:		1100	010L	ssss	kkkk	kkkk	kkkk	B010						
Description:		Subtract the zero-extended literal operand from the contents of the Working register, and place the result in the Working register Wn. If literal $\leq$ 31 and long operation is required, assemble as SUBLN instruction.												
		The 'L' and 'B' bits select operation data width.												
		The 's' bits select the Working register.												
		The 'k' bits specify the literal operand.												
I-Words:		1												
Cycles:		1												

<b>Example 1:</b>	<b>SUB.B #0x23, W0</b>	<b>; Sub. 0x23 from W0 (Byte mode)</b>
		<b>; Store result to W0</b>

Before Instruction			After Instruction		
W0	7804		W0	78E1	
SR	0000		SR	0008	(N = 1)

<b>Example 2:</b>	<b>SUB #0x108, W4</b>	<b>; Sub. 0x108 from W4 (Word mode)</b>
		<b>; Store result to W4</b>

Before Instruction			After Instruction		
W4	6234		W4	612C	
SR	0000		SR	0001	(C = 1)

SUBR	Subtract Wb from Ws							
Syntax:	{label:}	SUBR.b	Wb,	Ws,	Wd			
		SUBR.bz		[Ws],	[Wd]			
		SUBR{.w}		[Ws++],	[Wd++]			
		SUBR.I		[Ws--],	[Wd-]			
				[++Ws],	[++Wd]			
				[--Ws],	[--Wd]			
Operands:	$Wb \in [W0 \dots W15]; Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]$							
Operation:	$(Ws) - (Wb) \rightarrow Wd$							
Status Affected:	C, N, OV, Z							
Encoding:	S110	110L	dddd	ssss	pppq	qqww	wwUU	BU00
Description:	Subtract the contents of the base register Wb from the contents of the source register Ws, and place the result in the destination register Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.							
I-words:	1 or 0.5							
Cycles:	1							

Example 1:	SUBR.B W0, W1, W0	; Sub. W0 from W1 (Byte mode)
		; Store result to W0

Before Instruction		After Instruction	
W0	1732	W0	1712
W1	7844	W1	7844
SR	0000	SR	0001 (C = 1)

Example 2:	SUBR W7, [W8++], [W9++]	; Sub. W7 from [W8] (Word mode)
		; Store result to [W9]
		; Post-increment W8
		; Post-increment W9

Before Instruction		After Instruction	
W7	2450	W7	2450
W8	1808	W8	180A
W9	2022	W9	2024
Data 1808	92E4	Data 1808	92E4
Data 2022	A557	Data 2022	6E94
SR	0000	SR	0005 (OV, C = 1)

SUBR	Subtract Ws from Short Literal				
Syntax:	{label:}	SUBR.b	Ws,	lit7,	Wd
		SUBR.bz	[Ws],		[Wd]
		SUBR{.w} [Ws++],			[Wd++]
		SUBR.I	[Ws--],		[Wd-]
			[++Ws],		[++Wd]
			[~-Ws],		[~-Wd]
Operands:	$Ws \in [W0 \dots W15]; Wd \in [W0 \dots W15]; lit7 \in [1 \dots 63]$ (see note when lit7 = 0)				
Operation:	$lit7 - (Ws) \rightarrow Wd$ Note: The literal is zero-extended to the selected data size of the operation				
Status Affected:	C, N, OV, Z				
Encoding:	S111	101L	dddd	ssss	pppq
Description:	Subtract the contents of the source register Ws from the zero-extended unsigned literal, and place the result in the destination register Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 'k' bits provide the literal operand (MSb in op[2]). The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.				
I-Words:	1 or 0.5				
Cycles:	1				

SUB	Subtract Wn from f							
Syntax:	{label:}	SUB.b	f	Wn	{,WREG}			
		SUB.bz						
		SUB{.w}						
		SUB.I						
Operands:	$f \in [0 \dots 64KB]$ ; $Wn \in [W0 \dots W15]$							
Operation:	$(f) - (Wn) \rightarrow \text{destination designated by } D$							
Status Affected:	C, N, OV, Z							
Encoding:	1110	110L	ssss	ffff	ffff	ffff	ffff	BD01
Description:	<p>Subtract the contents of the Working register from the contents of the file register, and place the result in the destination designated by D. If the optional Wn is specified, D=0 and store the result in Wn; otherwise, D=1 and store the result in the file register.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>							
I-Words:	1							
Cycles:	1							

<b>SWAP</b>	<b>Word, Byte or Nibble Swap Wn</b>								
Syntax:	{label:}	SWAP.b	Wn						
		SWAP.bz							
		SWAP{.w}							
		SWAP.l							
Operands:		$Wn \in [W0 \dots W15]$							
Operation:		Byte mode: $Wn[7:4] \leftrightarrow Wn[3:0]$ Extended Byte mode: $Wn[7:4] \leftrightarrow Wn[3:0]; 24'h000000 \rightarrow Wn[31:8]$ Word mode: $Wn[15:8] \leftrightarrow Wn[7:0]$ Long Word mode: $Wn[31:24] \leftrightarrow Wn[7:0]; Wn[23:16] \leftrightarrow Wn[15:8]$							
Status Affected:		None							
Encoding:	1111	011L	UUUU	ssss	UUUU	UUUU	UUUU	B111	
Description:		If in Long Word mode, swap bytes in both the msw and lsw of Wn register. If in Word mode, swap bytes in the lsw of Wn register. Wn[31:16] is cleared. If in Extended Byte mode, swap nibbles in the LSb of Wn register. Wn[31:8] is cleared. If in Byte mode, swap nibbles in the LSb of Wn register. Wn[31:8] is unaffected. The 'L' and 'B' bits select operation data width. The 's' bits select the Working register. Long word byte swap is intended for data endian conversion.							
I-Words:	1								
Cycles:	1								

<b>TSTF</b>		<b>Test f and (optionally) Move to Wd</b>							
Syntax:	{label:} TST.b      f      {,Wnd}								
	TST.bz								
	TST{.w}								
	TST.I								
Operands:	$f \in [0 \dots 64KB]$ ; $Wnd \in [W0 \dots W14]$ ;								
Operation:	$(f) \rightarrow$ destination designated by D								
Status Affected:	Z, N								
Encoding:	1101	001L	dddd	ffff	ffff	ffff	ffff	BD01	
Description:	Read and test the contents of the file register. If the optional Wnd is specified (D=0), then store the data read in Wnd. Otherwise, D=1 and the only effect is to modify the status flags (i.e. test file register), and no write occurs.								
	The 'L' and 'B' bits select operation data width.								
	The 'D' bit selects if a data write occurs.								
	The 'f' bits select the address of the file register.								
I-Words:	1								
Cycles:	1								

ULNK	De-allocate Stack Frame			
Syntax:	{label:}	ULNK		
Operands:	None			
Operation:	W14 → W15 (W15) - 4 → W15 (TOS) → W14			
Status Affected:	None			
Encoding:	0111 001U UUUU 0011			
Description:	This instruction de-allocates a stack frame and adjusts the Stack Pointer and Frame Pointer.			
<b>Note:</b>	1. This instruction operates with long word aligned operands only.			
I-Words:	0.5			
Cycles:	1			

Example 1:	ULNK	;De-allocate stack frame	
		Before Instruction	After Instruction
		W14 0x5004	W14 0x5000
		W15 0x502C	W15 0x5000
		Data @0x5000 0x5000	Data @0x5000 0x5000
		SR[7:0] 8'b00000000	SR[7:0] 8'b00000001

<b>XOR</b>		<b>Exclusive or Wb and Ws</b>					
Syntax:	{label:}	XOR.b	Wb,	Ws,	Wd		
		XOR.bz		[Ws],	[Wd]		
		XOR{.w}		[Ws++],	[Wd++]		
		XOR.l		[Ws--],	[Wd-]		
				[++Ws],	[++Wd]		
				[--Ws],	[--Wd]		
				SR	SR		
Operands:		Wb ∈ [W0 ... W15]; Ws ∈ [W0 ... W15]; Wd ∈ [W0 ... W15]					
Operation:		(Wb).XOR.(Ws) → Wd					
Status Affected:		N, Z					
Encoding:		1111	000L	dddd	ssss	pppq	qqwwwwUU BU00
Description:		Compute exclusive OR of the contents of the source register Ws and the contents of the base register Wb, and place the result in the destination register Wd. The 'S' bit selects instruction size. The 'L' and 'B' bits select operation data width. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'q' bits select the destination addressing mode.					
Notes:		<ol style="list-style-type: none"> <li>When the SR is selected, SR data write will take priority over any SR update resulting from the XOR operation.</li> <li>.bz data size/mode is disallowed when writing to the SR.</li> </ol>					
I-Words:		1 or 0.5					
Cycles:		1					

<b>Example 1:</b>	<b>XOR.B W1, [W5++], [W9++]</b>	<b>; XOR W1 and [W5] (Byte mode)</b>
		<b>; Store result to [W9]</b>
		<b>; Post-increment W5 and W9</b>

Before Instruction	After Instruction
W1 AAAA	W1 AAAA
W5 2000	W5 2001
W9 2600	W9 2601
Data 2000 115A	Data 2000 115A
Data 2600 0000	Data 2600 00F0
SR 0000	SR 0008      (N = 1)

<b>Example 2:</b>	<b>XOR W1, W5, W9</b>	<b>; XOR W1 and W5 (Word mode)</b>
		<b>; Store the result to W9</b>

Before Instruction	After Instruction
W1 FEDC	W1 FEDC
W5 1234	W5 1234
W9 A34D	W9 ECE8
SR 0000	SR 0008      (N = 1)

XOR		Exclusive Or Ws and Short Literal						
Syntax:	{label:}	XOR.b	Ws,	lit7,	Wd			
		XOR.bz	[Ws],		[Wd]			
		XOR{.w}	[Ws++],		[Wd++]			
		XOR.l	[Ws-],		[Wd-]			
			[++Ws],		[++Wd]			
			[~-Ws],		[~-Wd]			
			SR		SR			
Operands:	$Ws \in [W0 \dots W15]; lit7 \in [0 \dots 127]; Wd \in [W0 \dots W15]$							
Operation:	$(Ws).XOR.lit7 \rightarrow Wd$							
<b>Note:</b> The literal is zero-extended to the selected data size of the operation								
Status Affected:	N, Z							
Encoding:	1111	000L	dddd	ssss	pppq	qqkk	kkkk	Bk10
Description:	Compute the exclusive OR of the contents of the source register Ws and the zero-extended literal operand, and place the result in the destination register Wd. The 'L' and 'B' bits select operation data width. The 'K' bits provide the literal operand (MSb in op[2]). The 'S' bits select the source register. The 'D' bits select the destination register. The 'P' bits select the source addressing mode. The 'Q' bits select the destination addressing mode.							
<b>Notes:</b>								
	<ol style="list-style-type: none"> <li>When the SR is selected, SR data write will take priority over any SR update resulting from the XOR operation.</li> <li>.bz data size/mode is disallowed when writing to the SR.</li> </ol>							
I-Words:	1							
Cycles:	1							

<b>XOR</b>		Exclusive Or Literal and Wn														
Syntax:	{label:}	XOR.b		lit16,	Wn											
		XOR.bz			SR											
		XOR{.w}														
		XOR.I														
Operands:	lit16 ∈ [0 ... 65535]; Wn ∈ [W0 ... W15]; Status Register (SR)															
Operation:	(Wn) .XOR. lit16 → Wn or (SR) .XOR. lit16 → SR															
Status Affected:	N, Z															
Encoding:	1101	000L	ssss	kkkk	kkkk	kkkk	kkkk	BT10								
Description:	Compute the exclusive OR of the zero-extended literal operand and the contents of the Working register Wn or SR, and place the result in the Working register Wn or SR. The 'L' and 'B' bits select operation data width.															
	The 's' bits select the Working register.															
	The 'k' bits specify the literal operand.															
	The 'T' bit selects between Ws (T = 0) and SR (T = 1) target registers.															
<b>Notes:</b>																
1.	When the SR is selected, SR data write will take priority over any SR update resulting from the XOR operation.															
2.	.bz data size/mode is disallowed when writing to the SR.															
I-Words:	1															
Cycles:	1															

<b>Example 1:</b>	XOR.B #0x23, W0	;	XOR 0x23 and W0 (Byte mode)
		;	Store result to W0
Before Instruction			After Instruction
W0 7804			W0 7827
SR 0000			SR 0000

<b>Example 2:</b>	XOR #0x108, W4	;	XOR 0x108 and W4 (Word mode)
		;	Store result to W4
Before Instruction			After Instruction
W4 6134			W4 603C
SR 0000			SR 0000

<b>XOR</b>		<b>Exclusive Or f and Wn</b>						
Syntax:	{label:}	XOR.b	f	,Wn	{,WREG}			
		XOR.bz						
		XOR{.w}						
		XOR.I						
Operands:		$f \in [0 \dots 64KB]; Wn \in [W0 \dots W15]$						
Operation:		$(f).XOR.(Wn) \rightarrow \text{destination designated by } D$						
Status Affected:		N, Z						
Encoding:		1111	000L	ssss	ffff	ffff	ffff	BD01
Description:		<p>Compute the XOR of the contents of the Working register and the contents of the file register, and place the result in the destination designated by D: If the optional Wn is specified, D=0 and store the result in Wn; otherwise, D=1 and store the result in the file register.</p> <p>The 'L' and 'B' bits select operation data width.</p> <p>The 'D' bit selects the destination.</p> <p>The 'f' bits select the address of the file register.</p> <p>The 's' bits select the Working register.</p>						
I-Words:		1						
Cycles:		1						

ZE	Zero Extend Ws							
Syntax:	{label:}	ZE.bz	Ws,	Wnd				
		ZE.w	[Ws],					
			[Ws++],					
			[Ws-],					
			[++Ws],					
			[--Ws],					
			[Ws+Wb],					
Operands:		Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W15]; Wnd ∈ [W0 ... W14]						
Operation:		If Byte mode: Ws[7:0] → Wd[7:0]; 0 → Wnd[31:8];						
		If Word mode: Ws[15:0] → Wd[15:0]; 0 → Wnd[31:16];						
Status Affected:	C, N, Z							
Encoding:	S111	111B	dddd	ssss	pppu	UUww	wwUU	UU00
Description:	Zero-extend an 8-bit or 16-bit signed value in Ws to a 32-bit value, then write the result back to Wnd. N is always cleared. C is always set.							
	The 'S' bit selects instruction size. The 'B' bit selects byte or word operation. The 's' bits select the source register. The 'd' bits select the destination register. The 'p' bits select the source addressing mode. The 'w' bits define the offset Wb.							
	<b>Note:</b> 1. This operation always writes a long word.							
I-Words:	1 or 0.5							
Cycles:	1							

Example 1:	ZE.bz	W0, W3	;Zero extend a byte in W0, store it in W3
		Before Instruction W0 0x12347879 W3 0x00000000 SR[7:0] 8'b00000000	After Instruction W0 0x12347879 W3 0x00000079 SR 8'b00000001 (C = 1)

Example 2:	ZE	W0, W3	;Zero extend a word in W0, store it in W3
		Before Instruction W0 0x12347879 W3 0x00000000 SR[7:0] 8'b00000000	After Instruction W0 0x12347879 W3 0x00007879 SR[7:0] 8'b00000001 (C = 1)

## 4.8. FPU Instruction Encoding and Opcode Field Description

**Table 4-11.** Symbols Used in FPU Instruction Encoding Field Description

Field	Description
lit16	16-bit unsigned literal $\{0...65535\}$
index	Literal address index into 32 entry SP and DP constant tables
.s	32-bit Single Precision selection
.d	64-bit Double Precision selection
Fd	One of 32 destination floating-point registers $\{F0..F31\}$ (Register Direct)
Fb	One of 32 source floating-point registers $\{F0..F31\}$ (Register Direct)
Fs	One of 32 source floating-point registers $\{F0..F31\}$ (Register Direct)
SUB, INF, FN, FZ, FNAN GT, LT, EQ, UN SUBO, HUGI, INX, UDF, OVF, DIV0, INVAL	FPU status bits (FSR) <b>Note:</b> Flags for exceptions that cannot be signaled by an instruction are always cleared by that instruction (sticky flags are unaffected)

**Table 4-12.** FPU Instruction Opcode Field Descriptions

Field	Description
P	Selects Single Precision (P=0) or Double Precision (P=1) operation
R	Selects between FPU coprocessor special registers (R=1) or F-reg (R=0)
U	Unused (don't care) instruction bit. Assembler to assign '0'
ddddd	R=0: FPU coprocessor Fd destination register select: 00000=F0; 11111=F31 R=1: FPU coprocessor special register select
eee	FPU rounding mode selection
kkkk kkkk kkkk kkkk	16-bit literal field, constant data
sssss	R=0: FPU coprocessor Fs source register select: 00000=F0; 11111=F31 R=1: FPU coprocessor special register select

## 4.9. Floating-Point Instruction Description

ABS	Absolute value of Fs						
Syntax:	{label:}	ABS.s	Fs,	Fd			
		ABS.d					
Operands (.s):		Fs ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]					
Operation:		ABS Fs → Fd					
Status Affected:		SUBO					
Exceptions Possible:		SUBO					
Encoding:	1000	100P	zzdd	ddds	ssss	UUUU	UUUU
Description:							
		Clear the sign of the contents of the source register Fs, and place the result in the destination register Fd.					
		The 'P' bit selects single (.s) or double precision (.d) operation.					
		The 'z' bits select the target coprocessor.					
		The 's' bits select the source register.					
		The 'd' bits select the destination register.					
I-Words:	1						
SP Execute Cycles:	1						
DP Execute Cycles:	1						
Repetition Rate:	1						

**Note:** This instruction only affects the sign bit and will not quit an sNaN input operand.

Example:1	ABS.s	F0, F3	;Find absolute value of a number in F0 and store result in F3	
			Before execution	
			F0	0x40800001 -1.0000001
			F3	0x00000000
			FSR[7:0]	8'b00000000
			After execution	
			F0	0x40800001
			F3	0x40800001
			FSR[7:0]	8'b00000000

Example:2	ABS.d F0, F2	;Find absolute of a number in F1:F0 and store result in F3:F2		
		Before execution		
			F1:F0	0x8000000000 0000001
			F3:F2	0x0000000000 0000000
			FSR[7:0]	8'b00000000
		After execution		
			F1:F0	0x8000000000000001
			F3:F2	0x0000000000000001
			FSR[7:0]	8'b01000000(SUBO) (SUBO)

ADD	Add Fb and Fs							
Syntax:	{label:}	ADD.s Fb,	Fs,	Fd				
		ADD.d						
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]						
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]						
Operation:		Fb + Fs → Fd						
Status Affected:		SUBO, INX, UDF, OVF, INVAL (see Note 1)						
Exceptions Possible:		SUBO, INX, UDF, OVF, INVAL						
Encoding:	1000	000P	zzdd	ddds	ssss	www	wuuu	0010
Description:	Add the contents of the source register Fb and the contents of the register Fs, then place the result in the destination register Fd. $(\infty + (-\infty))$ will result in distinguished qNaN and signal an INVAL exception. The 'P' bit selects single (.s) or double precision (.d) operation. The 'z' bits select the target coprocessor. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register.							
<b>Notes:</b>								
	1. If not already set, corresponding sticky exception status will also be set. 2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.							
I-Words:	1							
SP Execute Cycles:	2							
DP Execute Cycles:	2							
Repetition Rate:	1							

AND	AND Floating-Point Control/Status Register and Literal							
Syntax:	{label:}	AND	lit16,	FSR				
				FCR				
				FEAR				
Operand		lit16	$\in [0 \dots 65535]$					
Operation:		(FP special register[15:0]).AND.lit16	$\rightarrow$	FP special register[15:0]				
Status Affected:		None	(other than as a result of the instruction AND operation)					
Exceptions Possible:		None						
Encoding:	1001	000U	zzss	kkkk	kkkk	kkkk	kkkk	0010
Description:	Compute the AND of the lsw contents of the selected floating-point coprocessor special register and the literal operand, and write the result back into the lsw of the selected register. The 'z' bits select the target coprocessor.		The 's' bits select the floating-point coprocessor source register.					
			The 'k' bits specify the 16-bit literal operand.					
			<b>Notes:</b>					
			1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.					
			2. The effect of this instruction is solely to clear bits within the target register. No FPU interrupts can result from its execution.					
			3. FEAR destination includes the EACE control/status bit.					
I-Words:	1							
Execute Cycles:	1							
Repetition Rate:	1							

Example:1	AND	#FFF8, FCR	;FCR   0x7 -> To clear INVAL, DIV0 and OVM masks
	Before execution		After execution
	FCR	0x00000300	FCR
			0x00000307

Evaluate the Cosine of Fs							
COS							
Syntax:	{label:}	COS.s	Fs,	Fd			
Operands (.s):		Fs ∈ [F0 ... F31]	Fd ∈ [F0 ... F31]				
Operation:		cos(Fs) → Fd					
Status Affected:		SUBO, INX, UDF, INVAL (see Note 1)					
Exceptions Possible:		SUBO, INX, UDF, INVAL					
Encoding:	1000	110P	zzdd	ddds	ssss	UUUU	UUUU
Description:	Calculate the cosine of the contents of the source register Fs (where Fs = [2πk + θ rads]), and place the result in the destination register Fd. The 'z' bits select the target coprocessor.						
	The 's' bits select the source register						
	The 'd' bits select the destination register.						
<b>Notes:</b>							
1.	If not already set, corresponding sticky exception status will also be set.						
2.	This operation is only supported using single precision floating-point.						
3.	Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:	1						
SP Execute Cycles:	4						
Repetition Rate:	1						

Example:1	COS.s	F0, F3	;Find Cosine of radians in F0 and store result in F3
Before execution		After execution	
F0 0x4016DB6E		F0 0x4016DB6E	
F3 0x00000000		F3 0xBF3530E2	
FSR[7:0] 8'b00000000		FSR[7:0] 8'b00000000	
Example:2	COS.s	F0, F3	;Find Cosine of radians in F0 and store result in F3
Before execution		After execution	
F0 0xFF800000		F0 0xFF800000	
F3 0x00000000		F3 0x7FC00001 (distinguished qNaN)	
FSR[7:0] 8'b00000000		FSR[7:0] 8'b00000001 (INVAL)	

CPQ	Compare Fb with Fs, Quiet Signaling							
Syntax:	{label:}	CPQ.s	Fb,	Fs				
		CPQ.d						
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]						
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]						
Operation:		Fb - Fs (with IEEE signaling predicates)						
Status Affected:		LT, GT, EQ, UN, SUBO, INVAL (see Note 1)						
Exceptions Possible:		SUBO, INVAL						
Encoding:	1000	101P	zzUU	UUUs	ssss	wwWW	wUUU	1010
Description:	<p>Compare of the contents of the source registers Fb and Fs (Fb - Fs) then set flags, but do not store the result. An Invalid signaling exception will only be generated if either source operand is a sNaN. A qNaN source operand will not result in a signaling exception. If any source operand is a sNaN or qNaN, then set FSR.UN = 1.</p> <p>The LT, GT and EQ status bits are set depending on the result.</p> <p>The 'P' bit selects single (32-bit) or double precision (64-bit) operation.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p>							
	<b>Notes:</b>							
	<ol style="list-style-type: none"> <li>If not already set, corresponding sticky exception status will also be set.</li> <li>Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> <li>This instruction considers -0 and +0 as equivalent.</li> </ol>							
I-Words:	1							
SP Execute Cycles:	1							
DP Execute Cycles:	1							
Repetition Rate:	1							

Example:1	CPQ.s	F0, F3	;(Quiet Signaling)Compare F0 and F3; update FSR[23:16]
		Before execution	After execution
		F0 0xFF7FFFFB	F0 0xFF7FFFFB
		F3 0x80800001	F3 0x80800001
		FSR[23:16] 8'b00000000	FSR[23:16] 8'b000000100 (LT)

Example:2	CPQ.d	F0, F2	;(Quiet Signaling)Compare F1:F0 and F3:F2; update FSR[23:16]
		Before execution	After execution
		F1:F0 0xFFFF8000000000 000E	F1:F0 0xFFFF8000000000 000E
		F3:F2 0xFFFF8000000000 0008	F3:F2 0xFFFF8000000000 0008
		FSR[23:16] 8'b00000000	FSR[23:16] 8'b00000001 (UN)
		FSR[7:0] 8'b00000000	FSR[23:16] 8'b00000000

CPS	Compare Fb with Fs, Signaling							
Syntax:	{label:}	CPS.s	Fb,	Fs				
		CPS.d						
Operands (.s):		Fs	∈ [F0 ... F31];	Fb	∈ [F0 ... F31]			
Operands (.d):		Fs	∈ [F0, F2 ... F30];	Fb	∈ [F0, F2 ... F30]			
Operation:		Fb - Fs (with IEEE signaling predicates)						
Status Affected:		LT, GT, EQ, UN, SUBO, INVAL (see Note 1)						
Exceptions Possible:		SUBO, INVAL						
Encoding:	1000	101P	zzUU	UUUs	ssss	www	wUUU	1110
Description:	<p>Compare of the contents of the source registers Fb and Fs (Fb - Fs) then set flags, but do not store the result. If either source operand is a sNaN or a qNaN, an Invalid signaling exception will be generated.</p> <p>If any source operand is a sNaN or qNaN, then set FSR.UN = 1.</p> <p>The LT, GT and EQ status bits are set depending on the result.</p> <p>The 'P' bit selects single (32-bit) or double precision (64-bit) operation.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p>							
	<b>Notes:</b>							
	<ol style="list-style-type: none"> <li>1. If not already set, corresponding sticky exception status will also be set.</li> <li>2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> <li>3. This instruction consider -0 and +0 as equivalent.</li> </ol>							
I-Words:	1							
SP Execute Cycles:	1							
DP Execute Cycles:	1							
Repetition Rate:	1							

Example:1	CPS.s	F0, F3	;(Signaling)Compare F0 and F3 ; update FSR[23:16]
-----------	-------	--------	--

Before execution		After execution	
F0	0xFF7FFFFB	F0	0xFF7FFFFB
F3	0x80800001	F3	0x80800001
FSR[23:16]	8'b00000000	FSR[23:16]	8'b00000100 (LT)

Example:2	CPS.d	F0, F2	;(Signaling)Compare F1:F0 and F3:F2 ; update FSR[23:16]
-----------	-------	--------	--

Before execution		After execution	
F1:F0	0xFFFF80000000 000E	F1:F0	0xFFFF80000000 000E
F3:F2	0xFFFF80000000 0008	F3:F2	0xFFFF80000000 0008
FSR[23:16]	8'b00000000	FSR[23:16]	8'b00000001 (UN)
FSR[7:0]	8'b00000000	FSR[23:16]	8'b00000001 (INVAL)

DI2F Convert Double Word (64-bit) Integer to Floating Point							
Syntax:	{label:}	DI2F.s{rnd}	Fs,	Fd			
		DI2F.d{rnd}					
Operands (.s):		Fs ∈ [F0, F2 ... F30]; Fd ∈ [F0 ... F31]; rnd ∈ [e, z, p, n]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]; rnd ∈ [e, z, p, n]					
Operation:		Fs (integer) → Fd (float)					
Status Affected:		INX <sup>1</sup>					
Exceptions Possible:		INX					
Encoding:	1000	111P	zzdd	ddds	ssss	UUUU	Ueee
Description:	Convert the 64-bit integer contents of the source register Fs to floating point, round, then place the result in the destination register Fd. If rounding mode {rnd} is not specified, it will default to that defined by FCR.RND[1:0]. If the integer cannot be represented exactly as a floating point value, the Inexact signaling exception will be generated. The 'P' bit selects single (.s) or double precision (.d) operation. The 'e' bits define the rounding mode applied. The 'z' bits select the target coprocessor. The 's' bits select the source register. The 'd' bits select the destination register.						
<b>Notes:</b>	1. If not already set, corresponding sticky exception status will also be set. 2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:	1						
SP Execute Cycles:	2						
DP Execute Cycles:	2						
Repetition Rate:	1						

Example:1	DI2F.sp F0, F2 ;Convert 64-bit double long-word integer to a SP floating-point value (with rounding = pos)
-----------	--

Before execution		After execution	
F1:F0	0x7FFFFFFFFF FFFB	F1:F0	0x7FFFFFFFFF FFFB
F2	0x00000000	F2	0x5EFFFF
FSR[7:0]	8'b00000000	FSR[7:0]	8'b00010000 (INX)

Example:2 DI2F.dz F0, F2	;Convert 64-bit double long-word integer to a DP floating-point value (with rounding = zero)
--------------------------	--

Before execution		After execution	
F1:F0	0xFFFFFFFFFFFFFF	F1:F0	0xFFFFFFFFFFFFFF
F3:F2	0x0000000000000000	F3:F2	0xBFF0000000000000

DIV	Signed Floating-Point Divide						
Syntax:	{label:}	DIV.s	Fb	Fs,	Fd		
			DIV.d				
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]					
Operation:		Fb ÷ Fs → Fd					
Status Affected:		SUBO, INX, UDF, OVF, DIV0, INVAL (see Note 1)					
Exceptions Possible:		SUBO, INX, UDF, OVF, DIV0, INVAL					
Encoding:	1000	010P	zzdd	ddds	ssss	www	wuuu
Description:	Divide the contents of the source register Fb with the contents of the source register Fs, then place the result in the destination register Fd.						
	Result of (∞/0) is correctly signed infinity. Whereas, ((0/0) or (∞/∞)) will result in distinguished qNaN and signal an FPU INVAL exception.						
	The 'P' bit selects single (.s) or double precision (.d) operation.						
	The 'z' bits select the target coprocessor.						
	The 's' bits select the source register.						
	The 'w' bits select the base register.						
	The 'd' bits select the destination register.						
Notes:							
	1.	If not already set, corresponding sticky exception status will also be set.					
	2.	An attempt to divide by zero may or may not signal DIV0 depending upon operands.					
	3.	Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.					
I-Words:	1						
SP Execute Cycles:	11						
DP Execute Cycles:	32						

Example:1	DIV.s F0, F1, F3	;Divide the contents in F0 by F1 and store the result in F3
	Before execution	After execution
	F0 0x7F00000A	F0 0x7F00000A
	F1 0xFF0000AA	F1 0xFF0000AA
	F3 0x00000000	F3 0xBF7FFEC0
	FSR[7:0] 8'b00000000	FSR[7:0] 8'b00000000

Example:2	DIV.d F0, F4, F8	;Divide the contents in F1:F0 by F5:F4 and store the results in F9:F8
	Before execution	After execution
	F1:F0 0xFFEA36E2EB1 C432D	F1:F0 0xFFEA36E2EB1 C432D
	F5:F4 0x00000000000A 00000	F5:F4 0x00000000000A 00000
	F9:F8 0x000000000000 00000	F9:F8 0xFFFFFFF FFF
	FSR[7:0] 8'b00000000	FSR[7:0] 8'b01010100 (INX, OVF, SUBO)

FLIM	Force Signed Data Limit							
Syntax:	{label:}	FLIM.s	Fb,	Fs,	Fd			
		FLIM.d						
Operands (.s):		$Fs \in [F0 \dots F31]; Fb \in [F0 \dots F31]; Fd \in [F0 \dots F31]$						
Operands (.d):		$Fs \in [F0, F2 \dots F30]; Fb \in [F0, F2 \dots F30]; Fd \in [F0, F2 \dots F30]$						
Operation:		IF (Fd) > (Fs) THEN (Fs) → (Fd); IF (Fd) < (Fb) THEN (Fb) → (Fd);						
Status Affected:		SUBO, INVAL (see Note 1)						
Exceptions Possible:		SUBO, INVAL						
Encoding:	1001	001P	zzdd	ddds	ssss	www	wUUU	0010
Description:	<p>Simultaneously compare a signed floating-point value in Fd to a maximum signed floating-point value held in Fs and a minimum signed floating-point value held in Fb.</p> <p>If Fd is greater than Fs, set Fd to the limit value held in Fs.</p> <p>If Fd is less than Fb, set Fd to the limit value held in Fb.</p> <p>If Fd is less than or equal to the maximum limit value in Fs, and greater than or equal to the minimum limit value in Fb, Fd is not modified (i.e., data is within range and limits are not applied). See note 2.</p> <p>If any of the operands (limit values or value to be tested) are an sNaN or a qNaN, the instruction will return a qNaN as the result. In addition, if any of the operands are an sNaN, the INVAL status will be set.</p> <p>Should the (maximum) limit value Fs be inadvertently set to less than (minimum) limit value Fb, the instruction will return the distinguished qNaN as the result and set the INVAL status.</p> <p>This is a signed comparison operation. The limits may both be of the same sign.</p> <p>The 'P' bit selects single (32-bit) or double precision (64-bit) operation.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'd' bits select the destination register.</p> <p>If not already set, corresponding sticky exception status will also be set.</p> <p>Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</p>							
I-Words:	1							
SP Execute Cycles:	1							
DP Execute Cycles:	1							
Repetition Rate:	1							

Example:1	FLIM.s F3, F0, F1	;Limit the value in F3 to be within values in F0 and F1
		Before execution
		F0 0xFF669595
		F1 0xFF000000
		F3 0x44480000
		After execution
		F0 0xFF000000
		F1 0xFF000000
		F3 0x44480000
Example:2	FLIM.d F6, F4, F0	;Limit the value in F7:F6 to be within values in F5:F4 and F1:F0
		Before execution
		F1:F0 0xC02400000000 00000
		F5:F4 0x402400000000 00000
		F9:F8 0xC01400000000 00000
		After execution
		F1:F0 0xC01400000000 00000
		F5:F4 0x402400000000 00000
		F9:F8 0xC01400000000 00000

<b>F2DI</b>		Convert Floating-Point Fs to Double Word (64-bit) Integer						
Syntax:	{label:}	F2DI.s{rnd}	Fs,	Fd				
		F2DI.d{rnd}						
Operands (.s):		Fs ∈ [F0 ... F31]; Fd ∈ [F0, F2 ... F30]; rnd ∈ [e, z, p, n]						
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]; rnd ∈ [e, z, p, n]						
Operation:		Fs (float) → Fd (integer)						
Status Affected:		SUBO, HUGI, INX, INVAL (see Note 1)						
Exceptions Possible:		SUBO, HUGI, INX, INVAL						
Encoding:		1000      111P      zzdd      ddds      ssss      UUUU      Ueee      0110						
Description:		Convert the floating-point contents of the source register Fs to a rounded 64-bit integer, then place the result in the destination register Fd. If rounding mode {rnd} is not specified, it will default to that defined by FCR.RND[1:0].						
		If the source register contains ±NaN or ±∞, the INVAL exception will be signaled and the result will be the integer indefinite value of 0x8000_0000_0000_0000 (maximum negative integer).						
		If the source register contains a finite floating point number, but the result rounds to a number greater than $2^{63}$ .1 or less than $-2^{63}$ , the HUGI and INVAL exceptions will be signaled and the default result will be the integer indefinite value of 0x7FFF_FFFF_FFFF_FFFF or 0x8000_0000_0000_0000, respectively.						
		If the source register contains a subnormal value, the result will always be zero and the INX exception will be signaled.						
		The 'P' bit selects single (.s) or double precision (.d) operation.						
		The 'e' bits define the rounding mode applied.						
		The 'z' bits select the target coprocessor.						
		The 's' bits select the source register.						
		The 'd' bits select the destination register.						
		<b>Notes:</b>						
		1. If not already set, corresponding sticky exception status will also be set.						
		2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:		1						
SP Execute Cycles:		1 (subject to change)						
DP Execute Cycles:		2 (subject to change)						
Repetition Rate:		1						

**Example:1 F2DI.sp F0, F2 ;Convert SP float value to a 64-bit double long-word integer(with rounding = pos)**

Before execution		After execution	
F0	0xD7FFFFFF	F0	0xD7FFFFFF
F3:F2	0x0000000000000000	F3:F2	0xFFFFE000004000000
FSR[7:0]	8'b00000000	FSR[7:0]	8'b00010000 (INX)

**Example:2 F2DI.dz F0, F2 ;Convert DP float value to a 64-bit double long-word integer(with rounding = zero)**

Before execution		After execution	
F1:F0	0xC2E4F8B5877 B8000	F1:F0	0xC2E4F8B5877 B8000
F3:F2	0x000000000000 00000	F3:F2	0xFFFF583A53C 42400

F2LI								Convert Floating-Point Fs to Long Word (32-bit) Integer							
Syntax:	{label:}	F2LI.s{rnd}		Fs,		Fd									
		F2LI.d{rnd}													
Operands (.s):		Fs ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]; rnd ∈ [e, z, p, n]													
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fd ∈ [F0 ... F31]; rnd ∈ [e, z, p, n]													
Operation:		Fs (float) → Fd (integer)													
Status Affected:		SUBO, HUGI, INX, INVAL (see Note 1)													
Exceptions Possible:		SUBO, HUGI, INX, INVAL													
Encoding:	1000	111P	zzdd	ddds	ssss	UUUU	Ueee	0010							
Description:	<p>Convert the floating-point contents of the source register Fs to a rounded 32-bit integer, then place the result in the destination register Fd. If rounding mode {rnd} is not specified, it will default to that defined by FCR.RND[1:0].</p> <p>If the source register contains ±NaN or ±∞, the INVAL exception will be signaled and the result will be the integer indefinite value of 0x8000_0000 (maximum negative integer).</p> <p>If the source register contains a finite floating-point number but the result rounds to a number greater than <math>2^{31}</math>-1 or less than -<math>2^{31}</math>, the HUGI and INVAL exceptions will be signaled and the default result will be the integer indefinite value of 0x7FFF_FFFF or 0x8000_0000, respectively.</p> <p>If the source register contains a subnormal value, the result will always be zero and the INX exception will be signaled.</p> <p>The 'P' bit selects single (.s) or double precision (.d) operation.</p> <p>The 'e' bits define the rounding mode applied.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'd' bits select the destination register.</p>														
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. If not already set, corresponding sticky exception status will also be set.</li> <li>2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> </ol>														
I-Words:	1														
SP Execute Cycles:	1														
DP Execute Cycles:	2														
Repetition Rate:	1														

Example:1 F2LI.sp F0, F1 ;Convert SP float value to a long-word integer(with rounding = pos)	
Before execution	After execution
F0 0xCA0EF725	F0 0xCA0EF725
F1 0x00000000	F1 0xFFDC4237
FSR[7:0] 8'b00000000	FSR[7:0] 8'b00010000 (INX)

Example:2 F2LI.dz F0, F2 ;Convert DP float value to a long-word integer(with rounding = zero)	
Before execution	After execution
F1:F0 0xC09F40000000 00000	F1:F0 0xC09F400000000000
F2 0x00000000	F2 0xFFFFF830

FBRA CC	FPU Coprocessor Branch							
Syntax:	{label:}	FBRA	CC,	Expr				
Operands:	Expr is resolved by the linker to a signed word offset (slit20)							
Operation:	Condition (cc) = See Table x.y							
	If (cc) then {							
	If slit20 = 1 then skip next (16-bit) instruction							
	Else if (slit20 = 2 && next_op[31] = 1) then skip next (32-bit) instruction							
	Else (PC+4) + 2*slit20 → PC							
	}							
	Else no branch							
Status Affected:	None							
Encoding:	1011	bbbb	nnnn	nnnn	nnnn	nnnn	zz10	
	See <a href="#">Table 4-13</a> for the value of "bbbb"							

**Table 4-13.** FBRA Conditions (continued)

FBRA CC	FPU Coprocessor Branch
Description:	If the branch condition is met, then the instruction will either skip the next 16-bit or 32-bit instruction, or it will branch to any size of instruction with a forward or backward range of 1MB.
	If the 2's complement byte offset value '2*slit20' (the PC offset) equals two, the conditional branch will execute as a conditional skip of one 16-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.
	If the 2's complement byte offset value '2*slit20' (the PC offset) equals four and the instruction after the branch is a 32-bit opcode (see note), the conditional branch will execute as a conditional skip of that 32-bit instruction. This instruction will be speculatively executed as not skipped until such time that the branch decision can be determined.
	If requirements for a conditional skip are not met, the 2's complement byte offset value '2*slit20' is added to the PC to create a new PS word address. Since the PC will have incremented to fetch the next instruction, the new address will be (PC+4) + 2*slit20.
	Branch prediction is based on the direction of the branch. If the branch is backwards (negative), it will be predicted as taken. If the branch is forwards (positive), it will be predicted as not taken.
	The two instructions that follow the branch will then be speculatively executed in the predicted path until such time that the actual branch decision can be assessed. A correctly predicted branch will continue instruction execution in the same path unabated. An incorrectly predicted branch will abort the speculatively executed instructions and start execution from the correct path.
	The 'n' bits are a signed literal that specifies the number of PS words offset from (PC+4).
	The 'z' bits select the target coprocessor.
	FPU co-processor supports number of branch conditions (cc) as mentioned in following table.

**Table 4-13.** FBRA Conditions

cc	Conditions	Encoding (bbbb)
EQ	FSR.EQ	0000
UNE	(FSR.GT    FSR.LT    FSR.UN)	0001
NE	(FSR.GT    FSR.LT)	0010
UEQ	(FSR.EQ    FSR.UN)	0011
GT	FSR.GT	0100
ULE	(FSR.LT    FSR.EQ    FSR.UN)	0101
GE	(FSR.GT    FSR.EQ)	0110
ULT	(FSR.LT    FSR.UN)	0111
LT	FSR.LT	1000
UGE	(FSR.GT    FSR.EQ    FSR.UN)	1001
LE	(FSR.LT    FSR.EQ)	1010
UGT	(FSR.GT    FSR.UN)	1011
OR	(FSR.GT    FSR.LT    FSR.EQ)	1100
UN	FSR.UN	1101

**Notes:**

1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.
2. Should the byte offset equal four and instruction after the branch is a 16-bit opcode, a branch will occur (over the next two 16-bit ops) if the condition is met.
3. Branch conditions are evaluated within the coprocessor.

I-Words:

1

Cycles:

1 (2 or 3)

IOR	Inclusive OR Floating-Point Control/Status Register and Literal							
Syntax:	{label:}	IOR	lit16,	FSR				
				FCR				
				FEAR				
Operand	lit16 ∈ [0 ... 65535]							
Operation:	(FP special register[15:0]).IOR.lit16 → FP special register[15:0]							
Status Affected:	None (other than as a result of the instruction OR operation)							
Exceptions Possible:	None							
Encoding:	1001      000U      zzss      kkkk      kkkk      kkkk      kkkk      0110							
Description:	Compute the inclusive OR of the contents of the lsw of the selected floating-point coprocessor register and the literal operand, and write the result back into the lsw of the selected register. The 'z' bits select the target coprocessor.							
	The 's' bits select the floating-point coprocessor source register.							
	The 'k' bits specify the 16-bit literal operand.							
	<b>Notes:</b>							
	1. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.							
	2. The effect of this instruction is solely to set bits within the target register. No FPU interrupts can result from its execution.							
	3. FEAR destination includes the EACE control/status bit.							
I-Words:	1							
Execute Cycles:	1							
Repetition Rate:	1							

Example:1	IOR #7, FCR	:FCR   0x7 -> To set INVAL, DIV0 and OVM masks
	Before execution FCR      0x00000300	After execution FCR      0x00000307

LI2F Convert Long Word (32-bit) Integer to Floating-Point							
Syntax:	{label:}	LI2F.s{rnd}	Fs,	Fd			
		LI2F.d{rnd}					
Operands (.s):		Fs ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]; rnd ∈ [e, z, p, n]					
Operands (.d):		Fs ∈ [F0 ... F31]; Fd ∈ [F0, F2 ... F30]; rnd ∈ [e, z, p, n] <sup>3</sup>					
Operation:		Fs (integer) → Fd (float)					
Status Affected:		INX <sup>1,3</sup>					
Exceptions Possible:		INX					
Encoding:	1000	111P	zzdd	ddds	ssss	UUUU	Ueee
Description:		1010					
		Convert the 32-bit integer contents of the source register Fs to floating-point, round, and then place the result in the destination register Fd. If rounding mode {rnd} is not specified, it will default to that defined by FCR.RND[1:0]. If the integer cannot be represented exactly as a floating-point value, the Inexact signaling exception will be generated (see note 3). The 'P' bit selects single (.s) or double precision (.d) operation.					
		The 'e' bits define the rounding mode applied.					
		The 'z' bits select the target coprocessor.					
		The 's' bits select the source register.					
		The 'd' bits select the destination register.					
		<b>Notes:</b>					
		1. If not already set, corresponding sticky exception status will also be set.					
		2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.					
		3. Long integer to Double Precision float (LI2F.d) conversion are always exact, so Inexact (INX) status will never be set. Similarly, rounding modes will have no effect on LI2F.d result.					
I-Words:	1						
SP Execute Cycles:	1						
DP Execute Cycles:	1						
Repetition Rate:	1						

Example:1 LI2F.se F0, F1 ;Convert long-word integer to SP float value (with even rounding)							
Before execution				After execution			
F0	0xFFFFF830			F0	0xFFFFF830		
F1	0x00000000			F1	0xC4FA0000		

Example:2 LI2F.d F0, F2 ;Convert long-word integer to DP float value							
Before execution				After execution			
F0	0xAFF0BDBC			F0	0xAFF0BDBC		
F3:F2	0x0000000000000000			F3:F2	0xC1D403D091000000		

MAC Floating-Point Signed Multiply and Accumulate							
Syntax:	{label:}	MAC.s	Fb,	Fs,	Fd		
		MAC.d					
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]					
Operation:		(Fd) + ((Fb) * (Fs)) → Fd;					
Status Affected:		SUBO, INX, UDF, OVF, INVAL (see Note 1)					
Exceptions Possible:		SUBO, INX, UDF, OVF, INVAL					
Encoding:	1000	010P	zzdd	ddds	ssss	www	wuuu
Description:		Multiply the signed contents of the source register Fb with the signed contents of the source register Fs, then add the product to Fd.					
		The 'P' bit selects single (.s) or double precision (.d) operation.					
		The 'z' bits select the target coprocessor.					
		The 's' bits select the source register.					
		The 'w' bits select the base register.					
		The 'd' bits select the destination register.					
Notes:							
		1. If not already set, corresponding sticky exception status will also be set.					
		2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.					
I-Words:	1						
SP Execute Cycles:	3						
DP Execute Cycles:	4						
Repetition Rate:	1						

Example:1	MAC.s F3, F0, F1	;Multiply contents of F0, F1 and add the results with content in F3
	Before execution	After execution
	F0 0x80FFFFFF	F0 0x80FFFFFF
	F1 0x00000000	F1 0x00000000
	F3 0x80FFFFFE	F3 0x00000000
	FSR[7:0] 8'b00000000	FSR[7:0] 8'b00011000 (UDF, INX)
Example:2	MAC.d F6, F4, F0	;Multiply contents of F1:F0, F5:F4 and add the results with content in F7:F6
	Before execution	After execution
	F1:F0 0x406900000000 00000	F1:F0 0x406900000000 00000
	F5:F4 0x407F40000000 00000	F5:F4 0x407F40000000 00000
	F9:F8 0x402400000000 00000	F9:F8 0x40A388000000 00000
	FSR[7:0] 8'b00000000	FSR[7:0] 8'b00000000

<b>MAX</b>	Select the Signed Maximum of Fb and Fs {IEEE 754-2019 maximum(x,y)}						
Syntax:	{label:}	MAX.s Fb, MAX.d	Fs,	Fd			
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]					
Operation:		IF (Fs) >= (Fb) THEN (Fs) → (Fd); ELSE (Fb) → (Fd)					
Status Affected:		SUBO, INVAL (see Note 1)					
Exceptions Possible:		SUBO, INVAL					
Encoding:	1000 001P zzdd ddds ssss www wUUU 0010						
Description:	<p>Find the signed maximum of two floating-point values held in Fb and Fs. This instruction is compliant with the IEEE 754-2019 maximum() operation.</p> <p>Compare the contents of the source registers Fb and Fs, then place the maximum of the two values in the destination register Fd. This is a signed comparison operation (see note 3). If Fb = Fs (and of the same sign), Fd is loaded with the contents of Fs.</p> <p>When one (or both) of the input operands is a NaN, the instructions will return a qNaN.</p> <p>The 'P' bit selects single (32-bit) or double precision (64-bit) operation.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'd' bits select the destination register.</p>						
<b>Notes:</b>	<ol style="list-style-type: none"> <li>If not already set, corresponding sticky exception status will also be set.</li> <li>Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> <li>Operand value of -0 compares to less than +0.</li> </ol>						
I-Words:	1						
SP Execute Cycles:	1						
DP Execute Cycles:	1						
Repetition Rate:	1						

<b>Example:1</b>	MAX.s F0, F1, F3	;Returns maximum of F0, F1 to F3 (Follows IEEE 754-2019 maximum())
------------------	------------------	--

Before execution	After execution
F0 0x7F800000	F0 0x7F800000
F1 0xFF800000	F1 0xFF800000
F3 0x00000000	F3 0x7F800000
FSR[7:0] 8'b00000000	FSR[7:0] 8'b00000000

<b>Example:2</b>	MAX.d F0, F4, F8	;Returns maximum of F1:F0, F5:F4 to F9:F8 (Follows IEEE 754-2019 maximum())
------------------	------------------	---

Before execution	After execution
F1:F0 0x000000000000 0000A	F1:F0 0x000000000000 0000A
F5:F4 0xC00000000000 0000C	F5:F4 0xC00000000000 0000C
F9:F8 0x000000000000 000	F9:F8 0x000000000000 0000A
FSR[7:0] 8'b00000000	FSR[7:0] 8'b01000000 (SUBO)

<b>MAXNM</b>	<b>Select the Signed Maximum of Fb and Fs {IEEE 754-2019 maximumNumber(x,y)}</b>						
Syntax:	{label:}	MAXNM.s Fb, MAXNM.d	Fs,	Fd			
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]					
Operation:		IF (Fs) >= (Fb) THEN (Fs) → (Fd); ELSE (Fb) → (Fd)					
Status Affected:		SUBO, INVAL (see Note 1)					
Exceptions Possible:		SUBO, INVAL					
Encoding:	1000 001P zzdd ddds ssss www wUUU 0110						
Description:	<p>Find the signed maximum of two floating-point values held in Fb and Fs. This instruction is compliant with the IEEE 754-2019 maximumNumber() operation.</p> <p>Compare the contents of the source registers Fb and Fs, then place the maximum of the two values in the destination register Fd. If Fb = Fs, Fd is loaded with the contents of Fs. This is a signed comparison operation (see note 3). If Fb = Fs (and of the same sign), Fd is loaded with the contents of Fs.</p> <p>When one of the input operands is a NaN and the other input is a floating-point number (that is not a NaN), the instructions will return the floating-point number.</p> <p>If both input operands are a NaN, the instructions will return a qNaN.</p> <p>The 'P' bit selects single (32-bit) or double precision (64-bit) operation.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'd' bits select the destination register.</p>						
<b>Notes:</b>	<ol style="list-style-type: none"> <li>If not already set, corresponding sticky exception status will also be set.</li> <li>Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> <li>Operand value of -0 compares to less than +0.</li> </ol>						
I-Words:	1						
SP Execute Cycles:	1						
DP Execute Cycles:	1						
Repetition Rate:	1						

<b>Example:1</b>	<b>MAXNM.s F0, F1, F3</b>	<b>;Returns maximum of F0, F1 to F3</b>
Before execution		After execution
F0 0x00000000		F0 0x00000000
F1 0x80000000		F1 0x80000000
F3 0x00000000		F3 0x00000000
FSR[7:0] 8'b00000000		FSR[7:0] 8'b00000000

<b>Example:2</b>	<b>MAXNM.d F0, F4, F8</b>	<b>;Returns maximum of F1:F0, F5:F4 to F9:F8</b>
Before execution		After execution
F1:F0 0x800000000000 0000A		F1:F0 0x800000000000 0000A
F5:F4 0x3FF924920048 245C		F5:F4 0x3FF924920048 245C
F9:F8 0x000000000000 000		F9:F8 0x3FF924920048 245C
FSR[7:0] 8'b00000000		FSR[7:0] 8'b01000000 (SUBO)

<b>MIN</b>		Select the Signed Minimum of Fb and Fs {IEEE 754-2019 minimum(x,y)}						
Syntax:	{label:}	MIN.s	Fb,	Fs,	Fd			
		MIN.d						
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]						
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]						
Operation:		IF (Fs) <= (Fb) THEN (Fs) → (Fd); ELSE (Fb) → (Fd)						
Status Affected:		SUBO, INVAL (see note 1)						
Exceptions Possible:		SUBO, INVAL						
Encoding:		1000	001P	zzdd	ddds	ssss	www	wUUU
Description:		Find the signed minimum of two floating-point values held in Fb and Fs. This instruction is compliant with the IEEE 754-2019 minimum() operation. Compare the contents of the source registers Fb and Fs, then place the minimum of the two values in the destination register Fd. This is a signed comparison operation (see note 3). If Fb = Fs (and of the same sign), Fd is loaded with the contents of Fs. When one (or both) of the input operands is a NaN, the instructions will return a qNaN. The 'P' bit selects single (32-bit) or double precision (64-bit) operation. The 'z' bits select the target coprocessor. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register.						
		<b>Notes:</b>						
		1. If not already set, corresponding sticky exception status will also be set. 2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro. 3. Operand value of -0 compares to less than +0.						
I-Words:		1						
SP Execute Cycles:		1						
DP Execute Cycles:		1						
Repetition Rate:		1						

**Example:1** MIN.s F0, F1, F3 ;Returns minimum of F0, F1 to F3 (Follows IEEE 754-2019 minimum())

Before execution	After execution
F0 0x7F800000	F0 0x7F800000
F1 0xFF800000	F1 0xFF800000
F3 0x00000000	F3 0xFF800000
FSR[7:0] 8'b00000000	FSR[7:0] 8'b00000000

**Example:2** MIN.d F0, F4, F8 ;Returns minimum of F1:F0, F5:F4 to F9:F8 (Follows IEEE 754-2019 minimum())

Before execution	After execution
F1:F0 0x000000000000 0000A	F1:F0 0x000000000000 0000A
F5:F4 0xC00000000000 0000C	F5:F4 0xC00000000000 0000C
F9:F8 0x000000000000 000	F9:F8 0xC00000000000 0000C
FSR[7:0] 8'b00000000	FSR[7:0] 8'b01000000 (SUBO)

<b>MINNM</b>		Select the Signed Minimum of Fb and Fs [minimumNumber()]							
Syntax:	{label:}	MINNM.s	Fb,	Fs,	Fd				
			MINNM.d						
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]							
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]							
Operation:		IF (Fs) <= (Fb) THEN (Fs) → (Fd); ELSE THEN (Fb) → (Fd)							
Status Affected:		SUBO, INVAL (see Note 1)							
Exceptions Possible:		SUBO, INVAL							
Encoding:	1000 001P zzdd ddds ssss www wUUU 1110								
Description:	<p>Find the signed minimum of two floating-point values held in Fb and Fs. This instruction is compliant with the IEEE 754-2019 minimumNumber() operation.</p> <p>Compare the contents of the source registers Fb and Fs, then place the minimum of the two values in the destination register Fd. If Fb = Fs, Fd is loaded with the contents of Fs. This is a signed comparison operation (see note 3). If Fb = Fs (and of the same sign), Fd is loaded with the contents of Fs.</p> <p>When one of the input operands is a NaN and the other input is a floating-point number (that is not a NaN), the instructions will return the floating-point number.</p> <p>If both input operands are a NaN, the instructions will return a qNaN.</p> <p>The 'P' bit selects single (32-bit) or double precision (64-bit) operation.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'd' bits select the destination register.</p>								
<b>Notes:</b>	<ol style="list-style-type: none"> <li>If not already set, the corresponding sticky exception status will also be set.</li> <li>Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> <li>Operand value of -0 compares to less than +0.</li> </ol>								
I-Words:	1								
SP Execute Cycles:	1								
DP Execute Cycles:	1								
Repetition Rate:	1								

Example:1	MINNM.s F0, F1, F3	;Returns minimum of F0, F1 to F3										
		<table style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th style="text-align: center;">Before execution</th> <th style="text-align: center;">After execution</th> </tr> </thead> <tbody> <tr> <td style="text-align: center;">F0 0x00000000</td> <td style="text-align: center;">F0 0x00000000</td> </tr> <tr> <td style="text-align: center;">F1 0x80000000</td> <td style="text-align: center;">F1 0x80000000</td> </tr> <tr> <td style="text-align: center;">F3 0x00000000</td> <td style="text-align: center;">F3 0x80000000</td> </tr> <tr> <td style="text-align: center;">FSR[7:0] 8'b00000000</td> <td style="text-align: center;">FSR[7:0] 8'b00000000</td> </tr> </tbody> </table>	Before execution	After execution	F0 0x00000000	F0 0x00000000	F1 0x80000000	F1 0x80000000	F3 0x00000000	F3 0x80000000	FSR[7:0] 8'b00000000	FSR[7:0] 8'b00000000
Before execution	After execution											
F0 0x00000000	F0 0x00000000											
F1 0x80000000	F1 0x80000000											
F3 0x00000000	F3 0x80000000											
FSR[7:0] 8'b00000000	FSR[7:0] 8'b00000000											

Example:2	MINNM.d F0, F4, F8	;Returns minimum of F1:F0, F5:F4 to F9:F8																				
		<table style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th colspan="2" style="text-align: center;">Before execution</th> <th colspan="2" style="text-align: center;">After execution</th> </tr> </thead> <tbody> <tr> <td style="text-align: center;">F1:F0</td> <td style="text-align: center;">0x800000000000 0000A</td> <td style="text-align: center;">F1:F0</td> <td style="text-align: center;">0x800000000000 0000A</td> </tr> <tr> <td style="text-align: center;">F5:F4</td> <td style="text-align: center;">0x3FF924920048 245C</td> <td style="text-align: center;">F5:F4</td> <td style="text-align: center;">0x3FF924920048 245C</td> </tr> <tr> <td style="text-align: center;">F9:F8</td> <td style="text-align: center;">0x000000000000 000</td> <td style="text-align: center;">F9:F8</td> <td style="text-align: center;">0x800000000000 0000A</td> </tr> <tr> <td style="text-align: center;">FSR[7:0]</td> <td style="text-align: center;">8'b00000000</td> <td style="text-align: center;">FSR[7:0]</td> <td style="text-align: center;">8'b01000000 (SUBO)</td> </tr> </tbody> </table>	Before execution		After execution		F1:F0	0x800000000000 0000A	F1:F0	0x800000000000 0000A	F5:F4	0x3FF924920048 245C	F5:F4	0x3FF924920048 245C	F9:F8	0x000000000000 000	F9:F8	0x800000000000 0000A	FSR[7:0]	8'b00000000	FSR[7:0]	8'b01000000 (SUBO)
Before execution		After execution																				
F1:F0	0x800000000000 0000A	F1:F0	0x800000000000 0000A																			
F5:F4	0x3FF924920048 245C	F5:F4	0x3FF924920048 245C																			
F9:F8	0x000000000000 000	F9:F8	0x800000000000 0000A																			
FSR[7:0]	8'b00000000	FSR[7:0]	8'b01000000 (SUBO)																			



MOV							
Move Fs to Fd							
Syntax:	{label:}	MOV.s	Fs,	Fd			
		MOV.d					
Operands (.s):		Fs ∈ [F0 ... F31];	Fd ∈ [F0 ... F31]				
Operands (.d):		Fs ∈ [F0, F2 ... F30];	Fd ∈ [F0, F2 ... F30]				
Operation:		Fs → Fd					
Status Affected:		None					
Exceptions Possible:		None					
Encoding:	1000	011P	zzdd	ddds	ssss	UUUU	UUUU
Description:	Move (copy) the contents of the source register Fs to the destination register Fd. The 'P' bit selects single (.s) or double precision (.d) operation. The 'z' bits select the target coprocessor. The 's' bits select the source register. The 'd' bits select the destination register. <b>Note:</b> Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:	1						
SP Execute Cycles:	1						
DP Execute Cycles:	1						
Repetition Rate:	1						

Example:1      MOV.s F0, F1			;Move single precision content of F0 to F1		
Before execution			After execution		
F0            0x40C90FDB			F0            0x40C90FDB		
F1            0x00000000			F1            0x40C90FDB		
Example:2      MOV.d F0, F2			;Move double precision content of F1:F0 to F3:F2		
Before execution			After execution		
F1:F0        0x3FE62E42FEFA 39EF			F1:F0        0x3FE62E42FEFA 39EF		
F3:F2        0x000000000000 00000			F3:F2        0x3FE62E42FEFA 39EF		

<b>MOVC</b>		<b>Move Constant Value to Fd</b>						
Syntax:	{label:}	MOVC.s index, MOVC.d	Fd					
Operands (.s):		Fd ∈ [F0 ... F31]; index ∈ [0 ... 31]						
Operands (.d):		Fd ∈ [F0, F2 ... F30]; index ∈ [0 ... 31]						
Operation:		Constant table (index) → Fd						
Status Affected:		None						
Exceptions Possible:		None						
Encoding:	1000      011P      zzdd      dddU      UUUU      UUUk      kkkk      0110							
Description:	Move the constant value specified by the 5-bit index into the Fd register, as defined in the tables below. The 'P' bit selects single (.s) or double precision (.d) operation. The 'z' bits select the target coprocessor. The 'd' bits select the destination register. The 'k' bits select the index. <b>Note:</b> Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.							
I-Words:	1							
SP Execute Cycles:	1							
DP Execute Cycles:	1							
Repetition Rate:	1							
Single Precision:		Index	Data		Description			
		31	32'H7FFF_FFFF		QNAN - RANGE ENDS			
		30	32'H7FC0_0000		QNAN - RANGE STARTS			
		29	32'H7FBF_FFFF		SNAN - RANGE ENDS			
		28	32'H7F80_0001		SNAN - RANGE STARTS			
		27	32'H7F80_0000		+INF			
		26	32'H7F7F_FFFF		+MAX MAG NORMAL			
		25	32'H0080_0000		+MIN MAG NORMAL			
		24	32'H007F_FFFF		+MAX MAG SUB-NORMAL			
		23	32'H0000_0001		+MIN MAG SUB-NORMAL			
		22	32'H0000_0000		+0.0			
		21	32'H8000_0000		-0.0			
		20	32'H8000_0001		-MIN MAG SUB-NORMAL			
		19	32'H807F_FFFF		-MAX MAG SUB-NORMAL			
		18	32'H8080_0000		-MIN MAG NORMAL			
		17	32'HFF7F_FFFF		-MAX MAG NORMAL			
		16	32'HFF80_0000		-INF			
		15	32'HFF80_0001		SNAN - RANGE STARTS			
		14	32'HFFBF_FFFF		SNAN - RANGE ENDS			
		13	32'HFFC0_0000		QNAN - RANGE STARTS			
		12	32'FFFF_FFFF		QNAN - RANGE ENDS			
		11	32'H3F35_04F3		1/√ 2 = 0.70710...			
		10	32'H3FB5_04F3		√2 = 1.41421...			
		9	32'H4054_9A78		LOG2(10) = 3.321			
		8	32'H3F31_7218		LN(2) = 0.6931471 ...			
		7	32'H402D_F854		E = 2.71828...			
		6	32'H3F49_0FDB		Π/4 = 0.785398...			
		5	32'H3FC9_0FDB		Π/2 = 1.5707...			

**Floating-Point Instruction Description (continued)**

<b>MOVC</b>	<b>Move Constant Value to Fd</b>	
4	32'H4049_0FDB	$\Pi = 3.14159...$
3	32'H40C9_0FDB	$2\Pi = 6.283185...$
2	32'H4000_0000	2.0
1	32'H3F80_0000	1.0
0	32'H4120_0000	10.0
Double Precision:		
Index	Data	Description
31	64'H7FFF_FFFF_FFFF_FFFF	QNAN - RANGE ENDS
30	64'H7FF8_0000_0000_0000	QNAN RANGE STARTS
29	64'H7FF7_FFFF_FFFF_FFFF	SNAN RANGE ENDS
28	64'H7FF0_0000_0000_0001	SNAN RANGE STARTS
27	64'H7FF0_0000_0000_0000	+INF
26	64'H7FEF_FFFF_FFFF_FFFF	+MAX MAG NORMAL
25	64'H0010_0000_0000_0000	+MIN MAG NORMAL
24	64'H000F_FFFF_FFFF_FFFF	+MAX MAG SUB-NORMAL
23	64'H0000_0000_0000_0001	+MIN MAG SUB-NORMAL
22	64'H0000_0000_0000_0000	+0.0
21	64'H8000_0000_0000_0000	-0.0
20	64'H8000_0000_0000_0001	-MIN MAG SUB-NORMAL
19	64'H800F_FFFF_FFFF_FFFF	-MAX MAG SUB-NORMAL
18	64'H8010_0000_0000_0000	-MIN MAG NORMAL
17	64'HFFEF_FFFF_FFFF_FFFF	-MAX MAG NORMAL
16	64'HFFF0_0000_0000_0000	-INF
15	64'HFFF0_0000_0000_0001	SNAN RANGE STARTS
14	64'HFFF7_FFFF_FFFF_FFFF	SNAN RANGE ENDS
13	64'HFFF8_0000_0000_0000	QNAN RANGE STARTS
12	64'HFFFF_FFFF_FFFF_FFFF	QNAN RANGE ENDS
11	64'H3FE6_A09E_667F_3BCD	$1/\sqrt{2}=0.7071...$
10	64'H3FF6_A09E_667F_3BCD	$\sqrt{2}=1.41421..$
9	64'H400A_934F_0979_A371	$\text{LOG2}(10)= 3.321$
8	64'H3FE6_2E42_FEFA_39EF	$\text{LN}(2)=0.69314$
7	64'H4005_BF0A_8B14_5769	$E = 2.71828...$
6	64'H3FE9_21FB_5444_2D18	$\pi/4= 0.7853...$
5	64'H3FF9_21FB_5444_2D18	$\pi/2 = 1.5707..$
4	64'H4009_21FB_5444_2D18	$\pi = 3.14159...$
3	64'H4019_21FB_5444_2D18	$2\pi = 6.28...$
2	64'H4000_0000_0000_0000	2.0
1	64'H3FF0_0000_0000_0000	1.0
0	64'H4024_0000_0000_0000	10.0

<b>Example:1</b>	<b>MOVC.s #3, F0</b>	<b>;Move single precision constant <math>2\pi</math> to register F0</b>	
	Before execution F0	After execution F0	0x40C90FDB

<b>Example:2</b>	<b>MOVC.d #8, F0</b>	<b>;Move double precision constant <math>\text{Ln}(2)</math> to register F1:F0</b>	
	Before execution F0	After execution F0	0x40C90FDB

**Floating-Point Instruction Description (continued)**

Example:2	MOVC.d #8, F0	;Move double precision constant Ln(2) to register F1:F0	
		F1:F0	0x3FE62E42FEFA39EF 00000

MOV							Move Coprocessor Register to Wd							
Syntax:	{label:}	MOV.I	Fs, FSR, FSRH, FCR, FEAR,	Wd [Wd] [Wd++] [Wd--] [-Wd] [++Wd] [Wd+Wb]										
Operands:			$Fs \in [F0 \dots F31]; Wd \in [W0 \dots W15]; Wb \in [W0 \dots W14]$ (see Note 2)											
Operation:			$(Fs \text{ or } FSR \text{ or } FSRH \text{ or } FCR \text{ or } FEAR) \rightarrow (EAd)$											
Status Affected:			None											
Encoding:		S000	011s	ssss	dddd	UUUq	qqww	wwRU	zz00					
Description:		Move the contents of the source coprocessor data, status or control register into the destination register. Syntax shown is for the floating-point coprocessor registers. The 'S' bit selects instruction size. The 'z' bits select the target coprocessor. The 's' bits select the coprocessor source register. The 'd' bits select the destination register. The 'q' bits select the destination addressing mode. The 'w' bits define the offset Wb. The 'R' bit selects between the F-reg and FPU special registers.												
<b>Notes:</b>														
1. This instruction operates in Long Word mode only. 2. Stack Pointer (W15) limit checks are only performed for Wd = W15. Consequently, setting Wb = W15 is not permitted.														
I-Words:		1 or 0.5												
Cycles:		1												

<b>MOV</b>	<b>Move Wns with Signed Literal Offset to Coprocessor Register</b>													
Syntax:	{label:}	MOV.I	[Wns+Slit14], Fd											
Operands:	Wns $\in$ [W0 ... W15] Fd $\in$ [F0 ... F31]													
Operation:	[Wns+Slit14] $\rightarrow$ Fd													
Status Affected:	None													
Encoding:	1000	010d	dddd	ssss	kkkk	kkkk	kkkk	zz01						
Description:	<p>Moves contents of the source Effective Address into the destination register. Syntax shown is for the floating-point coprocessor F-reg. The contents of Wns are not modified by this operation.</p> <p>The instruction encoding includes space for a 12-bit literal which is scaled accordingly for long word data moves in order to generate the corresponding byte offset value.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'd' bits select the floating-point destination register.</p>													
<b>Notes:</b>	<ol style="list-style-type: none"> <li>1. This instruction only supported when a coprocessor is instantiated.</li> <li>2. This instruction operates in Long Word mode only.</li> </ol>													
I-Words:	1													
Cycles:	1													

MOV							Move Ws to Coprocessor Register						
Syntax:	{label:}	MOV.I	Ws, [Ws], [Ws++] [Ws-] [-Ws], [++Ws], [Ws+Wb],	Fd FSR FCR FEAR									
Operands:			Fd ∈ [F0 ... F31]; Ws ∈ [W0 ... W15]; Wb ∈ [W0 ... W14] (see Note 2)										
Operation:			(EAs) → Fd, FSR, FCR or FEAR										
Status Affected:			None										
Encoding:	S000	010d	dddd	ssss	pppU	UUww	wwRU	zz00					
Description:	Move the contents of the source register into the coprocessor destination data, status or control register. Syntax shown is for the floating-point coprocessor registers. The 'S' bit selects instruction size. The 'z' bits select the target coprocessor. The 's' bits select the source register. The 'd' bits select the coprocessor destination register. The 'p' bits select the source addressing mode. The 'w' bits define the offset Wb. The 'R' bit selects between the F-reg and FPU special registers.												
Notes:	<ol style="list-style-type: none"> <li>1. This instruction operates in Long Word mode only.</li> <li>2. Stack Pointer (W15) limit checks are only performed for Wd = W15. Consequently, setting Wb = W15 is not permitted.</li> </ol>												
I-Words:	1 or 0.5												
Cycles:	1												

<b>MOV</b>								
<b>Move Long Literal to Coprocessor Register</b>								
Syntax:	{label:}	MOV.I	lit32,	Fd				
				FSR				
				FCR				
Operands:	$\text{lit32} \in [0 \dots 4\text{GB}]$ ; $\text{Fd} \in [\text{F0} \dots \text{F31}]$							
Operation:	$\text{lit32} \rightarrow (\text{Fd or FSR or FCR})$							
Status Affected:	None							
Encoding:								
1st word	1000	001z	kkkk	kkkk	kkkk	kkkk	kkkk	z001
2nd word	1000	001d	dddd	Ukkk	kkkU	UUkk	kkkk	R101
Description:	<p>The literal 'K' is loaded into the coprocessor destination data register. Syntax shown is for the floating-point coprocessor register.</p> <p>The 'd' bits select the coprocessor destination register.</p> <p>The 'z' bits select the target coprocessor.</p> <p><math>\text{zz}[1:0] = \text{Word1}[24], \text{Word1}[3]</math></p> <p>The 'R' bit selects between the F-reg and FPU special registers.</p> <p>The 'K' bits specify the value of the 32-bit literal:</p> <p><math>\text{lit32}[31:0] = \text{Word1}[18:13], \text{Word2}[9:4], \text{Word1}[23:4]</math></p>							
I-Words:	2							
Cycles:	2							

<b>MUL</b>							
<b>Floating-Point Signed Multiply</b>							
Syntax:	{label:}	MUL.s	Fb,	Fs,	Fd		
			MUL.d				
Operands (.s):		Fs ∈ [F0 ... F31]; Fb ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fb ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]					
Operation:		(Fb) * (Fs) → Fd;					
Status Affected:		SUBO, INX, UDF, OVF, INVAL (see Note 1)					
Exceptions Possible:		SUBO, INX, UDF, OVF, INVAL					
Encoding:	1000	010P	zzdd	ddds	ssss	www	wuuu
Description:	Multiply the signed contents of the source register Fb with the signed contents of the source register Fs, then write the product to Fd. (0 * ∞) or (∞ * 0) will result in distinguished qNaN and signal an INVAL exception.						
	The 'P' bit selects single (.s) or double precision (.d) operation. The 'z' bits select the target coprocessor. The 's' bits select the source register. The 'w' bits select the base register. The 'd' bits select the destination register.						
Notes:							
	1. If not already set, the corresponding sticky exception status will also be set. 2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:	1						
SP Execute Cycles:	3						
DP Execute Cycles:	3						
Repetition Rate:	1						

<b>Example:1</b>	<b>MUL.s F0, F1, F3</b>	<b>;Multiply contents of F0 and F1 and store result in F3</b>
------------------	-------------------------	---

Before execution		After execution	
F0	0x41200000	F0	0x41200000
F1	0x41C80000	F1	0x41C80000
F3	0x00000000	F3	0x437A0000
FSR[7:0]	8'b00000000	FSR[7:0]	8'b00000000

<b>Example:2</b>	<b>MUL.d F0, F2, F6</b>	<b>;Multiply contents of F1:F0 and F3:F2 and store result in F7:F6</b>
------------------	-------------------------	--

Before execution		After execution	
F1:F0	0xFFE000000000 0030	F1:F0	0xFFE000000000 0030
F3:F2	0xFFE000000000 0030	F3:F2	0xFFE000000000 0030
F7:F6	0x000000000000 000	F7:F6	0x7FFFFFFFFF FFFF
FSR[7:0]	8'b00000000	FSR[7:0]	8'b00010100 (OVF, INX)

NEG	Negate Fs						
Syntax:	{label:}	NEG.s	Fs,	Fd			
			NEG.d				
Operands (.s):		Fs ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30]; Fd ∈ [F0, F2 ... F30]					
Operation:		-Fs → Fd					
Status Affected:		SUBO <sup>1</sup>					
Exceptions Possible:		SUBO					
Encoding:	1000	100P	zzdd	ddds	ssss	UUUU	UUUU
Description:	Negate the sign of the contents of the source register Fs and place the result in the destination register Fd.						
	The 'P' bit selects single (.s) or double precision (.d) operation.						
	The 'z' bits select the target coprocessor.						
	The 's' bits select the source register.						
	The 'd' bits select the destination register.						
	<b>Notes:</b>						
	1. If not already set, the corresponding sticky exception status will also be set.						
	2. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:	1						
SP Execute Cycles:	1						
DP Execute Cycles:	1						
Repetition Rate:	1						

**Note:** This instruction only affects the sign bit and will not quit an sNaN input operand.

Example:1	NEG.s F0, F3	;Find negate of a number in F0 and store result in F3
	Before execution	After execution
	F0 0x40800001	F0 0x40800001
	F3 0x00000000	F3 0xC0800001
	FSR[7:0] 8'b00000000	FSR[7:0] 8'b00000000
Example:2	NEG.d F0, F2	;Find negate of a number in F1:F0 and store result in F3:F2
	Before execution	After execution
	F1:F0 0x800000000000 00001	F1:F0 0x800000000000 00001
	F3:F2 0x000000000000 00000	F3:F2 0x000000000000 00001
	FSR[7:0] 8'b00000000	FSR[7:0] 8'b01000000 (SUBO)

POP					Pop FPU Coprocessor Register			
Syntax:	{label:}	POP.I	Fd					
			FSR					
			FCR					
			FEAR					
Operands:		Fd ∈ [F0 ... F31]						
Operation:		(W15)-4 → W15						
		(TOS) → (Fd or FSR or FCR or FEAR)						
Status Affected:		None						
Encoding:	0010	111d	dddd	zzRU				
Description:	Move the top of system stack to the coprocessor destination register. Use of the system Stack Pointer ([--W15]) is implied. Syntax shown is for the floating-point coprocessor registers. The 'd' bits select the floating point destination register.							
	The 'z' bits select the target coprocessor.							
	The 'R' bit selects between the F-reg and FPU special registers.							
<b>Notes:</b>								
	1. This instruction operates in Long Word mode only.							
	2. FSRH not supported as it would clear exception status in FSR[15:0].							
Words:	1							
Cycles:	1							

<b>PUSH</b>					<b>Push FPU Coprocessor Register</b>
Syntax:	{label:}	PUSH.I	Fs		
			FSR		
			FSRH		
			FCR		
			FEAR		
Operands:	$Fs \in [F0 \dots F31]$				
Operation:	$(Fs \text{ or } FSR \text{ or } FSRH \text{ or } FCR \text{ or } FEAR) \rightarrow (\text{TOS})$ $(W15)+4 \rightarrow W15$				
Status Affected:	None				
Encoding:	0010	011s	ssss	zzRU	
Description:	<p>Move the contents of the source floating point register to the top of system stack. Use of the system Stack Pointer ([W15++]) is implied. Syntax shown is for the floating-point coprocessor registers. When FSRH is selected, long word value {FSR[31:16], 16'b0} will be pushed onto the stack. FSR[31:16] comprises of FCPS/FCPQ and FTST instruction status only.</p> <p>The 's' bits select the floating point source register.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 'R' bit selects between the F-reg and FPU special registers.</p>				
<b>Note:</b>	<ol style="list-style-type: none"> <li>This instruction operates in Long Word mode only.</li> </ol>				
Words:	0.5				
Cycles:	1				

SIN	Evaluate the Sine of Fs							
Syntax:	{label:}	SIN.s	Fs,	Fd				
Operands (.s):		Fs ∈ [F0 ... F31]; Fd ∈ [F0 ... F31]						
Operation:		Sin(Fs) → Fd						
Status Affected:		SUBO, INX, UDF, INVAL (see Note 1)						
Exceptions Possible:		SUBO, INX, UDF, INVAL						
Encoding:	1000	110P	zzdd	ddds	ssss	UUUU	UUUU	0010
Description:		Calculate the sine of contents of the source register Fs (where Fs = [2πk + θ rads]), and place the result in the destination register Fd. The 'z' bits select the target coprocessor.						
		The 's' bits select the source register						
		The 'd' bits select the destination register.						
		<b>Notes:</b>						
		1. If not already set, the corresponding sticky exception status will also be set.						
		2. This operation is only supported using single precision floating-point.						
		3. Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:	1							
SP Execute Cycles:	4							
Repetition Rate:	1							

Example:1	SIN.s F0, F3	;Find Sine of radians in F0 and store result in F3		
		Before execution	After execution	
		F0            0x3F492492	F0            0x3F492492	
		F3            0x00000000	F3            0x3F351398	
		FSR[7:0]    8'b00000000	FSR[7:0]    8'b00000000	
Example:2	SIN.s F0, F3	;Find Sine of radians in F0 and store result in F3		
		Before execution	After execution	
		F0            0x7F800000	F0            0x7F800000	
		F3            0x00000000	F3            0x7FC00001	(distinguished qNaN)
		FSR[7:0]    8'b00000000	FSR[7:0]    8'b00000001	(INVAL)

SQRT	Square Root Fs, result in Fd							
Syntax:	{label:}	SQRT.s	Fs,	Fd				
		SQRT.d						
Operands (.s):		Fs ∈ [F0 ... F31];	Fd ∈ [F0 ... F31]					
Operands (.d):		Fs ∈ [F0, F2 ... F30];	Fd ∈ [F0, F2 ... F30]					
Operation:		$\sqrt{Fs} \rightarrow Fd$						
Status Affected:		SUBO, INX, INVAL (see Note 1)						
Exceptions Possible:		SUBO, INX, INVAL						
Encoding:	1000	100P	zzdd	ddds	ssss	UUUU	UUUU	1010
Description:	Take the square root of contents of the source register Fs, and place the result in the destination register Fd. If Fs < 0, destination will be written with distinguished qNaN and signal an FPU INVAL exception. The 'P' bit selects single (.s) or double precision (.d) operation. The 'z' bits select the target coprocessor. The 's' bits select the source register. The 'd' bits select the destination register.							
Notes:	<ol style="list-style-type: none"> <li>If not already set, the corresponding sticky exception status will also be set.</li> <li>Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> </ol>							
I-Words:	1							
SP Execute Cycles:	10							
DP Execute Cycles:	13							
Repetition Rate:	TBD							

Example:1	SQRT.s F0, F3	;Find Square root of number in F0 and store result in F3			
		Before execution		After execution	
	F0	0x00000030		F0	0x00000030
	F3	0x00000000		F3	0x1B9CC470
	FSR[7:0]	8'b00000000		FSR[7:0]	8'b00000000

Example:2	SQRT.d F0, F2	;Find Square root of number in F1:F0 and store result in F3:F2			
		Before execution		After execution	
	F1:F0	0x405000000000 00000		F1:F0	0x405000000000 00000
	F3:F2	0x000000000000 00000		F3:F2	0x402000000000 00000
	FSR[7:0]	8'b00000000		FSR[7:0]	8'b00010000 (INX)

<b>SUB</b>		Subtract Fs from Fb						
Syntax:	{label:}	SUB.s	Fb,	Fs,	Fd			
			SUB.d					
Operands (.s):		Fs ∈ [F0 ... F31];	Fb ∈ [F0 ... F31];	Fd ∈ [F0 ... F31]				
Operands (.d):		Fs ∈ [F0, F2 ... F30];	Fb ∈ [F0, F2 ... F30];	Fd ∈ [F0, F2 ... F30]				
Operation:		Fb - Fs → Fd						
Status Affected:		SUBO, INX, UDF, OVF, INVAL (see Note 1)						
Exceptions Possible:		SUBO, INX, UDF, OVF, INVAL						
Encoding:	1000	000P	zzdd	ddds	ssss	www	wuuu	0110
Description:	<p>Subtract the contents of the source register Fs from the contents of the base register Fb, then place the result in the destination register Fd.</p> <p>(∞ - ∞) will result in distinguished qNaN and signal an FPU INVAL exception.</p> <p>The 'P' bit selects single (.s) or double precision (.d) operation.</p> <p>The 'z' bits select the target coprocessor.</p> <p>The 's' bits select the source register.</p> <p>The 'w' bits select the base register.</p> <p>The 'd' bits select the destination register.</p>							
<b>Notes:</b>	<ol style="list-style-type: none"> <li>If not already set, the corresponding sticky exception status will also be set.</li> <li>Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.</li> </ol>							
I-Words:	1							
SP Execute Cycles:	2							
DP Execute Cycles:	2							
Repetition Rate:	1							

Example:1	SUB.s F0, F1, F3	;Subtracts F1 from F0 and stores the result in F3
		Before execution
	F0	0x40800001
	F1	0xC0F800F1
	F3	0x00000000
	FSR[7:0]	8'b00000000
		After execution
	F0	0x40800001
	F1	0xC0F800F1
	F3	0xC07001E0
	FSR[7:0]	8'b00000000

Example:2	SUB.d F0, F2, F6	;Subtracts F3:F2 from F1:F0 and stores the result in F7:F6
		Before execution
	F1:F0	0x000FFFFFFFFF FFFF
	F3:F2	0x000FFFFFFFFF FFFF
	F7:F6	0x001FFFFFFFFF FFFE
	FSR[7:0]	8'b00000000
		After execution
	F1:F0	0x000FFFFFFFFF FFFF
	F3:F2	0x000FFFFFFFFF FFFF
	F7:F6	0x001FFFFFFFFF FFFE
	FSR[7:0]	8'b01000000 (SUBO)

TST	Test Fs							
Syntax:	{label:}	TST.s	Fs					
			TST.d					
Operands (.s):		Fs ∈ [F0 ... F31]						
Operands (.d):		Fs ∈ [F0, F2 ... F30]						
Operation:		If Fs is +/-Subnormal FSR.SUB = 1 else FSR.SUB = 0 If Fs is +/- Infinity FSR.INF = 1 else FSR.INF = 0 If Fs is Negative FSR.FN = 1 else FSR.FN = 0 If Fs is +/- Zero FSR.FZ = 1 else FSR.FZ = 0 If Fs is qNaN or sNaN FSR.FNAN = 1 else FSR.FNAN = 0						
Status Affected:		SUB, INF, FN, FZ, FNAN						
Exceptions Possible:		None						
Encoding:	1001	010P	zzUU	UUUs	ssss	UUUU	UUUU	0010
Description:		Inspect the contents of Fs and set the FSR status accordingly. The 'P' bit selects single (.s) or double precision (.d) operation. The 'z' bits select the target coprocessor. The 's' bits select the source register.						
		<b>Note:</b> Coprocessor select bit field zz = 2'b00 for floating-point coprocessor macro.						
I-Words:	1							
SP Execute Cycles:	1							
DP Execute Cycles:	1							
Repetition Rate:	1							

Example:1	TST.s F0	;Inspects contents of F0 register and updates FSR
		Before execution
	F0	0xFFC0000A (a -ve NaN number)
	FSR[31:24]	8'b00000000
		After execution
		F0 0xFFC0000A
		FSR[31:24] 8'b000000101 (FN, NAN)
Example:2	TST.d F0	;Inspects contents of F1:F0 register and updates FSR
		Before execution
	F1:F0	0xFFFF00000000 0000 (a -ve Inf number)
	FSR[31:24]	8'b00000000
		After execution
		F1:F0 0xFFFF00000000 0000
		FSR[31:24] 8'b00001100 (FN, INF)

## 5. Built-In Functions

### 5.1. Introduction

Built-in functions give the C programmer access to assembler operators or machine instructions that are currently only accessible using in-line assembly but are sufficiently useful that they are applicable in a broad range of applications. Built-in functions are coded in C source files syntactically like function calls, but they are compiled to assembly code that directly implements the function and does not involve function calls or library routines.

The built-in functions are preferred by programmers more than in-line assembly, for the following reasons :

1. Providing built-in functions for specific purposes simplifies coding.
2. Certain optimizations are disabled when in-line assembly is used as compared to built-in functions.
3. For machine instructions that use dedicated registers, coding with in-line assembly needs care to avoid register allocation errors . The built-in functions make this process simpler as user need not be concerned with the particular register requirements for each individual machine instruction.

The built-in functions are listed below followed by their individual detailed descriptions.

- [\\_builtin\\_addab](#)
- [\\_builtin\\_add](#)
- [\\_builtin\\_btg](#)
- [\\_builtin\\_clr](#)
- [\\_builtin\\_divf](#)
- [\\_builtin\\_divmodsd](#)
- [\\_builtin\\_divmodud](#)
- [\\_builtin\\_divsd](#)
- [\\_builtin\\_divud](#)
- [\\_builtin\\_dmaoffset](#)
- [\\_builtin\\_ed](#)
- [\\_builtin\\_edac](#)
- [\\_builtin\\_fbcl](#)
- [\\_builtin\\_lac](#)
- [\\_builtin\\_mac](#)
- [\\_builtin\\_modsd](#)
- [\\_builtin\\_modud](#)
- [\\_builtin\\_mpy](#)
- [\\_builtin\\_mpyn](#)
- [\\_builtin\\_msc](#)
- [\\_builtin\\_mulss](#)
- [\\_builtin\\_mulsu](#)
- [\\_builtin\\_mulus](#)
- [\\_builtin\\_muluu](#)

- `_builtin_nop`
- `_builtin_readsfr`
- `_builtin_return_address`
- `_builtin_sac`
- `_builtin_sacr`
- `_builtin_sftac`
- `_builtin_subab`

The compiler provides additional built-in functions for operations, such as writing to Flash program memory and changing the oscillator settings. Refer to the *MPLAB® XC32 C Compiler User's Guide for PIC32A MCU* (DS50003831) for a complete list of compiler built-in functions.

## 5.2. Built-In Functions List

This section describes the programmer interface to the compiler built-in functions. Since the functions are “built-in”, there are no header files associated with them. Similarly, there are no command-line switches associated with the built-in functions; they are always available. The built-in function names are chosen such that they belong to the compiler’s namespace (they all have the prefix: `_builtin_`), so they will not conflict with function or variable names in the programmer’s namespace.

### `_builtin_addab`

#### Description:

Adds Accumulators A and B with the result written back to the specified accumulator. For example:

```
register int result asm("A");
register int B asm("A");
result = _builtin_addab(result,B);
```

will generate:

```
add A
```

#### Prototype:

```
int _builtin_addab(int Accum_a, int Accum_b);
```

#### Argument:

`Accum_a` First accumulator to add.

`Accum_b` Second accumulator to add.

#### Return Value:

Returns the addition result to an accumulator.

#### Assembler Operator/Machine Instruction:

`add`

#### Error Messages:

An error message appears if the result is not an Accumulator register.

### `_builtin_add`

#### Description:

Adds `value` to the accumulator specified by `result` with a shift specified by literal shift. For example:

```
register int result asm("A");
int      value;
result = _builtin_add(result,value,0);
```

If `value` is held in `w0`, the following will be generated:

```
add w0, #0, A
```

**Built-In Functions List (continued)****\_builtin\_add****Prototype:**

```
int _builtin_add(int Accum, int value,
const int shift);
```

**Argument:**

*Accum* Accumulator to add.

*value* Integer number to add to accumulator value.

*shift* Amount to shift resultant accumulator value.

**Return Value:**

Returns the shifted addition result to an accumulator.

**Assembler Operator/Machine Instruction:**

add

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- Argument 0 is not an accumulator
- The shift value is not a literal within range

**\_builtin\_btg****Description:**

This function will generate a `btg` machine instruction. Some examples include:

```
int i; /* near by default */
int l __attribute__((far));
struct foo {
    int bit1:1;
} barbits;
int bar;
void some_bit toggles() {
    register int j asm("w9");
    int k;
    k = i;
    _builtin_btg(&i,1);
    _builtin_btg(&j,3);
    _builtin_btg(&k,4);
    _builtin_btg(&l,11);
    return j+k;
}
```

Note that taking the address of a variable in a register will produce a warning by the compiler and cause the register to be saved onto the stack (so that its address may be taken); this form is not recommended. This caution only applies to variables explicitly placed in registers by the programmer.

**Prototype:**

```
void _builtin_btg(unsigned int *, unsigned int Oxn);
```

**Argument:**

\* A pointer to the data item for which a bit should be toggled.

*Oxn* A literal value in the range of 0 to 15.

**Return Value:**

Returns a `btg` machine instruction.

**Assembler Operator/Machine Instruction:**

btg

**Error Messages:**

An error message appears if the parameter values are not within range.

**\_builtin\_clr****Description:**

Clears the specified accumulator. For example:

```
register int result asm("A");
result = _builtin_clr();
```

will generate:

```
clr A
```

**Prototype:**

```
int _builtin_clr(void);
```

**Argument:**

None.

**Return Value:**

Returns the cleared value result to an accumulator.

**Assembler Operator/Machine Instruction:**

```
clr
```

**Error Messages:**

An error message appears if the result is not an Accumulator register.

**\_builtin\_divf****Description:**

Computes the quotient:  $num / den$ . A math error exception occurs if  $den$  is zero. Function arguments are unsigned, as is the function result.

**Prototype:**

```
unsigned int _builtin_divf(unsigned int num,
                           unsigned int den);
```

**Argument:**

*num* Numerator.

*den* Denominator.

**Return Value:**

Returns the unsigned integer value of the quotient:  $num / den$ .

**Assembler Operator/Machine Instruction:**

```
div.f
```

**\_builtin\_divmodsd****Description:**

Issues the 16-bit architecture's native signed divide support. Notably, if the quotient does not fit into a 16-bit result, the results (including remainder) are unexpected. This form of the built-in function will capture both the quotient and remainder.

**Prototype:**

```
signed int _builtin_divmodsd(
    signed long dividend, signed int divisor,
    signed int *remainder);
```

**Argument:**

*dividend* Number to be divided.

*divisor* Number to divide by.

*remainder* Pointer to remainder.

**Return Value:**

Quotient and remainder.

**Built-In Functions List (continued)****\_builtin\_divmodsd****Assembler Operator/Machine Instruction:**

divmodsd

**Error Messages:**

None.

**\_builtin\_divmodud****Description:**

Issues the 16-bit architecture's native unsigned divide support. Notably, if the quotient does not fit into a 16-bit result, the results (including remainder) are unexpected. This form of the built-in function will capture both the quotient and remainder.

**Prototype:**

```
unsigned int _builtin_divmodud(
    unsigned long dividend, unsigned int divisor,
    unsigned int *remainder);
```

**Argument:***dividend* Number to be divided.*divisor* Number to divide by.*remainder* Pointer to remainder.**Return Value:**

Quotient and remainder.

**Assembler Operator/Machine Instruction:**

divmodud

**Error Messages:**

None.

**\_builtin\_divsd****Description:**

Computes the quotient: *num / den*. A math error exception occurs if *den* is zero. Function arguments are signed, as is the function result. The command-line option, `-Wconversions`, can be used to detect unexpected sign conversions.

**Prototype:**

```
int _builtin_divsd(const long num, const int den);
```

**Argument:***num* Numerator.*den* Denominator.**Return Value:**Returns the signed integer value of the quotient: *num / den*.**Assembler Operator/Machine Instruction:**

div.sd

**\_builtin\_divud****Description:**

Computes the quotient: *num / den*. A math error exception occurs if *den* is zero. Function arguments are unsigned, as is the function result. The command-line option, `-Wconversions`, can be used to detect unexpected sign conversions.

**Prototype:**

```
unsigned int _builtin_divud(const unsigned
    long num, const unsigned int den);
```

**Built-In Functions List (continued)****\_builtin\_divud****Argument:***num* Numerator.*den* Denominator.**Return Value:**Returns the unsigned integer value of the quotient: *num / den*.**Assembler Operator/Machine Instruction:**

div.ud

**\_builtin\_dmaoffset****Description:**

Obtains the offset of a symbol within DMA memory.

For example:

```
unsigned int result;
char buffer[256] __attribute__((space(dma)));
result = _builtin_dmaoffset(&buffer);
```

May generate:

```
mov #dmaoffset(buffer), w0
```

**Prototype:**

```
unsigned int _builtin_dmaoffset(const void *p);
```

**Argument:***\*p* Pointer to DMA address value.**Return Value:**

Returns the offset to a variable located in DMA memory.

**Assembler Operator/Machine Instruction:**

dmaoffset

**Error Messages:**

An error message appears if the parameter is not the address of a global symbol.

**\_builtin\_ed****Description:**Squares *sqr*, returning it as the result. Also prefetches data for future square operation by computing *\*\*xptr - \*\*yptr* and storing the result in *\*distance*.*xincr* and *yincr* may be the literal values: -6, -4, -2, 0, 2, 4, 6 or an integer value.

For example:

```
register int result asm("A");
int *xmemory, *ymemory;
int distance;
result = _builtin_ed(distance,
                     &xmemory, 2,
                     &ymemory, 2,
                     &distance);
```

May generate:

```
ed w4*w4, A, [w8]+=2, [W10]+=2, w4
```

**Built-In Functions List (continued)****\_builtin\_ed****Prototype:**

```
int _builtin_ed(int sqr, int **xptr, int xincr,
int **yptr, int yincr, int *distance);
```

**Argument:**

*sqr* Integer squared value.

*xptr* Integer Pointer to pointer to X prefetch.

*xincr* Integer increment value of X prefetch.

*yptr* Integer Pointer to pointer to Y prefetch.

*yincr* Integer increment value of Y prefetch.

*distance* Integer Pointer to distance.

The arguments, *xptr* and *yptr*, must point to the arrays located in the X data memory and Y data memory, respectively.

**Return Value:**

Returns the squared result to an accumulator.

**Assembler Operator/Machine Instruction:**

```
ed
```

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- *xptr* is null
- *yptr* is null
- *distance* is null

**\_builtin\_edac****Description:**

Squares *sqr* and sums with the nominated Accumulator register, returning it as the result. Also prefetches data for future square operation by computing *\*\*xptr - \*\*yptr* and storing the result in *\*distance*.

*xincr* and *yincr* may be the literal values: -6, -4, -2, 0, 2, 4, 6 or an integer value.

For example:

```
register int result asm("A");
int      *xmemory, *ymemory;
int      distance;
result = _builtin_ed(result, distance,
                     &xmemory, 2,
                     &ymemory, 2,
                     &distance);
```

May generate:

```
edac w4*w4, A, [w8]+=2, [W10]+=2, w4
```

**Prototype:**

```
int _builtin_edac(int Accum, int sqr,
int **xptr, int xincr, int **yptr, int yincr,
int *distance);
```

**Built-In Functions List (continued)****\_builtin\_edac****Argument:**

*Accum* Accumulator to sum.

*sqr* Integer squared value.

*xptr* Integer Pointer to pointer to X prefetch.

*xincr* Integer increment value of X prefetch.

*yptr* Integer Pointer to pointer to Y prefetch.

*yincr* Integer increment value of Y prefetch.

*distance* Integer Pointer to distance.

The arguments, *xptr* and *yptr*, must point to the arrays located in the X data memory and Y data memory, respectively.

**Return Value:**

Returns the squared result to specified accumulator.

**Assembler Operator/Machine Instruction:**

edac

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- *Accum* is not an Accumulator register
- *xptr* is null
- *yptr* is null
- *distance* is null

**\_builtin\_fbcl****Description:**

Finds the first bit change from left in value. This is useful for dynamic scaling of fixed-point data. For example:

```
int result, value;
result = _builtin_fbcl(value);
```

May generate:

```
fbcl w4, w5
```

**Prototype:**

```
int _builtin_fbcl(int value);
```

**Argument:**

*value* Integer number of first bit change.

**Return Value:**

Returns the shifted addition result to an accumulator.

**Assembler Operator/Machine Instruction:**

fbcl

**Error Messages:**

An error message appears if the result is not an Accumulator register.

**\_builtin\_lac****Description:**

Shifts value by *shift* (a literal between -8 and 7) and returns the value to be stored into the Accumulator register. For example:

```
register int result asm("A");
int      value;
result = _builtin_lac(value, 3);
```

May generate:

```
lac w4, #3, A
```

**Prototype:**

```
int _builtin_lac(int value, int shift);
```

**Argument:**

*value* Integer number to be shifted.

*shift* Literal amount to shift.

**Return Value:**

Returns the shifted addition result to an accumulator.

**Assembler Operator/Machine Instruction:**

```
lac
```

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- The shift value is not a literal within range

**\_builtin\_mac****Description:**

Computes  $a \times b$  and sums with the accumulator; it also prefetches data ready for a future MAC operation.

*xptr* may be null to signify no X prefetch to be performed; in which case, the values of *xincr* and *xval* are ignored but required.

*yptr* may be null to signify no Y prefetch to be performed; in which case, the values of *yincr* and *yval* are ignored but required.

*xval* and *yval* nominate the address of a C variable where the prefetched value will be stored.

*xincr* and *yincr* may be the literal values: -6, -4, -2, 0, 2, 4, 6 or an integer value.

If *AWB* is non-null, the other accumulator will be written back into the referenced variable.

For example:

```
register int result asm("A");
register int B asm("B");
int      *xmemory;
int      *ymemory;
int      xVal, yVal;
result = _builtin_mac(result, xVal, yVal,
                      &xmemory, &xVal, 2,
                      &ymemory, &yVal, 2, 0, B);
```

May generate:

```
mac w4*w5, A, [w8]+=2, w4, [w10]+=2, w5
```

**Prototype:**

```
int _builtin_mac(int Accum, int a, int b,
int **xptr, int *xval, int xincr,
int **yptr, int *yval, int yincr, int *AWB,
int AWB_accum);
```

**Built-In Functions List (continued)****\_builtin\_mac****Argument:**

*Accum* Accumulator to sum.

*a* Integer multiplicand.

*b* Integer multiplier.

*xptr* Integer Pointer to pointer to X prefetch.

*xval* Integer Pointer to value of X prefetch.

*xincr* Integer increment value of X prefetch.

*yptr* Integer Pointer to pointer to Y prefetch.

*yval* Integer Pointer to value of Y prefetch.

*yincr* Integer increment value of Y prefetch.

*AWB* Accumulator Write-Back location.

*AWB\_accum* Accumulator to Write-Back.

The arguments, *xptr* and *yptr*, must point to the arrays located in the X data memory and Y data memory, respectively.

**Return Value:**

Returns the cleared value result to an accumulator.

**Assembler Operator/Machine Instruction:**

```
mac
```

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- *Accum* is not an Accumulator register
- *xval* is a null value but *xptr* is not null
- *yval* is a null value but *yptr* is not null
- *AWB\_accum* is not an Accumulator register and *AWB* is not null

**\_builtin\_modsd****Description:**

Issues the 16-bit architecture's native signed divide support. Notably, if the quotient does not fit into a 16-bit result, the results (including remainder) are unexpected. This form of the built-in function will capture only the remainder.

**Prototype:**

```
signed int _builtin_modsd(signed long dividend,
                           signed int divisor);
```

**Argument:**

*dividend* Number to be divided.

*divisor* Number to divide by.

**Return Value:**

Remainder.

**Assembler Operator/Machine Instruction:**

```
modsd
```

**Error Messages:**

None.

**\_builtin\_modud****Description:**

Issues the 16-bit architecture's native unsigned divide support. Notably, if the quotient does not fit into a 16-bit result, the results (including remainder) are unexpected. This form of the built-in function will capture only the remainder.

**Built-In Functions List (continued)****\_builtin\_modud****Prototype:**

```
unsigned int _builtin_modud(unsigned long dividend,
                           unsigned int divisor);
```

**Argument:**

*dividend* Number to be divided.

*divisor* Number to divide by.

**Return Value:**

Remainder.

**Assembler Operator/Machine Instruction:**

```
modud
```

**Error Messages:**

None.

**\_builtin\_mpy****Description:**

Computes  $a \times b$ ; also, prefetches data ready for a future MAC operation.

*xptr* may be null to signify no X prefetch to be performed; in which case, the values of *xincr* and *xval* are ignored but required.

*yptr* may be null to signify no Y prefetch to be performed; in which case, the values of *yincr* and *yval* are ignored but required.

*xval* and *yval* nominate the address of a C variable where the prefetched value will be stored.

*xincr* and *yincr* may be the literal values: -6, -4, -2, 0, 2, 4, 6 or an integer value.

For example:

```
register int result asm("A");
int      *xmemory;
int      *ymemory;
int      xVal, yVal;
result = _builtin_mpy(xVal, yVal,
                      &xmemory, &xVal, 2,
                      &ymemory, &yVal, 2);
```

May generate:

```
mac w4*w5, A, [w8]+=2, w4, [w10]+=2, w5
```

**Prototype:**

```
int _builtin_mpy(int a, int b,
                 int **xptr, int *xval, int xincr,
                 int **yptr, int *yval, int yincr);
```

**Argument:**

*a* Integer multiplicand.

*b* Integer multiplier.

*xptr* Integer Pointer to pointer to X prefetch.

*xval* Integer Pointer to value of X prefetch.

*xincr* Integer increment value of X prefetch.

*yptr* Integer Pointer to pointer to Y prefetch.

*yval* Integer Pointer to value of Y prefetch.

*yincr* Integer increment value of Y prefetch.

*AWB* Integer Pointer to accumulator selection.

The arguments *xptr* and *yptr* must point to the arrays located in the X data memory and the Y data memory, respectively.

**Built-In Functions List (continued)****\_builtin\_mpy****Return Value:**

Returns the cleared value result to an accumulator.

**Assembler Operator/Machine Instruction:**

```
mpy
```

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- *xval* is a null value but *xptr* is not null
- *yval* is a null value but *yptr* is not null

**\_builtin\_mpyn****Description:**

Computes  $-a \times b$ ; also, prefetches data ready for a future MAC operation.

*xptr* may be null to signify no X prefetch to be performed; in which case, the values of *xincr* and *xval* are ignored but required.

*yptr* may be null to signify no Y prefetch to be performed; in which case, the values of *yincr* and *yval* are ignored but required.

*xval* and *yval* nominate the address of a C variable where the prefetched value will be stored.

*xincr* and *yincr* may be the literal values: -6, -4, -2, 0, 2, 4, 6 or an integer value.

For example:

```
register int result asm("A");
int      *xmemory;
int      *ymemory;
int      xVal, yVal;
result = _builtin_mpy(xVal, yVal,
                      &xmemory, &xVal, 2,
                      &ymemory, &yVal, 2);
```

May generate:

```
mac w4*w5, A, [w8]+=2, w4, [w10]+=2, w5
```

**Prototype:**

```
int _builtin_mpyn(int a, int b,
                  int **xptr, int *xval, int xincr,
                  int **yptr, int *yval, int yincr);
```

**Argument:**

*a* Integer multiplicand.

*b* Integer multiplier.

*xptr* Integer Pointer to pointer to X prefetch.

*xval* Integer Pointer to value of X prefetch.

*xincr* Integer increment value of X prefetch.

*yptr* Integer Pointer to pointer to Y prefetch.

*yval* Integer Pointer to value of Y prefetch.

*yincr* Integer increment value of Y prefetch.

*AWB* Integer Pointer to accumulator selection.

The arguments *xptr* and *yptr* must point to the arrays located in the X data memory and the Y data memory, respectively.

**Return Value:**

Returns the cleared value result to an accumulator.

**Built-In Functions List (continued)****\_builtin\_mbyn****Assembler Operator/Machine Instruction:**

mpyn

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- *xval* is a null value but *xptr* is not null
- *yval* is a null value but *yptr* is not null

**\_builtin\_msc****Description:**Computes *a* *x* *b* and subtracts from accumulator; also, prefetches data ready for a future *MAC* operation.*xptr* may be null to signify no X prefetch to be performed; in which case, the values of *xincr* and *xval* are ignored but required.*yptr* may be null to signify no Y prefetch to be performed; in which case, the values of *yincr* and *yval* are ignored but required.*xval* and *yval* nominate the address of a C variable where the prefetched value will be stored.*xincr* and *yincr* may be the literal values: -6, -4, -2, 0, 2, 4, 6 or an integer value.If *AWB* is non-null, the other accumulator will be written back into the referenced variable.

For example:

```
register int result asm("A");
int      *xmemory;
int      *ymemory;
int      xVal, yVal;
result = _builtin_msc(result, xVal, yVal,
                      &xmemory, &xVal, 2,
                      &ymemory, &yVal, 2, 0, 0);
```

May generate:

```
msc w4*w5, A, [w8]+=2, w4, [w10]+=2, w5
```

**Prototype:**

```
int _builtin_msc(int Accum, int a, int b,
int **xptr, int *xval, int xincr,
int **yptr, int *yval, int yincr, int *AWB,
int AWB_accum);
```

**Argument:***Accum* Accumulator to sum.*a* Integer multiplicand.*b* Integer multiplier.*xptr* Integer Pointer to pointer to X prefetch.*xval* Integer Pointer to value of X prefetch.*xincr* Integer increment value of X prefetch.*yptr* Integer Pointer to pointer to Y prefetch.*yval* Integer Pointer to value of Y prefetch.*yincr* Integer increment value of Y prefetch.*AWB* Accumulator Write-Back location.*AWB\_accum* Accumulator to Write-Back.The arguments *xptr* and *yptr* must point to the arrays located in the X data memory and the Y data memory, respectively.**Return Value:**

Returns the cleared value result to an accumulator.

**Built-In Functions List (continued)****\_builtin\_msc****Assembler Operator/Machine Instruction:**

msc

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- *Accum* is not an Accumulator register
- *xval* is a null value but *xptr* is not null
- *yval* is a null value but *yptr* is not null
- *AWB\_accum* is not an Accumulator register and *AWB* is not null

**\_builtin\_mulss****Description:**

Computes the product *p0*  $\times$  *p1*. Function arguments are signed integers and the function result is a signed long integer. The command-line option, *-Wconversions*, can be used to detect unexpected sign conversions.

**Prototype:**

```
signed long _builtin_mulss(const signed int p0, const signed int p1);
```

**Argument:**

*p0* Multiplicand.

*p1* Multiplier.

**Return Value:**

Returns the signed long integer value of the product *p0*  $\times$  *p1*.

**Assembler Operator/Machine Instruction:**

mul.ss

**\_builtin\_mulsu****Description:**

Computes the product *p0*  $\times$  *p1*. Function arguments are integers with mixed signs and the function result is a signed long integer. The command-line option, *-Wconversions*, can be used to detect unexpected sign conversions. This function supports the full range of addressing modes of the instruction, including Immediate mode for operand *p1*.

**Prototype:**

```
signed long _builtin_mulsu(const signed int p0, const unsigned int p1);
```

**Argument:**

*p0* Multiplicand.

*p1* Multiplier.

**Return Value:**

Returns the signed long integer value of the product *p0*  $\times$  *p1*.

**Assembler Operator/Machine Instruction:**

mul.su

**\_builtin\_mulus****Description:**

Computes the product *p0*  $\times$  *p1*. Function arguments are integers with mixed signs and the function result is a signed long integer. The command-line option, *-Wconversions*, can be used to detect unexpected sign conversions. This function supports the full range of addressing modes of the instruction.

**Built-In Functions List (continued)****\_builtin\_mulus****Prototype:**

```
signed long _builtin_mulus(const unsigned int p0, const signed int p1);
```

**Argument:**

*p0* Multiplicand.

*p1* Multiplier.

**Return Value:**

Returns the signed long integer value of the product  $p0 \times p1$ .

**Assembler Operator/Machine Instruction:**

```
mul.us
```

**\_builtin\_muluu****Description:**

Computes the product  $p0 \times p1$ . Function arguments are unsigned integers and the function result is an unsigned long integer. The command-line option, `-Wconversions`, can be used to detect unexpected sign conversions. This function supports the full range of addressing modes of the instruction, including Immediate mode for operand *p1*.

**Prototype:**

```
unsigned long _builtin_muluu(const unsigned int p0, const unsigned int p1);
```

**Argument:**

*p0* Multiplicand.

*p1* Multiplier.

**Return Value:**

Returns the signed long integer value of the product  $p0 \times p1$ .

**Assembler Operator/Machine Instruction:**

```
mul.uu
```

**\_builtin\_nop****Description:**

Generates a NOP instruction.

**Prototype:**

```
void _builtin_nop(void);
```

**Argument:**

None.

**Return Value:**

Returns a no operation (NOP).

**Assembler Operator/Machine Instruction:**

```
NOP
```

**\_builtin\_readsfr****Description:**

Reads the SFR.

**Prototype:**

```
unsigned int _builtin_readsfr(const void *p);
```

**Built-In Functions List (continued)****\_builtin\_readsfr****Argument:***p* Object address.**Return Value:**

Returns the SFR.

**Assembler Operator/Machine Instruction:**

readsfr

**\_builtin\_return\_address****Description:**

Returns the return address of the current function or of one of its callers. For the *level* argument, a value of zero yields the return address of the current function, a value of one yields the return address of the caller of the current function, and so forth. When *level* exceeds the current stack depth, zero will be returned. This function should only be used with a non-zero argument for debugging purposes.

**Prototype:**

int \_builtin\_return\_address (const int level);

**Argument:***level* Number of frames to scan up the call stack.**Return Value:**

Returns the return address of the current function or of one of its callers.

**Assembler Operator/Machine Instruction:**

return\_address

**\_builtin\_sac****Description:**Shifts value by *shift* (a literal between -8 and 7) and returns the value.

For example:

```
register int value asm("A");
int result;
result = _builtin_sac(value,3);
```

May generate:

sac A, #3, w0

**Prototype:**

int \_builtin\_sac(int value, int shift);

**Argument:***value* Integer number to be shifted.*shift* Literal amount to shift.**Return Value:**

Returns the shifted result to an accumulator.

**Assembler Operator/Machine Instruction:**

sac

**Built-In Functions List (continued)****\_builtin\_sac****Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- The shift value is not a literal within range

**\_builtin\_sacr****Description:**

Shifts value by *shift* (a literal between -8 and 7) and returns the value, which is rounded using the Rounding mode determined by the RND (CORCON<1>) control bit.

For example:

```
register int value asm("A");
int      result;
result = _builtin_sacr(value, 3);
```

May generate:

```
sacr A, #3, w0
```

**Prototype:**

```
int _builtin_sacr(int value, int shift);
```

**Argument:**

*value* Integer number to be shifted.

*shift* Literal amount to shift.

**Return Value:**

Returns the shifted result to the CORCON register.

**Assembler Operator/Machine Instruction:**

```
sacr
```

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- The shift value is not a literal within range

**\_builtin\_sftac****Description:**

Shifts accumulator by *shift*. The valid shift range is -16 to 16.

For example:

```
register int result asm("A");
int      i;
result = _builtin_sftac(result,i);
```

May generate:

```
sftac A, w0
```

**Prototype:**

```
int _builtin_sftac(int Accum, int shift);
```

**Argument:**

*Accum* Accumulator to shift.

*shift* Amount to shift.

**Built-In Functions List (continued)****\_builtin\_sftac****Return Value:**

Returns the shifted result to an accumulator.

**Assembler Operator/Machine Instruction:**

```
sftac
```

**Error Messages:**

An error message appears if:

- The result is not an Accumulator register
- *Accum* is not an Accumulator register
- The shift value is not a literal within range

**\_builtin\_subab****Description:**

Subtracts accumulators A and B with the result written back to the specified accumulator. For example:

```
register int result asm("A");
register int B asm("B");
result = _builtin_subab(result,B);
```

Will generate:

```
sub A
```

**Prototype:**

```
int _builtin_subab(int Accum_a, int Accum_b);
```

**Argument:**

*Accum\_a* Accumulator from which to subtract.

*Accum\_b* Accumulator to subtract.

**Return Value:**

Returns the subtraction result to an accumulator.

**Assembler Operator/Machine Instruction:**

```
sub
```

**Error Messages:**

An error message appears if the result is not an Accumulator register.

## 6. Reference

### 6.1. Instruction Bit Map

The CPU instruction opcode map for the PIC32A class of devices is shown in [Figure 6-1](#). The map comprises 4 x 64 slot “Quadrants,” each selected by a 2-bit opcode extension bit field within the (32-bit) instruction. Each instruction has a 6-bit opcode, selecting one of 64 opcodes within a Quadrant. The 16-bit instructions do not have the extension bit field but assume use of the default Quadrant 0. New instructions are highlighted in blue text, while unused slots are grayed out. The opcode map shows the instructions defined for the FPU coprocessor in green text.

**Note:** The complete opcode for each instruction can be determined by the instruction descriptions in [Instruction Descriptions](#).

**Figure 6-1. Instruction Encoding**

Col Ref →		Opcode Size	Opcode Extension	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15											
Row Ref ↓		Opcode Size	Opcode Extension	0000	0001	0010	0011	0100	0101	0110	0111	1000	1001	1010	1011	1100	1101	1110	1111											
a	16/32-bit	00	NEOP (16-bit only)	MOV MOVW (W → W)	MOV MOVW (W → W)	GP Co-processor Move Ops (16-bit)			MOV (W15-H17) W → (W14-H17)	MOV (W15-H17) W → (W14-H17) → W	MOV/ IIS		MULSS MULUS	MULUU MULSU			POP (16-bit only)	PUSH (16-bit only)												
			ASR by Short Lit	LSR by Short Lit	SL	PUSH Co-processor	MOV (F → W)	ASR	LSR	SL	POP Co-processor	RLC	RNC	RRC	RRNC	CP #13	CPB #13	CP	CPB											
			BQR	BSET	BTG	BTSF	BTSTS	BSW	BTST Wn with Wb																					
			ADD	ADDC	SUB	SUBB	AND	IOR	SUBR	SUBBR	XOR	Inherent	ADD RS	SUB IIS	COM	SUB#17	SE	ZE												
e	32-bit	01	MOVJ lit32	GP Co-processor Move Ops (32-bit)			MOVJ Wn+Slit14 → Fd Pn → (Wn+Slit14)	MOVJ Wn+Slit14 → Fd Pn → (Wn+Slit14)			MOVJ/w/J literal MOVS			MOVJ/w/J	MOVJ/w/J			MOVJ/bz	MOVJ/bz											
f			ASR	LSR	SL	CLR	CPO	CP	CPB	SET	RLC	RNC	RRC	RRNC		DTB	MUL	PUSH												
g			Instruction with File Register			Instruction with File Register			Instruction with File Register			Instruction with File Register			Instruction with File Register			Instruction with File Register												
h			BQR	BSET	BTG	BTSF	BTSTS	BRNS lit16	BRNS BFEXT		TST	INC	DEC	COM	NEGf	JNZ2	DEC2													
i	32-bit	10	ADD lit16 and Wn			SUB lit16 and Wn CP (lit16)			SUBB Ws and lit7			SUBR Ws and lit7			XORWF			Reserved for Co-processor Ops												
j			FADD FSUB	FMAX FMX2	FMAC FMUL FDIV	FMOV FMOVc	FABS FNEG FSQRT	FOPS FQFQ	FSN FQFS	RF21 RF22 RF23 RF24	FAND FIOR	FRIM	FIST	Reserved	Reserved	Reserved	Reserved	Reserved												
k			ADD lit16 and Wn			SUB lit16 and Wn CP (lit16)			AND lit16 and Wn			RET lit16 and Wn			XOR lit16 ad Wn			Reserved for Co-processor Branches												
l			ADD Ws and lit7			ADD Ws and lit7			SUB Ws and lit7			AND Ws and lit7			SUBR Ws and lit7			RECALL REPEAT LINK												
m	11	00																												
n																														
o			LAC ADD, SUB			SAC SACR, SAC SACR, SAC			MIN, SRC MAX, NORM			SFTAC, SFTAC, SIS LAC			MAC/SQAC MSQAC, SRC			MPY/SQR MPY/SQNN												
p			ASRM LSRM SLM by lit5	ASRM LSRM SLM by Wb	FIM MIN MAX	FF1L FF1R FBCL	DIVL DIVU DIVF	DIVSL DIVUL DIVLF			MOV/L Move f (extended) to Wd	MOV/L Move Wd to f (extended)	GOTO CALL I22	MOV/F SWAP	BFINS BFEXT	RESET BOOTSWAP	MOV [Wn + Slit4], Wind	NOPR (32-bit)												

### 6.2. Instruction Set Summary Table

The complete 16-bit DSC device instruction set is summarized in [Table 6-1](#). This table contains an alphabetized listing of the instruction set. It includes instruction assembly syntax, description, size (in 24-bit words), execution time (in instruction cycles), affected Status bits and the page number where the detailed description can be found. The PIC32A CPU Instruction Status Flag Operations section identifies the symbols that are used in the [Table 6-1](#).

**Table 6-1.** Instruction Set Summary Table

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
ADD	f + Wn or Wn = f + Wn	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wn = Wn + lit5	0.5	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wn = Wn + lit16	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = Wb + Ws	0.5/1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = Wb + lit7 (literal zero-extended)	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Add Accumulators	0.5	1	—	—	—	—	↔	↔	↑	↑	↔	↑
	16-bit Signed Add to Accumulator	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
ADDC	f = f + Wn + (C) or Wn = f + Wn + (C)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Wn = Wn + lit16 + (C)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Wd = Wb + Ws + (C)	0.5/1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Wd = Wb + lit7 + (C) (literal zero-extended)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
AND	f = f .AND. Wn or Wn = f .AND. Wn	1	1	—	↔	—	↔	—	—	—	—	—	—
	Wn = Wn .AND. lit16	1	1	—	↔	—	↔	—	—	—	—	—	—
	Wd = Wb .AND. Ws	0.5/1	1	—	↔	—	↔	—	—	—	—	—	—
	Wd = Wb .AND. lit7 (literal zero-extended)	1	1	—	↔	—	↔	—	—	—	—	—	—
AND1	Wd = Wb .AND. lit7 (literal ones-extended)	1	1	—	↔	—	↔	—	—	—	—	—	—
ASR	f = Arithmetic Right Shift f by 1 or Wn = Arithmetic Right Shift f by 1	1	1	↔	↔	↔	—	—	—	—	—	—	—
	Wd = Arithmetic Right Shift Ws by 1	0.5/1	1	↔	↔	↔	—	—	—	—	—	—	—
	Wnd = Arithmetic Right Shift Ws by Wb	0.5/1	1	↔	↔	↔	—	—	—	—	—	—	—
	Wnd = Arithmetic Right Shift Ws by lit5	0.5/1	1	↔	↔	↔	—	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↔ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
ASRM	Wnd = Arithmetic Right Shift Ws by lit5, then logically OR with next LS-word	1	2	—	↑↓	—	↓	—	—	—	—	—	—
	Wnd = Arithmetic Right Shift Ws by Wb, then logically OR with next LS-word	1	2	—	↑↓	—	↓	—	—	—	—	—	—
BCLR	Bit Clear f	1	1	—	—	—	—	—	—	—	—	—	—
	Bit Clear Ws	0.5/1	1	—	—	—	—	—	—	—	—	—	—
BFEXT	Bit Field Extract from Ws to Wb	1	1	—	—	—	—	—	—	—	—	—	—
	Bit Field Extract from f to Wb	2	2	—	—	—	—	—	—	—	—	—	—
BFINS	Bit Field Insert from Wb into Ws	1	1	—	—	—	—	—	—	—	—	—	—
	Bit Field Insert from Wb into f	2	2	—	—	—	—	—	—	—	—	—	—
	Bit Field Insert lit8 into Ws	2	2	—	—	—	—	—	—	—	—	—	—
BOOTSWP	Swap Active and Inactive address space	1	2	—	—	—	↑↓	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↑ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
BRA	Branch Unconditionally	1	1	—	—	—	—	—	—	—	—	—	—
	Computed Branch	1	2	—	—	—	—	—	—	—	—	—	—
	Branch if Carry	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if greater than or equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if unsigned greater than or equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if greater than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if unsigned greater than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if less than or equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if unsigned less than or equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if less than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if unsigned less than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if Negative	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if Not Carry	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if Not Negative	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if Not Overflow	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if Not Zero	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
BREAK	Branch if accumulator A overflow	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if accumulator B overflow	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if Overflow	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if accumulator A saturated	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if accumulator B saturated	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Branch if Zero	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
BREAK	Stop user code execution	1	1	—	—	—	—	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↑ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
BSET	Bit Set f	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	Bit Set Ws	1	1	—	—	—	—	—	—	—	—	—	—
BSW	Write C or Z bit to Ws<Wb>	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	Write C or Z bit to Ws<Wb>	1	1	—	—	—	—	—	—	—	—	—	—
BTG	Bit Toggle f	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	Bit Toggle Ws	1	1	—	—	—	—	—	—	—	—	—	—
BTST	Bit Test f	0.5/1	1	‡‡	—	—	‡‡	—	—	—	—	—	—
	Bit Test Ws to C	1	1	‡‡	—	—	‡‡	—	—	—	—	—	—
	Bit Test Ws to Z	0.5/1	1	‡‡	—	—	‡‡	—	—	—	—	—	—
	Bit Test Ws<Wb> to C	1	1	‡‡	—	—	‡‡	—	—	—	—	—	—
	Bit Test Ws<Wb> to Z	1	1	‡‡	—	—	‡‡	—	—	—	—	—	—
BTSTS	Bit Test then Set f	0.5/1	1	‡‡	—	—	‡‡	—	—	—	—	—	—
	Bit Test Ws to C then Set	1	1	‡‡	—	—	‡‡	—	—	—	—	—	—
	Bit Test Ws to Z then Set			‡‡	—	—	‡‡	—	—	—	—	—	—
CALL	Call subroutine (label < ~16MB)	1	1	—	—	—	—	—	—	—	—	—	—
	Call subroutine (label > ~16MB)	2	2	—	—	—	—	—	—	—	—	—	—
	Call indirect subroutine at address [W11]	1	2	—	—	—	—	—	—	—	—	—	—
CLR	f = 0x0000	1	1										
	Wd = 0x0000	1	1										
	Clear Accumulator	0.5	1										
CLRWDT	Clear Watchdog Timer	0.5	1										
COM	f = f or Wd = f	1	1	—	‡‡	—	‡‡	—	—	—	—	—	—
	Wd = Ws	0.5/1	1	—	‡‡	—	‡‡	—	—	—	—	—	—
CP	Compare f with Ws	1	1	‡‡	‡‡	‡‡	‡‡	—	—	—	—	—	—
	Compare Ws with lit13 (literal zero-extended)	1	1	‡‡	‡‡	‡‡	‡‡	—	—	—	—	—	—
	Compare Wb with lit16 (literal zero-extended)	1	1	‡‡	‡‡	‡‡	‡‡	—	—	—	—	—	—
	Compare Wb with Ws	0.5/1	1	‡‡	‡‡	‡‡	‡‡	—	—	—	—	—	—

<sup>1</sup> Legend:

- ‡‡ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ‡‡ may be cleared but never set (sticky), †† may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
CP0	Compare f with 0x0000	1	1	1	↔	0	↔	—	—	—	—	—	—
	Compare Ws with 0x0000	1	1	↔	↔	↔	↔	—	—	—	—	—	—
CPB	Compare f with Ws, with borrow	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Compare Wb with lit13, with borrow(literal zero-extended)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Compare Wb with lit16, with borrow(literal zero-extended)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Compare Borrow Wb with Ws	0.5/1	1	↔	↔	↔	↓	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↔ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
CBRA	Floating Point Branch if Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Not Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Greater Than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Greater Than or Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Less Than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Less Than or Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Ordered	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Unordered or Not Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Unordered or Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Unordered or Less Than or Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Unordered or Less Than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Unordered or Greater Than or Equal	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Unordered or Greater Than	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
	Floating Point Branch if Unordered	1	1(2/3)	—	—	—	—	—	—	—	—	—	—
CTXTSWP	Swap to CPU register context defined in lit3	0.5	2	—	—	—	—	—	—	—	—	—	—
	Swap to CPU register context defined in Wn[2:0]	1	2	—	—	—	—	—	—	—	—	—	—
DTB	Decrement Wn, then branch if not zero	1	1(2/3)	—	—	—	—	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↑ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
DEC	f = f - 1 or Wnd = f - 1	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
	Wd = Ws - 1 (substitute with SUB Ws,#1,Wd)	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
DEC2	f = f - 2 or Wnd = f - 2	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
	Wd = Ws - 2 (substitute with SUB Ws,#2,Wd)	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
DISICTL	Disable interrupts at IPL <= lit3 Optionally save prior IPL threshold to Wd	1	1	—	—	—	—	—	—	—	—	—	—
	Disable interrupts at IPL <= Wns[2:0] Optionally save prior IPL threshold to Wd	1	1	—	—	—	—	—	—	—	—	—	—
DIVF	Interruptible Signed 16/16 or 32/16 Fractional Divide	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
DIVFL	Interruptible Signed 32/32 Fractional Divide	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
DIVS	Interruptible Signed 16/16-bit Integer Divide	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
	Interruptible Signed 32/16-bit Integer Divide	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
DIVSL	Interruptible Signed 32/32 Integer Divide	1	1	⊕	⊕	⊕	⊕	—	—	—	—	—	—
DIVU	Interruptible Unsigned 16/16-bit Integer Divide	1	1	⊕	0	0	⊕	—	—	—	—	—	—
	Interruptible Unsigned 32/16-bit Integer Divide	1	1	⊕	0	⊕	⊕	—	—	—	—	—	—
DIVUL	Interruptible Unsigned 32/32 Integer Divide	1	1	⊕	0	0	⊕	—	—	—	—	—	—
ED	Euclidean Distance	1	2	—	—	—	—	⊕	⊕	↑	↑	⊕	↑
EDAC	Euclidean Distance Accumulate	1	2	—	—	—	—	⊕	⊕	↑	↑	⊕	↑
EXCH	Swap Wns with Wnd	1	2	—	—	—	—	—	—	—	—	—	—
FBCL	Find Bit Change from Left (MSb) Side	1	1	⊕	—	—	—	—	—	—	—	—	—
FF1L	Find First One from Left (MSb) Side	1	1	⊕	—	—	—	—	—	—	—	—	—

<sup>1</sup> Legend:

- ⊕ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ⊖ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
FF1R	Find First One from Right (Lsb) Side	1	1	↔	—	—	—	—	—	—	—	—	—
FLIM	Force Data (Upper and Lower) Range Limit without Limit Excess Result	1	1	—	↔	0	↔	—	—	—	—	—	—
	Force Data (Upper and Lower) Range Limit with Limit Excess Flag (Wd = -1)	1	2	—	↔	0	↔	—	—	—	—	—	—
	Force Data (Upper and Lower) Range Limit with Limit Excess Result	1	2	—	↔	0	↔	—	—	—	—	—	—
GOTO	Go to address (address < ~16MB)	1	1	—	—	—	—	—	—	—	—	—	—
	Go to address (address > ~16MB)	2	2	—	—	—	—	—	—	—	—	—	—
	Go to indirect address at [W11]	1	2	—	—	—	—	—	—	—	—	—	—
INC	f = f + 1 or Wnd = f + 1	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = Ws + 1 (substitute with ADD Ws,#1,Wd)	1	1	↔	↔	↔	↔	—	—	—	—	—	—
INC2	f = f + 2 or Wnd = f + 2	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = Ws +2 (Substitute with ADD Ws,#2,Wd)	1	1	↔	↔	↔	↔	—	—	—	—	—	—
IOR	f = f .IOR. Wn or Wn = f .IOR. Wn	1	1	—	↔	—	↔	—	—	—	—	—	—
	Wn = Wn .IOR. lit16	1	1	—	↔	—	↔	—	—	—	—	—	—
	Wd = Wb .IOR. Ws	0.5/1	1	—	↔	—	↔	—	—	—	—	—	—
	Wd = Wb .IOR. lit7	1	1	—	↔	—	↔	—	—	—	—	—	—
LAC	Load Accumulator (16/32-bit), literal shift	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
LLAC	Load Lower (LS-word of) Accumulator (32-bit), literal shift	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
LUAC	Load Upper (LS-byte) of Accumulator (32-bit), literal shift	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑

<sup>1</sup> Legend:

- ↔ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
LNK	Link frame pointer	1	1	—	—	—	—	—	—	—	—	—	—
	Link frame pointer (literal < 128)	0.5	1	—	—	—	—	—	—	—	—	—	—
LSR	f = Logical Right Shift f by 1 or Wd = Logical Right Shift f by 1	1	1	↔	0	—	↔	—	—	—	—	—	—
	Wd = Logical Right Shift Ws by 1	0.5/1	1	↔	0	—	↔	—	—	—	—	—	—
	Wnd = Logical Right Shift Ws by Wns	0.5/1	1	↔	0	—	↔	—	—	—	—	—	—
	Wnd = Logical Right Shift Ws by lit5	0.5/1	1	↔	0	—	↔	—	—	—	—	—	—
LSRM	Wnd = Logical Right Shift Ws by lit5, then logically OR with next LS-word	1	2	↔	0	—	↔	—	—	—	—	—	—
	Wnd = Logical Right Shift Ws by Wb, then logically OR with next LS-word	1	2	↔	0	—	↔	—	—	—	—	—	—
MAC	Multiply and Accumulate	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
MAX	Force Data Maximum Range Limit	1	1	—	0	0	↔	—	—	—	—	—	—
	Force Data Maximum Range Limit	0.5	1	—	0	0	↔	—	—	—	—	—	—
	Force Data Maximum Range Limit with Result	1	2	—	0	0	↔	—	—	—	—	—	—
MIN	Force Data Minimum Range Limit	1	1	—	↔	0	↔	—	—	—	—	—	—
	Force Data Minimum Range Limit	0.5	1	—	↔	0	↔	—	—	—	—	—	—
	Force Data Minimum Range Limit with Result	1	2	—	↔	0	↔	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↔ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
MOV	Move Ws to Wd	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	Move coprocessor register to Wd	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	Move Ws to coprocessor register	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	Move 32-bit unsigned literal to Wnd	2	2	—	—	—	—	—	—	—	—	—	—
	Move 8/16/24-bit unsigned literal to Wnd; 0 extend to 32-bits	1	1	—	—	—	—	—	—	—	—	—	—
	Move 32-bit unsigned literal to co-processor register	2	2	—	—	—	—	—	—	—	—	—	—
	Move from system stack with literal offset to Wnd using SP or FP	0.5	1	—	—	—	—	—	—	—	—	—	—
	Move from Wns to system stack with literal offset using SP or FP	0.5	1	—	—	—	—	—	—	—	—	—	—
	Move f to Wnd (Word or Long Word only) (f < ~1MB)	1	1	—	—	—	—	—	—	—	—	—	—
	Move f to Wnd (f > ~1MB)	2	2	—	—	—	—	—	—	—	—	—	—
	Move f to Wnd (Byte)	1	1	—	—	—	—	—	—	—	—	—	—
	Move Wns to f (Word or Long Word only) (f < ~1MB)	1	1	—	—	—	—	—	—	—	—	—	—
	Move Wns to f (f > ~1MB)	2	2	—	—	—	—	—	—	—	—	—	—
	Move Wns to f (Byte)	1	1	—	—	—	—	—	—	—	—	—	—
	Move [Wns+Slit12/13/14] to Wnd (Slitxx size varies with data width)	1	1	—	—	—	—	—	—	—	—	—	—
	Move [Wns+Slit14] to Fd	1	1	—	—	—	—	—	—	—	—	—	—
	Move Wns to [Wnd+Slit12/13/14] (Slitxx size varies with data width)	1	1	—	—	—	—	—	—	—	—	—	—
	Move Fs to [Wnd+Slit14]	1	1	—	—	—	—	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↑ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
MOVIF	If SR.Z = 1 Move Wb to Wd Else Move Wns to Wd	1	1	—	—	—	—	—	—	—	—	—	—
MOVR	Move Ws to Wd with destination Bit Reversed addressing	1	1	—	—	—	—	—	—	—	—	—	—
MOVS.I	Move signed extended 16-bit literal to Wd	1	1	—	—	—	—	—	—	—	—	—	—
MOVS.w	Move 16-bit literal to Wd; sign extend to 32-bits if register direct mode.	1	1	—	—	—	—	—	—	—	—	—	—
MOVS.b	Move 8-bit literal to Wd; no extension.	1	1	—	—	—	—	—	—	—	—	—	—
MPY	Multiply Wm by Wn to Accumulator	1	1	—	—	—	—	↑↑	↑↑	↑↑	↑↑	↑↑	↑↑
MPYN	-(Multiply Wm by Wn) to Accumulator	1	1	—	—	—	—	↑↑	↑↑	↑↑	↑↑	↑↑	↑↑
MSC	Multiply and Subtract from Accumulator	1	1	—	—	—	—	↑↑	↑↑	↑↑	↑↑	↑↑	↑↑
MUL	W2 = f * Wn	1	1	—	—	—	—	—	—	—	—	—	—
MULFSS	Fractional: Acc(A or B) = signed(Wb) * signed(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
MULFSU	Fractional: Acc(A or B) = signed(Wb) * unsigned(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
MULFUS	Fractional: Acc(A or B) = unsigned(Wb) * signed(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
MULFUU	Fractional: Acc(A or B) = unsigned(Wb) * unsigned(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
MULISS	Integer: Acc(A or B) = signed(Wb) * signed(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
	Integer: Acc(A or B) = signed(Wb) * signed(slit8)	1	1	—	—	—	—	—	—	—	—	—	—
MULISU	Integer: Acc(A or B) = signed(Wb) * unsigned(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
	Integer: Acc(A or B) = signed(Wb) * unsigned(lit8)	1	1	—	—	—	—	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↑↑ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓↓ may be cleared but never set (sticky), ↑↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
MULIUS	Integer: Acc(A or B) = unsigned(Wb) * signed(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
	Integer: Acc(A or B) = unsigned(Wb) * signed(slit8)	1	1	—	—	—	—	—	—	—	—	—	—
MULIUU	Integer: Acc(A or B) = unsigned(Wb) * unsigned(Ws)	1	1	—	—	—	—	—	—	—	—	—	—
	Integer: Acc(A or B) = unsigned(Wb) * unsigned(lit8)	1	1	—	—	—	—	—	—	—	—	—	—
MULSS	{Wd}=signed(Wb) * signed(Ws)	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	{Wd}=signed(Wb) * signed(slit8)	1	1	—	—	—	—	—	—	—	—	—	—
MULSU	{Wd}=signed(Wb) * unsigned(Ws)	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	{Wd}=signed(Wb) * unsigned(lit8)	1	1	—	—	—	—	—	—	—	—	—	—
MULUS	{Wd}=unsigned(Wb) * signed(Ws)	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	{Wd}=unsigned(Wb) * signed(slit8)	1	1	—	—	—	—	—	—	—	—	—	—
MULUU	{Wd}=unsigned(Wb) * unsigned(Ws)	0.5/1	1	—	—	—	—	—	—	—	—	—	—
	{Wd}=unsigned(Wb) * unsigned(lit8)	1	1	—	—	—	—	—	—	—	—	—	—
NEG	Negate Accumulator	0.5	1	—	—	—	—	↑↑	↑↑	↑↑	↑↑	↑↑	↑↑
	f = $\bar{f} + 1$ or Wd = $\bar{f} + 1$	1	1	↑↑	0	↑↑	↑↑	—	—	—	—	—	—
	Wd = $\bar{W_s} + 1$	0.5/1	1	↑↑	0	↑↑	↑↑	—	—	—	—	—	—
NEOP	None executable NOP (16-bit instruction pad)	0.5	0	—	—	—	—	—	—	—	—	—	—
NOP	No Operation	1	1	—	—	—	—	—	—	—	—	—	—
	No Operation	1	1	—	—	—	—	—	—	—	—	—	—
NORM	Normalize Accumulator	1	1	—	↑↑	—	↑↑	0	0	—	—	—	—
POP	Pop f from top of stack (TOS)	1	1	—	—	—	—	—	—	—	—	—	—
	Pop Wnd Register from system stack.	0.5	1	—	—	—	—	—	—	—	—	—	—
	Pop Fd Register from system stack.	0.5	1	—	—	—	—	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↑↑ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓↓ may be cleared but never set (sticky), ↑↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
PUSH	Push f to top of stack (TOS)	1	1	—	—	—	—	—	—	—	—	—	—
	Push Wns Register to system stack	0.5	1	—	—	—	—	—	—	—	—	—	—
	Push Fs Register to system stack	0.5	1	—	—	—	—	—	—	—	—	—	—
PWRSAV	Go into standby mode	0.5	2	—	—	—	—	—	—	—	—	—	—
RCALL	Relative Call	1	1	—	—	—	—	—	—	—	—	—	—
	Computed Call	1	2	—	—	—	—	—	—	—	—	—	—
REPEAT	Repeat Next Instruction lit15+1 times	1	1	—	—	—	—	—	—	—	—	—	—
	Repeat Next Instruction lit5+1 times	0.5	1	—	—	—	—	—	—	—	—	—	—
	Repeat Next Instruction (Wn)+1 times	1	1	—	—	—	—	—	—	—	—	—	—
RESET	Software device RESET	1	1	—	—	—	—	—	—	—	—	—	—
RETFIE	Return from interrupt enable	0.5	4	↕	↕	↕	↕	—	—	—	—	—	—
RETLW	Return from Subroutine with literal in Wn	1	3	—	—	—	—	—	—	—	—	—	—
RETURN	Return from Subroutine	0.5	3	—	—	—	—	—	—	—	—	—	—
RLC	f = Rotate Left through Carry f or Wd = Rotate Left through Carry f	1	1	↕	↕	—	↕	—	—	—	—	—	—
	Wd = Rotate Left through Carry Ws	0.5/1	1	↕	↕	—	↕	—	—	—	—	—	—
RLNC	f = Rotate Left (No Carry) f or Wd = Rotate Left (No Carry) f	1	1	—	↕	—	↕	—	—	—	—	—	—
	Wd = Rotate Left (No Carry) Ws	0.5/1	1	—	↕	—	↕	—	—	—	—	—	—
RRC	f = Rotate Right through Carry f or Wd = Rotate Right through Carry f	1	1	↕	↕	—	↕	—	—	—	—	—	—
	Wd = Rotate Right through Carry Ws	0.5/1	1	↕	↕	—	↕	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↕ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
RRNC	f = Rotate Right (No Carry) f or Wd = Rotate Right (No Carry) f	1	1	—	↑↓	—	↑↓	—	—	—	—	—	—
	Wd = Rotate Right (No Carry) Ws	0.5/1	1	—	↑↓	—	↑↓	—	—	—	—	—	—
SAC	Store Accumulator (16/32-bit)	1	1	—	—	—	—	—	—	—	—	—	—
SACR	Store Rounded Accumulator (16/32-bit), literal shift	1	1	—	—	—	—	—	—	—	—	—	—
	Store Rounded Accumulator (16/32-bit), Wb shift	1	1	—	—	—	—	—	—	—	—	—	—
SLAC	Store Lower (LS-Word of) Accumulator (32-bit), literal shift	1	1	—	—	—	—	—	—	—	—	—	—
SUAC	Store sign extended Upper (MS-Byte of) Accumulator (32-bit), literal shift	1	1	—	—	—	—	—	—	—	—	—	—
SE	Wd = sign-extended Ws	0.5/1	1	↑↓	↑↓	—	↑↓	—	—	—	—	—	—
SETM	f = 0xFFFF	1	1	—	—	—	—	—	—	—	—	—	—
	Wd = 0xFFFF (substitute with MOVS #0xFFFF, Wd)	1	1	—	—	—	—	—	—	—	—	—	—
SFTAC	Arithmetic Shift Accumulator by (Wn)	1	1	—	—	—	—	↑↓	↑↓	↑	↑	↑↓	↑
	Arithmetic Shift Accumulator by Slit7	1	1	—	—	—	—	↑↓	↑↓	↑	↑	↑↓	↑
SL	f = Left Shift f by 1 or Wd = Left Shift f by 1	1	1	↑↓	↑↓	—	↑↓	—	—	—	—	—	—
	Wd = Left Shift Ws by 1	0.5/1	1	↑↓	↑↓	—	↑↓	—	—	—	—	—	—
	Wnd = Left Shift Wb by Wns	0.5/1	1	↑↓	↑↓	—	↑↓	—	—	—	—	—	—
	Wnd = Left Shift Ws by lit5	0.5/1	1	↑↓	↑↓	—	↑↓	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↑↓ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
SLM	Wnd = Left Shift Wb by lit5, then logically OR with next MS-word	1	2	—	—	—	↓	—	—	—	—	—	—
	Wnd = Left Shift Wb by Wb, then logically OR with next MS-word	1	2	—	—	—	↓	—	—	—	—	—	—
SQR	Square to Accumulator	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
SQRAC	Square and Accumulate	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
SQRN	Negated Square to Accumulator	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
SQRSC	Square and Subtract from Accumulator	1	1	—	—	—	—	↔	↔	↑	↑	↔	↑
SSTEP	Single step	1	5	—	—	—	—	—	—	—	—	—	—
SUB	Subtract Accumulators	0.5	1	↔	↔	↔	↔	—	—	—	—	—	—
	f = f - Wn or Wn = f - Wn	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wn = Wn - lit5	0.5	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wn = Wn - lit16	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = Wb - Ws	0.5/1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = Ws - lit7 (literal zero-extended)	1	1	↔	↔	↔	↔	—	—	—	—	—	—
SUBB	16-bit Signed Subtract from Accumulator	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	f = f - Wn - (C) or Wn = f - Wn - (C)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Wn = Wn - lit16 - (C)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	Wd = Wb - Ws - (C)	0.5/1	1	↔	↔	↔	↓	—	—	—	—	—	—
SUBR	Wd = Ws - lit7 - (C) (literal zero-extended)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	f = Wn - f or Wn = Wn - f	1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = Ws - Wb	0.5/1	1	↔	↔	↔	↔	—	—	—	—	—	—
	Wd = lit7 - Ws (literal zero-extended)	0.5/1	1	↔	↔	↔	↔	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↔ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-1.** Instruction Set Summary Table (continued)

Assembly Mnemonic <sup>1</sup>	Description	Word s	Cycle s	C	N	OV	Z	OA	OB	SA	SB(1)	OAB	SAB(1 )
SUBBR	$f = Wn - f - (C)$ or $Wn = Wn - f - (C)$	1	1	↔	↔	↔	↓	—	—	—	—	—	—
	$Wd = Ws - Wb - (C)$	0.5/1	1	↔	↔	↔	↓	—	—	—	—	—	—
	$Wd = \text{lit7} - Ws - (C)$ (literal zero-extended)	1	1	↔	↔	↔	↓	—	—	—	—	—	—
SWAP	$Wn = \text{Word or byte swap } Wn$	1	1	—	—	—	—	—	—	—	—	—	—
TST	Test f	1	1	—	↔	—	↔	—	—	—	—	—	—
	Test f and move f to Wnd	1	1	—	↔	—	↔	—	—	—	—	—	—
UFEX	Execute op(s) in FEX2:FEX in Mission	1	5	—	—	—	—	—	—	—	—	—	—
ULNK	Unlink frame pointer	0.5	1	—	—	—	—	—	—	—	—	—	—
URUN	Switch from debug PC to user PC	1	3	—	—	—	—	—	—	—	—	—	—
XOR	$f = f .XOR. Wn$ or $Wn = f .XOR. Wn$	1	1	—	↔	—	↔	—	—	—	—	—	—
	$Wn = Wn .XOR. \text{lit16}$	1	1	—	↔	—	↔	—	—	—	—	—	—
	$Wd = Wb .XOR. Ws$	0.5/1	1	—	↔	—	↔	—	—	—	—	—	—
	$Wd = Wb .XOR. \text{lit7}$ (literal zero-extended)	1	1	—	↔	—	↔	—	—	—	—	—	—
ZE	$Wd = \text{Zero-extend } Ws$	0.5/1	1	1	0	—	↔	—	—	—	—	—	—

<sup>1</sup> Legend:

- ↔ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-2.** FPU Instruction Set Summary

Instruction <sup>2</sup> : Description	Cycle Count	FPU Status Register (FSR)															
		FSR[28:24]					FSR[19:16]				FSR[6:0]						
		SUB	INF	FN	FZ	FNA N	GT	LT	EQ	UN	SUBO	HUGI	INX	UDF	OVF	DIV0	INVAL
<b>Move instructions</b>																	
MOV: Move (Wns+offset) to Fd	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV: Move Fs to (Wnd+offset)	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV: Move (Ws) to Fd or FSR or FCR or FEAR	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV: Move Fs or FSR or FCR or FEAR to (Wd)	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV: Move 32-bit literal value into Fd or FSR or FCR	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
PUSH: Push Fs or FSR or FCR or FSRH to system stack	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
POP: Pop Fd or FSR or FCR from system stack	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV{.s, .d}: Move Fs to Fd	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOVC{.s, .d}: Load constant value into Fd	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
<b>Status Bit Set/Clear/Update Instructions</b>																	
AND: Clear bit(s) in active FP special reg, bits specified by literal value.	1	—	—	—	—	—	—	—	—	—	↓	↓	↓	↓	↓	↓	↓
IOR: Set bit(s) in active FP special reg, bits specified by literal value.	1	—	—	—	—	—	—	—	—	—	↑	↑	↑	↑	↑	↑	↑
TST{.s, .d}: Test (inspect) Fs and update FSR accordingly.	1	↑	↑	↑	↑	↑	—	—	—	—	—	—	—	—	—	—	—
<b>Conversion Instructions</b>																	
LI2F{.s, .d}: Convert 32-bit Integer Fs to FP in Fd	1	—	—	—	—	—	—	—	—	—	0	0	↑	0	0	0	0
DI2F{.s, .d}: Convert 64-bit Integer Fs to FP in Fd	1	—	—	—	—	—	—	—	—	—	0	0	↑	0	0	0	0
F2LI{.s, .d}: Convert FP Fs to 32-bit Integer in Fd	1 for SP 2 for DP	—	—	—	—	—	—	—	—	—	↑	↑	↑	0	0	0	↑
F2DI{.s, .d}: Convert FP Fs to 64-bit Integer in Fd	1 for SP 2 for DP	—	—	—	—	—	—	—	—	—	↑	↑	↑	0	0	0	↑
<b>Comparison Instructions</b>																	
CPS{.s, .d}: FP Compare (Fb - Fs) with signaling exceptions	1	—	—	—	—	—	↑	↑	↑	↑	↑	↑	↑	0	0	0	0
CPQ{.s, .d}: FP Compare (Fb - Fs) with quiet signaling exceptions	1	—	—	—	—	—	↑	↑	↑	↑	↑	↑	↑	0	0	0	↑
FLIM{.s, .d}: Clamp Fd to limits specified by Fb and Fs	1	—	—	—	—	—	—	—	—	—	↑	0	0	0	0	0	↑

<sup>2</sup> Legend:

- ↑ set or cleared, '1' always set, '0' always cleared, — unchanged

- ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)

**Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-2.** FPU Instruction Set Summary (continued)

Instruction <sup>2</sup> : Description	Cycle Count	FPU Status Register (FSR)															
		FSR[28:24]					FSR[19:16]				FSR[6:0]						
		SUB	INF	FN	FZ	FNA N	GT	LT	EQ	UN	SUB0	HUGI	INX	UDF	OVF	DIV0	INVAL
MAX{.s, .d}: Maximum of Fb or Fs moved to Fd: IEEE 754-2019 maximumNumber(x,y)	1	—	—	—	—	—	—	—	—	—	↑↑	0	0	0	0	0	↑↑
MAXNUM{.s, .d}: Maximum of Fb or Fs moved to Fd: IEEE 754-2019 maximum(x,y)	1	—	—	—	—	—	—	—	—	—	↑↑	0	0	0	0	0	↑↑
MIN{.s, .d}: Minimum of Fb or Fs moved to Fd: IEEE 754-2019 minimumNumber(x,y)	1	—	—	—	—	—	—	—	—	—	↑↑	0	0	0	0	0	↑↑
MINNUM{.s, .d}: Minimum of Fb or Fs moved to Fd: IEEE 754-2019 minimum(x,y)	1	—	—	—	—	—	—	—	—	—	↑↑	0	0	0	0	0	↑↑
<b>Math Instructions</b>																	
ADD{.s, .d}: Add: Fb + Fs,	2	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	↑↑	↑↑	0	↑↑
SUB{.s, .d}: Subtract: Fb - Fs,	2	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	↑↑	↑↑	0	↑↑
NEG{.s, .d}: Negate: -1 x Fs,	1	—	—	—	—	—	—	—	—	—	↑↑	—	—	—	—	—	—
ABS{.s, .d}: Absolute value of Fs,	1	—	—	—	—	—	—	—	—	—	↑↑	—	—	—	—	—	—
MUL{.s, .d}: Multiply: Fb x Fs	3	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	↑↑	↑↑	0	↑↑
MAC{.s, .d}: Multiply and Accumulate: Fd = Fd + Fb x Fs	3 for SP 4 for DP	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	↑↑	↑↑	0	↑↑
DIV{.s, .d}: Divide: Fb ÷ Fs,	11 for SP 32 for DP	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	↑↑	↑↑	↑↑	↑↑
SQRT{.s, .d}: Square Root of Fs,	10 for SP 13 for DP	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	0	0	0	↑↑
SIN{.s}: Calculate Sine of Fs (radians),	4	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	↑↑	0	0	↑↑
COS{.s}: Calculate Cosine of Fs (radians), result in Fd	4	—	—	—	—	—	—	—	—	—	↑↑	0	↑↑	↑↑	0	0	↑↑
<b>Branch Instructions</b>																	
FBRA cc: PC branch based on FPU compare status bits.	1 for backward branch (2 or 3) for forward branch	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—

<sup>2</sup> Legend:

- ↑↑ set or cleared, '1' always set, '0' always cleared, — unchanged
  - ↓↓ may be cleared but never set (sticky), ↑↑ may be set but never cleared (sticky)
- Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

**Table 6-3.** FPU Instruction Set Summary

Instruction <sup>3</sup> : Description	Cycle Count	FPU Status Register (FSR)														
		FSR[28:24]					FSR[19:16]				FSR[6:0]					
		SUB	INF	FN	FZ	FNA N	GT	LT	EQ	UN	SUBO	HUGI	INX	UDF	OVF	DIV0
<b>Move instructions</b>																
MOV: Move (Wns+offset) to Fd	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV: Move Fs to (Wnd+offset)	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV	Move (Ws) to Fd or FSR or FCR or FEAR	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV	Move Fs or FSR or FCR or FEAR to (Wd)	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV	Move 32-bit literal value into Fd or FSR or FCR	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—
PUSH	Push Fs or FSR or FCR or FSRH to system stack	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—
POP	Pop Fd or FSR or FCR from system stack	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOV{.s, .d}	Move Fs to Fd	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—
MOVC{.s, .d}	Load constant value into Fd	1	—	—	—	—	—	—	—	—	—	—	—	—	—	—
<b>Status Bit Set/Clear/Update Instructions</b>																
AND	Clear bit(s) in active FP special reg, bits specified by literal value.	1	—	—	—	—	—	—	—	—	↓	↓	↓	↓	↓	↓
IOR	Set bit(s) in active FP special reg, bits specified by literal value.	1	—	—	—	—	—	—	—	—	↑	↑	↑	↑	↑	↑
TST{.s, .d}	Test (inspect) Fs and update FSR accordingly.	1	↔	↔	↔	↔	↔	—	—	—	—	—	—	—	—	—
<b>Conversion Instructions</b>																
LI2F{.s, .d}	Convert 32-bit Integer Fs to FP in Fd	1	—	—	—	—	—	—	—	—	0	0	0	0	0	0
DI2F{.s, .d}	Convert 64-bit Integer Fs to FP in Fd	1	—	—	—	—	—	—	—	—	0	0	0	0	0	0
F2LI{.s, .d}	Convert FP Fs to 32-bit Integer in Fd	1 for SP 2 for DP	—	—	—	—	—	—	—	—	↔	↔	↔	0	0	↔
F2DI{.s, .d}	Convert FP Fs to 64-bit Integer in Fd	1 for SP 2 for DP	—	—	—	—	—	—	—	—	↔	↔	↔	0	0	↔
<b>Comparison Instructions</b>																
CPS{.s, .d}	FP Compare (Fb - Fs) with signaling exceptions	1	—	—	—	—	—	↔	↔	↔	↔	↔	0	0	0	0
CPQ{.s, .d}	FP Compare (Fb - Fs) with quiet signaling exceptions	1	—	—	—	—	—	↔	↔	↔	↔	↔	0	0	0	↔

<sup>3</sup> Legend:

- ↔ set or cleared, '1' always set, '0' always cleared, — unchanged
- ↓ may be cleared but never set (sticky), ↑ may be set but never cleared (sticky)

**Table 6-3.** FPU Instruction Set Summary (continued)

Instruction <sup>3</sup> : Description		Cycle Count	FPU Status Register (FSR)															
			FSR[28:24]				FSR[19:16]				FSR[6:0]							
			SUB	INF	FN	FZ	FNA N	GT	LT	EQ	UN	SUBO	HUGI	INX	UDF	OVF	DIV0	INVAL
FLIM{.s, .d}	Clamp Fd to limits specified by Fb and Fs	1	—	—	—	—	—	—	—	—	—	‡	0	0	0	0	0	‡
MAX{.s, .d}	Maximum of Fb or Fs moved to Fd: IEEE 754-2019 maximumNumber(x,y)	1	—	—	—	—	—	—	—	—	—	‡	0	0	0	0	0	‡
MAXNUM{.s, .d}	Maximum of Fb or Fs moved to Fd: IEEE 754-2019 maximum(x,y)	1	—	—	—	—	—	—	—	—	—	‡	0	0	0	0	0	‡
MIN{.s, .d}	Minimum of Fb or Fs moved to Fd: IEEE 754-2019 minimumNumber(x,y)	1	—	—	—	—	—	—	—	—	—	‡	0	0	0	0	0	‡
MINNUM{.s, .d}	Minimum of Fb or Fs moved to Fd: IEEE 754-2019 minimum(x,y)	1	—	—	—	—	—	—	—	—	—	‡	0	0	0	0	0	‡
<b>Math Instructions</b>																		
ADD{.s, .d}	Add: Fb + Fs,	2	—	—	—	—	—	—	—	—	—	‡	0	‡	‡	‡	0	‡
SUB{.s, .d}	Subtract: Fb - Fs,	2	—	—	—	—	—	—	—	—	—	‡	0	‡	‡	‡	0	‡
NEG{.s, .d}	Negate: -1 x Fs,	1	—	—	—	—	—	—	—	—	—	‡	—	—	—	—	—	—
ABS{.s, .d}	Absolute value of Fs,	1	—	—	—	—	—	—	—	—	—	‡	—	—	—	—	—	—
MUL{.s, .d}	Multiply: Fb x Fs	3	—	—	—	—	—	—	—	—	—	‡	0	‡	‡	‡	0	‡
MAC{.s, .d}	Multiply and Accumulate: Fd = Fd + Fb x Fs	3 for SP 4 for DP	—	—	—	—	—	—	—	—	—	‡	0	‡	‡	‡	0	‡
DIV{.s, .d}	Divide: Fb ÷ Fs,	11 for SP 32 for DP	—	—	—	—	—	—	—	—	—	‡	0	‡	‡	‡	‡	‡
SQRT{.s, .d}	Square Root of Fs,	10 for SP 13 for DP	—	—	—	—	—	—	—	—	—	‡	0	‡	0	0	0	‡
SIN{.s}	Calculate Sine of Fs (radians),	4	—	—	—	—	—	—	—	—	—	‡	0	‡	‡	0	0	‡
COS{.s}	Calculate Cosine of Fs (radians), result in Fd	4	—	—	—	—	—	—	—	—	—	‡	0	‡	‡	0	0	‡
<b>Branch Instructions</b>																		
FBRA cc	PC branch based on FPU compare status bits.	1 for backward branch (2 or 3) for forward branch	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	

**Note:** SA and/or SB only modified if corresponding saturation bit is enabled, otherwise unchanged.

## 7. Revision History

The revision history describes the changes that were implemented in the document. The changes are listed by revision, starting with the most current publication.

Revision	Date	Description
A	May 2025	This is the initial release of this document. Originally combined with the dsPIC33A Programmer's Reference Manual.

# Microchip Information

## Trademarks

The "Microchip" name and logo, the "M" logo, and other names, logos, and brands are registered and unregistered trademarks of Microchip Technology Incorporated or its affiliates and/or subsidiaries in the United States and/or other countries ("Microchip Trademarks"). Information regarding Microchip Trademarks can be found at <https://www.microchip.com/en-us/about/legal-information/microchip-trademarks>.

ISBN: 979-8-3371-1269-5

## Legal Notice

This publication and the information herein may be used only with Microchip products, including to design, test, and integrate Microchip products with your application. Use of this information in any other manner violates these terms. Information regarding device applications is provided only for your convenience and may be superseded by updates. It is your responsibility to ensure that your application meets with your specifications. Contact your local Microchip sales office for additional support or, obtain additional support at [www.microchip.com/en-us/support/design-help/client-support-services](http://www.microchip.com/en-us/support/design-help/client-support-services).

THIS INFORMATION IS PROVIDED BY MICROCHIP "AS IS". MICROCHIP MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WHETHER EXPRESS OR IMPLIED, WRITTEN OR ORAL, STATUTORY OR OTHERWISE, RELATED TO THE INFORMATION INCLUDING BUT NOT LIMITED TO ANY IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE, OR WARRANTIES RELATED TO ITS CONDITION, QUALITY, OR PERFORMANCE.

IN NO EVENT WILL MICROCHIP BE LIABLE FOR ANY INDIRECT, SPECIAL, PUNITIVE, INCIDENTAL, OR CONSEQUENTIAL LOSS, DAMAGE, COST, OR EXPENSE OF ANY KIND WHATSOEVER RELATED TO THE INFORMATION OR ITS USE, HOWEVER CAUSED, EVEN IF MICROCHIP HAS BEEN ADVISED OF THE POSSIBILITY OR THE DAMAGES ARE FORESEEABLE. TO THE FULLEST EXTENT ALLOWED BY LAW, MICROCHIP'S TOTAL LIABILITY ON ALL CLAIMS IN ANY WAY RELATED TO THE INFORMATION OR ITS USE WILL NOT EXCEED THE AMOUNT OF FEES, IF ANY, THAT YOU HAVE PAID DIRECTLY TO MICROCHIP FOR THE INFORMATION.

Use of Microchip devices in life support and/or safety applications is entirely at the buyer's risk, and the buyer agrees to defend, indemnify and hold harmless Microchip from any and all damages, claims, suits, or expenses resulting from such use. No licenses are conveyed, implicitly or otherwise, under any Microchip intellectual property rights unless otherwise stated.

## Microchip Devices Code Protection Feature

Note the following details of the code protection feature on Microchip products:

- Microchip products meet the specifications contained in their particular Microchip Data Sheet.
- Microchip believes that its family of products is secure when used in the intended manner, within operating specifications, and under normal conditions.
- Microchip values and aggressively protects its intellectual property rights. Attempts to breach the code protection features of Microchip products are strictly prohibited and may violate the Digital Millennium Copyright Act.
- Neither Microchip nor any other semiconductor manufacturer can guarantee the security of its code. Code protection does not mean that we are guaranteeing the product is "unbreakable". Code protection is constantly evolving. Microchip is committed to continuously improving the code protection features of our products.